



- Expert Verified, Online, **Free**.

Custom View Settings

Topic 1 - Single Topic

Question #1

Every employee of your company has a Google account. Your operational team needs to manage a large number of instances on Compute Engine. Each member of this team needs only administrative access to the servers. Your security team wants to ensure that the deployment of credentials is operationally efficient and must be able to determine who accessed a given instance. What should you do?

- A. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key in the metadata of each instance.
- B. Ask each member of the team to generate a new SSH key pair and to send you their public key. Use a configuration management tool to deploy those keys on each instance.
- C. Ask each member of the team to generate a new SSH key pair and to add the public key to their Google account. Grant the `compute.osAdminLogin` role to the Google group corresponding to this team.
- D. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key as a project-wide public SSH key in your Cloud Platform project and allow project-wide public SSH keys on each instance.

Correct Answer: D

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

Community vote distribution

C (100%)

✉  **dan80**  2 years, 5 months ago

C is correct - <https://cloud.google.com/compute/docs/instances/managing-instance-access>
upvoted 50 times

✉  **adedj99** 2 years ago

We recommend collecting users with the same responsibilities into groups and assigning IAM roles to the groups rather than to individual users. For example, you can create a "data scientist" group and assign appropriate roles to enable interaction with BigQuery and Cloud Storage. When a new data scientist joins your team, you can simply add them to the group and they will inherit the defined permissions. You can create and manage groups through the Admin Console.

upvoted 6 times

✉  **zakhili**  2 years, 5 months ago

Send private key to users is not safe, i think it's C
upvoted 18 times

✉  **anandvpm** 2 months, 1 week ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my ananddrajagopal@gmail.com address. Thanks in advance!
upvoted 2 times

✉  **Nipun25** 3 days, 23 hours ago

Have you got the whole question set? if yes please mail to nipunbanga972@gmail.com
upvoted 1 times

✉  **jyotimishra1**  1 day, 21 hours ago

I am preparing for GCP Associate engineer exam, if anyone has the entire questions please share them to my jyotimishra.jec@gmail.com address.
upvoted 1 times

✉  **Edideng** 4 days, 21 hours ago

I am preparing for GCP Associate cloud engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my email kaldereta55@gmail.com . Thanks in advance!
upvoted 1 times

✉  **gyani_06** 6 days, 1 hour ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my email-gautamnitisha@gmail.com address. Thanks in advance!
upvoted 1 times

✉  **Nipun25** 3 days, 23 hours ago

Have you got the whole question set? if yes please mail to nipunbanga972@gmail.com
upvoted 1 times

✉  **Vismaya** 1 week, 5 days ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my shinesree.m@gmail.com address. Thanks in advance!

upvoted 1 times

 **Samsom6** 2 weeks ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please kindly share them to my gkadeoti@yahoo.com address. Thanks in advance!

upvoted 1 times

 **thenmozhi12** 2 weeks ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my thenmozhi.maniskm@gmail.com address. Thanks in advance!

upvoted 1 times

 **leogor** 2 weeks, 3 days ago

C is correct, cuz it doesn't need to grant access to specific user / instance

upvoted 1 times

 **mudassirsmk27** 3 weeks ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my mudassirsmk27@gmail.com address. Thanks in advance!

upvoted 1 times

 **Bothell** 3 weeks, 6 days ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my ashishwakhare1234@yahoo.com address. Thanks in advance!

upvoted 1 times

 **vamosspoli76** 1 month ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my a.polisini@libero.it address. Thanks in advance!

upvoted 1 times

 **davidsalomon** 1 month ago

if the answer is C, why does it still display D as the right answer?

upvoted 1 times

 **sheffy** 1 month ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my sheffybatra29@gmail.com address. Thanks in advance!

upvoted 2 times

 **rwl555** 1 month, 1 week ago

I was passing exam last week - almost the same questions as in pdf.

upvoted 1 times

 **Aras123** 1 month, 1 week ago

Hi...I have access to 92 questions only, Please send the PDF to prathu83@gmail.com

upvoted 1 times

 **IDpka** 1 month, 1 week ago

Hi all, I am preparing for the GCP exam , I am unable to access all the questions here, if anyone has the entire questions pdf, please send it to my narasimhak950@gmail.com thanks in advance

upvoted 1 times

Question #2

You need to create a custom VPC with a single subnet. The subnet's range must be as large as possible. Which range should you use?

- A. 0.0.0.0/0
- B. 10.0.0.0/8
- C. 172.16.0.0/12
- D. 192.168.0.0/16

Correct Answer: A*Community vote distribution*

B (94%)	6%
---------	----

 **CarIS** Highly Voted 2 years, 7 months ago

I got this question in one of the trainings I did in udemy: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?referralCode=810D02D4A159FC3E36CC>

B is correct. Pay attention to the question, is talking about custom VPC subnet and is not mentioning you will use automatic subnet mode creation. If you set subnet to custom, the minimum size is /8.

upvoted 53 times

 **passnow** Highly Voted 2 years, 7 months ago

B is legit as fuck!

upvoted 24 times

 **FunkyB** 4 months, 3 weeks ago

Language. :-(

upvoted 2 times

 **hems4all** 2 years ago

B is correct

Use 10.0.0.0/8 CIDR range. is the right answer.

The private network range is defined by IETF (Ref: <https://tools.ietf.org/html/rfc1918>) and adhered to by all cloud providers. The supported internal IP Address ranges are

1. 24-bit block 10.0.0.0/8 (16777216 IP Addresses)
2. 20-bit block 172.16.0.0/12 (1048576 IP Addresses)
3. 16-bit block 192.168.0.0/16 (65536 IP Addresses)

10.0.0.0/8 gives you the most extensive range - 16777216 IP Addresses.

upvoted 11 times

 **dennydream** Most Recent 1 week, 5 days ago

I'd think A, but it's really B. A is not allowed on GCP

upvoted 1 times

 **leogor** 2 weeks, 3 days ago

0.0.0.0/0 -_-

Ans is B

upvoted 1 times

 **hiromi** 2 weeks, 6 days ago

Selected Answer: B

B is correct:

<https://cloud.google.com/vpc/docs/subnets#valid-ranges>

upvoted 1 times

 **Sgegmund** 1 month, 1 week ago

Selected Answer: B

0.0.0.0/0 ??? Really ? B!

upvoted 1 times

 **Sgegmund** 1 month, 1 week ago

0.0.0.0/0 ??? Who selected the answers to these questions ?

B

upvoted 1 times

 **ashit44244** 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

 **adarsh4503** 1 month, 2 weeks ago

Type "examtopics gcp ace question __", and the number of the question you want in the blank space, you can access all the questions and discussion like that.

upvoted 2 times

 **braj1234** 2 months, 1 week ago

Selected Answer: B

B is Correct

upvoted 1 times

 **kkrzakk** 2 months, 1 week ago

Selected Answer: B

Even checked in console to ensure that A is wrong

upvoted 1 times

 **Arulkumar** 2 months, 1 week ago

Google doesn't allow subnets with ip range 0.0.0.0/0 Refer section https://cloud.google.com/vpc/docs/subnetsProhibited_IPv4_subnet_ranges

upvoted 1 times

 **pkmdb66** 3 months ago

Selected Answer: B

It's B

upvoted 1 times

 **imperadoor** 3 months, 1 week ago

Selected Answer: B

B is correct, 10.0.0.0/8

Ref: RFC 1918

upvoted 1 times

 **imperadoor** 3 months, 1 week ago

B. 10.0.0.0/8 its correct

Ref: RFC 1918

upvoted 1 times

 **abirroy** 3 months, 2 weeks ago

Selected Answer: B

Answer is B

upvoted 1 times

 **kecy992** 4 months ago

Selected Answer: B

B is correct

upvoted 1 times

 **orious** 4 months, 1 week ago

Selected Answer: A

A is correct.

Pay attention to the question. You need the maximum IPs possible! There is any IP class constraint in the question.

While 10.0.0.0/8 is per standard (RFC1918), the minimal CIDR for private subnetworks, It is possible to customize your IP range to 0.0.0.0/0 (non-RFC1918). However, any egress traffic will remain within the VPC.

upvoted 1 times

Question #3

You want to select and configure a cost-effective solution for relational data on Google Cloud Platform. You are working with a small set of operational data in one geographic location. You need to support point-in-time recovery. What should you do?

- A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.
- B. Select Cloud SQL (MySQL). Select the create failover replicas option.
- C. Select Cloud Spanner. Set up your instance with 2 nodes.
- D. Select Cloud Spanner. Set up your instance as multi-regional.

Correct Answer: A

Reference:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/restore>*Community vote distribution*

A (100%)

 **YashBindlish** Highly Voted 2 years, 6 months ago

A is Correct. You must enable binary logging to use point-in-time recovery. Enabling binary logging causes a slight reduction in write performance. <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 40 times

 **ryumada** 3 months, 2 weeks ago

In this link below, the docs explains clearly that point-in-time recovery requires binary logging.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr#disk-usage>

upvoted 2 times

 **Bharathy** Highly Voted 2 years, 7 months ago

A is correct, as Binary Logging enables Point in Recovery in Cloud SQL

upvoted 10 times

 **leogor** Most Recent 2 weeks, 3 days ago

A, small amount of operation data doesnt need cloud spanner

upvoted 1 times

 **eledu1985** 2 weeks, 5 days ago

the most accurate link is:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr#disk-usage>

upvoted 1 times

 **Sgegmund** 1 month, 1 week ago

Selected Answer: A

PiT Recovery : A

upvoted 1 times

 **GayuSundar** 2 months, 1 week ago

Answer A: small set of operational data in one geographic location - CloudSQL (Cloud Spanner generally to hold large data and global) Point-in-time recovery means binary logging should be enabled.

upvoted 1 times

 **HILMY** 3 months ago

-cost effective (eliminating point C and D)
-point-in-time recovery (Binary Logging)

So the answer is A

upvoted 1 times

 **abirroy** 3 months, 2 weeks ago

Selected Answer: A

Answer is A

upvoted 1 times

 **kecy992** 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **patashish** 4 months ago

A is correct.
Two ways
1) using timestamps
2) by providing a specific binary log position in a binary log file.
Ref : <https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr>

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is correct
upvoted 1 times

 **haroldbenites** 5 months, 3 weeks ago

go for A.
upvoted 1 times

 **ebrulato** 8 months, 2 weeks ago

Selected Answer: A

A gustas
upvoted 1 times

 **harishbohra089** 9 months, 1 week ago

A is correct
upvoted 1 times

 **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

 **SleepyHitman** 10 months ago

https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr#perform_the_point-in-time_recovery_using_binary_log_positions
upvoted 1 times

 **koxow74213** 11 months, 2 weeks ago

Selected Answer: A

One location - MySQL no doubt.
Between A & B, Point in time recovery needs binary logs. So A
upvoted 5 times

Question #4

You want to configure autohealing for network load balancing for a group of Compute Engine instances that run in multiple zones, using the fewest possible steps.

You need to configure re-creation of VMs if they are unresponsive after 3 attempts of 10 seconds each. What should you do?

- A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)
- B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.
- C. Create a managed instance group. Set the Autohealing health check to healthy (HTTP)
- D. Create a managed instance group. Verify that the autoscaling setting is on.

Correct Answer: D*Community vote distribution*

C (100%)

 **ReyBan** Highly Voted 2 years, 6 months ago

C, Agreed

reference : <https://cloud.google.com/compute/docs/tutorials/high-availability-autohealing>

Pro Tip: Use separate health checks for load balancing and for autohealing. Health checks for load balancing detect unresponsive instances and direct traffic away from them. Health checks for autohealing detect and recreate failed instances, so they should be less aggressive than load balancing health checks. Using the same health check for these services would remove the distinction between unresponsive instances and failed instances, causing unnecessary latency and unavailability for your users.

upvoted 64 times

 **ashrafh** 1 year, 2 months ago

I also vote for C

go to gcp console create a httpa load balancer and in the health check settings take your mouse to question mark it says

""Ensures that requests are sent only to instances that are up and running""

so its not recreating, if the vm not working it redirect to one which work.

go to gpc console create MIG and check the questions mark of Autohealing health check settings it says

"""Autohealing allows recreating VM instances when needed. You can use a health check to recreate a VM instance if the health check finds it unresponsive. If you don't select a health check, Compute Engine will recreate VM instances only when they're not running."""

hope this help :)

upvoted 13 times

 **bryanchew** Highly Voted 2 years, 7 months ago

A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy(HTTP)

This is a possible answer. This answer assumes that the existing backend is configured correctly.

B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.

This is a possible answer. This answer assumes that the existing backend is configured correctly. This answer adds an additional step over answer A.

C. Create a managed instance group. Set the Autohealing health check to healthy(HTTP)

This is only a partial solution. The default configuration is auto scaling enabled. You still need to create the HTTP Load Balancer.

D. Create a managed instance group. Verify that the auto scaling setting is on.

This is only a partial solution. Creating a Managed Instance Group with Auto Scaling is required, but you still need to create the HTTP Load Balancer.

Therefore the best answer is A in my opinion.

upvoted 19 times

 **Ridhanya** 11 months, 2 weeks ago

it cannot be option A because as you said, load balancer with the health check is already present and now the problem is simply auto healing. so we need to focus only on recreation which can happen only if option C is correct

upvoted 1 times

 **koniec** 2 years ago

It's A.

Managed group already exists so create a LB with health checks.

If you go for C, you will have to create a LB anyway so it's more steps to achieve the goal

upvoted 1 times

tavva_prudhvi 1 year, 7 months ago

https://www.youtube.com/watch?v=dT7xDEtALPQ&list=PLlivdWyY5sqljj_cglNUHZDMnGjVx3rx&index=36

step-1: go to the instance group

step-2: click edit

step-3: scroll down you will see auto-healing off by default change to ON

step-4: create a health check saying 10 seconds as CHECK INTERVAL and UNHEALTHY THRESHOLD=3

upvoted 5 times

pYWORLD 1 year, 3 months ago

I hope you saw that in the questions is stated: "network load balancing"

upvoted 1 times

DickDastardly 1 year, 8 months ago

It can't be A as a load balancer does not re-create unhealthy instances, as per the requirement.

Has to be C

<https://cloud.google.com/compute/docs/instance-groups>

upvoted 6 times

anirban7172 [Most Recent] 2 days, 21 hours ago

The correct Answer is C;

upvoted 1 times

eledu1985 2 weeks, 5 days ago

Selected Answer: C

The correct answer is C as the autoscaling is on by default but if you want to recreate an instance you will need in any case a Health check, for every other option.

upvoted 1 times

Sgegmund 1 month, 1 week ago

C is correct

upvoted 1 times

ashit44244 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

gcp_world123 2 months, 3 weeks ago

C is correct

upvoted 1 times

imperadoor 3 months, 1 week ago

Selected Answer: C

The Answer correct is C, is the only one with the AUTOHEALING option

upvoted 1 times

SaiSaiA 3 months, 3 weeks ago

C is the only one with the AUTOHEALING option, but it is not really correct. Remember, the GIVEN information are "a NETWORK load balancer and a group of Compute Engine Instances that run in multiple zones" which gives us an idea that the existing configuration is a target pool-based network lb.

If we are to use the existing group of VMs, we need to choose UNMANAGED Instance Group, UNMANAGED Instance Group does not have Autohealing, only a health check. Health check only checks if VMs are responsive or not but does not re-create instances as what Autohealing and Autoscaling do.

You can also try re-creating the scenario or check this

<https://cloud.google.com/load-balancing/docs/network/transition-to-backend-services#console>

So, if a MANAGED INSTANCE group is to be used, then you need to create an instance template and use it for your MIG. Ofc, you cannot use the existing VMs, then you create a new load balancer. Ofc, the existing group of VMs mentioned in the question will no longer be used but rather a new set of VMs based on the instance template will be created. The choices should be updated.

upvoted 2 times

AzureDP900 4 months, 2 weeks ago

C is the correct answer

upvoted 1 times

Himadhar1997 5 months, 2 weeks ago

C is correct option

upvoted 1 times

 **haroldbenites** 5 months, 3 weeks ago

go for C.

<https://cloud.google.com/compute/docs/tutorials/high-availability-autohealing>

upvoted 2 times

 **kiemkhach** 6 months, 2 weeks ago

C is correct. Autohealing is not from Load Balancer. D is for autoscaling which when memory, cpu, etc get to a certain level and create new or drop instances.

upvoted 1 times

 **Tx4free** 7 months, 4 weeks ago

Selected Answer: C

Is the correct option

upvoted 1 times

 **JohnPi** 8 months, 1 week ago

Selected Answer: C

it should be C

upvoted 1 times

 **Eminenza22** 9 months ago

C is correct

<https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs>

upvoted 1 times

 **Eminenza22** 9 months ago

Create a health check for autohealing that is more conservative than a load balancing health check.

upvoted 1 times

 **Williehk** 9 months, 1 week ago

A, not C

Question stats "for network load balancing for a group of Compute Engine instances"

upvoted 1 times

Question #5

You are using multiple configurations for gcloud. You want to review the configured Kubernetes Engine cluster of an inactive configuration using the fewest possible steps. What should you do?

- A. Use gcloud config configurations describe to review the output.
- B. Use gcloud config configurations activate and gcloud config list to review the output.
- C. Use kubectl config get-contexts to review the output.
- D. Use kubectl config use-context and kubectl config view to review the output.

Correct Answer: D

Reference:

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

Community vote distribution

D (100%)

✉  **zukko78** Highly Voted 2 years, 6 months ago

D is correct

upvoted 35 times

✉  **nhusain** 1 year, 6 months ago

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

explains it well

upvoted 6 times

✉  **poogcp** Highly Voted 2 years, 5 months ago

C is correct , Use kubectl config get-contexts to review the output : shows the clusters and the configurations and based on the output we can identify the inactive configurations

upvoted 19 times

✉  **Gurnoor** 2 years, 4 months ago

This is wrong get-contexts does not show clusters it only shows contexts.

upvoted 8 times

✉  **jilly** 2 years, 4 months ago

True .

Will give only below results

kubectl config get-contexts

CURRENT NAME CLUSTER AUTHINFO NAMESPACE

* white white dazwilkin

black black dazwilkin

upvoted 1 times

✉  **abirroy** Most Recent 3 months, 2 weeks ago

Selected Answer: D

Right answer is D

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

D is right

upvoted 1 times

✉  **guruanytime** 5 months ago

D is right

upvoted 1 times

✉  **kiemkhach** 6 months, 2 weeks ago

D is correct. A lot details mentioned in this group. Here I only say about eliminating answers. As we go down to between C and D. The question is want to review a inactive config. So, to me, C is viewing info about all config while D is apply a specific config and viewing it. So I eliminate C to go with D.

upvoted 4 times

✉  **janukaz** 7 months, 4 weeks ago

Selected Answer: D

Answer bouncing among the C and D. Here I explain original results from C code,

If we executed "kubectl config get-contexts" code, Output will be,
CURRENT NAME | NAMESPACE | CLUSTER | AUTHINFO

upvoted 1 times

pranavdhopey 9 months, 2 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

NYF 9 months, 4 weeks ago

Selected Answer: D

```
kubectl config view -o jsonpath='{.users[0].name}' # display the first user
kubectl config view -o jsonpath='{.users[*].name}' # get a list of users
kubectl config get-contexts # display list of contexts
kubectl config current-context # display the current-context
kubectl config use-context my-cluster-name # set the default context to my-cluster-name
```

<https://kubernetes.io/docs/reference/kubectl/cheatsheet/>

upvoted 4 times

jabrrJ68w02ond1 11 months, 2 weeks ago

Selected Answer: D

"get-contexts" shows us our Kubernetes cluster contexts, that's right. But the question says that you want to review the cluster itself, so you need to use-context to get into the cluster.

upvoted 4 times

shawnkkk 11 months, 3 weeks ago

D is correct,

upvoted 3 times

vishnukumartr 11 months, 3 weeks ago

D. Use kubectl config use-context and kubectl config view to review the output.

upvoted 1 times

Jaira1256 11 months, 3 weeks ago

Selected Answer: D

And -D

upvoted 1 times

Bukowski 1 year ago

why not A?

<https://cloud.google.com/sdk/gcloud/reference/config/configurations/describe>

SYNOPSIS

gcloud config configurations describe CONFIGURATION_NAME [--all] [GCLOUD_WIDE_FLAG ...]

DESCRIPTION

Describes a named configuration by listing its properties.

upvoted 4 times

Pravallika123 1 year, 1 month ago

D is correct

upvoted 3 times

JustJack21 1 year, 2 months ago

The question says “multiple gcloud configurations” . That has nothing to do with Kubernetes. This question could have been for ANYTHING. The answer is “B” to activate the correct configuration in gloud SDK.

upvoted 4 times

Shruti_Pal 1 year, 2 months ago

D is right

upvoted 1 times

Question #6

Your company uses Cloud Storage to store application backup files for disaster recovery purposes. You want to follow Google's recommended practices. Which storage option should you use?

- A. Multi-Regional Storage
- B. Regional Storage
- C. Nearline Storage
- D. Coldline Storage

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/storage-classes#nearline>

Community vote distribution

D (86%)	14%
---------	-----

✉  **poogcp** Highly Voted 2 years, 5 months ago

Best Answer is " Archive Storage "

<https://cloud.google.com/storage/docs/storage-classes>

But as per the given option next best solution is " Coldline Storage"

upvoted 44 times

✉  **Mutune** 1 year, 8 months ago

Perfectly stated

upvoted 5 times

✉  **zukko78** Highly Voted 2 years, 6 months ago

D is correct,

Coldline Storage COLDLINE 90 days

99.95% in multi-regions and dual-regions

99.9% in regions

upvoted 9 times

✉  **dasun_eranga** Most Recent 3 days, 8 hours ago

correct answer is Archive Storage

watch this video - https://www.cloudskillsboost.google/course_sessions/1672384/video/326664

upvoted 1 times

✉  **leogor** 2 weeks, 3 days ago

D, coldline is used if there is no option for archive

upvoted 1 times

✉  **GaneshSurwase** 3 weeks, 4 days ago

correct ans is D as per te new doc announcement

<https://cloud.google.com/blog/products/gcp/introducing-coldline-and-a-unified-platform-for-data-storage>

upvoted 1 times

✉  **GaneshSurwase** 3 weeks, 4 days ago

but doc says - Archive storage is the lowest-cost, highly durable storage service for data archiving, online backup, and disaster recovery.

<https://cloud.google.com/storage/docs/storage-classes#:~:text=Archive%20storage%20is%20the%20lowest,online%20backup%2C%20and%20disaster%20recovery.>

upvoted 1 times

✉  **Sgegmund** 1 month, 1 week ago

Selected Answer: D
Coldline Storage for Disaster Recovery

upvoted 1 times

✉  **abirroy** 3 months ago

Selected Answer: D
Cloud Storage Coldline: a low-latency storage class for long-term archiving

Coldline is a new Cloud Storage class designed for long-term archival and disaster recovery. Coldline is perfect for the archival needs of big data or multimedia content, allowing businesses to archive years of data. Coldline provides fast and instant (millisecond) access to data and changes the way that companies think about storing and accessing their cold data.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Since there is no archive mentioned cold line storage is best option.

upvoted 1 times

 **Miwado** 4 months, 4 weeks ago

D coldline next to archive

upvoted 1 times

 **huwaku** 7 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **Ery** 6 months, 3 weeks ago

why u chose a?

upvoted 1 times

 **Tx4free** 7 months, 4 weeks ago

Selected Answer: D

Coldline is the best option

upvoted 1 times

 **PLUTONIUM54** 9 months, 2 weeks ago

D is correct

upvoted 1 times

 **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

 **[Removed]** 11 months ago

D: coldline

upvoted 1 times

 **FernandoJ** 11 months, 2 weeks ago

Selected Answer: D

The recommended option is Archive, but the next best option in this list is Coldline

<https://cloud.google.com/storage/docs/storage-classes#:~:text=Archive%20Storage%20is%20the%20lowest%2Dcost%2C%20highly%20durable%20storage%20service%20for%20data%20archiving%2C%20online%20backup%2C%20and%20disaster%20recovery>

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

D. Coldline Storage

upvoted 1 times

Question #7

Several employees at your company have been creating projects with Cloud Platform and paying for it with their personal credit cards, which the company reimburses. The company wants to centralize all these projects under a single, new billing account. What should you do?

- A. Contact cloud-billing@google.com with your bank account details and request a corporate billing account for your company.
- B. Create a ticket with Google Support and wait for their call to share your credit card details over the phone.
- C. In the Google Platform Console, go to the Resource Manager and move all projects to the root Organization.
- D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

Correct Answer: D

Reference:

<https://www.whizlabs.com/blog/google-cloud-interview-questions/>

Community vote distribution

D (69%)	C (31%)
---------	---------

✉  **samcat84** Highly Voted 2 years, 5 months ago

C is incomplete. Moving projects under an organization doesn't change their linked billing project.
<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>

Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet.

D is incomplete as well, after setting the billing account in the organization you need to link the projects to the new billing account.
upvoted 42 times

✉  **sarahf** 1 year, 10 months ago

I agree that neither C or D is correct. I did the cert a month ago and this question was not on it. Although a similar question about how to change the payment method from your own card in your project to the company's "card". So they might have removed this one.
upvoted 2 times

✉  **ehizren** 1 year, 9 months ago

What's was the answer you chose for your particular exam question?
upvoted 7 times

✉  **GokulVelusaamy** 3 days, 5 hours ago

We need to add a new payment method and need to set that as Primary, post that we need to remove the previous one
"If you want to remove a payment method, you should add a new payment method first."

Refer : <https://cloud.google.com/billing/docs/how-to/payment-methods>

upvoted 1 times

✉  **RazOr** 9 months, 2 weeks ago

The given answers make D the only possible solution. C can not be right, you all need to look it up here: https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account

upvoted 2 times

✉  **ryumada** 3 months, 1 week ago

This link explains clearly that move a project won't affect billing.
<https://cloud.google.com/resource-manager/docs/project-migration#permissions-billing>
upvoted 3 times

✉  **poogcp** Highly Voted 2 years, 5 months ago

C is correct Answer. there will be 1 billing account for the organization and all projects under that organization are linked to single billing account.
upvoted 13 times

✉  **arathefu** 7 months, 3 weeks ago

https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account
"Moving a project from one organization to another won't impact billing, and charges will continue against the old billing account."
upvoted 5 times

✉  **leogor** Most Recent 2 weeks, 3 days ago

D, as other options are not correct
upvoted 1 times

✉  **hiromi** 2 weeks, 6 days ago

Selected Answer: D

D is more complete than C

upvoted 1 times

✉ **AwesomeGCP** 1 month ago

Selected Answer: D

It's pretty straight-forward: you should establish a new billing account with a company-based payment method. Then set all the projects to use that new billing account.

upvoted 1 times

✉ **Rajagopal** 1 month, 1 week ago

I think "D" is the right answer (https://cloud.google.com/resource-manager/docs/project-migration#change_billing_account)

upvoted 2 times

✉ **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

✉ **ale_brd_** 2 months, 1 week ago

Selected Answer: D

D is the correct one, even tho could be written in a better way.

upvoted 1 times

✉ **Qureshizaid64** 2 months, 1 week ago

Selected Answer: C

If we move all projects under the root organization hierarchy, they still need to modify to use a billing account within the organization.

Ref: <https://cloud.google.com/resource-manager/docs/migrating-projects-billing#top of page> Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization, they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet.

But in this option, all projects are in the organization resource hierarchy so the organization can uniformly apply organization policies to all its projects which is a Google recommended practice. So this is the better of the two options.

Ref: <https://cloud.google.com/billing/docs/concepts>

upvoted 2 times

✉ **akg001** 2 months, 4 weeks ago

Selected Answer: D

D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

upvoted 1 times

✉ **habros** 3 months, 1 week ago

D

<https://cloud.google.com/billing/docs/concepts>

upvoted 1 times

✉ **pedroyrosa5700** 3 months, 3 weeks ago

I took the exam twice, it's this type of page, and I haven't passed it XD hahaha

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

In the GCP Console, navigate to the Resource Manage section and move all projects to the root Organization. C is right

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

C is right .. o In the GCP Console, navigate to the Resource Manage section and move all projects to the root Organization. Tutorial Dojo practice questions having the similar question.

upvoted 1 times

✉ **Paulv82003** 3 months, 1 week ago

Tutorial Dojo's question mentions Organization in the question, here no mention of Org, assuming org does not exist. Just link billing account to each project.

upvoted 1 times

✉ **dinesh198728** 7 months, 3 weeks ago

Selected Answer: D

Catch the word : new billing account

upvoted 1 times

✉ **janukaz** 7 months, 4 weeks ago

The answer is D.

Carefully read this sentence "Several employees at your company have been creating projects with Cloud Platform and paying for it with their personal credit cards". all projects are already in the root org. no need to move especially. but the thing is they are paying individually. So need to solve this should create a new account and move it to all projects.

upvoted 3 times

 **tigerbaer** 9 months ago

Selected Answer: D

<https://cloud.google.com/resource-manager/docs/project-migration>

upvoted 1 times

Question #8

You have an application that looks for its licensing server on the IP 10.0.3.21. You need to deploy the licensing server on Compute Engine. You do not want to change the configuration of the application and want the application to be able to reach the licensing server. What should you do?

- A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.
- B. Reserve the IP 10.0.3.21 as a static public IP address using gcloud and assign it to the licensing server.
- C. Use the IP 10.0.3.21 as a custom ephemeral IP address and assign it to the licensing server.
- D. Start the licensing server with an automatic ephemeral IP address, and then promote it to a static internal IP address.

Correct Answer: A*Community vote distribution*

A (100%)

 **zukko78** Highly Voted 2 years, 6 months ago

it's obvious, A
upvoted 22 times

 **Khaled_Rashwan** Highly Voted 2 years, 5 months ago

A
IP 10.0.3.21 is internal by default, and to ensure that it will be static non-changing it should be selected as static internal ip address.
upvoted 20 times

 **ashit44244** Most Recent 1 month, 3 weeks ago

Its not mentioned anywhere that the application is deployed in GCP or its external
upvoted 2 times

 **Chris_91_DE** 2 weeks, 5 days ago

I went into the same trap.
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Static IP point to internal .. Option A is right
upvoted 1 times

 **rajsharma2108** 5 months, 2 weeks ago

A
obvious answer
upvoted 1 times

 **akshaychavan7** 5 months, 3 weeks ago

If they haven't mentioned that the application has been deployed internally(inside GCP) or externally then how can we determine if the server IP address should be internal or external? -_-
upvoted 4 times

 **knsujith** 5 months, 3 weeks ago

Selected Answer: A
A, is correct option
upvoted 1 times

 **Alejondri** 6 months, 2 weeks ago

Selected Answer: A
A, obvio
upvoted 1 times

 **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: A
A is correct
upvoted 1 times

 **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: A
A is correct
upvoted 1 times

 **Abdul88** 9 months, 3 weeks ago

Selected Answer: A

A 100%

upvoted 1 times

  **jabrrJ68w02ond1** 11 months, 2 weeks ago

A is correct. By RFC 1918, IP range 10.0.0.0/8 is reserved for internal usage.

upvoted 4 times

  **shawnkkk** 11 months, 3 weeks ago

A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

upvoted 1 times

  **vishnukumartr** 11 months, 3 weeks ago

A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

upvoted 1 times

  **Jaira1256** 11 months, 3 weeks ago

Ans - A

upvoted 1 times

  **Ajay_Katkade** 1 year, 1 month ago

A IS RIGHT

upvoted 1 times

  **Finger41** 1 year, 6 months ago

Its A, if its in the same subnet you can use the internal IP, which you would want as a static IP. There is no point of connecting to the server via the internet.

upvoted 3 times

Question #9

You are deploying an application to App Engine. You want the number of instances to scale based on request rate. You need at least 3 unoccupied instances at all times. Which scaling type should you use?

- A. Manual Scaling with 3 instances.
- B. Basic Scaling with min_instances set to 3.
- C. Basic Scaling with max_instances set to 3.
- D. Automatic Scaling with min_idle_instances set to 3.

Correct Answer: D

Reference:

<https://cloud.google.com/appengine/docs/standard/python/how-instances-are-managed>

Community vote distribution

D (100%)

 **zukko78** Highly Voted 2 years, 6 months ago

D is correct.

App Engine supports the following scaling types, which controls how and when instances are created:

Automatic

Basic

Manual

You specify the scaling type in your app's app.yaml.

Automatic scaling

Automatic scaling creates instances based on request rate, response latencies, and other application metrics. You can specify thresholds for each of these metrics, as well as a minimum number instances to keep running at all times.

upvoted 42 times

 **Finger41** Highly Voted 1 year, 6 months ago

D is correct : <https://cloud.google.com/appengine/docs/standard/go/config/appref>

"App Engine calculates the number of instances necessary to serve your current application traffic based on scaling settings such as target_cpu_utilization and target_throughput_utilization. Setting min_idle_instances specifies the number of instances to run in addition to this calculated number. For example, if App Engine calculates that 5 instances are necessary to serve traffic, and min_idle_instances is set to 2, App Engine will run 7 instances (5, calculated based on traffic, plus 2 additional per min_idle_instances)."

upvoted 12 times

 **sanju032** Most Recent 2 weeks, 1 day ago

If any one had all these set of questions to prepare GCP ACE please send to sanjeevkumartm254@gmail.com

upvoted 1 times

 **leogor** 2 weeks, 3 days ago

D, basic and manual scaling are not wise actions

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is correct.

upvoted 1 times

 **ace_2021_1** 1 month, 2 weeks ago

why D.....pls tell the reason...

upvoted 1 times

 **haroldbenites** 5 months, 3 weeks ago

Go for D

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

D. Automatic Scaling with min_idle_instances set to 3.

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Automatic Scaling with min_idle_instances set to 3.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Selected Answer: D

Ans - D

upvoted 1 times

  **beginnercloud** 1 year, 2 months ago

D is correct

upvoted 2 times

  **pacman_user** 1 year, 4 months ago

D is correct.

upvoted 2 times

  **sidharthwader** 1 year, 4 months ago

Obviously Auto-scaling with min no of instances as 3

upvoted 3 times

  **vmart** 1 year, 4 months ago

D is correct

upvoted 1 times

  **mcaromit** 1 year, 6 months ago

D is correct

upvoted 1 times

  **[Removed]** 1 year, 7 months ago

D is correct. Automatic Scaling with min_idle_instances set to 3.

upvoted 2 times

  **pondai** 1 year, 7 months ago

D is right

B is wrong <https://cloud.google.com/compute/docs/disks/create-snapshots?hl=zh-tw>

upvoted 2 times

  **EABDAJA** 1 year, 8 months ago

D is the Answer

upvoted 1 times

Question #10

You have a development project with appropriate IAM roles defined. You are creating a production project and want to have the same IAM roles on the new project, using the fewest possible steps. What should you do?

- A. Use gcloud iam roles copy and specify the production project as the destination project.
- B. Use gcloud iam roles copy and specify your organization as the destination organization.
- C. In the Google Cloud Platform Console, use the 'create role from role' functionality.
- D. In the Google Cloud Platform Console, use the 'create role' functionality and select all applicable permissions.

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>*Community vote distribution*

A (100%)

✉ **coldpar** Highly Voted 2 years, 7 months ago

Correct Answer is A not B

upvoted 35 times

✉ **Agents89** Highly Voted 2 years, 6 months ago

Correct answer is A

upvoted 19 times

✉ **dasun_eranga** Most Recent 1 day, 6 hours ago

correct answer is A

upvoted 1 times

✉ **dennydream** 1 week, 5 days ago

A makes sense. B doesn't. If you copied to the Org all projects underneath inherit those policies?

upvoted 1 times

✉ **leogor** 2 weeks, 3 days ago

Correct answer is A.

upvoted 1 times

✉ **handras80** 3 weeks, 3 days ago

Selected Answer: A

Correct answer is A.

upvoted 1 times

✉ **Shane_trt** 1 month, 1 week ago

Definitely A

upvoted 1 times

✉ **Sgegmund** 1 month, 1 week ago

Selected Answer: A

A not B

upvoted 1 times

✉ **ssahu975** 1 month, 3 weeks ago

Correct Ans is A

upvoted 1 times

✉ **its_another_samsun** 3 months ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **12234** 3 months ago

Selected Answer: A

Correct answer A

upvoted 1 times

✉ **abirroy** 3 months, 2 weeks ago

Selected Answer: A

Correct answer is A

upvoted 1 times

  **abiso** 4 months, 1 week ago

A. Link in the answer clearly says this

To create a copy of an existing role spanner.databaseAdmin into a project with PROJECT_ID, run:

```
gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomSpannerDbAdmin --dest-project=PROJECT_ID
```

upvoted 2 times

  **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

  **orious** 4 months, 1 week ago**Selected Answer: A**

The goal is to copy a role from a project into another project (not an organization).

upvoted 4 times

  **AzureDP900** 4 months, 2 weeks ago

going with A

upvoted 1 times

  **utkan** 5 months, 1 week ago**Selected Answer: A**

must be a

upvoted 1 times

Question #11

You need a dynamic way of provisioning VMs on Compute Engine. The exact specifications will be in a dedicated configuration file. You want to follow Google's recommended practices. Which method should you use?

- A. Deployment Manager
- B. Cloud Composer
- C. Managed Instance Group
- D. Unmanaged Instance Group

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instances/>*Community vote distribution*

A (100%)

leogor 2 weeks, 3 days ago

A. deployment manager,
instance template / group is wrong

upvoted 1 times

priyanka2209 3 weeks, 4 days ago

Hi Chris. can you please share all questions on priyankasmaill@gmail.com

upvoted 1 times

alexander_reznicek 1 month, 1 week ago

Hello, if possible, please email me the questions at alex.reznicek@live.com. I have access to only a limited number as well. Thank you.

upvoted 1 times

akshaydoifode88 3 weeks ago

Hey Alex, if you got the questions, can you send it to me at aksrocks11@gmail.com

upvoted 1 times

Arnaud_6789 1 month, 2 weeks ago

Please, if someone has all the questions, share them with me at arnaud.noutch@gmail.com. I am preparing for the ACE exam. I was able to access only 92 questions.

Thank you

upvoted 1 times

b_script 1 month, 1 week ago

Open every new page in Incognito Window

upvoted 2 times

STARSON 1 month, 2 weeks ago

If anyone have the entire questions , Could you please share it to the mail snehastarson19@gmail.com.Thanks in Advance

upvoted 1 times

Saibaba13 2 months ago

I am preparing for GCP Associate engineer exam. I was able to access 92 questions only, if anyone has the entire questions please share them to my saibabaareti13@gmail.com address. Thanks in advance!

upvoted 1 times

anirban7172 1 week, 4 days ago

A contributor access only allows to see all the questions. or some one can share it please

upvoted 1 times

Laxman72 1 month, 2 weeks ago

If got all kindly share with me rachalwar98@outlook.com

upvoted 1 times

Adhithiyan 1 month, 3 weeks ago

if anyone have the entire questions share it to the mail adhivip02@gmail.com

upvoted 1 times

Revupvn 1 month, 3 weeks ago

If you receive all the questions please share it with me also revu.pvn@gmail.com

Even i am preparing for GCP exam next week.

upvoted 1 times

 **RAVI321** 2 months, 3 weeks ago

Hey guys anybody have the whole 197 question of this site

upvoted 4 times

 **Adhithiyan** 1 month, 3 weeks ago

if anyone have the entire questions share it to the mail adhivip02@gmail.com

upvoted 1 times

 **isatemelci** 1 month, 4 weeks ago

Could you please share with me? tmlc@duck.com

upvoted 1 times

 **Letahrgicbeagle** 1 month, 4 weeks ago

Did anyone get that from Chris?

upvoted 1 times

 **Chris10X** 2 months, 1 week ago

I do... yes

upvoted 3 times

 **Aras123** 3 weeks, 3 days ago

Hi Chris ..Can you please send them to prathu83@gmail.com?

upvoted 1 times

 **Irenia111** 4 weeks ago

hi Chris Can you please send them to 15000616693@162.com

upvoted 1 times

 **DarkRyuk** 4 weeks, 1 day ago

hey Chris, could you please share the pdf to me as well, Thanks buddy!

(mail : alf141616@gmail.com)

upvoted 1 times

 **12234** 3 months ago

Selected Answer: A

A correct answer

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **orious** 4 months, 1 week ago

Selected Answer: A

A is correct.

Managed Instance Groups don't support Configuration file in order to provision VM instances.

upvoted 3 times

 **AzureDP900** 4 months, 2 weeks ago

Deployment Manager is same like AWS Cloud Formation.. A is right answer

upvoted 1 times

 **alex000** 5 months, 2 weeks ago

Selected Answer: A

https://cloud.google.com/deployment-manager/docs/fundamentals?hl=en&_ga=2.254774664.-1447792053.1630588214

upvoted 3 times

 **haroldbenites** 5 months, 3 weeks ago

Go for A

upvoted 1 times

 **Krmnpi** 6 months ago

Selected Answer: A

It should be A

upvoted 1 times

 **rsuresh27** 6 months, 2 weeks ago

Selected Answer: A

Managed instance groups use instance templates. C is not correct. The correct answer is A.

upvoted 2 times

 **Chile** 6 months, 2 weeks ago

The answer should be A

upvoted 1 times

 **DesmoMike74** 6 months, 2 weeks ago

pls see in this article <https://cloud.google.com/deployment-manager/docs/configuration>

upvoted 2 times

 **darkroomie** 1 month, 1 week ago

So you're saying that "B" is correct?

upvoted 1 times

Question #12

You have a Dockerfile that you need to deploy on Kubernetes Engine. What should you do?

- A. Use kubectl app deploy <dockerfilename>.
- B. Use gcloud app deploy <dockerfilename>.
- C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.
- D. Create a docker image from the Dockerfile and upload it to Cloud Storage. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

Correct Answer: C

Reference -

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

Community vote distribution

C (94%)	6%
---------	----

 **Agents89** Highly Voted 2 years, 6 months ago

C is correct

upvoted 35 times

 **NC4F** Highly Voted 10 months, 3 weeks ago

Selected Answer: C

C is correct (21/12/2021 in my exam)

upvoted 6 times

 **leogor** Most Recent 2 weeks, 3 days ago

c, u need to upload to Container Registry before deploying

upvoted 1 times

 **hiromi** 2 weeks, 6 days ago

Selected Answer: C

C is correct, cuz you need to upload to Container Registry at first

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right

upvoted 1 times

 **Biksthrow** 4 months, 3 weeks ago

C is correct answer

upvoted 1 times

 **wjtb** 5 months, 1 week ago

Selected Answer: C

C is correct.

A can be eliminated because kubectl app * is not a valid command

B can be eliminated because gcloud app deploy deploys on app engine, not on kubernetes (also it still requires a config file pointing to the image).

D is not correct, since you cannot deploy a container image directly from GCS

upvoted 4 times

 **haroldbenites** 5 months, 3 weeks ago

Go for C

upvoted 1 times

 **Prassanth** 5 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

 **Chile** 6 months, 2 weeks ago

C is correct

upvoted 1 times

 **pgb54** 8 months, 2 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

 **HansKloss611** 9 months, 2 weeks ago

Selected Answer: C

Answer: C

We cannot deploy directly dockerfile on k8s. We need create docker image first

upvoted 3 times

 **pnVino27** 11 months ago

Selected Answer: A

Answer: A

A. Use kubectl app deploy <dockerfilename>.

Deploying on K8's Engine we use Kubectl command

upvoted 1 times

 **wh1t4k3r** 11 months ago

Create the image prior to deployment:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

You cant deploy a docker file

upvoted 3 times

 **shawnkkk** 11 months, 3 weeks ago

C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

upvoted 3 times

 **vishnukumartr** 11 months, 3 weeks ago

C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

upvoted 2 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - C

upvoted 2 times

Question #13

Your development team needs a new Jenkins server for their project. You need to deploy the server using the fewest steps possible. What should you do?

- A. Download and deploy the Jenkins Java WAR to App Engine Standard.
- B. Create a new Compute Engine instance and install Jenkins through the command line interface.
- C. Create a Kubernetes cluster on Compute Engine and create a deployment with the Jenkins Docker image.
- D. Use GCP Marketplace to launch the Jenkins solution.

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine>*Community vote distribution*

D (100%)

⊕  **Agents89** Highly Voted 2 years, 6 months ago

D is correct

upvoted 31 times

⊕  **leogor** Most Recent 2 weeks, 3 days ago

D, marketplace has do it for u

upvoted 1 times

⊕  **PKookNN** 4 weeks, 1 day ago

Selected Answer: D

D is correct one

upvoted 2 times

⊕  **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

⊕  **krnagg** 3 months, 3 weeks ago

D is the correct

upvoted 1 times

⊕  **RanjithK** 4 months, 1 week ago

Answer is D

upvoted 1 times

⊕  **Pritesh28** 5 months ago

D. Use GCP Marketplace to launch the Jenkins solution.

upvoted 1 times

⊕  **Denis_TKM** 5 months, 1 week ago

D is the good answer

upvoted 1 times

⊕  **ybaqqiou** 5 months, 2 weeks ago

D is correct the Market Place is the quicker

upvoted 1 times

⊕  **haroldbenites** 5 months, 3 weeks ago

Go for D

upvoted 1 times

⊕  **Chile** 6 months, 2 weeks ago

D is correct

upvoted 1 times

⊕  **dapdommaG** 6 months, 3 weeks ago

D is Correct

upvoted 1 times

⊕  **kchandank** 7 months, 3 weeks ago

D is correct

upvoted 1 times

 **[Removed]** 10 months, 3 weeks ago

Market Place is the place to be (just click and go)

upvoted 3 times

 **AugustoKras011111** 7 months, 1 week ago

YES, I think Qwiklabs have a lab about this, Easypeasy

upvoted 1 times

 **onyb** 11 months ago

MarketPlace is the place to be

upvoted 4 times

 **Skyone** 11 months, 2 weeks ago

Selected Answer: D

D is correct

upvoted 3 times

 **shawnkkk** 11 months, 3 weeks ago

D. Use GCP Marketplace to launch the Jenkins solution.

upvoted 2 times

Question #14

You need to update a deployment in Deployment Manager without any resource downtime in the deployment. Which command should you use?

- A. gcloud deployment-manager deployments create --config <deployment-config-path>
- B. gcloud deployment-manager deployments update --config <deployment-config-path>
- C. gcloud deployment-manager resources create --config <deployment-config-path>
- D. gcloud deployment-manager resources update --config <deployment-config-path>

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/deployments/update>*Community vote distribution*

B (100%)

 **Agents89** Highly Voted 2 years, 6 months ago

B is correct

upvoted 38 times

 **[Removed]** Highly Voted 2 years, 1 month ago

B is correct Additional tip, update and create resource is not even a command under deployment management service.

upvoted 35 times

 **leogor** Most Recent 2 weeks, 3 days ago

should not use 'create', so B or D is correct

and keyword 'resource' is invalid, so B is the ans

upvoted 1 times

 **PKookNN** 4 weeks, 1 day ago

Selected Answer: B

B is correct

upvoted 2 times

 **eloyus** 1 month ago

Selected Answer: B

B is correct

upvoted 1 times

 **krnagg** 3 months, 3 weeks ago

B is correct. C and D are invalid commands and A is to create a new deployment.

upvoted 2 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right to update the deployment.

upvoted 1 times

 **Baan** 5 months ago

B is correct

upvoted 1 times

 **haroldbenites** 5 months, 3 weeks ago

Go for B

upvoted 1 times

 **OlayemiO** 6 months, 3 weeks ago

Selected Answer: B

B is the correct answer.

upvoted 2 times

 **kchandank** 7 months, 3 weeks ago

B is correct one

upvoted 1 times

 **nooneknows** 8 months, 2 weeks ago

Selected Answer: B

B is correct!

upvoted 1 times

 **Zufair** 10 months, 2 weeks ago

B is correct

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

B. gcloud deployment-manager deployments update --config <deployment-config-path>

upvoted 3 times

 **vishnukumartr** 11 months, 3 weeks ago

B. gcloud deployment-manager deployments update --config <deployment-config-path>

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - B

upvoted 1 times

Question #15

You need to run an important query in BigQuery but expect it to return a lot of records. You want to find out how much it will cost to run the query. You are using on-demand pricing. What should you do?

- A. Arrange to switch to Flat-Rate pricing for this query, then move back to on-demand.
- B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.
- C. Use the command line to run a dry run query to estimate the number of bytes returned. Then convert that bytes estimate to dollars using the Pricing Calculator.
- D. Run a select count (*) to get an idea of how many records your query will look through. Then convert that number of rows to dollars using the Pricing Calculator.

Correct Answer: B

Reference:

<https://cloud.google.com/bigquery/docs/estimate-costs>*Community vote distribution*

B (100%)

  **Agents89** Highly Voted 2 years, 6 months ago

B is Correct

upvoted 34 times

  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answers is (B):

On-demand pricing

Under on-demand pricing, BigQuery charges for queries by using one metric: the number of bytes processed (also referred to as bytes read). You are charged for the number of bytes processed whether the data is stored in BigQuery or in an external data source such as Cloud Storage, Drive, or Cloud Bigtable. On-demand pricing is based solely on usage.

https://cloud.google.com/bigquery/pricing#on_demand_pricing

upvoted 25 times

  **leogor** Most Recent 2 weeks, 3 days ago

B, READ but not RETURNED

upvoted 1 times

  **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 2 times

  **AzureDP900** 4 months, 2 weeks ago

B is right . number bytes read instead of number of byte returned.

upvoted 2 times

  **haroldbenites** 5 months, 3 weeks ago

Go for B

upvoted 2 times

  **Alejondri** 6 months, 2 weeks agoSelected Answer: B

B Is correct

upvoted 2 times

  **DanielB96** 6 months, 2 weeks agoSelected Answer: B

B is correct :)

upvoted 1 times

  **naresh136** 8 months, 4 weeks ago

B is correct answer.

upvoted 1 times

  **Durgesh1997** 9 months, 2 weeks ago

B is correct

upvoted 2 times

 **CK389** 11 months, 2 weeks ago

B is correct

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.

upvoted 2 times

 **Jaira1256** 11 months, 3 weeks ago

B is correct

upvoted 1 times

 **KunK** 1 year, 1 month ago

B, it's read instead of return

upvoted 1 times

 **vikram____** 1 year, 2 months ago

Correct answer is B !

upvoted 1 times

 **Chotebhaisahab** 1 year, 2 months ago

agree B is the Correct option.

upvoted 1 times

Question #16

You have a single binary application that you want to run on Google Cloud Platform. You decided to automatically scale the application based on underlying infrastructure CPU usage. Your organizational policies require you to use virtual machines directly. You need to ensure that the application scaling is operationally efficient and completed as quickly as possible. What should you do?

- A. Create a Google Kubernetes Engine cluster, and use horizontal pod autoscaling to scale the application.
- B. Create an instance template, and use the template in a managed instance group with autoscaling configured.
- C. Create an instance template, and use the template in a managed instance group that scales up and down based on the time of day.
- D. Use a set of third-party tools to build automation around scaling the application up and down, based on Stackdriver CPU usage monitoring.

Correct Answer: B*Community vote distribution*

B (100%)

 **coldpar** Highly Voted 2 years, 7 months ago

correct is B as you have to use VM instances directly.

upvoted 50 times

 **mohdafiuddin** Highly Voted 1 year, 10 months ago

I'll take a simple and logical approach for answering this.

Let's first break down the question into key requirements -

1. automatically scale the application based on underlying infrastructure CPU usage.
2. use virtual machines directly.

A. Not feasible because VMs are not used directly here.

B. This is the correct answer.

C. Time of Day... Easy elimination because this does not scale on CPU usage and time of day is mentioned NOWHERE.

D. Third Party Tools.... Nobody would use GCP if they needed third party tools to do something as simple as scaling based on CPU usage. all popular cloud providers have native solutions for this including GCP.

upvoted 43 times

 **kopper2019** 1 year, 7 months ago

and also D is out because why would I use a third party tool when is a GCP exam

upvoted 9 times

 **RMO000** 1 year, 1 month ago

If the resource/solution is not available. It's a possibility.

upvoted 3 times

 **leogor** Most Recent 2 weeks, 3 days ago

B, managed instance group (VM instances) with autoscaling

upvoted 1 times

 **gcp_world123** 2 months, 3 weeks ago

Our requirements are as per the question

1. Use Virtual Machines directly (i.e. not container-based)
2. Scale Automatically
3. Scaling is efficient & is quick

B is correct

Managed instance groups offer autoscaling capabilities that let you automatically add or delete instances from a managed instance group based on increases or decreases in load (CPU Utilization in this case). Autoscaling helps your apps gracefully handle increases in traffic and reduce costs when the need for resources is lower. You define the autoscaling policy and the autoscaler performs automatic scaling based on the measured load (CPU Utilization in this case). Autoscaling works by adding more instances to your instance group when there is more load (upscale), and deleting instances when the need for instances is lowered (downscale).

Ref: <https://cloud.google.com/compute/docs/autoscaler>

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is perfect for given scenario.

upvoted 1 times

 **jagan_cloud** 5 months, 2 weeks ago

option B is correct

upvoted 1 times

 **haroldbenites** 5 months, 3 weeks ago

Go for B

upvoted 1 times

 **sid0127** 11 months, 2 weeks ago

Selected Answer: B

correct is B

upvoted 5 times

 **shawnkkk** 11 months, 3 weeks ago

B. Create an instance template, and use the template in a managed instance group with autoscaling configured.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Create an instance template, and use the template in a managed instance group with autoscaling configured.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

B is correct

upvoted 1 times

 **KunK** 1 year, 1 month ago

B, since managed instance groups involve the use of VM and can enable autoscale

upvoted 1 times

 **ankatsu2010** 1 year, 1 month ago

B is correct. The Key word is 'to use virtual machines directly'.

upvoted 1 times

 **AdarshTrivedi** 1 year, 1 month ago

B is correct

upvoted 1 times

 **beginnercloud** 1 year, 2 months ago

B is correct answers

upvoted 1 times

 **AB15** 1 year, 4 months ago

B is correct.

upvoted 1 times

Question #17

You are analyzing Google Cloud Platform service costs from three separate projects. You want to use this information to create service cost estimates by service type, daily and monthly, for the next six months using standard query syntax. What should you do?

- A. Export your bill to a Cloud Storage bucket, and then import into Cloud Bigtable for analysis.
- B. Export your bill to a Cloud Storage bucket, and then import into Google Sheets for analysis.
- C. Export your transactions to a local file, and perform analysis with a desktop tool.
- D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

Correct Answer: D

Community vote distribution

D (100%)

 **mohdafiuddin** Highly Voted 1 year, 10 months ago

Solving this by first eliminating the options that don't suit us. By breaking down the question into the key requirements-

1. Analyzing Google Cloud Platform service costs from three separate projects.
2. Using standard query syntax. -> (Relational data and SQL)

- A. 'Cloud Storage bucket'.....'Cloud Bigtable'. Not feasible, mainly because cloud BigTable is not good for Structured Data (or Relational Data on which we can run SQL queries as per the question's requirements). BigTable is better suited for Semi Structured data and NoSQL data.
- B. 'Cloud Storage bucket'.....'Google Sheets'. Not Feasible because there is no use of SQL in this option, which is one of the requirements.
- C. Local file, external tools... this is automatically eliminated because the operation we need is simple, and there has to be a GCP native solution for this. We shouldn't need to rely on going out of the cloud for such a simple thing.
- D. 'BigQuery'.....'SQL queries' -> This is the right answer.

upvoted 66 times

 **ryumada** 3 months, 1 week ago

Cloud billing data can only be exported to a JSON local file and to Bigquery. So, using Cloud Storage to export cloud billing data is not possible to do.

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 3 times

 **cesar7816** Highly Voted 2 years, 8 months ago

Agreed, BigQuery

upvoted 17 times

 **yurstev** 1 year, 10 months ago

the key is standard query syntax

upvoted 3 times

 **Kopy** Most Recent 3 days, 2 hours ago

Selected Answer: D

Agreed, BigQuery

upvoted 1 times

 **leogor** 2 weeks, 3 days ago

D, using standard query syntax so use BQ

upvoted 1 times

 **Haarish** 3 months, 3 weeks ago

Agreed, BigQuery

Ans: D

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is the right answer.

upvoted 1 times

 **denkyira** 4 months, 4 weeks ago

I agreed, BigQuery Answer D is correct

upvoted 1 times

✉ **jagan_cloud** 5 months, 2 weeks ago

answer is Option D

upvoted 1 times

✉ **haroldbenites** 5 months, 3 weeks ago

Go for D

upvoted 1 times

✉ **LaxmanTiwari** 5 months, 3 weeks ago

Agreed , answer is D

upvoted 1 times

✉ **shawnkkk** 11 months, 3 weeks ago

D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

upvoted 2 times

✉ **vishnukumartr** 11 months, 3 weeks ago

D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

upvoted 1 times

✉ **Jaira1256** 11 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 3 times

✉ **Chotebhaisahab** 1 year, 2 months ago

agree D is the right choice.

upvoted 1 times

✉ **YAS007** 1 year, 3 months ago

D is correct : "...we recommend that you enable Cloud Billing data export to BigQuery at the same time that you create a Cloud Billing account. "
<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

<https://medium.com/google-cloud/analyzing-google-cloud-billing-data-with-big-query-30bae1c2aae4>

upvoted 3 times

✉ **HogwartsTrue** 1 year, 5 months ago

Ans:D BigQuery is for Analytical Purpose

upvoted 1 times

Question #18

You need to set up a policy so that videos stored in a specific Cloud Storage Regional bucket are moved to Coldline after 90 days, and then deleted after one year from their creation. How should you set up the policy?

- A. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 275 days (365 - 90)
- B. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 365 days.
- C. Use gsutil rewrite and set the Delete action to 275 days (365 - 90).
- D. Use gsutil rewrite and set the Delete action to 365 days.

Correct Answer: A*Community vote distribution*

B (73%)

A (27%)

 **Sammigbo** Highly Voted 2 years, 5 months ago

Answer is B. There should be no reason to recalculate the time needed to delete after a year.

upvoted 45 times

 **JKRowlings** 1 year, 11 months ago

The correct ans is A.

upvoted 4 times

 **yvinisiupacuando** 1 year, 6 months ago

Right answer is clearly B, "A" does not make any sense.

upvoted 11 times

 **cloudenthu01** Highly Voted 2 years, 4 months ago

Correct is B.

You only re-calculate expiry date when objects are re-written using re-write option to another storage class in which case creation date is reset. But in this case objects is moveed to Coldline class after 90 days and then we want to delete the object after 365 days.

upvoted 29 times

 **T_T_M** 2 years, 2 months ago

You can change the storage class of an existing object either by rewriting the object or by using Object Lifecycle Management...Since Object Life cycle management was used there was no need to recalculate the expiration date and delete action still remains 365 days.

<https://cloud.google.com/storage/docs/storage-classes>

upvoted 9 times

 **Baba_Voss** Most Recent 1 week, 4 days ago

Selected Answer: A

A is the correct Answer

upvoted 1 times

 **leogor** 2 weeks, 3 days ago

B, no need to minus 90 days

upvoted 1 times

 **vijay456** 4 weeks ago

Age is counted from when an object was uploaded to the current bucket, even if it moved from another

upvoted 1 times

 **PKookNN** 4 weeks, 1 day ago

Selected Answer: B

The key part is "deleted after one year from their creation" - B is correct at 365 days. If A is chosen then the object would be deleted BEFORE 1 year.

upvoted 2 times

 **darkroomie** 1 month, 1 week ago

Selected Answer: A

A is the correct answer. It clearly specifies that videos are moved to coldline after 90 days and deleted after "A YEAR OF THE CREATION" (that includes the 90 days that counts for moving to coldline. So, videos are created one day, then moved to coldline after 90 days (+90) and deleted after a year of creation (90+275=365)

A IS CORRECT....!!!

upvoted 3 times

 **darkroomie** 1 month, 1 week ago

A is the correct answer. It clearly specifies that videos are moved to coldline after 90 days and deleted after "A YEAR OF THE CREATION" (that includes the 90 days that counts for moving to coldline. So, videos are created one day, then moved to coldline after 90 days (+90) and deleted after a year of creation (90+275=365)

A IS CORRECT....!!!

upvoted 3 times

 **kadc** 1 month, 3 weeks ago

Selected Answer: B

Some suggested that rewritten changes the creation date, but here it stated that

"Unlike changing an object's storage class manually, using SetStorageClass does not rewrite an object. This gives Object Lifecycle Management certain pricing advantages."

<https://cloud.google.com/storage/docs/lifecycle#setstorageclass-cost>

With other scenarios stated at lower section, it doesn't look like creation date would change

upvoted 1 times

 **shreykul** 3 months ago

Why not A, I mean question states 365 days from the date of creation, and the data will be in cold-line for 90 days, $365-90=275$ -> which should be set in lifeline? Am I wrong? or is the 90 days(Cold line period) is inclusive in the 365 days when creating life cycle management?

upvoted 2 times

 **gandor** 3 months, 1 week ago

Selected Answer: A

A is the answer.

The first object version should still be there after 90 days even if it was moved to cold line storage.

upvoted 3 times

 **gandor** 3 months, 1 week ago

A is the answer.

The first object version should still be there after 90 days even if it was moved to cold line storage.

upvoted 2 times

 **SaiSaiA** 3 months, 2 weeks ago

Answer B, the date counted is when the object is uploaded into the Cloud Storage. It doesn't matter if it changed to a different storage class, it is still in the same bucket.

upvoted 2 times

 **jayaen** 3 months, 2 weeks ago

A is correct

upvoted 3 times

 **Karthik_SIDDE** 3 months, 3 weeks ago

Selected Answer: A

Answer is A. The videos are stored in GCS and moved to coldline after 90 days, then deleted after one year of their creation

upvoted 3 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right , This is also in Tutorials Dojo with detailed explanations.

upvoted 1 times

Question #19

You have a Linux VM that must connect to Cloud SQL. You created a service account with the appropriate access rights. You want to make sure that the VM uses this service account instead of the default Compute Engine service account. What should you do?

- A. When creating the VM via the web console, specify the service account under the 'Identity and API Access' section.
- B. Download a JSON Private Key for the service account. On the Project Metadata, add that JSON as the value for the key compute-engine-service-account.
- C. Download a JSON Private Key for the service account. On the Custom Metadata of the VM, add that JSON as the value for the key compute-engine-service-account.
- D. Download a JSON Private Key for the service account. After creating the VM, ssh into the VM and save the JSON under `~/.gcloud/compute-engine-service-account.json`.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Community vote distribution

A (67%)

C (33%)

✉️  **Agents89** Highly Voted 2 years, 6 months ago

A is correct

upvoted 36 times

✉️  **ready2rock** 1 year, 5 months ago

How can this be? It says you HAVE a VM, meaning it's already created. A cannot be the solution.

upvoted 6 times

✉️  **jiniguez** 11 months ago

As the comment says:

"To change an instance's service account and access scopes, the instance must be temporarily stopped ... After changing the service account or access scopes, remember to restart the instance." So we can stop the instance, change the service account, then start it up again.

upvoted 2 times

✉️  **boof** 1 year, 1 month ago

A seems legit, the answer is worded poorly but is the most correct.

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#changeserviceaccountandscopes>

"To change an instance's service account and access scopes, the instance must be temporarily stopped ... After changing the service account or access scopes, remember to restart the instance." So we can stop the instance, change the service account, then start it up again.

upvoted 2 times

✉️  **ashrafh** 1 year, 2 months ago

I vote A

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Changing the service account and access scopes for an instance

If you want to run the VM as a different identity, or you determine that the instance needs a different set of scopes to call the required APIs, you can change the service account and the access scopes of an existing instance. For example, you can change access scopes to grant access to a new API, or change an instance so that it runs as a service account that you created, instead of the Compute Engine default service account. However, Google recommends that you use the fine-grained IAM policies instead of relying on access scopes to control resource access for the service account.

To change an instance's service account and access scopes, the instance must be temporarily stopped. To stop your instance, read the documentation for Stopping an instance. After changing the service account or access scopes, remember to restart the instance. Use one of the following methods to the change service account or access scopes of the stopped instance.

Hope this helps :)

upvoted 12 times

✉️  **jabrrJ68w02ond1** Highly Voted 11 months ago

Either the question or the answers are wrong. The question says that we HAVE a Linux VM, so we should strike all the answers that include "when creating the VM.." - on the other hand, adding JSON Tokens to VM metadata is terrible because it's readable in clear-text for everyone. So, what do we need to do here?

upvoted 6 times

 **leogor** Most Recent 2 weeks, 3 days ago

A is enough for the setting
upvoted 1 times

 **Untamables** 2 weeks, 5 days ago

Selected Answer: C

There is no correct answer in all options, but option C is the closest because a VM instance has meta data related to service accounts and you need to change the service account of the VM.

<https://cloud.google.com/compute/docs/metadata/default-metadata-values>

Refer to the below Google document to ensure the true answer.

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#changeserviceaccountandscopes>

upvoted 1 times

 **Allaman** 2 months, 1 week ago

Selected Answer: A

Go to the Create an instance page.

1-Go to Create an instance

2-Specify the VM details.

3-In the Identity and API access section, choose the service account you want to use from the drop-down list.

Continue with the VM creation process.

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 2 times

 **amindbesideitself** 6 months ago

Selected Answer: A

A is correct here

upvoted 2 times

 **DanielB96** 6 months, 2 weeks ago

Selected Answer: A

A.

All seems wrong for some reasons but A is the most applicable here.

upvoted 1 times

 **jjweckl** 7 months ago

aprender a interpretar un texto indios de mierda

upvoted 2 times

 **gielda211** 7 months, 1 week ago

Selected Answer: C

You have a Linux VM that must connect to Cloud SQL.

upvoted 1 times

 **wh1t4k3r** 11 months ago

My insights on this:

A: It's wrong, but it is the most applicable here. Here is why:

B: Talks about project wide, it's too broad and should be ruled out.

C: I did vote for this one at first, but after some research, I've found this:

The json file downloaded is a private key, and as per documentation, the service-account metadata of a VM should be used for tokens, unique to the VMs, not the service account:

<https://cloud.google.com/compute/docs/metadata/default-metadata-values>

Now, even considering this to be a key scenario, wouldn't the SQL need to be configured with the private key of the service account for this to work? And adding this private key alone to the VMs does not guarantee that the default service account won't be used (which is stated by the question)

D: This is for gcloud tool/SDK, not VMs.

Conclusion: All these are wrong lol, but A is the less wrong

upvoted 4 times

 **shawnkkk** 11 months, 3 weeks ago

A. When creating the VM via the web console, specify the service account under the Identity and API Access™ section.

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

A. When creating the VM via the web console, specify the service account under the **Identity and API Access**™ section.
upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. When creating the VM via the web console, specify the service account under the **Identity and API Access**™ section.
upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

A is correct
upvoted 1 times

 **peter77** 1 year, 1 month ago

It has to be A - You just don't put a private key in the metadata for everybody to see!
upvoted 3 times

Question #20

You created an instance of SQL Server 2017 on Compute Engine to test features in the new version. You want to connect to this instance using the fewest number of steps. What should you do?

- A. Install a RDP client on your desktop. Verify that a firewall rule for port 3389 exists.
- B. Install a RDP client in your desktop. Set a Windows username and password in the GCP Console. Use the credentials to log in to the instance.
- C. Set a Windows password in the GCP Console. Verify that a firewall rule for port 22 exists. Click the RDP button in the GCP Console and supply the credentials to log in.
- D. Set a Windows username and password in the GCP Console. Verify that a firewall rule for port 3389 exists. Click the RDP button in the GCP Console, and supply the credentials to log in.

Correct Answer: B

Reference:

<https://medium.com/falafel-software/sql-server-in-the-google-cloud-a17e8a1f11ce>

Community vote distribution

B (59%)

D (41%)

 **vnxt** Highly Voted 2 years, 6 months ago

I would say B is correct. RDP is enabled by default when you create a Windows instance (no need to check on it). Just make sure you install an RDP client (chrome ext or RDP) and set windows password.

upvoted 39 times

 **ryumada** 3 months, 1 week ago

The firewall rule for port 3389 is created by default if you create windows server on Compute Engine. So, no need to verify it.

<https://cloud.google.com/compute/docs/instances/connecting-to-windows#before-you-begin>

upvoted 3 times

 **eledu1985** 2 weeks, 2 days ago

In the link you provided the first step is to verify firewall rule was created, even if it is the default option!, so D is the most accurate even by the link.

upvoted 1 times

 **pas77** 1 year, 4 months ago

Obviously, B is not the answer because you have to install an RDP client which is an extra step. D is the answer because you can connect directly using the RDP button in the GCP console.

upvoted 6 times

 **jabrrJ68w02ond1** 11 months, 2 weeks ago

Tested it myself. At least on my Machines, I was asked to First install a RDP Client.

upvoted 8 times

 **UtsavDM** 1 year, 2 months ago

No, we can't connect using RDP directly in the GCP console. When we click on it, it asks us to install RDP client. So ultimately, B is more accurate.

upvoted 9 times

 **ankit89** Highly Voted 2 years, 6 months ago

D seems more correct

upvoted 31 times

 **obeythefist** 8 months, 2 weeks ago

I tested this on on Compute Engine today by deploying a new instance. D is not correct. When you click the RDP button, you are asked to install a client or use the Windows RDP client if you are running Windows. There is no option to enter credentials or get an RDP session through the web interface.

upvoted 6 times

 **anolive** Most Recent 1 day, 13 hours ago

Selected Answer: D

In my simulation was closest D

upvoted 1 times

 **xxz** 1 week, 1 day ago

Selected Answer: D

D is the correct answer.

<https://cloud.google.com/compute/docs/instances/connecting-to-windows#remote-desktop-connection-app>
<https://cloud.google.com/compute/docs/instances/windows/generating-credentials>

upvoted 1 times

eledu1985 2 weeks, 2 days ago**Selected Answer: D**

<https://cloud.google.com/compute/docs/instances/connecting-to-windows#before-you-begin>

upvoted 1 times

leogor 2 weeks, 3 days ago

B is the answer

upvoted 1 times

hiromi 2 weeks, 6 days ago**Selected Answer: B**

B is correct. From GCP console:

"Use a remote desktop protocol (RDP) client to connect to this instance. If you are running Windows on your local machine, use Remote Desktop Connection. Other operating systems might require you to use third-party software. The first time you connect, enter the username and password that you provided when you created the instance.

Note: You must configure the network firewall to open TCP port 3389 to enable RDP access"

upvoted 2 times

darkroomie 1 month, 1 week ago**Selected Answer: B**

Correct answer is B

upvoted 2 times

adil114 1 month, 2 weeks ago

B- IS A CORRECT ANSWER

upvoted 1 times

iadars 1 month, 3 weeks ago**Selected Answer: D**

D is correct because, before login to the instance we must ensure first that the firewall rule is allowing traffic through port 3369, otherwise we will not able to do RDP even we have the credentials.

upvoted 1 times

ale_brd_ 2 months ago**Selected Answer: B**

Option 2 is correct because:

Using the internal IP address is always preferred to going through the Internet using the Public one.

Through Internal IP option:

Chrome Remote Desktop is a service that lets you remotely access another computer by using a web browser. Chrome Remote Desktop works on Windows, macOS, and Linux and does not require the VM instance to have a public IP address.

Before you connect by using Chrome Remote Desktop, make sure that the following prerequisites are met:

You've created a Windows account and password on the VM instance.

You've installed the Chrome Remote Desktop service on the VM instance.

<https://cloud.google.com/compute/docs/instances/connecting-to-windows#chrome-remote-desktop>

upvoted 2 times

pedro_ 2 months ago

B and D have incorrect statements.

Firewall rule 3389 MUST be explicitly allowed. (B incorrect)

and

You required either the RDP Chrome extension OR a third party RDP client to use it (D incorrect)

upvoted 2 times

on2it 4 months ago**Selected Answer: B**

I know it

upvoted 2 times

RanjithK 4 months, 1 week ago

Answer is B. Tried and tested

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **patashish** 4 months, 1 week ago

Option D is correct.

Difference between B and D is only firewall setting.

With default network , firewall rules will come in built to allow access on 3389

But in answer , they have not mentioned either default or custom network.

We do not know if its changed.

Better to cross check rules , add/edit if required

Hence , as per my view D is correct answer.

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

Bis the correct answer.

upvoted 2 times

Question #21

You have one GCP account running in your default region and zone and another account running in a non-default region and zone. You want to start a new

Compute Engine instance in these two Google Cloud Platform accounts using the command line interface. What should you do?

- A. Create two configurations using gcloud config configurations create [NAME]. Run gcloud config configurations activate [NAME] to switch between accounts when running the commands to start the Compute Engine instances.
- B. Create two configurations using gcloud config configurations create [NAME]. Run gcloud configurations list to start the Compute Engine instances.
- C. Activate two configurations using gcloud configurations activate [NAME]. Run gcloud config list to start the Compute Engine instances.
- D. Activate two configurations using gcloud configurations activate [NAME]. Run gcloud configurations list to start the Compute Engine instances.

Correct Answer: A

Community vote distribution

A (90%) 10%

 **coldpar** Highly Voted 2 years, 7 months ago

A is the correct option

upvoted 28 times

 **leba** Highly Voted 2 years, 6 months ago

Correct answer is A as you can create different configurations for each account and create compute instances in each account by activating the respective account. Refer GCP documentation - Configurations Create & Activate Options B, C & D are wrong as gcloud config configurations list does not help create instances. It would only lists existing named configurations.

upvoted 25 times

 **PSS387** Most Recent 2 weeks, 5 days ago

A is correct answer

upvoted 1 times

 **PKookNN** 4 weeks, 1 day ago

Selected Answer: A

a is correct

upvoted 1 times

 **VarunGo** 1 month ago

Selected Answer: A

A is correct bro!

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **Alok17** 4 months, 2 weeks ago

A is corect

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right, similar question is part of Tutorials Dojo questions.

upvoted 1 times

 **ybaqqiou** 5 months, 2 weeks ago

A is correct

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

upvoted 1 times

 **theBestStudent** 6 months, 3 weeks ago

Obviously A is correct. All the other options don't make any sense when day say "Run gcloud configurations list to start the Compute Engine instances". How the heck are you expecting to "start" GCE instances doing "configuration list". Obviously B,C,D don't make any sense.

upvoted 4 times

 **j mugambi** 9 months, 2 weeks ago

Selected Answer: A

A is the only option.

By elimination: B/C/D have "gcloud config" or 'gcloud configurations' as a way to create compute instances which is wrong. Right command begins with 'gcloud compute'

Correct method: create configurations -> create instances in active config -> switch configuration -> create instances in newly active config
upvoted 2 times

 **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **blackthorn_01** 11 months ago

Selected Answer: C

<https://cloud.google.com/sdk/gcloud/reference/config/configurations/activate>

upvoted 1 times

 **Skyone** 11 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 2 times

 **test157** 11 months, 3 weeks ago

Selected Answer: A

A is the correct answer!

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

A. Create two configurations using gcloud config configurations create [NAME]. Run gcloud config configurations activate [NAME] to switch between accounts when running the commands to start the Compute Engine instances.

upvoted 1 times

Question #22

You significantly changed a complex Deployment Manager template and want to confirm that the dependencies of all defined resources are properly met before committing it to the project. You want the most rapid feedback on your changes. What should you do?

- A. Use granular logging statements within a Deployment Manager template authored in Python.
- B. Monitor activity of the Deployment Manager execution on the Stackdriver Logging page of the GCP Console.
- C. Execute the Deployment Manager template against a separate project with the same configuration, and monitor for failures.
- D. Execute the Deployment Manager template using the “-preview” option in the same project, and observe the state of interdependent resources.

Correct Answer: D

Reference:

<https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments>

Community vote distribution

D (100%)

✉  **YashBindlish** Highly Voted 2 years, 6 months ago

Correct answer is D as Deployment Manager provides the preview feature to check on what resources would be created
upvoted 37 times

✉  **Kopy** Most Recent 3 days, 17 hours ago

Selected Answer: D

deployment Manager provides the preview feature
upvoted 1 times

✉  **leogor** 2 weeks, 3 days ago

D Preview mode
upvoted 1 times

✉  **RanjithK** 4 months, 1 week ago

Answer is D
upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

D is right.. Preview mode is best to verification
upvoted 2 times

✉  **jagopi9381** 4 months, 3 weeks ago

Selected Answer: D
D FOR SURE
upvoted 1 times

✉  **haroldbenites** 5 months, 2 weeks ago

Go for D
upvoted 1 times

✉  **tigerbaer** 9 months ago

Selected Answer: D
Preview feature is available
upvoted 2 times

✉  **vishnukumartr** 11 months, 3 weeks ago

D. Execute the Deployment Manager template using the “-preview” option in the same project, and observe the state of interdependent resources.
upvoted 1 times

✉  **Jaira1256** 11 months, 3 weeks ago

D is correct
upvoted 1 times

✉  **beginnercloud** 1 year, 2 months ago

D is correct
upvoted 1 times

 **blan_ak** 1 year, 2 months ago

D is correct.

upvoted 1 times

 **vmart** 1 year, 5 months ago

yes D is the correct answer

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

D correct option

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

D is correct. Execute the Deployment Manager template using the "-preview" option in the same project, and observe the state of interdependent resources.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

D is correct

upvoted 1 times

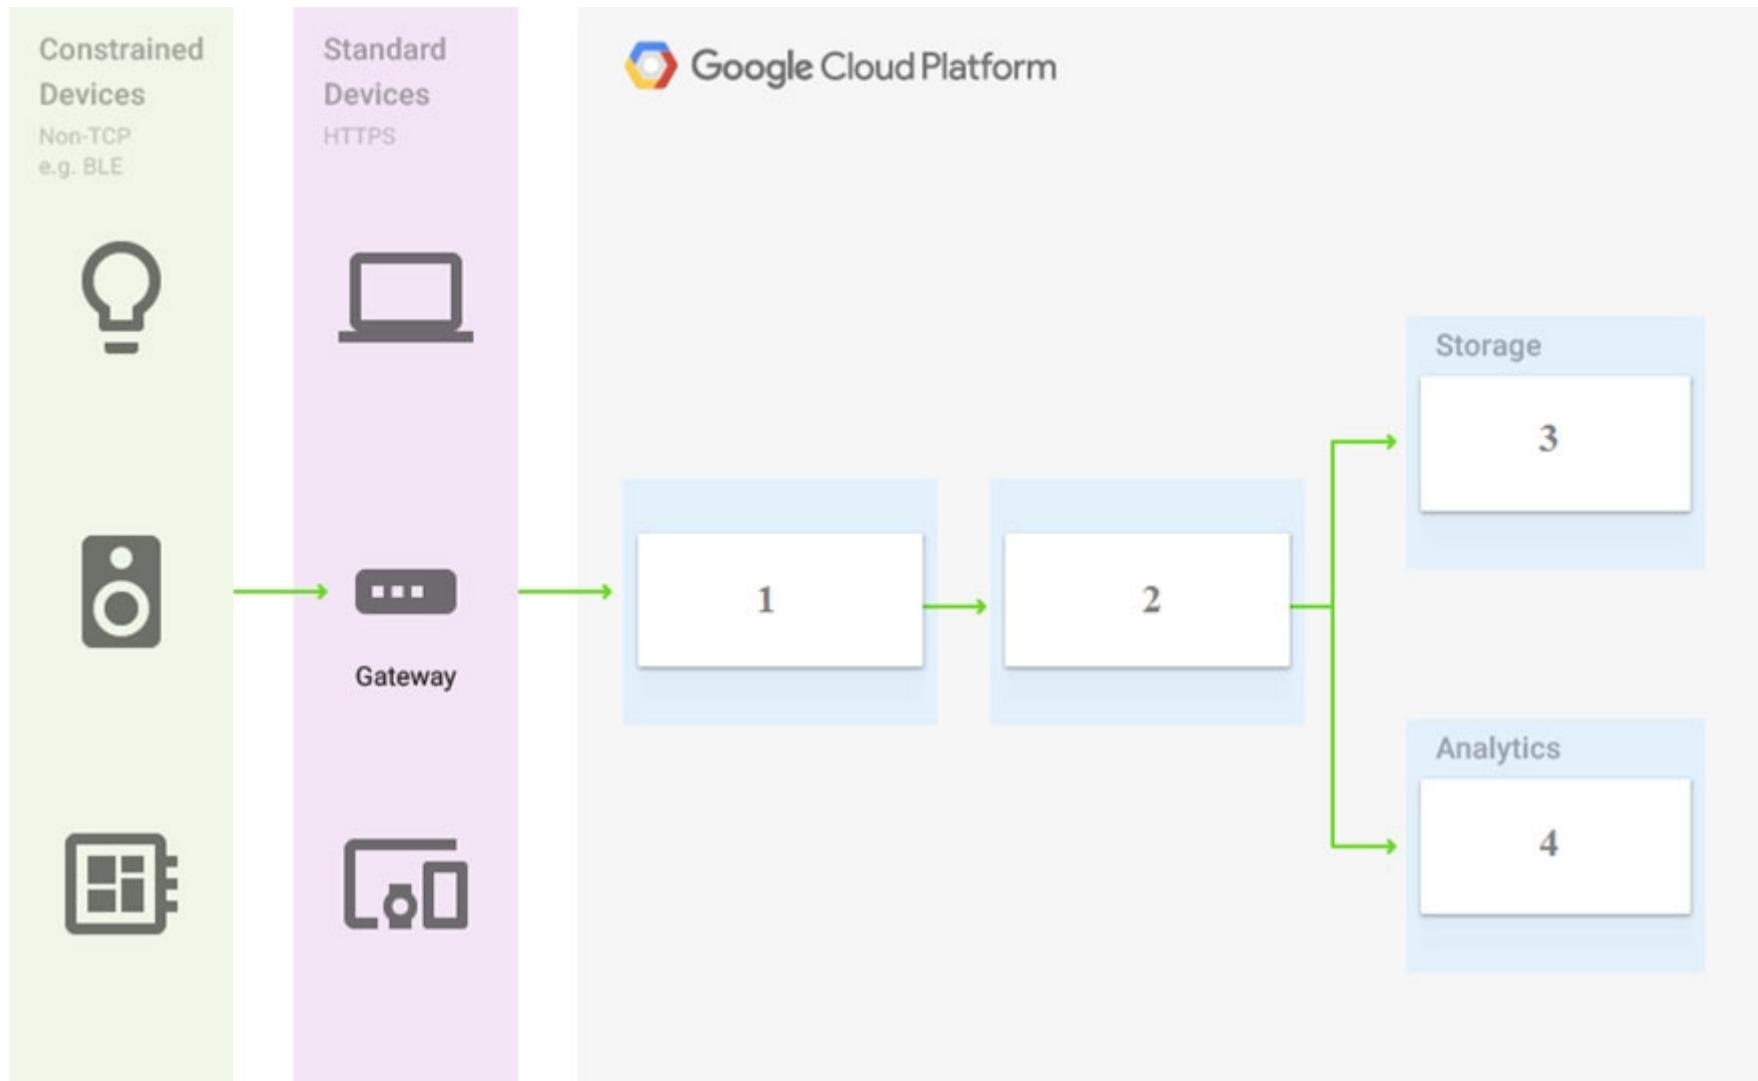
 **GCP_Student1** 1 year, 8 months ago

D. Execute the Deployment Manager template using the "-preview" option in the same project, and observe the state of interdependent resources.

upvoted 1 times

Question #23

You are building a pipeline to process time-series data. Which Google Cloud Platform services should you put in boxes 1,2,3, and 4?



- A. Cloud Pub/Sub, Cloud Dataflow, Cloud Datastore, BigQuery
- B. Firebase Messages, Cloud Pub/Sub, Cloud Spanner, BigQuery
- C. Cloud Pub/Sub, Cloud Storage, BigQuery, Cloud Bigtable
- D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/correlating-time-series-dataflow>

Community vote distribution

D (100%)

 **cludenthu01** Highly Voted 2 years, 4 months ago

Without a doubt D.

Whenever we want to process timeseries data look for BigTable.

Also you want to perform analytics in Box 4 ..look for BigQuery

Only D provides this option.

upvoted 56 times

 **vodia** 2 years, 4 months ago

Speaker also looks like an IoT device so D not A

upvoted 2 times

 **adedj99** 1 year, 11 months ago

are we considering bigtable as storage in here , since they expecting some storage

upvoted 2 times

 **shuvo_123** Highly Voted 2 years, 2 months ago

Option D:

See this reference : <https://cloud.google.com/blog/products/databases/getting-started-with-time-series-trend-predictions-using-gcp>

upvoted 9 times

 **dennydream** Most Recent 1 week, 5 days ago

I know the answer is D, but it's misleading. "Storage" would suggest cloud storage, not BigTable.

upvoted 1 times

✉ ale_brd_ 2 months ago

Selected Answer: D

whenever IOT is present go with DataStore

upvoted 1 times

✉ RanjithK 4 months, 1 week ago

Answer is D

upvoted 1 times

✉ AzureDP900 4 months, 2 weeks ago

D is right

upvoted 1 times

✉ Pleaseletmein 4 months, 2 weeks ago

D is the right answer

upvoted 1 times

✉ bamosk 4 months, 3 weeks ago

Selected Answer: D

Look the picture posted on this page here

<https://cloud.google.com/dataflow#section-7>

upvoted 1 times

✉ bamosk 4 months, 3 weeks ago

IoT = Unstructured data -> eliminated Datastore + Timeseries requirement = BigTable

Ingestion point = Pub/Sub (Firebase messaging is a push notification service to client)

Process data + realtime or batch = Data flow

Analytics = BigQuery

Hence D.

upvoted 3 times

✉ mplibunao 5 months ago

Selected Answer: D

Answer is D

upvoted 1 times

✉ haroldbenites 5 months, 2 weeks ago

Go for D

upvoted 1 times

✉ vishnukumartr 11 months, 3 weeks ago

D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

upvoted 1 times

✉ Jaira1256 11 months, 3 weeks ago

D is correct

upvoted 1 times

✉ bubblegumbeach 1 year, 1 month ago

why is it dataflow and not cloud storage

upvoted 2 times

✉ Vishal_b 9 months, 1 week ago

dataflow need to push message from pub/sub to any storage using pipeline

upvoted 1 times

✉ theBestStudent 6 months, 3 weeks ago

one detail Vishal_b, pull (and not push) messages from a pub/sub subscription in this case. The ones pushing messages (to a pub/sub topic) are the ones in the "standard devices" layer.

upvoted 1 times

✉ beginnercloud 1 year, 2 months ago

In this case...D is correct

upvoted 1 times

✉ akshay_jadhav 1 year, 2 months ago

D ..is the ans

upvoted 1 times

✉ Agraved 1 year, 5 months ago

awesome question with graphical representation, clearing a lot concepts in one Diagram. Answer has to be D

upvoted 2 times

Question #24

You have a project for your App Engine application that serves a development environment. The required testing has succeeded and you want to create a new project to serve as your production environment. What should you do?

- A. Use gcloud to create the new project, and then deploy your application to the new project.
- B. Use gcloud to create the new project and to copy the deployed application to the new project.
- C. Create a Deployment Manager configuration file that copies the current App Engine deployment into a new project.
- D. Deploy your application again using gcloud and specify the project parameter with the new project name to create the new project.

Correct Answer: A*Community vote distribution*

A (100%)

 **coldpar** Highly Voted 2 years, 7 months ago

Correct is A.

Option B is wrong as the option to use gcloud app cp does not exist.

Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage

upvoted 54 times

 **manu202020** 2 years, 4 months ago

you're missing one thing. D isn't about using deployment manager to copy the configuration, instead, using the configuration file to copy the configuration from test project.

upvoted 6 times

 **leba** Highly Voted 2 years, 6 months ago

Correct answer is A as gcloud can be used to create a new project and the gcloud app deploy can point to the new project. Refer GCP documentation - GCloud App Deploy.

Option B is wrong as the option to use gcloud app cp does not exist

. Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage.

upvoted 9 times

 **leogor** Most Recent 2 weeks, 3 days ago

A, one project can only have one deployment

upvoted 1 times

 **ale_brd_** 2 months ago

Selected Answer: A

cloud can be used to create a new project and the gcloud app deploy can point to the new project.

upvoted 1 times

 **gcp_world123** 2 months, 3 weeks ago

A is correct

You can deploy to a different project by using –project flag.

By default, the service is deployed the current project configured via:

\$ gcloud config set core/project PROJECT

To override this value for a single deployment, use the –project flag:

\$ gcloud app deploy ~/my_app/app.yaml –project=PROJECT

Ref: <https://cloud.google.com/sdk/gcloud/reference/app/deploy>

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

go for A

upvoted 1 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: A

Think A also.

upvoted 1 times

✉  **onyb** 11 months ago

Although the answer is not worded correctly, I think C is the right answer . Command is gcloud app deploy ~/my_app/app.yaml --project=PROJECT (<https://cloud.google.com/sdk/gcloud/reference/app/deploy>)

upvoted 3 times

✉  **shawnkkk** 11 months, 3 weeks ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

✉  **vishnukumartr** 11 months, 3 weeks ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

✉  **Jaira1256** 11 months, 3 weeks ago

Correct is A

upvoted 1 times

✉  **vmart** 1 year, 5 months ago

A is correct answer

upvoted 1 times

✉  **Umesh09** 1 year, 5 months ago

I would have gone with D if the project was an existing project then option D is the best choice since it would have been one step process. Now since the new project has to be created Option A seems to be better choice with the two step process. First create a new project and then deploy your application hence my choice Option A

upvoted 3 times

✉  **didilmfs** 1 year, 1 month ago

1 project only have 1 app engine. 1 app engine/project. the rule. cmiw

upvoted 1 times

✉  **mcaromit** 1 year, 6 months ago

A is correct option

upvoted 1 times

✉  **[Removed]** 1 year, 7 months ago

A is correct. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 2 times

✉  **EABDAJA** 1 year, 8 months ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

Question #25

You need to configure IAM access audit logging in BigQuery for external auditors. You want to follow Google-recommended practices. What should you do?

- A. Add the auditors group to the 'logging.viewer' and 'bigQuery.dataViewer' predefined IAM roles.
- B. Add the auditors group to two new custom IAM roles.
- C. Add the auditor user accounts to the 'logging.viewer' and 'bigQuery.dataViewer' predefined IAM roles.
- D. Add the auditor user accounts to two new custom IAM roles.

Correct Answer: C

Reference:

<https://cloud.google.com/iam/docs/roles-audit-logging>*Community vote distribution*

A (96%)	4%
---------	----

 **coldpar** Highly Voted 2 years, 7 months ago

Correct is A.

As per google best practices it is recommended to use predefined roles and create groups to control access to multiple users with same responsibility

upvoted 57 times

 **droogie** 2 years, 4 months ago

You assume Auditors Group = External Auditors only. Auditors Group may contain both Internal and External Auditors.

upvoted 4 times

 **robor97** 1 year, 11 months ago

The question literally says - External Auditors

upvoted 11 times

 **adeice** 1 year, 7 months ago

I can create External group and Internal group Auditors

upvoted 2 times

 **JavierCorrea** Highly Voted 2 years, 3 months ago

Correct answer is A as per:

https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors

upvoted 27 times

 **vijay456** Most Recent 1 week, 4 days ago

Selected Answer: A<https://cloud.google.com/iam/docs/job-functions/auditing>

upvoted 1 times

 **vijay456** 1 week, 4 days ago

Selected Answer: C

there is no group created and no option says create group too, so provided option C is suitable best answer

upvoted 1 times

 **leogor** 2 weeks, 2 days ago

A, auditors group

upvoted 1 times

 **PSS387** 2 weeks, 5 days ago

Selected Answer: A

Google recommends to group users to allow permissions

upvoted 1 times

 **PKookNN** 4 weeks, 1 day ago

Selected Answer: A

best practices recommend going with group instead of individual users so A is more correct than C

upvoted 2 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

✉ **rixson** 1 month ago

what's the answer?

upvoted 1 times

✉ **DjayTest21** 1 month, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **iadarsh** 1 month, 3 weeks ago

Selected Answer: A

A is Correct

Because if you directly add users to the IAM roles, then if any users left the organization then you have to remove the users from multiple places and need to revoke his/her access from multiple places.

But, if you put a user into a group then its very easy to manage these type of situations. Now, if any user left then you just need to remove the user from the group and all the access got revoked

upvoted 1 times

✉ **12234** 3 months ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

A is right.. don't fall on trap with C

upvoted 2 times

✉ **haroldbenites** 5 months, 2 weeks ago

Go for A

<https://cloud.google.com/iam/docs/job-functions/auditing>

<https://cloud.google.com/iam/docs/job-functions/auditing>

upvoted 1 times

✉ **haroldbenites** 5 months, 2 weeks ago

Scenario: External auditors

The organization creates a Google group for these external auditors and adds the current auditor to the group. This group is monitored and is typically granted access to the dashboard application.

During normal access, the auditors' Google group is only granted access to view the historic logs stored in BigQuery. If any anomalies are discovered, the group is granted permission to view the actual Cloud Logging Admin Activity logs via the dashboard's elevated access mode. At the end of each audit period, the group's access is then revoked.

Data is redacted using Cloud DLP before being made accessible for viewing via the dashboard application.

The table below explains IAM logging roles that an Organization Administrator can grant to the service account used by the dashboard, as well as the resource level at which the role is granted.

upvoted 1 times

✉ **YaaElon** 6 months ago

Selected Answer: A

Best practice is to add to groups

upvoted 2 times

✉ **jgnogueira** 6 months, 2 weeks ago

Selected Answer: A

The main best practice in any attribution of permission is to create a group

upvoted 1 times

✉ **theBestStudent** 6 months, 3 weeks ago

Answer is A. Google recommended practices is to use PREDEFINED roles (As in this case we need granularity to give specific roles to the group). And also use groups

upvoted 1 times

Question #26

You need to set up permissions for a set of Compute Engine instances to enable them to write data into a particular Cloud Storage bucket. You want to follow Google-recommended practices. What should you do?

- A. Create a service account with an access scope. Use the access scope 'https://www.googleapis.com/auth/devstorage.write_only'.
- B. Create a service account with an access scope. Use the access scope '<https://www.googleapis.com/auth/cloud-platform>'.
- C. Create a service account and add it to the IAM role 'storage.objectCreator' for that bucket.
- D. Create a service account and add it to the IAM role 'storage.objectAdmin' for that bucket.

Correct Answer: D*Community vote distribution*

C (100%)

 **coldpar** Highly Voted 2 years, 7 months ago

As per as the least privilege recommended by google, C is the correct Option, A is incorrect because the scope doesn't exist. B incorrect because it will give him full control

upvoted 43 times

 **johnconnor** 4 months, 2 weeks ago

Check here, it is A-> <https://cloud.google.com/storage/docs/authentication>
<https://cloud.google.com/storage/docs/authentication>

upvoted 1 times

 **robor97** 1 year, 11 months ago

The scope does exist - <https://download.huihoo.com/google/gdgdevkit/DVD1/developers.google.com/compute/docs/api/how-tos/authorization.html>

upvoted 1 times

 **peter77** 1 year, 1 month ago

No it doesn't. You have read-only, read-write, full-control and others... but "write-only" is not a thing.

<https://cloud.google.com/storage/docs/authentication>

upvoted 3 times

 **gielda211** 7 months, 1 week ago

it doesn't exist. show us this on official google website

upvoted 2 times

 **XRiddlerX** Highly Voted 2 years, 4 months ago

In reviewing this, it looks to be a multiple answer question. According to Best Practices in this Google Doc (https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices) you grant the instance the scope and the permissions are determined by the IAM roles of the service account. In this case, you would grant the instance the scope and the role (storage.objectCreator) to the service account.

Ans B and C

Role from GCP Console:

ID = roles/storage.objectCreator

Role launch stage = General Availability

Description = Access to create objects in GCS.

3 assigned permissions

resourcemanager.projects.get

resourcemanager.projects.list

storage.objects.create

upvoted 14 times

 **nickyshil** 3 months, 1 week ago

There are many access scopes available to choose from, but a best practice is to set the cloud-platform access scope, which is an OAuth scope for most Google Cloud services, and then control the service account's access by granting it IAM roles..you have an app that reads and writes files on Cloud Storage, it must first authenticate to the Cloud Storage API. You can create an instance with the cloud-platform scope and attach a service account to the instance

<https://cloud.google.com/compute/docs/service-accounts>

upvoted 1 times

 **ryumada** 3 months, 1 week ago

Reading the second point of the best practice. You should grant your VM the <https://www.googleapis.com/auth/cloud-platform> scope to allow access to most of Google Cloud APIs.

So, that the IAM permissions are completely determined by the IAM roles you granted to the service account.

The conclusion is you should not mess up with the VM scopes to grant access to Google Services, instead you should grant the access via IAM roles of the service account you attached to the VM.

https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices

upvoted 1 times

 **dennydream** (Most Recent) 1 week, 5 days ago

How is the right answer D? It should be C. You just need creator rights to write a file there. No way should admin be granted.

upvoted 1 times

 **leogor** 2 weeks, 2 days ago

C, storage.objectCreator

upvoted 1 times

 **Untamables** 2 weeks, 5 days ago

Selected Answer: C

C is the best choice.

https://cloud.google.com/iam/docs/understanding-service-accounts#using_service_accounts_with_compute_engine

upvoted 1 times

 **Sgegmund** 1 month, 1 week ago

Selected Answer: C

Least Privileges

upvoted 1 times

 **adil114** 1 month, 2 weeks ago

ANSWER IS D

upvoted 1 times

 **ale_brd_** 2 months ago

Selected Answer: C

access scope are legacy way and not anymore the best practice even though they can still be used.

upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: C

C is the correct answer

upvoted 1 times

 **gandor** 3 months, 1 week ago

Selected Answer: C

qweqweqweqweqweqw

upvoted 2 times

 **AllMaaya** 3 months, 1 week ago

D must be correct , Since its gives ONLY Object.admin , not Storage Admin

upvoted 1 times

objectCreator has this description:

Allows users to create objects. Does not give permission to view, delete, or replace objects.

objectAdmin has this description:

Grants full control over objects, including listing, creating, viewing, and deleting objects.

The **objectAdmin** has unnecessary permissions that doesn't needed by the question context.

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 3 times

 **polarbear** 4 months, 1 week ago

objectCreator role does not gives permission to write the data and the question says they need to write the data, so I guess Answer is D

upvoted 1 times

ryumada 3 months, 1 week ago

Object Creator has the ability to create data which called object in this description. Object Storage is a storage type that called their stuffs as objects. **objectCreator** has this ability to create this object or what you call it data. A data is an object in the Object Storage type.

<https://cloud.google.com/learn/what-is-object-storage>

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 1 times

✉️ **RanjithK** 4 months, 1 week ago

Answer is D

upvoted 1 times

✉️ **AzureDP900** 4 months, 2 weeks ago

C is the correct answer

upvoted 1 times

✉️ **Ranjith2709** 5 months ago

Selected Answer: C

C is the correct answer

upvoted 2 times

✉️ **haroldbenites** 5 months, 2 weeks ago

Go for C

<https://cloud.google.com/storage/docs/access-control/iam-roles>

upvoted 2 times

✉️ **Akash7** 6 months, 2 weeks ago

Answer is C, access scope are a legacy way....

upvoted 2 times

Question #27

You have sensitive data stored in three Cloud Storage buckets and have enabled data access logging. You want to verify activities for a particular user for these buckets, using the fewest possible steps. You need to verify the addition of metadata labels and which files have been viewed from those buckets. What should you do?

- A. Using the GCP Console, filter the Activity log to view the information.
- B. Using the GCP Console, filter the Stackdriver log to view the information.
- C. View the bucket in the Storage section of the GCP Console.
- D. Create a trace in Stackdriver to view the information.

Correct Answer: A*Community vote distribution*

A (53%) B (47%)

 **iamgcp** Highly Voted 2 years, 5 months ago

A is correct. As mentioned in the question, data access logging is enabled. I tried to download a file from a bucket and was able to view this information in Activity tab in console

upvoted 31 times

 **MEHDIGRB** 1 month ago

activity log is deprecated:

<https://cloud.google.com/compute/docs/logging/activity-logs>

upvoted 2 times

 **vito9630** 2 years, 5 months ago

data access logging don't provide information about addition of metadata, so B is correct

upvoted 17 times

 **RegisFTM** 10 months, 3 weeks ago

I did all the configuration enabling data access logging but I still not able to see the logs when uploading or downloading a file. Does someone here has done it with a different result?

upvoted 1 times

 **ryumada** 3 months, 1 week ago

I agree with liyux21 and vito9630. In this reference link below says:

In the Activity page, where the identity performing logged actions is redacted from the audit log entry, User (anonymized) is displayed.

Beacause of this, I think you can't verify the addition of metadata labels through Activity Logs.

<https://cloud.google.com/logging/docs/audit#view-activity>

upvoted 1 times

 **eliteone11** Highly Voted 1 year, 11 months ago

Answer is A. Activity log does indeed show information about metadata.

I agree with Eshkrkrkr based on <https://cloud.google.com/storage/docs/audit-logs> Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

upvoted 11 times

 **injarapu** 8 months, 1 week ago

'Admin activity logs' capture metadata modification, but its different from 'Data Access logging', right ?

upvoted 2 times

 **Stephanyjempot** Most Recent 17 hours ago

Selected Answer: B

B IS CORRECT

upvoted 1 times

 **Untamables** 2 weeks, 5 days ago

Selected Answer: B

B is correct.

In this scenario, you need to select data access audit logs in Cloud Logging.

<https://cloud.google.com/storage/docs/audit-logging>

Note the Stackdriver Logging is now named Cloud Logging.

upvoted 2 times

 **Tanvinitish** 3 weeks, 1 day ago

B is the right ans. Activity log is not for data access audit logs, use log explorer for this purpose
upvoted 1 times

 **AwesomeGCP** 4 weeks, 1 day ago

Selected Answer: A

A. Using the GCP Console, filter the Activity log to view the information.
upvoted 1 times

 **anolive** 1 month, 3 weeks ago

A is correct based on <https://cloud.google.com/storage/docs/audit-logs>
upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: A

Activity logs option A
upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Activity log does have data access option but it doesn't show any result when I download a couple of objects from my GCS bucket. Infact, I tried looking into the stackdriver logs as well but haven't found anything. Does it usually take some time to show ?
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is correct
upvoted 1 times

 **bamosk** 4 months, 3 weeks ago

Selected Answer: B

Link: <https://cloud.google.com/logging/docs/audit#user-id>

If you're viewing audit logs using the Google Cloud console Activity page, User (anonymized) is displayed for any log entries where identity is redacted or empty.

<https://cloud.google.com/logging/docs/audit#view-activity>

Even though Data access audit log type is enabled, in order to obtain Updating bucket metadata - Admin activity audit log type needs to be filtered which can be done in Stackdriver and not in "Activity page".

Best practice suggests that for CloudStorage "Audit log" is better option than user or storage log.

Audit logs are enabled by default <https://cloud.google.com/storage/docs/audit-logging#settings>

Dataaccess logs needs to be explicitly enabled.

<https://cloud.google.com/blog/products/it-ops/best-practices-for-working-with-google-cloud-audit-logging>

You can view audit logs from two places in the GCP Console: via the activity feed, which provides summary entries, and via the Stackdriver Logs viewer page, which gives full entries.

So answer is B. Not A

upvoted 3 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

<https://cloud.google.com/storage/docs/audit-logging>

Note: Cloud Audit Logs is separate from Cloud Storage usage logs. Both features provide information about who accesses your Cloud Storage resources, but Cloud Audit Logs is the recommended method for doing so.

upvoted 1 times

 **zjitz** 7 months, 2 weeks ago

Selected Answer: A

Audit logs meet the requirement

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object.

upvoted 2 times

 **ErenYeager** 8 months, 1 week ago

Selected Answer: A

Compute Engine activity logs are deprecated. Which is now audit logs is more suitable for verifying the addition of metadata labels as described here:

https://cloud.google.com/compute/docs/logging/audit-logging#audited_operations

upvoted 2 times

 **somenick** 8 months, 1 week ago

Selected Answer: A

<https://cloud.google.com/migrate/compute-engine/docs/4.8/how-to/monitoring/viewing-stackdriver-logs>

<https://cloud.google.com/products/operations>

<https://cloud.google.com/logging/docs/audit/configure-data-access>

upvoted 1 times

 **tigerbaer** 9 months ago

Selected Answer: B

B is incorrect, in the activity log, user information is redacted , so you can not verify activities for a particular user
upvoted 1 times

 **likki8798** 9 months, 1 week ago

Selected Answer: A

Its A sure
upvoted 1 times

Question #28

You are the project owner of a GCP project and want to delegate control to colleagues to manage buckets and files in Cloud Storage. You want to follow Google-recommended practices. Which IAM roles should you grant your colleagues?

- A. Project Editor
- B. Storage Admin
- C. Storage Object Admin
- D. Storage Object Creator

Correct Answer: B*Community vote distribution*

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Storage Admin (roles/storage.admin) Grants full control of buckets and objects.

When applied to an individual bucket, control applies only to the specified bucket and objects within the bucket.

```
firebase.projects.get
resourcemanager.projects.get
resourcemanager.projects.list
storage.buckets.*
storage.objects.*
```

upvoted 41 times

 **iambatmanadarkknight** 1 year ago

why not storage object admin?

upvoted 3 times

 **TenshiID** 1 year ago

Because the objet admin don't have control over buckets and you need it

upvoted 12 times

 **RazOr** 10 months ago

Exactly, you want to give someone right to edit storages not just objects. Google does this kind of answers to confuse us.

upvoted 2 times

 **dang1986** 8 months, 2 weeks ago

Question states "Buckets and Objects"

upvoted 2 times

 **Agents89** Highly Voted 2 years, 6 months ago

B is correct

upvoted 14 times

 **leogor** Most Recent 2 weeks, 2 days ago

B. Storage Admin

upvoted 1 times

 **Jesse20323645** 3 weeks ago

Selected Answer: B

According to the question, your colleagues need to manage "buckets" in Cloud Storage(storage.objects.* permission), so (B) is correct. (C) doesn't have control over the buckets.

Cloud document:

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

-->Storage Object Admin (Cannot find "storage.buckets.*" Permission)

---->Storage Admin (Has "storage.buckets.*" Permission)

upvoted 1 times

 **leogor** 1 month, 2 weeks ago

Selected Answer: B

the ans is definitively B

upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: B

Storage Admin option A - As we are supposed to create buckets as well
upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is B
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Storage Admin is right.. B is correct
upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for B
upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

<https://cloud.google.com/storage/docs/access-control/iam-roles>
upvoted 1 times

 **gielda211** 7 months, 1 week ago

Selected Answer: B

<https://cloud.google.com/storage/docs/access-control/iam-roles>
upvoted 3 times

 **gielda211** 7 months, 1 week ago

Selected Answer: B

B. Storage Admin
upvoted 2 times

 **likki8798** 9 months, 1 week ago

Selected Answer: B

B for sure
upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

B. Storage Admin
upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Storage Admin
upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

B is correct
upvoted 1 times

 **Student7** 1 year, 3 months ago

B. Storage Admin
upvoted 1 times

 **Furqon** 1 year, 5 months ago

B is the correct one
upvoted 1 times

Question #29

You have an object in a Cloud Storage bucket that you want to share with an external company. The object contains sensitive data. You want access to the content to be removed after four hours. The external company does not have a Google account to which you can grant specific user-based access privileges. You want to use the most secure method that requires the fewest steps. What should you do?

- A. Create a signed URL with a four-hour expiration and share the URL with the company.
- B. Set object access to 'public' and use object lifecycle management to remove the object after four hours.
- C. Configure the storage bucket as a static website and furnish the object's URL to the company. Delete the object from the storage bucket after four hours.
- D. Create a new Cloud Storage bucket specifically for the external company to access. Copy the object to that bucket. Delete the bucket after four hours have passed.

Correct Answer: A*Community vote distribution*

A (100%)

 **JJ_ME** Highly Voted 2 years, 1 month ago

A.

Signed URLs are used to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account.

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 31 times

 **Agents89** Highly Voted 2 years, 6 months ago

A is correct

upvoted 9 times

 **leogor** Most Recent 2 weeks, 2 days ago

A. signed URL

upvoted 1 times

 **Jesse20323645** 3 weeks ago

Selected Answer: A

Go for (A)

Anyone in possession of the signed URL can use it while it's active, regardless of whether they have a Google account. It gives time-limited access to a specific Cloud Storage resource.

Reference:

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 1 times

 **Vidhey** 3 weeks, 2 days ago

Selected Answer: A

A is correct

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **sandeshsegu** 4 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

 **Alok17** 4 months, 2 weeks ago

A is correct

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right.

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

go for A

upvoted 1 times

 **Himadhar1997** 5 months, 2 weeks ago

A IS CORRECT. signed URLs give time-limited access to a specific Cloud Storage resource.

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

upvoted 1 times

 **gielda211** 7 months, 1 week ago

Selected Answer: A

A IS CORRECT. signed URLs give time-limited access to a specific Cloud Storage resource. Anyone in possession of the signed URL can use it while it's active, regardless of whether they have a Google account

upvoted 2 times

 **PawanS** 11 months ago

Selected Answer: A

Signed URLs are very effective to give a time limited access to anyone

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

A. Create a signed URL with a four-hour expiration and share the URL with the company.

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Create a signed URL with a four-hour expiration and share the URL with the company.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

A is correct

upvoted 1 times

Question #30

You are creating a Google Kubernetes Engine (GKE) cluster with a cluster autoscaler feature enabled. You need to make sure that each node of the cluster will run a monitoring pod that sends container metrics to a third-party monitoring solution. What should you do?

- A. Deploy the monitoring pod in a StatefulSet object.
- B. Deploy the monitoring pod in a DaemonSet object.
- C. Reference the monitoring pod in a Deployment object.
- D. Reference the monitoring pod in a cluster initializer at the GKE cluster creation time.

Correct Answer: B*Community vote distribution*

B (100%)

 **JackGlemins** Highly Voted 1 year, 8 months ago

B is right: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>
Some typical uses of a DaemonSet are:

running a cluster storage daemon on every node
running a logs collection daemon on every node
running a node monitoring daemon on every node
upvoted 27 times

 **Agents89** Highly Voted 2 years, 6 months ago

B is correct
upvoted 24 times

 **leogor** Most Recent 2 weeks, 2 days ago

B DaemonSet
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct
upvoted 1 times

 **pfabio** 5 months, 1 week ago

Selected Answer: B
DaemonSets attempt to adhere to a one-Pod-per-node model, either across the entire cluster or a subset of nodes. As you add nodes to a node pool, DaemonSets automatically add Pods to the new nodes as needed.
upvoted 5 times

 **haroldbenites** 5 months, 2 weeks ago

go for B
upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for B
upvoted 1 times

 **adityachowta** 6 months, 1 week ago

Selected Answer: B
B is correct
upvoted 1 times

 **obeythefist** 8 months, 1 week ago

"Every Node" is the keyword here, which is what DaemonSet is used for
upvoted 1 times

 **ionutr** 9 months, 2 weeks ago

Selected Answer: B
B 2000%
upvoted 1 times

 **Skyerh** 11 months, 1 week ago

Selected Answer: B
I will vote B

upvoted 3 times

 **shawnkkk** 11 months, 3 weeks ago

B. Deploy the monitoring pod in a DaemonSet object.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Deploy the monitoring pod in a DaemonSet object.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

B is correct

upvoted 1 times

 **vikram___** 1 year, 2 months ago

B is right <https://cloud.google.com/kubernetes-engine/docs/concepts/daemonset>

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

B is correct answer

upvoted 1 times

 **max94** 1 year, 6 months ago

B is correct https://cloud.google.com/kubernetes-engine/docs/concepts/daemonset#usage_patterns

upvoted 1 times

Question #31

You want to send and consume Cloud Pub/Sub messages from your App Engine application. The Cloud Pub/Sub API is currently disabled. You will use a service account to authenticate your application to the API. You want to make sure your application can use Cloud Pub/Sub. What should you do?

- A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
- B. Rely on the automatic enablement of the Cloud Pub/Sub API when the Service Account accesses it.
- C. Use Deployment Manager to deploy your application. Rely on the automatic enablement of all APIs used by the application being deployed.
- D. Grant the App Engine Default service account the role of Cloud Pub/Sub Admin. Have your application enable the API on the first connection to Cloud Pub/Sub.

Correct Answer: A

Community vote distribution

A (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A)

Quickstart: using the Google Cloud Console

This page shows you how to perform basic tasks in Pub/Sub using the Google Cloud Console.

Note: If you are new to Pub/Sub, we recommend that you start with the interactive tutorial.

Before you begin

Set up a Cloud Console project.

Set up a project

Click to:

Create or select a project.

Enable the Pub/Sub API for that project.

You can view and manage these resources at any time in the Cloud Console.

Install and initialize the Cloud SDK.

Note: You can run the gcloud tool in the Cloud Console without installing the Cloud SDK. To run the gcloud tool in the Cloud Console, use Cloud Shell .

<https://cloud.google.com/pubsub/docs/quickstart-console>

upvoted 25 times

 **Bharathy** Highly Voted 2 years, 5 months ago

We need to enable the pub/sub API, if we are going to use it in your project... then APP engine can able to access it with required ServiceAccount
upvoted 23 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

A is right

upvoted 1 times

 **sedado77** 5 months, 1 week ago

Selected Answer: A

Yup its A

upvoted 2 times

 **haroldbenites** 5 months, 2 weeks ago

go for A

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

upvoted 1 times

 **Schladde** 5 months, 3 weeks ago

Selected Answer: A

A is correct, although there are options to activate APIs programmatically (gcloud, curl) <https://cloud.google.com/service-usage/docs/enable-disable>

B and C are incorrect because "rely on ..." is never a good option

D could be possible, but too much permissions are given to the app

upvoted 2 times

✉ **gielda211** 7 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **flashlight01** 7 months, 3 weeks ago

Selected Answer: A

Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

upvoted 1 times

✉ **obeythefist** 8 months, 1 week ago

This question seems to be asking "Can you automatically enable the API or do you have to do it manually?" and I think the answer is that there's no automatic enablement of the API.

upvoted 1 times

✉ **shawnkkk** 11 months, 3 weeks ago

A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

upvoted 2 times

✉ **vishnukumartr** 11 months, 3 weeks ago

A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

upvoted 1 times

✉ **Jaira1256** 11 months, 3 weeks ago

A is correct

upvoted 1 times

✉ **vikram___** 1 year, 2 months ago

Enabling API is the first step before using it, Ans A

upvoted 1 times

✉ **Kaushik_Mallick** 1 year, 3 months ago

Does anyone has latest questions & answers. Also if you guide me the reading material/blog/portal to pass this exam.

upvoted 4 times

✉ **BAGCP** 1 year, 3 months ago

Even i also want this information

upvoted 1 times

✉ **Nick334** 1 year, 3 months ago

I am also looking for it.

upvoted 1 times

✉ **vishnukumartr** 1 year, 3 months ago

I'm also friends...your help will be much appreciable (vishnutr98@gmail.com)

upvoted 1 times

✉ **mcaromit** 1 year, 6 months ago

As a first step A is the correct answer

upvoted 2 times

✉ **gspb** 1 year, 6 months ago

In the question - "You want to make sure your application can use Cloud Pub/Sub"

Since we want to make sure that our application uses pub/sub, enabling it manually is the best option.

Hence, option A is correct

upvoted 1 times

Question #32

You need to monitor resources that are distributed over different projects in Google Cloud Platform. You want to consolidate reporting under the same Stackdriver.

Monitoring dashboard. What should you do?

- A. Use Shared VPC to connect all projects, and link Stackdriver to one of the projects.
- B. For each project, create a Stackdriver account. In each project, create a service account for that project and grant it the role of Stackdriver Account Editor in all other projects.
- C. Configure a single Stackdriver account, and link all projects to the same account.
- D. Configure a single Stackdriver account for one of the projects. In Stackdriver, create a Group and add the other project names as criteria for that Group.

Correct Answer: D*Community vote distribution*

C (79%)

D (21%)

 **sahedge** Highly Voted  1 year, 10 months ago

First of all D is incorrect, Groups are used to define alerts on set of resources(such as VM instances, databases, and load balancers). FYI tried adding Two projects into a group it did not allowed me as the "AND"/"OR" criteria for the group failed with this combination of resources.

C is correct because,

When you intially click on Monitoring(Stackdriver Monitoring) it creates a workspace(a stackdriver account) linked to the ACTIVE(CURRENT) Project from which it was clicked.

Now if you change the project and again click onto Monitoring it would create an another workspace(a stackdriver account) linked to the changed ACTIVE(CURRENT) Project, we don't want this as this would not consolidate our result into a single dashboard(workspace/stackdriver account).

If you have accidentally created two diff workspaces merge them under Monitoring > Settings > Merge Workspaces > MERGE.

If we have only one workspace and two projects we can simply add other GCP Project under Monitoring > Settings > GCP Projects > Add GCP Projects.

In both of these cases we did not create a GROUP, we just linked GCP Project to the workspace(stackdriver account).

upvoted 67 times

 **coldpar** Highly Voted  2 years, 7 months ago

C is correct not D

upvoted 32 times

 **Sanay** Most Recent  3 days, 12 hours ago

If anyone is the contributor please share the dumps on shah.sana@northeastern.edu

upvoted 1 times

 **leogor** 2 weeks, 2 days ago

C. Configure a single Stackdriver account, and link all projects to the same account.

upvoted 1 times

 **KravitejaK** 2 weeks, 3 days ago

if any one have contributor access please share all the questions pdf to raviteja922.kotha@outlook.com. I have scheduled the exam tomorrow

upvoted 1 times

 **Priam** 1 month, 1 week ago

if anyone has contributor access please share the pdf of all questions with answers to - devappy07@gmail.com

upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: C

Option C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is correct

upvoted 1 times

 **madadis** 4 months, 3 weeks ago

Selected Answer: C

C is Correct

upvoted 1 times

  **Tirthankar17** 5 months ago**Selected Answer: C**

C is correct.

upvoted 1 times

  **ravenor** 5 months, 2 weeks ago

"View metrics for multiple Cloud projects:

C.

<https://cloud.google.com/monitoring/settings/multiple-projects>

Nothing about groups

<https://cloud.google.com/monitoring/settings?hl=en>

upvoted 1 times

  **haroldbenites** 5 months, 2 weeks ago

Go for C

upvoted 1 times

  **haroldbenites** 5 months, 2 weeks ago

Go for C.

upvoted 1 times

  **HansKloss611** 9 months, 2 weeks ago**Selected Answer: C**

C is correct

upvoted 3 times

  **alshad** 10 months, 2 weeks ago**Selected Answer: C**

C is correct

upvoted 1 times

  **rafsrod** 10 months, 3 weeks ago**Selected Answer: C**

C is correct, not D

upvoted 1 times

  **arapin** 10 months, 4 weeks ago**Selected Answer: C**

The correct answer is C.

Workspaces can track your existing GCP projects, which form the basis for managing permissions and resources within GCP, as well as any Amazon Web Services (AWS) accounts that you want to monitor.

<https://cloud.google.com/blog/products/management-tools/using-stackdriver-workspaces-help-manage-your-hybrid-and-multicloud-environment>

D is incorrect because:

Groups provide a mechanism for alerting on the behavior of a set of resources, rather than on individual resources.

<https://cloud.google.com/monitoring/groups/>

A and B are not even worth mentioning.

upvoted 1 times

Question #33

You are deploying an application to a Compute Engine VM in a managed instance group. The application must be running at all times, but only a single instance of the VM should run per GCP project. How should you configure the instance group?

- A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- B. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- C. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 2.
- D. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 2.

Correct Answer: A*Community vote distribution*

A (100%)

 **XRiddlerX** Highly Voted 2 years, 4 months ago

In my GCP console, I created a managed instance group for each answer. For each answer I deleted the instance that was created as a simple test to prove or disprove each answer.

In answer A, another instance was created after I deleted the instance

In answer B, no other instance was created after I deleted the instance

In answer C, another instance was created after I deleted the instance

In answer D, no other instance was created after I deleted the instance

My observation is A is the correct Answer.

A - Correct - It correctly solves the problem with only a single instance at one time

B - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created

C - Incorrect - It creates another instance after the delete HOWEVER it 2 VM's could be created even if the target is exceeded

D - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created

upvoted 108 times

 **Verve** 2 years, 3 months ago

Thanks for confirming this for us.

upvoted 1 times

 **ESP_SAP** 2 years, 2 months ago

thanks so much for your time to test every option.

This kind of collaboration really clarify the doubts!

upvoted 8 times

 **prasaru** 2 years, 1 month ago

thanks for confirming

upvoted 1 times

 **soukumar369** 1 year, 12 months ago

Thank you very much for this detail testing and explaination.

upvoted 1 times

 **khaliqn** Highly Voted 2 years, 4 months ago

We want the application running at all times. If the VM crashes due to any underlying hardware failure, we want another instance to be added so we need autoscaling ON

Correct answer is A

upvoted 23 times

 **Gurnoor** 2 years, 4 months ago

A is wrong - What you are talking about is Autohealing. Autoscale will not rebuild the VM on the crash.

upvoted 4 times

 **Dips_** 1 year, 10 months ago

won't auto healing take time to recreate that VM? and there should be one VM running all time

upvoted 1 times

 **Ozymandiasx** 1 year, 10 months ago

As said in other coment, you can fix this with Autohealing. autoscaling means more machines, autohealing means re creating of VM's

upvoted 1 times

 **GCP_Student1** 1 year, 9 months ago

Even if I agree with your response, it still does not meet the requirements asked in the question, which is "the application should be running all the time", because in your case when the VM instance for whatever reason stops, that mean the application will experience a downtime.
upvoted 2 times

 **[Removed]** 1 year, 7 months ago

Read this warning message: "The minimum number of instances is equal to maximum number of instances. This means the autoscaler cannot add or remove instances from the instance group. Make sure this is the correct setting."
So when minimum is equal to maximum, it does not matter whether autoscaling is on or off. So now the question is who takes care of running the MINIMUM instances: MIG itself.
upvoted 2 times

 **leogor** (Most Recent) 2 weeks, 2 days ago

A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
upvoted 1 times

 **SaiSaiA** 3 months, 4 weeks ago

Selected Answer: A

I think the answer should be A. People who mentioned about autohealing but it is not indicated in the provided choices. As per GCP <https://cloud.google.com/compute/docs/autoscaler#specifications>

Autoscaling works independently from autohealing. If you configure autohealing for your group and an instance fails the health check, the autohealer attempts to recreate the instance. Recreating an instance can cause the number of instances in the group to fall below the autoscaling threshold (minNumReplicas) that you specify.

You can configure a MIG without Autohealing and if we stick to the choices that are provided, then it should be A.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right answer.

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for A

upvoted 1 times

 **Teraflow** 7 months ago

Selected Answer: A

its obvioussss

upvoted 1 times

 **zjitz** 7 months, 2 weeks ago

Autoscaling setting is required to make sure MIG is able to create a new instance.

With Autoscaling off it doesn't create/delete instances. so Answer is A

upvoted 1 times

 **sai1621** 10 months, 1 week ago

Answer is B

B worked for me. I was able to recreate the scenario.

ASG - ON is required only if we want to scale up / down based on any instance metric. Hence the decision.

upvoted 1 times

 **arapin** 10 months, 4 weeks ago

Selected Answer: A

If I assume that there is a mistake in the answers, and it is supposed to be AUTOHEALING instead of AUTOSCALING, than the answer under A makes sense. Otherwise autoscaling makes no difference when there must be only a single instance at all times.

So... I'll pick A as an answer with the assumption from above.

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

upvoted 1 times

 **vishnukumart** 11 months, 3 weeks ago

A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Correct Answer is A

upvoted 1 times

 **BobLoblawsLawBlog** 1 year, 1 month ago

Going with B here after reviewing this for a long time.

I've seen questions regarding B such as 'What happens if the instance stops for some reason?".

First: "Autoscaling works independently from autohealing. If you configure autohealing for your group and an instance fails the health check, the autohealer attempts to recreate the instance." - see <https://cloud.google.com/compute/docs/autoscaler>

Second: "If you configure an application-based health check and the autohealer determines that your application isn't responding, the MIG automatically recreates that VM. For a preemptible VM, the group recreates the VM when the necessary resources become available again." - see <https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs>

upvoted 1 times

 **BinaryBatman** 1 year, 1 month ago

Answer should be B, high availability is feature of MIGs that says "If a VM in the group stops, crashes, or is deleted by an action other than an instance group management command (for example, an intentional scale in), the MIG automatically recreates that VM".

upvoted 1 times

Question #34

You want to verify the IAM users and roles assigned within a GCP project named my-project. What should you do?

- A. Run gcloud iam roles list. Review the output section.
- B. Run gcloud iam service-accounts list. Review the output section.
- C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.
- D. Navigate to the project and then to the Roles section in the GCP Console. Review the roles and status.

Correct Answer: C

Community vote distribution

C (100%)

 **coldpar** Highly Voted 2 years, 7 months ago

Correct answer is C as IAM section provides the list of both Members and Roles. Option A is wrong as it would provide information about the roles only. Option B is wrong as it would provide only the service accounts. Option D is wrong as it would provide information about the roles only.
upvoted 59 times

 **Agents89** Highly Voted 2 years, 6 months ago

C is the correct answer

upvoted 10 times

 **leogor** Most Recent 2 weeks, 2 days ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 1 times

 **SaiSaiA** 3 months, 4 weeks ago

Selected Answer: C

C is the only logical answer. If you go IAM & Admin > IAM: You can see Principals and Roles. Users, groups, service accounts are Principals
upvoted 1 times

 **NOOGLer** 4 months ago

Selected Answer: C

C is correct

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right! A & B doesn't satisfy the requirement.

upvoted 1 times

 **Himadhar1997** 5 months, 2 weeks ago

Correct answer is C as IAM section provides the list of both Members and Roles. Option A is wrong as it would provide information about the roles only. Option B is wrong as it would provide only the service accounts. Option D is wrong

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

go for C

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for C

upvoted 1 times

 **HarshaKumar** 6 months, 2 weeks ago

Selected Answer: C

C is correct

upvoted 2 times

 **obeythefist** 8 months, 1 week ago

I chose C after experimenting in the console.

A. Wrong, this just shows you all the roles, not users etc. It's not useful.

B. Wrong, This will just show service accounts, not users and roles like the question asks

C. Correct. I logged onto console and followed the steps and was able to see all the assigned users and roles.

D. Wrong. We need to see how the roles are used. I couldn't even see a "Roles" option directly, you need to browse to IAM then to the Roles subsection, which is not useful anyway.

upvoted 4 times

 **shawnkkk** 11 months, 3 weeks ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 2 times

 **Jaira1256** 11 months, 3 weeks ago

C is correct

upvoted 2 times

 **sidharthwader** 1 year, 4 months ago

A lists all the roles in gcp

B lists all the service account

C is right answer

upvoted 3 times

 **HogwartsTrue** 1 year, 5 months ago

Ans is C

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

C is correct

upvoted 1 times

Question #35

You need to create a new billing account and then link it with an existing Google Cloud Platform project. What should you do?

- A. Verify that you are Project Billing Manager for the GCP project. Update the existing project to link it to the existing billing account.
- B. Verify that you are Project Billing Manager for the GCP project. Create a new billing account and link the new billing account to the existing project.
- C. Verify that you are Billing Administrator for the billing account. Create a new project and link the new project to the existing billing account.
- D. Verify that you are Billing Administrator for the billing account. Update the existing project to link it to the existing billing account.

Correct Answer: B*Community vote distribution*

B (72%)

A (25%)

 **coldpar** Highly Voted 2 years, 7 months ago

D is correct as the project is already created
upvoted 41 times

 **Shubham0072** 9 months, 1 week ago

The Question clearly says "Create a new Billing account" and a Bill Administrator cannot create A billing account so according to the question the clear Answer is B.

upvoted 12 times

 **spudleymcdudley** 2 years, 4 months ago

Listen to this guy. I spent 10 minutes and came to the same. <https://cloud.google.com/billing/docs/how-to/billing-access>
upvoted 7 times

 **lxgywil** 1 year, 6 months ago

This documentation is literally saying:

"Billing Account Administrator (roles/billing.admin) - Manage billing accounts (but not create them). "

This role has nothing to do with what is required, while Project Billing Manager allows to link/unlink the project to/from a billing account.

upvoted 8 times

 **UtsavDM** 1 year, 2 months ago

It says Billing Account Admin can also link/ unlink projects from the billing account.

Here -

Billing Account Administrator:

This role is an owner role for a billing account. Use it to manage payment instruments, configure billing exports, view cost information, link and unlink projects and manage other user roles on the billing account.

upvoted 2 times

 **bhavik7899** 1 year, 4 months ago

billing administration can not create billing account
so A

upvoted 2 times

 **afooh** 9 months, 3 weeks ago

B is the correct answer. Existing project with new billing account.

upvoted 5 times

 **Agents89** Highly Voted 2 years, 6 months ago

Project Billing Manager cannot create a billing account, there is nothing like Billing Administrator it is Billing Account Administrator. Both Project Billing Manager and Billing Account Administrator cannot create a billing account. A is the only answer that make sense. We have to assume the the billing account is already created

upvoted 25 times

 **jcloud965** 1 year, 3 months ago

Wrong because "You need to create a new billing account". You can't assume the billing account is already created.

upvoted 7 times

 **spudleymcdudley** 2 years, 4 months ago

'Verify that you are Billing Administrator for the billing account' - aka Billing Account Administrator. A and B do not have privilege to make a new billing account (<https://cloud.google.com/billing/docs/how-to/billing-access>). C is correct, this has been incorrectly up voted

upvoted 4 times

 **nightflyer** 1 year, 10 months ago

The answer is A, yes

upvoted 5 times

✉ **ankatsu2010** 1 year, 1 month ago

Project Billing Manager can't create a billing account, still, A is the only feasible answer.

upvoted 1 times

✉ **kamranthomas** [Most Recent] 4 days, 18 hours ago

THB, this one got me a little confused, but I think I figured out the logic:

I suppose in this question you need to assume you have permissions to create a new billing account. Since both Billing Administrator and Project Billing Manager roles, do not have enough permissions to create a billing account. Therefore, by following google cloud best practices, more specifically the principle of least privilege, we have eliminated the options with Billing Administrator (C and D). Also the question says you need to create a new Billing account (so it does not exist) which eliminates option A, therefore option B is the most correct answer imo.

Hope this helps!

upvoted 1 times

✉ **leogor** 2 weeks, 2 days ago

Selected Answer: B

Ans is obviously B

- A. Update the existing project to "link it to the existing billing account"
- B. Create a new billing account and "link the new billing account to the existing project"

upvoted 1 times

✉ **PKookNN** 4 weeks, 1 day ago

Selected Answer: B

B is correct per: "create a new billing account then link to existing project" Option A has existing billing acc which does not meet the first condition.

upvoted 1 times

✉ **Priyanka109** 1 month ago

It's mentioned in the question itself that you have to create a new billing account so B.

upvoted 1 times

✉ **VaneA** 1 month, 1 week ago

Selected Answer: B

B is correct

upvoted 1 times

✉ **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

✉ **gcpviewer** 1 month, 2 weeks ago

Selected Answer: B

Answer is B. Billing Administrators can not create a new billing account, and the project is presumably already created. Project Billing Manager allows you to link the created billing account to the project. It is vague on how the billing account gets created but by process of elimination, I believe B to be the correct answer.

upvoted 1 times

✉ **ashit44244** 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

✉ **anolive** 1 month, 3 weeks ago

To me B makes more sense, if the question is asking for create a new billing account and then link it with an existing Google Cloud Platform project.

upvoted 1 times

✉ **learn_GCP** 1 month, 3 weeks ago

Selected Answer: B

B is correct as its mentioned to create new billing account.
in A ur not creating new billing account.

upvoted 1 times

✉ **sandipk91** 2 months, 1 week ago

Selected Answer: B

Option B

upvoted 1 times

✉ **GayuSundar** 2 months, 1 week ago

B is correct. Why A is not correct, it is saying linking existing billing account to existing project, but we need to create new billing account.

upvoted 1 times

✉ **habros** 3 months ago

I'm sticking with B. A is non-existent as there is no existing billing accounts to attach to!

upvoted 2 times

✉ **12234** 3 months ago

The answer can be both A or D as neither billing administrator nor Project Billing Manager can create new billing account we can just assume the account was already created.

upvoted 1 times

✉ **ryumada** 3 months ago

I don't know what is the right answer. Project Billing Manager and Billing Administrator are the role that can assign billing account to a project. They both have this permission:

`resourcemanager.projects.createBillingAssignment`

The only role that has the permission to create a billing account is Billing Account Creator. The role has this permission in order to create a billing account:

`billing.accounts.create`

That permission is neither in the Project Billing Manager nor the Billing Administrator role.

This make the answer in B and C can be eliminated. But this also makes the answer in A and D Equivalent.

Ref:

- https://cloud.google.com/billing/docs/how-to/custom-roles#billing_account_hierarchy
- https://cloud.google.com/billing/docs/how-to/custom-roles#resource_associations
- <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 2 times

Question #36

You have one project called proj-sa where you manage all your service accounts. You want to be able to use a service account from this project to take snapshots of VMs running in another project called proj-vm. What should you do?

- A. Download the private key from the service account, and add it to each VM's custom metadata.
- B. Download the private key from the service account, and add the private key to each VM's SSH keys.
- C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
- D. When creating the VMs, set the service account's API scope for Compute Engine to read/write.

Correct Answer: C*Community vote distribution*

C (100%)

 **jackdbd** Highly Voted 1 year, 3 months ago

C is the correct answer.

It took me a while to figure it out because I didn't understand how service accounts work across project. This article made it clear for me. <https://gtseries.medium.com/using-service-accounts-across-projects-in-gcp-cf9473fef8f0>

You create the service account in proj-sa and take note of the service account email, then you go to proj-vm in IAM > ADD and add the service account's email as new member and give it the Compute Storage Admin role.

upvoted 25 times

 **JelloMan** 7 months, 2 weeks ago

As of now, service accounts may be impersonated (new-term). AKA, you can create a service account in one project and then impersonate it in others. Essentially, it involves the same steps as what the medium article suggests (create a service account in the principal (main) project and then add the email of the main project to the project you want to impersonate) <https://cloud.google.com/iam/docs/impersonating-service-accounts#impersonate-sa-level>

upvoted 1 times

 **SaiSaiA** 3 months, 4 weeks ago

I have tried C, it doesn't work. Also, this refers to a different Principal (user) impersonating a Service Account which is a different case from what is in the question.

upvoted 1 times

 **kishoredeena** Highly Voted 2 years, 4 months ago

Option C is the right one

upvoted 22 times

 **leogor** Most Recent 2 weeks, 2 days ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

 **habros** 2 months, 3 weeks ago

Safe to eliminate any options that demand transferring of private keys. NOT SAFE

Hence, C.

upvoted 2 times

 **theBestStudent** 1 month, 1 week ago

highly agree with this thoughts! transferring private keys is a big no no here.

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C. is the correct answer

Compute Storage Admin
(roles/compute.storageAdmin)
Permissions to create, modify, and delete disks, images, and snapshots.

For example, if your company has someone who manages project images and you don't want them to have the editor role on the project, then grant this role to their account on the project.

Lowest-level resources where you can grant this role:

Disk
Image
Snapshot Beta
upvoted 3 times

✉ **haroldbenites** 5 months, 2 weeks ago

go for C
upvoted 1 times

✉ **somenick** 8 months, 1 week ago

Selected Answer: C

<https://cloud.google.com/compute/docs/access/iam#compute.storageAdmin>

upvoted 2 times

✉ **RealEL40** 11 months ago

When a service account is in one project, and it accesses a resource in another project, you usually must enable the API for that resource in both projects. For example, if you have a service account in the project my-service-accounts and a Cloud SQL instance in the project my-application, you must enable the Cloud SQL API in both my-service-accounts and my-application.

upvoted 2 times

✉ **shawnkkk** 11 months, 3 weeks ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
upvoted 1 times

✉ **vishnukumartr** 11 months, 3 weeks ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
upvoted 1 times

✉ **Jaira1256** 11 months, 3 weeks ago

C is correct
upvoted 1 times

✉ **cuongnd** 1 year, 4 months ago

C. Compute Storage Admin role has this: compute.snapshots.*
upvoted 4 times

✉ **mcaromit** 1 year, 6 months ago

C is the most apt from the given options, though the question is unclear
upvoted 2 times

✉ **[Removed]** 1 year, 7 months ago

C is correct. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
upvoted 2 times

✉ **pondai** 1 year, 7 months ago

This quest has problem
upvoted 1 times

✉ **pondai** 1 year, 7 months ago

I'll take it back. I try this solution. It success. C is correct.
upvoted 1 times

✉ **yushi69** 1 year, 7 months ago

C is correct
upvoted 1 times

Question #37

You created a Google Cloud Platform project with an App Engine application inside the project. You initially configured the application to be served from the us-central region. Now you want the application to be served from the asia-northeast1 region. What should you do?

- A. Change the default region property setting in the existing GCP project to asia-northeast1.
- B. Change the region property setting in the existing App Engine application from us-central to asia-northeast1.
- C. Create a second App Engine application in the existing GCP project and specify asia-northeast1 as the region to serve your application.
- D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

Correct Answer: C*Community vote distribution*

D (100%)

 **Bharathy** Highly Voted 2 years, 7 months ago

Option D is correct, as there can be only one App Engine application inside a project . C is incorrect, as GCP can't have two app engine applications..

upvoted 71 times

 **jcloud965** 1 year, 3 months ago

Yes, and you can't change an App Engine application region once created

upvoted 13 times

 **saurabh1805** Highly Voted 2 years, 5 months ago

Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

<https://cloud.google.com/appengine/docs/flexible/nodejs/managing-projects-apps-billing#create>

upvoted 30 times

 **Ratna_15** Most Recent 1 day, 6 hours ago

I am getting confuse which answer i need to choose correct answer or community vote distribution

upvoted 1 times

 **Ratna_15** 1 day, 6 hours ago

For few questions we have option c or d then which to choose correct answer or community vote distribution

upvoted 1 times

 **leogor** 2 weeks, 2 days ago

D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application

upvoted 1 times

 **Vidhey** 3 weeks, 1 day ago

Selected Answer: D

D is correct as you can't modify the region of the existing app-engine app.

upvoted 1 times

 **666Amitava666** 1 month ago

Selected Answer: D

Here it's clearly mentioned

<https://cloud.google.com/appengine/docs/flexible/managing-projects-apps-billing#:~:text=Each%20Cloud%20project%20can%20contain%20only%20a%20single%20App%20Engine%20application%2C%20and%20once%20created%20you%20cannot%20change%20the%20location%20of%20your%20App%20Engine%20application>.

upvoted 1 times

 **theBestStudent** 1 month, 1 week ago

Selected Answer: D

Correct answer is D. Also Two App engine can't be running on the same project: you can check this easy diagram for more info:

https://cloud.google.com/appengine/docs/standard/an-overview-of-app-engine#components_of_an_application

And you can't change location after setting it for your app Engine. <https://cloud.google.com/appengine/docs/standard/locations>

upvoted 1 times

 **iadarsh** 1 month, 3 weeks ago

Selected Answer: D

D is correct

You cannot change the region or zone of the deployed App Engine Service(Application)

upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: D

Option D is the correct answer

upvoted 1 times

 **Allaman** 2 months, 1 week ago

Selected Answer: D

The App Engine application is a resource that you must create to enable App Engine. The App Engine application resource defines the region in which you want to run your web service and store its related settings, credentials, and metadata. Only a single App Engine application resource can be created per Cloud project and the specified region cannot be changed after those resources are created.

<https://cloud.google.com/appengine/docs/standard/python3/building-app/creating-gcp-project>

upvoted 2 times

 **ajkrizz** 1 month, 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my ajkriswork@gmail.com address. Thanks in advance!

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is right

upvoted 1 times

 **Iodusia** 6 months, 3 weeks ago

Selected Answer: D

Option D is correct, as there can be only one App Engine application inside a project.

upvoted 4 times

 **Madj** 8 months, 1 week ago

Important: Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

Option D is correct

upvoted 5 times

 **sharmaishu** 8 months, 3 weeks ago

Selected Answer: D

<https://cloud.google.com/appengine/docs/standard/python/console>

Important: Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

upvoted 1 times

 **Charlie7569** 8 months, 3 weeks ago

D is the correct answer

upvoted 1 times

 **shanz910** 9 months, 3 weeks ago

D . you cannot create the location of app once created

<https://cloud.google.com/appengine/docs/standard/python/console#:~:text=Important%3A%20Each%20Cloud%20project%20can,of%20your%20App%20Engine%20application.>

upvoted 1 times

Question #38

You need to grant access for three users so that they can view and edit table data on a Cloud Spanner instance. What should you do?

- A. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to the role.
- B. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to a new group. Add the group to the role.
- C. Run gcloud iam roles describe roles/spanner.viewer - --project my-project. Add the users to the role.
- D. Run gcloud iam roles describe roles/spanner.viewer - --project my-project. Add the users to a new group. Add the group to the role.

Correct Answer: A*Community vote distribution*

B (96%)	4%
---------	----

 **yasu** Highly Voted 2 years, 7 months ago

I think it should be B, setup a group first are suggested way from Google.
upvoted 66 times

 **Agents89** Highly Voted 2 years, 6 months ago

B is the correct option
upvoted 24 times

 **PKookNN** Most Recent 3 weeks, 1 day ago

Selected Answer: B

B is the most correct based on this
<https://cloud.google.com/spanner/docs/iam#spanner.databaseUser>
upvoted 1 times

 **GaneshSurwase** 4 weeks ago

Selected Answer: B

Ans is B
upvoted 1 times

 **Robertolo** 1 month ago

Selected Answer: A

Yeah, option B if we involve best practices. But the question does not say anything about any group. Thus option A is the correct one
upvoted 1 times

 **danny19g** 2 months ago

Option A and B are eliminated because it talks about looking for viewer role and we want view + edit. Option B is correct as we should create a new group and add the group to the role.
upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: B

Correct answer is option B - as per IAM best practices
upvoted 2 times

 **RockingRohit6** 3 months, 3 weeks ago

The predefined role roles/spanner.databaseUser contains the permissions spanner.databases.read and spanner.database.write. There are two types of predefined roles for Cloud Spanner. B is correct
upvoted 2 times

 **SaiSaiA** 3 months, 4 weeks ago

It can be A or B.
If the project is not part of an organization, then you cannot create Groups.
Hence, it should be A. <https://cloud.google.com/iam/docs/groups-in-cloud-console#group-permissions>

But if the project is part of the organization, then let us stick to the best practice which is B. Create a new group and users to it and then assign a role to the group.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right.
Using the gcloud tool, execute the gcloud iam roles describe roles/spanner.databaseUser command on Cloud Shell. Attach the users to a newly created Google group and add the group to the role.
upvoted 2 times

 **sedado77** 5 months, 1 week ago

Selected Answer: B

B is correct! A is wrong, as groups are best practice

upvoted 3 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

go for B

upvoted 1 times

 **rfmartinezv** 6 months, 1 week ago

Selected Answer: B

B is correct.

upvoted 1 times

 **gielda211** 7 months, 1 week ago

Selected Answer: B

B is correct.

<https://cloud.google.com/spanner/docs/iam>
upvoted 1 times

 **test1234ff** 11 months ago

Selected Answer: B

b IS CORRECT

upvoted 7 times

 **ACEngineer** 11 months ago

Selected Answer: B

As recommended practice, for 3 users creating a group and assigning role to that is the correct method. Hence B is right option.

upvoted 6 times

Question #39

You create a new Google Kubernetes Engine (GKE) cluster and want to make sure that it always runs a supported and stable version of Kubernetes. What should you do?

- A. Enable the Node Auto-Repair feature for your GKE cluster.
- B. Enable the Node Auto-Upgrades feature for your GKE cluster.
- C. Select the latest available cluster version for your GKE cluster.
- D. Select Container-Optimized OS (cos) as a node image for your GKE cluster.

Correct Answer: B*Community vote distribution*

B (100%)

 **Lush** Highly Voted 2 years, 6 months ago

The answer is B
<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>
 upvoted 35 times

 **4bsolut** Highly Voted 2 years, 4 months ago

"Creating or upgrading a cluster by specifying the version as <latest> does not provide automatic upgrades. Enable automatic node upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version." --source: <https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>

-Correct answer: B
 upvoted 25 times

 **leogor** Most Recent 2 weeks, 2 days ago

Selected Answer: B
 B. Enable the Node Auto-Upgrades
 upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct.
 Creating or upgrading a cluster by specifying the version as latest does not provide automatic upgrades. Enable node auto-upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version.
 upvoted 2 times

 **pfabio** 5 months, 1 week ago

Selected Answer: B
 Node auto-upgrades help you keep the nodes in your cluster up-to-date with the cluster control plane version when your control plane is updated on your behalf. When you create a new cluster or node pool with Google Cloud console or the gcloud command, node auto-upgrade is enabled by default.
 upvoted 2 times

 **Harbeeb** 6 months ago

Selected Answer: B
<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>
 upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.
 upvoted 4 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.
 upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans is B
 upvoted 1 times

 **Rameshfmf** 1 year ago

Ans B- With auto-upgrades, GKE automatically ensures that security updates are applied and kept up to date. Ease of use: Provides a simple way to keep your nodes up to date with the latest Kubernetes features.
 upvoted 2 times

 **viswanand** 1 year, 6 months ago

I was trying to create GKE but I dont see that option in Console
upvoted 1 times

 **mcaromit** 1 year, 6 months ago

B is correct
upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Enable the Node Auto-Upgrades feature for your GKE cluster.
upvoted 1 times

 **yubi69** 1 year, 7 months ago

answer is B . <https://cloud.google.com/kubernetes-engine/versioning> in 'specifying cluster version' section
upvoted 1 times

 **sumanshu** 1 year, 7 months ago

Vote for B
upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

B is correct
upvoted 1 times

 **JackGlemins** 1 year, 8 months ago

B is right: <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-upgrades>
upvoted 1 times

Question #40

You have an instance group that you want to load balance. You want the load balancer to terminate the client SSL session. The instance group is used to serve a public web application over HTTPS. You want to follow Google-recommended practices. What should you do?

- A. Configure an HTTP(S) load balancer.
- B. Configure an internal TCP load balancer.
- C. Configure an external SSL proxy load balancer.
- D. Configure an external TCP proxy load balancer.

Correct Answer: A

Reference:

<https://cloud.google.com/load-balancing/docs/https/>*Community vote distribution*

A (100%)

Gini Highly Voted 2 years, 7 months ago

According to the documentation of SSL Proxy Load Balancing on Google, "SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing." in my opinion A should be the most suitable choice.

upvoted 46 times

yvinisiupacuando 1 year, 6 months ago

Agree with you but, A is not the most suitable choice, it is the only choice, as the other Load Balancers cannot route HTTP(S) traffic.
upvoted 12 times

Agents89 Highly Voted 2 years, 6 months ago

For HTTP(s) Load balancer, the client SSL session terminates at the load balancer. A is the correct option.
upvoted 24 times

leogor Most Recent 2 weeks, 2 days ago

Selected Answer: A
HTTP(S) load balancer.
upvoted 1 times

AzureDP900 4 months, 2 weeks ago

I will go with A
upvoted 1 times

Akash7 5 months ago

A is correct.
According to this guide for setting up an HTTP (S) load balancer in GCP: The client SSL session terminates at the load balancer. Sessions between the load balancer and the instance can either be HTTPS (recommended) or HTTP.
upvoted 1 times

haroldbenites 5 months, 2 weeks ago

Go for C.
It dont say Global balancer.
upvoted 1 times

Imttt 6 months, 1 week ago

Selected Answer: A
Correct A
upvoted 2 times

karim_chaara 11 months, 1 week ago

Correct answer is A - HTTP LB
for SSL proxy, its doesnt support HTTP traffic
"SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing."
<https://cloud.google.com/load-balancing/docs/ssl>
upvoted 5 times

shawnkkk 11 months, 3 weeks ago

A. Configure an HTTP(S) load balancer.
upvoted 1 times

✉ **vishnukumartr** 11 months, 3 weeks ago

A. Configure an HTTP(S) load balancer.
upvoted 1 times

✉ **Jaira1256** 11 months, 3 weeks ago

Ans is A
upvoted 1 times

✉ **Rameshfmf** 1 year ago

A- please refer the this link - <https://cloud.google.com/load-balancing/docs/choosing-load-balancer>
upvoted 2 times

✉ **sidharthwader** 1 year, 4 months ago

No doubt its answer is (A) in Azure its application gateway. These both are layer 7 traffic load balancer
upvoted 2 times

✉ **contaexamtopics** 1 year, 5 months ago

A or c?
upvoted 1 times

✉ **mcaromit** 1 year, 6 months ago

A is the correct option. The lb can have target-http-proxy in the forwarding-rule
upvoted 1 times

✉ **meh009** 1 year, 6 months ago

GCP Best practices
"SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing."
<https://cloud.google.com/load-balancing/docs/ssl>

Answer is A in my opinion

upvoted 3 times

✉ **phony2** 1 year, 7 months ago

The answer should be A
SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

Question #41

You have 32 GB of data in a single file that you need to upload to a Nearline Storage bucket. The WAN connection you are using is rated at 1 Gbps, and you are the only one on the connection. You want to use as much of the rated 1 Gbps as possible to transfer the file rapidly. How should you upload the file?

- A. Use the GCP Console to transfer the file instead of gsutil.
- B. Enable parallel composite uploads using gsutil on the file transfer.
- C. Decrease the TCP window size on the machine initiating the transfer.
- D. Change the storage class of the bucket from Nearline to Multi-Regional.

Correct Answer: B*Community vote distribution*

B (100%)

 **leba** Highly Voted 2 years, 6 months ago

Correct answer is B as the bandwidth is good and its a single file, gsutil parallel composite uploads can be used to split the large file and upload in parallel. Refer GCP documentation - Transferring Data to GCP &

upvoted 37 times

 **berezinsn** Highly Voted 2 years, 5 months ago

Truly B is absolutely correct

upvoted 16 times

 **leogor** Most Recent 2 weeks, 2 days ago

Selected Answer: B

parallel composite uploads

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

<https://cloud.google.com/storage/docs/parallel-composite-uploads>

upvoted 3 times

 **lmttt** 6 months, 1 week ago

Selected Answer: B

b is right

upvoted 2 times

 **Ery** 7 months, 1 week ago

Selected Answer: B

youtube.com/watch?v=NlevtGlo-E0 slice upload elephant file

upvoted 3 times

 **alaahakim** 11 months, 3 weeks ago

Ans : B

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

. Enable parallel composite uploads using gsutil on the file transfer.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Enable parallel composite uploads using gsutil on the file transfer.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

B is correct

upvoted 1 times

 **NizarIrfani** 1 year, 5 months ago

Play online games on <http://159.89.187.75/link-alternatif-m88/>

upvoted 2 times

 **NamiKaze** 1 year, 3 months ago

Yess You Smart

upvoted 3 times

 **djgodzilla** 1 year, 5 months ago

B is correct :

i.e :

```
gsutil -o GSUtil:parallel_composite_upload_threshold=200M,parallel_composite_upload_component_size=100M cp ./localbigfile gs://your-bucket
```

upvoted 5 times

 **mcaromit** 1 year, 6 months ago

B is the obvious choice

upvoted 1 times

 **ApaMokus** 1 year, 6 months ago

Although parallel composite transfer is not optimal for Nearline storage, option B is correct as the question focuses on quicker upload and not costs.

upvoted 2 times

 **kopper2019** 1 year, 7 months ago

it's similar to AWS in order to upload faster using a parallel way so B

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Enable parallel composite uploads using gsutil on the file transfer.

upvoted 2 times

 **yuv169** 1 year, 7 months ago

answer is B.

<https://cloud.google.com/storage/docs/uploads-downloads#parallel-composite-uploads>

upvoted 2 times

Question #42

You've deployed a microservice called myapp1 to a Google Kubernetes Engine cluster using the YAML file specified below:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp1-deployment
spec:
  selector:
    matchLabels:
      app: myapp1
  replicas: 2
  template:
    metadata:
      labels:
        app: myapp1
    spec:
      containers:
        - name: main-container
          image: gcr.io/my-company-repo/myapp1:1.4
          env:
            - name: DB_PASSWORD
              value: "t0ugh2guess!"
        ports:
          - containerPort: 8080
```

You need to refactor this configuration so that the database password is not stored in plain text. You want to follow Google-recommended practices. What should you do?

- A. Store the database password inside the Docker image of the container, not in the YAML file.
- B. Store the database password inside a Secret object. Modify the YAML file to populate the DB_PASSWORD environment variable from the Secret.
- C. Store the database password inside a ConfigMap object. Modify the YAML file to populate the DB_PASSWORD environment variable from the ConfigMap.
- D. Store the database password in a file inside a Kubernetes persistent volume, and use a persistent volume claim to mount the volume to the container.

Correct Answer: C*Community vote distribution*

B (100%)

 **rramani7** Highly Voted 2 years, 5 months ago

it is good practice to use Secrets for confidential data (like API keys) and ConfigMaps for non-confidential data (like port numbers). B is correct.

upvoted 59 times

 **saurabh1805** Highly Voted 2 years, 5 months ago

B is correct answer

<https://cloud.google.com/kubernetes-engine/docs/concepts/secret>

upvoted 34 times

 **hjyhf** 1 year, 3 months ago

"Storing sensitive data in Secrets is more secure than in plaintext ConfigMaps or in Pod specifications"

upvoted 7 times

 **leogor** Most Recent 2 weeks, 2 days ago

Selected Answer: B

Secret object

upvoted 1 times

 **Robertolo** 1 month ago

Selected Answer: B

Passwords should ALWAYS be saved as k8s secrets. There's no doubt

upvoted 1 times

 **raghu09** 1 month, 3 weeks ago

Selected Answer: B

Sensitive content is always gets stored in Secrets.

upvoted 1 times

 **iadarsh** 1 month, 3 weeks ago

Selected Answer: B

B is correct

In kubernetes we always use secrets to keep sensitive and confidential data.

upvoted 1 times

 **Allaman** 2 months, 1 week ago

Selected Answer: B

B for sure. secret object used to store sensitive data such as passwords, key and token

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

This is no brainer question if you already worked on Kubernetes. Secrets is the way to keep passwords safe.. B is right

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for B

upvoted 1 times

 **badrik** 6 months ago

Selected Answer: B

"B" 100%

upvoted 1 times

 **pulse008** 7 months, 2 weeks ago

B is correct

upvoted 2 times

 **Durgesh1997** 9 months, 1 week ago

B is correct answer as secretes are confidential data

upvoted 1 times

 **Rakesh_SAP** 10 months, 1 week ago

Selected Answer: B

<https://cloud.google.com/config-connector/docs/how-to/secrets#gcloud>

upvoted 3 times

 **wHencut** 10 months, 2 weeks ago

Selected Answer: B

B IS CORRECT

upvoted 3 times

 **karim_chaara** 11 months, 1 week ago

Selected Answer: B

ref:<https://cloud.google.com/kubernetes-engine/docs/concepts/secret>

upvoted 4 times

Question #43

Topic 1

You are running an application on multiple virtual machines within a managed instance group and have autoscaling enabled. The autoscaling policy is configured so that additional instances are added to the group if the CPU utilization of instances goes above 80%. VMs are added until the instance group reaches its maximum limit of five VMs or until CPU utilization of instances lowers to 80%. The initial delay for HTTP health checks against the instances is set to 30 seconds.

The virtual machine instances take around three minutes to become available for users. You observe that when the instance group autoscales, it adds more instances than necessary to support the levels of end-user traffic. You want to properly maintain instance group sizes when autoscaling. What should you do?

- A. Set the maximum number of instances to 1.
- B. Decrease the maximum number of instances to 3.
- C. Use a TCP health check instead of an HTTP health check.
- D. Increase the initial delay of the HTTP health check to 200 seconds.

Correct Answer: D*Community vote distribution*

D (100%)

 **berezinsn** Highly Voted 2 years, 5 months ago

D is correct answer.
upvoted 21 times

 **professor** Highly Voted 2 years, 4 months ago

Ans is D

The virtual machine instances take around three minutes to become available for users.

upvoted 18 times

 **Abalagu** 1 year, 5 months ago

D is correct. The reason is that when you do health check, you want the VM to be working. Do the first check after initial setup time of 3 mins = 180 s < 200 s is reasonable.
upvoted 16 times

 **leogor** Most Recent 2 weeks, 2 days ago

Selected Answer: D
Increase the initial delay
upvoted 1 times

 **ale_brd_** 1 month, 3 weeks ago

Selected Answer: D
D is correct answer.
upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: D
Increase initial delay
upvoted 1 times

 **GouseBasha** 4 months, 1 week ago

D is correct one
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is more appropriate for this question
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D
upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for D
upvoted 1 times

 **Ram_krishna** 7 months, 3 weeks ago

Selected Answer: D

logically first 3 options are out of context, so d is right

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Increase the initial delay of the HTTP health check to 200 seconds.

upvoted 2 times

 **Jaira1256** 11 months, 3 weeks ago

Ans is D

upvoted 1 times

 **vikram___** 1 year, 2 months ago

Ans D so that it waits for the instance to be booted n available

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

D is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

D is correct. Increase the initial delay of the HTTP health check to 200 seconds.

upvoted 2 times

 **yuvি69** 1 year, 7 months ago

D is correct

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

D is Correct.

upvoted 1 times

Question #44

You need to select and configure compute resources for a set of batch processing jobs. These jobs take around 2 hours to complete and are run nightly. You want to minimize service costs. What should you do?

- A. Select Google Kubernetes Engine. Use a single-node cluster with a small instance type.
- B. Select Google Kubernetes Engine. Use a three-node cluster with micro instance types.
- C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.
- D. Select Compute Engine. Use VM instance types that support micro bursting.

Correct Answer: C*Community vote distribution*

C (100%)

 **gcper** Highly Voted 2 years ago

As everyone has said the answer is C but here is the source for the information. "For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances."

srouce: <https://cloud.google.com/compute/docs/instances/preemptible>

upvoted 39 times

 **vnxt** Highly Voted 2 years, 6 months ago

I woud say C is the correct answer

upvoted 28 times

 **RAVI321** Most Recent 2 months ago

batch processing jobs can run on preemptible instances. if some of those instances stop during processing, the job slows but does not completely stop. preemptible instances complete your batch processing tasks without placing additional worklods on your existing instances and without requring you to pay full price for additional normal instances"

upvoted 1 times

 **RAVI321** 2 months, 3 weeks ago

hey guys tell me one important thing i am learning GCP but did not get anything i mean whatever you guys are discussing in this forum

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right .

If your apps are fault-tolerant and can withstand possible instance preemptions, then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances.

upvoted 3 times

 **sedado77** 5 months, 1 week ago

Selected Answer: C
Yup, C for batch and cost

upvoted 3 times

 **haroldbenites** 5 months, 1 week ago

Go for C

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for C

upvoted 1 times

 **ruben82** 10 months ago

I read that C is the right answer, but the question doesn't say that batch can be stopped and restarted.

upvoted 3 times

 **VijaySrinivasan** 9 months ago

Batch will not be stopped and load will be shifted to another instances.

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.
upvoted 3 times

 **Jaira1256** 11 months, 3 weeks ago

Ans is C
upvoted 3 times

 **beginnercloud** 1 year, 2 months ago

C - corectly
upvoted 2 times

 **dc2016bte0006** 1 year, 4 months ago

Option C is the correct answer
upvoted 3 times

 **mcaromit** 1 year, 6 months ago

C is correct...For cost-saving & not immediate fault-tolerant workloads like batch jobs use Preemptible VM instances
upvoted 1 times

 **Faizan2991** 1 year, 6 months ago

Might be D , because preemptible VM can be terminate at any time with short notice
upvoted 1 times

 **VijaySrinivasan** 9 months ago

Even if the VM terminates, load will be shifted to another VM since the batch required resources
upvoted 1 times

 **[Removed]** 1 year, 7 months ago

C is correct. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.
upvoted 2 times

 **yuv169** 1 year, 7 months ago

option C is correct
upvoted 1 times

Question #45

You recently deployed a new version of an application to App Engine and then discovered a bug in the release. You need to immediately revert to the prior version of the application. What should you do?

- A. Run gcloud app restore.
- B. On the App Engine page of the GCP Console, select the application that needs to be reverted and click Revert.
- C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.
- D. Deploy the original version as a separate application. Then go to App Engine settings and split traffic between applications so that the original version serves 100% of the requests.

Correct Answer: D

Reference:

<https://medium.com/google-cloud/app-engine-project-cleanup-9647296e796a>

Community vote distribution

C (100%)

 **coldpar** Highly Voted 2 years, 7 months ago

correct is C NOT D.

Option A is wrong as gcloud app restore was used for backup and restore and has been deprecated. Option B is wrong as there is no application revert functionality available. Option D is wrong as App Engine maintains version and need not be redeployed.

upvoted 59 times

 **Bharathy** Highly Voted 2 years, 7 months ago

App engine maintains versions and to revert back to previous version, traffic can be set to 100% for the prior version.. hence correct answer is C

upvoted 25 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: C

route 100% of the traffic to the previous version

upvoted 2 times

 **danny19g** 2 months ago

Correct option is C.

AppEngine already creates a version for you. Also you do not create a application as one project is associated with one AppEngine application.

upvoted 2 times

 **joeMP** 2 months, 4 weeks ago

Selected Answer: C

App engine is a version control tool for your running app

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for C

upvoted 2 times

 **pfabio** 5 months, 1 week ago

Selected Answer: C

D - is wrong because there is only one app engine per project.

C is correct, you have to split traffic

upvoted 5 times

 **legen** 4 months, 2 weeks ago

thanks for your reminder. i constantly forgot that only ONE app engine per project is allowed, and answer D means they'll be 2 app engine in the project

upvoted 1 times

 **haroldbenites** 5 months, 2 weeks ago

Go for C

upvoted 1 times

 **Charlie7569** 8 months, 3 weeks ago

C is the correct answer, D might be correct but that's kinda old and C is better than that

upvoted 1 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: C

App Engine supports versioning for this exact same reason

upvoted 3 times

 **ravan281099** 10 months, 1 week ago

C is the correct NOT D

upvoted 2 times

 **pratheek** 11 months ago

Selected Answer: C

C is legit

upvoted 2 times

 **jaffarali** 11 months ago

Selected Answer: C

Correct Answer is C

upvoted 2 times

 **test1234ff** 11 months ago

Selected Answer: C

CORRECT ANSWER IS C

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Ans: C

upvoted 2 times

Question #46

You deployed an App Engine application using gcloud app deploy, but it did not deploy to the intended project. You want to find out why this happened and where the application deployed. What should you do?

- A. Check the app.yaml file for your application and check project settings.
- B. Check the web-application.xml file for your application and check project settings.
- C. Go to Deployment Manager and review settings for deployment of applications.
- D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.

Correct Answer: A

Reference:

<https://cloud.google.com/endpoints/docs/openapi/troubleshoot-aeflex-deployment>

Community vote distribution

D (82%)

A (18%)

✉  **Bharathy**  2 years, 7 months ago

I would opt option D : as it would help to check the config details and Option A is not correct, as app.yaml would have only the runtime and script to run parameters and not the Project details

upvoted 47 times

✉  **alejandrombc** 1 year, 1 month ago

Why would you choose Cloud Shell if its not even mention on the question? (what if the person did the command on its own computer?, this would not work)

upvoted 1 times

✉  **zaxxon** 1 year ago

gcloud app deploy means sdk

upvoted 9 times

✉  **csrazdan** 4 months ago

Regardless if you use your computer or cloud shell, you have to use SDK for gcloud command-line interface. gcloud uses a configuration file which contains default project, region and zone details so that command line can omit these parameters and use default.

upvoted 1 times

✉  **ahmed812**  2 years, 7 months ago

Option D - The config list will give the name of the project

C:\GCP\appeng>gcloud config list

[core]

account = xxx@gmail.com

disable_usage_reporting = False

project = my-first-demo-xxxx

upvoted 35 times

✉  **leogor**  2 weeks, 1 day ago

Selected Answer: D

D. check project setting by gcloud config list

upvoted 1 times

✉  **danny19g** 2 months ago

Both A and D seem correct.

Lots of people mentioned that app.yaml does not contain project id. That is an incorrect statement. Project is contained in app.YAML for "Standard" app engine application but gcloud config list has to be used for "FLEXIBLE" application. Since the questions does not inform us whether it is a standard or flexible app, Option D is correct. Look at the link here.

<https://cloud.google.com/appengine/docs/flexible/python/reference/app-yaml#app-id>.

The doc states:

"In some App Engine standard environment runtimes, you might have specified the Cloud Platform project ID (sometimes called "app ID") in the project's app.yaml file.

However, in the flexible environment, the project ID (app ID) is specified either:

By using gcloud init when you install the Google Cloud CLI. To view the default project ID of the gcloud CLI, run gcloud config list.

By using the gcloud config set project [YOUR_PROJECT_ID] command to set the default project ID of the gcloud CLI.

By using the --project flag when you deploy your app, for example: gcloud app deploy --project [YOUR_PROJECT_ID]

"

upvoted 4 times

✉️ **RAVI321** 2 months, 3 weeks ago

Why C is incorrect
upvoted 2 times

✉️ **Haz993** 1 month, 2 weeks ago

App Engine Does not use deployment manager
upvoted 1 times

✉️ **anallava** 3 months ago

D
<https://cloud.google.com/endpoints/docs/openapi/troubleshoot-gce-deployment>
upvoted 1 times

✉️ **RanjithK** 4 months, 1 week ago

Answer is D
upvoted 1 times

✉️ **AzureDP900** 4 months, 2 weeks ago

Dis correct
To list the project property in the core section, run:

gcloud config list project
upvoted 1 times

✉️ **haroldbenites** 5 months, 1 week ago

Go for D
upvoted 1 times

✉️ **haroldbenites** 5 months, 2 weeks ago

Go for D
upvoted 1 times

✉️ **BAPII** 8 months, 1 week ago

Selected Answer: D
Correct answer will be D. app.yaml doesn't content the projects information.
upvoted 3 times

✉️ **SleepyHitman** 10 months ago

Selected Answer: A
https://cloud.google.com/endpoints/docs/openapi/troubleshoot-aeflex-deployment#error_messages_about_rollout_strategy
upvoted 2 times

✉️ **svsilence** 6 months ago

app.yaml not include project information. correct anwser is D
upvoted 1 times

✉️ **rafsrod** 10 months, 3 weeks ago

Selected Answer: D
D is correct.
upvoted 1 times

✉️ **hellof20** 10 months, 3 weeks ago

Selected Answer: D
D is correct
upvoted 1 times

✉️ **sid0127** 11 months, 2 weeks ago

Selected Answer: D
D is correct
upvoted 3 times

✉️ **Ridhanya** 11 months, 3 weeks ago

The recommended approach is to remove the application element from your app.yaml file and instead, use a command-line flag to specify your application ID (which is the project id). So A is possible too. although D seems more right
upvoted 1 times

✉️ **vishnukumartr** 11 months, 3 weeks ago

D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.
upvoted 1 times

Question #47

You want to configure 10 Compute Engine instances for availability when maintenance occurs. Your requirements state that these instances should attempt to automatically restart if they crash. Also, the instances should be highly available including during system maintenance. What should you do?

- A. Create an instance template for the instances. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance. Add the instance template to an instance group.
- B. Create an instance template for the instances. Set 'Automatic Restart' to off. Set 'On-host maintenance' to Terminate VM instances. Add the instance template to an instance group.
- C. Create an instance group for the instances. Set the 'Autohealing' health check to healthy (HTTP).
- D. Create an instance group for the instance. Verify that the 'Advanced creation options' setting for 'do not retry machine creation' is set to off.

Correct Answer: B*Community vote distribution*

A (100%)

 **lio123** Highly Voted 1 year, 8 months ago

A

<https://cloud.google.com/compute/docs/instances/setting-instance-scheduling-options>

onHostMaintenance: Determines the behavior when a maintenance event occurs that might cause your instance to reboot.

[Default] MIGRATE, which causes Compute Engine to live migrate an instance when there is a maintenance event.
TERMINATE, which stops an instance instead of migrating it.

automaticRestart: Determines the behavior when an instance crashes or is stopped by the system.

[Default] true, so Compute Engine restarts an instance if the instance crashes or is stopped.
false, so Compute Engine does not restart an instance if the instance crashes or is stopped.

upvoted 31 times

 **Imdeepak12** Highly Voted 1 year ago

Seems like it was a very obvious option i.e. A...Who selected B, I want to know his/her location?

upvoted 14 times

 **Kickbuttonowski_** 11 months, 3 weeks ago

Nikki singh.

upvoted 4 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: A

A. Set the 'Automatic Restart' to on. Set the 'On-host maintenance' to Migrate VM instance

upvoted 1 times

 **Vidhey** 3 weeks, 1 day ago

Selected Answer: A

A is correct

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **Yolex** 4 months, 2 weeks ago

Option C!!

<https://cloud.google.com/compute/docs/tutorials/high-availability-autohealing>

Using autohealing for highly available apps

This interactive tutorial shows how to use autohealing to build highly available apps on Compute Engine.

Highly available apps are designed to serve clients with minimal latency and downtime. Availability is compromised when an app crashes or freezes. Clients of a compromised app can experience high latency or downtime.

Autohealing lets you automatically restart apps that are compromised. It promptly detects failed instances and recreates them automatically, so clients can be served again. With autohealing, you no longer need to manually bring an app back to service after a failure.

upvoted 2 times

 **ace_2021_1** 1 month, 2 weeks ago

"You want to configure 10 Compute Engine instances for availability when maintenance occurs." when maintenance occurs ...is the key word here.....A is the best option

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **Schladde** 5 months, 3 weeks ago

I would rank D before A then C ..

A: Why need to set 'Automatic Restart' and 'On-host maintenance'? It's set by default

B: obviously false

C: Also not wrong but for the question unnecessary

D: Except there is maybe an 's' missing for instances also not wrong. 'do not retry machine creation' 'off' is also default, but action is verify only

upvoted 1 times

 **kmkavitha** 8 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 2 times

 **ArunTaneja** 9 months, 2 weeks ago

Selected Answer: A

A. Create an instance template for the instances. Set the "Automatic Restart" to on. Set the "On-host maintenance" to Migrate VM instance. Add the instance template to an instance group.

upvoted 3 times

 **[Removed]** 10 months, 1 week ago

D

app.yaml doesn't have info about gcp project

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

Ans: A

upvoted 3 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Create an instance template for the instances. Set the "Automatic Restart" to on. Set the "On-host maintenance" to Migrate VM instance. Add the instance template to an instance group.

upvoted 4 times

 **DoctorSly** 11 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 3 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - A

upvoted 3 times

 **TenshiD** 11 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 2 times

 **vijai_nimmala** 1 year, 2 months ago

Many of the answers are wrong. Why it has not been corrected

upvoted 4 times

Question #48

You host a static website on Cloud Storage. Recently, you began to include links to PDF files on this site. Currently, when users click on the links to these PDF files, their browsers prompt them to save the file onto their local system. Instead, you want the clicked PDF files to be displayed within the browser window directly, without prompting the user to save the file locally. What should you do?

- A. Enable Cloud CDN on the website frontend.
- B. Enable 'Share publicly' on the PDF file objects.
- C. Set Content-Type metadata to application/pdf on the PDF file objects.
- D. Add a label to the storage bucket with a key of Content-Type and value of application/pdf.

Correct Answer: C*Community vote distribution*

C (100%)

 **berezinsn** Highly Voted 2 years, 5 months ago

C is correct

upvoted 23 times

 **yas_cloud** Highly Voted 2 years, 2 months ago

C is correct here. tested

upvoted 12 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: C

C. Set Content-Type metadata to application/pdf

upvoted 1 times

 **Untamables** 2 weeks, 4 days ago

Selected Answer: C

FYI

Importance of setting the correct MIME type

https://developer.mozilla.org/en-US/docs/Web/HTTP/Basics_of_HTTP/MIME_Types#importance_of_setting_the_correct_mime_type

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is correct. Edit the PDF objects in Cloud Storage and reconfigure their Content-Type metadata into application/pdf.

upvoted 2 times

 **BBB_ACE** 5 months, 2 weeks ago

Selected Answer: C

Ans: C

upvoted 2 times

 **PLUTONIUM54** 9 months, 1 week ago

C IS CORRECT

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - C

upvoted 1 times

 **sunilw** 1 year, 4 months ago

C is correct.

<https://cloud.google.com/storage/docs/metadata#content-type>

upvoted 9 times

 **vmart** 1 year, 5 months ago

c is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

C is correct. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **yuvি69** 1 year, 7 months ago

correct option is C

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **nliaustemac** 1 year, 8 months ago

C. is the correct answer:

Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

 **victory108** 1 year, 9 months ago

C - Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times

Question #49

You have a virtual machine that is currently configured with 2 vCPUs and 4 GB of memory. It is running out of memory. You want to upgrade the virtual machine to have 8 GB of memory. What should you do?

- A. Rely on live migration to move the workload to a machine with more memory.
- B. Use gcloud to add metadata to the VM. Set the key to required-memory-size and the value to 8 GB.
- C. Stop the VM, change the machine type to n1-standard-8, and start the VM.
- D. Stop the VM, increase the memory to 8 GB, and start the VM.

Correct Answer: D*Community vote distribution*

D (100%)

-  **cesar7816** Highly Voted  2 years, 7 months ago
coldpar, why are you getting the people confused? you need to stop teh VM and modify the RAM, that's all
upvoted 58 times
-  **iambatmanadarkknight** 1 year ago
who is coldpar
upvoted 5 times
-  **spatidar2711** 3 months, 2 weeks ago
He deleted his comment
upvoted 1 times
-  **Carls** Highly Voted  2 years, 7 months ago
D is correct. If you pay attention to the question, option C mentions n1-standard-8. That instance type has 8vCPUs and 30 GB RAM, and we only need 8GB. On top of that, it is possible to use custom machine type to adjust current VM RAM to the value we need. Got the answer from this course I did to prepare the exam: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?couponCode=21CDE6A4C2B95F79BD97>
good luck!
upvoted 38 times
-  **Veera_Venkata_Satyanarayana** 4 months, 2 weeks ago
How to use coupon code carls
upvoted 1 times
-  **leogor** Most Recent  2 weeks, 1 day ago
Selected Answer: D
C is incorrect because the 8 in "n1-standard-8" means 8 cpus instead of 8 gb of ram
upvoted 1 times
-  **santhu1039** 3 months, 1 week ago
Selected Answer: D
D is right one
upvoted 1 times
-  **haroldbenites** 5 months, 1 week ago
Go for D
upvoted 1 times
-  **ChiefArch** 7 months, 1 week ago
The reason A is not correct is because live migration simply moves an existing VM between hosts, no attributes or properties are changed otherwise. Hence, you cannot live migrate from 1 VM type to another.
<https://cloud.google.com/compute/docs/instances/live-migration>
upvoted 3 times
-  **mchaconr** 10 months, 3 weeks ago
To be exact... There is no "n1" option that has 2vCPU and 4GB. The closest would be n1-standard-2 with 2vCPU and 7.50GB RAM. This machine is already custom. So, option D.
upvoted 2 times
-  **vishnukumartr** 11 months, 3 weeks ago
D. Stop the VM, increase the memory to 8 GB, and start the VM.
upvoted 5 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - D

upvoted 1 times

 **Gpjha** 1 year, 3 months ago

We do not have the option to increase memory in GCP VM directly. Instead, machine type needs to be changed. However, given the choices, D is the best answer here.

upvoted 2 times

 **YAS007** 1 year, 3 months ago

D is correct

https://cloud.google.com/compute/docs/instances/creating-instance-with-custom-machine-type#add_extended_memory_to_an_existing_vm_instance

upvoted 2 times

 **Umesh09** 1 year, 5 months ago

why do we want to stop the VM when live migration is an option I would go for Option A

upvoted 3 times

 **YAS007** 1 year, 3 months ago

I don't find how to manually initiate a live migration, its apparently a google managed process

upvoted 1 times

 **Umesh09** 1 year, 5 months ago

Why not A GCP does give a live migration option.

upvoted 3 times

 **mcaromit** 1 year, 6 months ago

D is correct

upvoted 1 times

 **ri_unhou119** 1 year, 6 months ago

Option A [Live migration] Why is it not applicable?

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

D is correct. Stop the VM, increase the memory to 8 GB, and start the VM.

upvoted 2 times

 **yuv169** 1 year, 7 months ago

correct option is D

upvoted 2 times

Question #50

You have production and test workloads that you want to deploy on Compute Engine. Production VMs need to be in a different subnet than the test VMs. All the VMs must be able to reach each other over Internal IP without creating additional routes. You need to set up VPC and the 2 subnets. Which configuration meets these requirements?

- A. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.
- B. Create a single custom VPC with 2 subnets. Create each subnet in the same region and with the same CIDR range.
- C. Create 2 custom VPCs, each with a single subnet. Create each subnet in a different region and with a different CIDR range.
- D. Create 2 custom VPCs, each with a single subnet. Create each subnet in the same region and with the same CIDR range.

Correct Answer: A*Community vote distribution*

A (82%)

B (18%)

 **JamesBond** Highly Voted 2 years, 7 months ago

A is correct

upvoted 33 times

 **nwk** Highly Voted 2 years ago

Vote A

<https://cloud.google.com/vpc/docs/using-vpc#subnet-rules>

Primary and secondary ranges for subnets cannot overlap with any allocated range, any primary or secondary range of another subnet in the same network, or any IP ranges of subnets in peered networks.

upvoted 19 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: A

A, same VPC network with different CIDR range

upvoted 1 times

 **PKookNN** 3 weeks, 1 day ago

A is correct

upvoted 1 times

 **Robertolo** 3 weeks, 5 days ago

Selected Answer: A

It has to be A: one VPC, two subnets in different regions and different CIDR range. (It would also be valid to have both subnets in the same region.)

What about option B? We know that "each primary or secondary IPv4 range for all subnets in a VPC network must be a unique valid CIDR block" (read here <https://cloud.google.com/vpc/docs/subnets#ipv4-ranges>). Thus, prod and test subnets cannot overlap -> option B is not valid.

Options C and D are not valid neither, because "all the VMs must be able to reach each other" - this will not happen if we distribute the VMs across two VPC.

upvoted 1 times

 **manukoli1986** 1 month ago

Leutenant_Ololo, I test and checked. It is B answer

upvoted 1 times

 **Akhi2022** 3 months ago

Option A is incorrect, see what us says "create each subnet in a different region and with a different CIDR range" in this case routing is required due different CIDR range.

Option B is correct, If you create 2 subnets in same CIDR you can communicate over internal IP without additional routing.

upvoted 4 times

 **Leutenant_Ololo** 1 month ago

routers will be created automatically. Just go create a new VPC with 2 subnets and then check routes.

<https://cloud.google.com/vpc/docs/routes#subnet-routes>

upvoted 2 times

 **73173v2** 3 months, 2 weeks ago

Selected Answer: A

Different regions is something odd, but the main reason why its A is cause the CIDR range.
CIDR is the short for Classless Inter-Domain Routing. So, if we have 2 subnets, they CAN NOT BE the same CIDR.

IPv4 subnet ranges

"Each primary or secondary IPv4 range for all subnets in a VPC network must be a unique valid CIDR block. Refer to the per network limits for the number of secondary IP ranges you can define."

<https://cloud.google.com/vpc/docs/vpc>

upvoted 4 times

✉ **ryumada** 3 months ago

You should understand that the VPC is a Global Resource. You can create a VPC and having subnets across regions. These subnets could communicate to each other through their Private IP.

<https://cloud.google.com/sql/docs/mysql/private-ip#overview>

upvoted 3 times

✉ **csrazdan** 3 months, 3 weeks ago

A is correct.

B - is incorrect because 2 subnets cannot have the same CIDR Range.

C and D - are incorrect because creating 2 VPC's will require parsing and exchange routes for communication.

upvoted 3 times

✉ **AzureDP900** 4 months, 2 weeks ago

A is right, Tutorials Dojo have same question!

upvoted 2 times

✉ **Tirthankar17** 5 months ago

Selected Answer: A

A is correct. Why on earth would it be B.

upvoted 1 times

✉ **Schladde** 5 months, 3 weeks ago

Selected Answer: A

Two subnets with same CIDR are not possible

upvoted 2 times

✉ **PAUGURU** 6 months, 1 week ago

Selected Answer: A

B is wrong, try it yourself, create a new VPC and create 2 subnets in the same region. It won't let you use the same parameters(IP class, netmask) for the 2 subnets. The error was "Operation type [insert] failed with message "Invalid IPCidrRange: 10.10.10.0/24 conflicts with existing subnetwork 'net1-2' in region 'us-central1'." It is obvious anyway, using the same CIDR on two subnets would require bridging capabilities that are not implemented on GCP. Networks are routed at layer 3 not bridged at layer 2 on GCP. Answer is A.

upvoted 6 times

✉ **mrmccoy007** 6 months ago

Agreed. Just tried this and got the same error. A is correct.

upvoted 1 times

✉ **PAUGURU** 6 months, 1 week ago

VPC of course not VPN

upvoted 1 times

✉ **VivekSolutionArchitect** 6 months, 2 weeks ago

I would say both A and B are correct. One VPC can span across the regions and we can have subnets segregated between 2 regions. It is best practice (suggested by all Cloud providers) to keep workloads separate, hence A could be the answer.

upvoted 1 times

✉ **DanielB96** 6 months, 2 weeks ago

Selected Answer: B

Voting for B.

No additional routing is key.

upvoted 1 times

✉ **Leutenant_Ololo** 1 month ago

just try to do it ;-)

upvoted 1 times

✉ **Bableves** 7 months ago

Selected Answer: B

It is fascinating how you convince each other for A, it is B for sure, the link provided for subnet rules have nothing to do with the question, there is nothing about primary and secondary subnets, peering is outer domain routing connecting different VPCs and why should you deploy in different region? Makes no sense at all. Polatik is right, and, if you have different subnets you NEED TO CONFIGURE routing! Moreover:
https://cloud.google.com/vpc/docs/vpc#intra_vpc_reqs

upvoted 3 times

 **badrik** 6 months ago

You'll have to understand the question properly. There is a requirement to let the environments communicate internally using the Private IP. Given this consideration it is impossible to have this requirement met when two subnets have same CIDR ranges and the answer narrows down to A.

upvoted 1 times

 **Bableves** 7 months ago

CIDR - classless inter domain routing, how can you have two domains and communicate between if you don't have additional routing set up?

upvoted 1 times

 **Paulv82003** 3 months, 2 weeks ago

You obviously have no idea about routing and subnetting. How in the world will you assign same CIDR range to 2 subnets in same VPC? How prey tell will you route between 2 subnets with same CIDR range, you can't even peer 2 VPCs with same range.

upvoted 2 times

 **VidyaJi** 11 months, 1 week ago

Selected Answer: A

A is perfect

upvoted 3 times

Question #51

You need to create an autoscaling managed instance group for an HTTPS web application. You want to make sure that unhealthy VMs are recreated. What should you do?

- A. Create a health check on port 443 and use that when creating the Managed Instance Group.
- B. Select Multi-Zone instead of Single-Zone when creating the Managed Instance Group.
- C. In the Instance Template, add the label 'health-check'.
- D. In the Instance Template, add a startup script that sends a heartbeat to the metadata server.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>*Community vote distribution*

A (80%) C (20%)

✉ **cesar7816** Highly Voted 2 years, 7 months ago

I'll go with A, MIGs support autohealing, load balancing, autoscaling, and auto-updating. no the Images templates, this is set up in the MIG
upvoted 52 times

✉ **tanito83** Highly Voted 1 year, 5 months ago

The correct answer is A. Please, modify it.
upvoted 11 times

✉ **dennydream** Most Recent 1 week, 5 days ago

This shows the importance of reading exactly what it says. The label can throw you off. A is the answer.
upvoted 1 times

✉ **leogor** 2 weeks, 1 day ago

Selected Answer: A

A obviously, port 443 for HTTPS application
upvoted 1 times

✉ **Untamables** 2 weeks, 4 days ago

A is correct.
https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs#setting_up_an_autohealing_policy
upvoted 1 times

✉ **ravip12345** 1 month, 4 weeks ago

Answer is A. Just adding label doesn't do any thing
upvoted 1 times

✉ **RanjithK** 4 months, 1 week ago

Answer is A
upvoted 2 times

✉ **Yolex** 4 months, 2 weeks ago

Selected Answer: A

Adding label won't effect to health chech!!
upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

A is right answer
upvoted 1 times

✉ **Tirthankar17** 5 months ago

Selected Answer: A

A is right
upvoted 1 times

✉ **haroldbenites** 5 months, 1 week ago

Go for A
upvoted 1 times

✉ **feg2421** 8 months, 2 weeks ago

Selected Answer: C

If the health check's protocol is SSL, HTTPS, or HTTP2, it uses --port=443.
upvoted 5 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: A

A

C is absurd. what will adding a label do?
upvoted 5 times

 **Vidyaji** 11 months, 1 week ago

A perfect

upvoted 3 times

 **itznotme** 11 months, 1 week ago

Selected Answer: A

Create a health check on port 443 and use that when creating the Managed Instance Group.
upvoted 5 times

 **PradeepPen** 11 months, 2 weeks ago

Selected Answer: A

Correct answer is A.

upvoted 3 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Create a health check on port 443 and use that when creating the Managed Instance Group.

upvoted 1 times

Question #52

Your company has a Google Cloud Platform project that uses BigQuery for data warehousing. Your data science team changes frequently and has few members.

You need to allow members of this team to perform queries. You want to follow Google-recommended practices. What should you do?

- A. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery jobUser role to the group.
- B. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery dataViewer user role to the group.
- C. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery jobUser role to the group.
- D. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery dataViewer user role to the group.

Correct Answer: D

Reference:

<https://cloud.google.com/bigquery/docs/cloud-sql-federated-queries>

Community vote distribution

C (67%)

D (33%)

 **Gini**  2 years, 6 months ago

C is correct because dataViewer does not allow user to perform queries. jobUser can.

upvoted 83 times

 **rbrto** 2 years, 6 months ago

dataviewer can perform queries as well. D is correct

upvoted 6 times

 **nightflyer** 1 year, 10 months ago

incorrect

upvoted 2 times

 **Dips_** 1 year, 10 months ago

no it does not.

BigQuery Job User Provides permissions to run jobs, including queries, within the project.

upvoted 1 times

 **kabu_shawarib** 2 years, 5 months ago

I just checked IAM roles, it doesn't have job.create like jobUser.

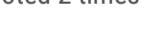
upvoted 5 times

 **RazOr** 9 months, 2 weeks ago

How about this?

<https://cloud.google.com/bigquery/docs/access-control#:~:text=Provides%20permissions%20to%20run%20jobs%2C%20including%20queries%2C%20within%20the%20project>.

upvoted 2 times

 **zukko78**  2 years, 6 months ago

C is correct, doc's said: When applied to a dataset, dataViewer provides permissions to:

Read the dataset's metadata and to list tables in the dataset.

Read data and metadata from the dataset's tables.

When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 43 times

 **spudleymcdudley** 2 years, 4 months ago

Ref 'D' data viewer = When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 4 times

 **spudleymcdudley** 2 years, 4 months ago

listen to this guy, he's right

upvoted 2 times

 **dartharlinton** 4 months ago

according to the principle of least privilege that Google instills on its IAM offerings, answer c would be correct!

upvoted 1 times

□  **shubhamsaurabh** Most Recent 2 weeks, 4 days ago

C is correct
BigQuery Data Viewer
(roles/bigquery.dataViewer)

When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.
This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.
Read data and metadata from the dataset's tables.
When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

Lowest-level resources where you can grant this role:

Table
View

BigQuery Job User
(roles/bigquery.jobUser)

Provides permissions to run jobs, including queries, within the project.

Lowest-level resources where you can grant this role:

Project
upvoted 1 times

□  **Souvik_79** 2 weeks, 6 days ago

Selected Answer: C
Data viewer role doesn't allow to query tables/views
upvoted 1 times

□  **Souvik_79** 2 weeks, 6 days ago

Selected Answer: C
Job viewer role doesn't allow to query data
upvoted 1 times

□  **anolive** 3 weeks ago

Selected Answer: D
i think is d
dataViewer, dataEditor, and dataOwner roles
are similar to the primitive roles except
can be assigned for individual datasets
don't give users permission to run jobs or queries
upvoted 1 times

□  **anolive** 3 weeks ago

reading again, i think is c;
<https://jayendrapatil.com/google-cloud-bigquery-security/>
upvoted 1 times

□  **PKookNN** 3 weeks, 1 day ago

Selected Answer: C
Anwer is C : to run jobs
<https://cloud.google.com/bigquery/docs/access-control#bigquery.jobUser>
databaseUser needs additional role permission to run jobs
<https://cloud.google.com/spanner/docs/iam#spanner.databaseUser>
upvoted 1 times

□  **GaneshSurwase** 3 weeks, 4 days ago

C is correct ans , with the bigqueryjobuser role user can run query
upvoted 1 times

□  **PKookNN** 3 weeks, 6 days ago

Selected Answer: C
I would go for C since data viewer can't run a job but jobUser can
<https://cloud.google.com/bigquery/docs/access-control#bigquery.filteredDataViewer>
<https://cloud.google.com/bigquery/docs/access-control#bigquery.jobUser>
upvoted 1 times

□  **Leutenant_Ololo** 1 month ago

I have just tested it, ladies and gentlemen :-) I have created a fresh new user in my project with just one role assigned.

Role bigQuery.dataViewer allows me to open BQ and access the dataset and table and even to open Query editio, but as soon as I hit the "Run" button, GCP Console gives the "Access Denied: Project ***-sandbox: User does not have bigquery.jobs.create permission in project ***- sandbox."

Role bigQuery.jobUser allows me to run query with no issues.

However, I can't see the neither dataset name, nor table name (I just need to obtain the exact dataset/table names before I run the query).

The answer is C :-) "Group with the bigQuery.jobUser"

upvoted 4 times

 **gcpBeginner** 1 month, 1 week ago

C is correct. user should be added to a group and need to perform work.

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

 **ace_2021_1** 1 month, 2 weeks ago

roles/bigquery.dataViewer role has a permission of "bigquery.tables.getData"....as per google doc bigquery.tables.getData, permission is required for querying table data.

C should be the best choice.

upvoted 1 times

 **ace_2021_1** 1 month, 2 weeks ago

I mean D

upvoted 1 times

 **gcpreviewer** 1 month, 2 weeks ago

Selected Answer: C

Many others have explained it well. A and B can be discarded per IAM best practice to create groups. D is not correct because a viewer is not able to query the data. C is able to run queries.

upvoted 1 times

 **ale_brd_** 1 month, 3 weeks ago

Selected Answer: C

Correct one is C.

BigQuery Job User
(roles/bigquery.jobUser)

Provides permissions to run jobs, including queries, within the project.

Lowest-level resources where you can grant this role:

Project

bigquery.config.get

bigquery.jobs.create

resourcemanager.projects.get

resourcemanager.projects.list

One thing is being able to see the table (Headers, Columns names, Schema etc) as with BigQuery Data Viewer, one thing is being able to create a Job with it, as with roles/bigquery.jobUser, Jobs are also queries.

=====
upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: C

Option C is the correct answer

upvoted 1 times

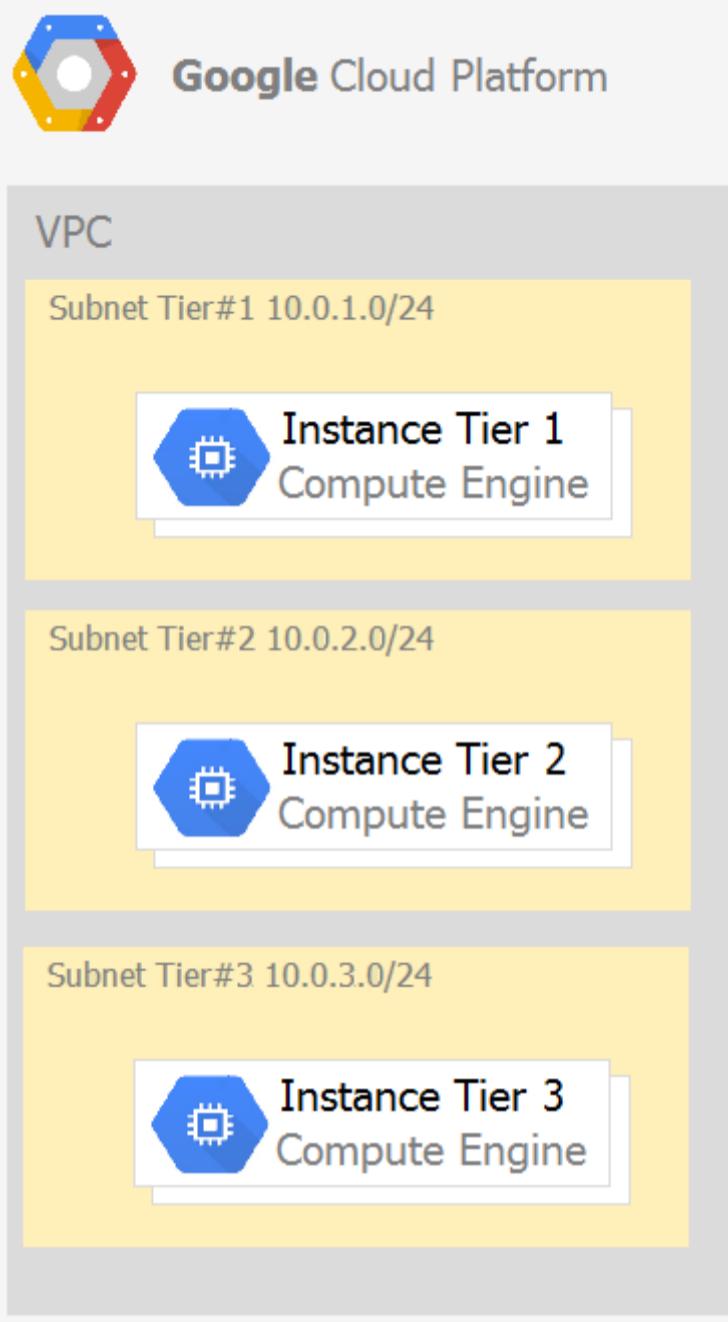
 **canbadroit** 2 months, 1 week ago

C is correct . I created a role with just dataviewer role and tried querying the table it throws permission error. I later added jobuser role and the query worked.

upvoted 1 times

Question #53

Your company has a 3-tier solution running on Compute Engine. The configuration of the current infrastructure is shown below.



Each tier has a service account that is associated with all instances within it. You need to enable communication on TCP port 8080 between tiers as follows:

- * Instances in tier #1 must communicate with tier #2.
- * Instances in tier #2 must communicate with tier #3.

What should you do?

- A. 1. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow all
- B. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow TCP: 8080
- C. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow all
- D. 1. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow TCP: 8080 2. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow TCP: 8080

Correct Answer: B

Community vote distribution

B (100%)

kopper2019 Highly Voted 1 year, 7 months ago

if you see closely, port 8080 and service account is required so B is the answer without reading all answers
upvoted 22 times

 **GCPjunkie** 9 months, 2 weeks ago

Love they way you think, drill down to the important details.
upvoted 2 times

 **obeythefist**  7 months, 4 weeks ago

This question is designed to waste your time during the exam by making you read all those long answers. Remember that part of exam technique is not about knowing the product at all, but understanding multiple choice questions.

For example when two answers are very similar to each other, this can increase the likelihood that the correct answer is one of those two.

In this case it's an easy process of elimination as all answers are similar, we just need to filter out the wrong ones (and whacking the wrong answer in an exam is sometimes the best way to find the right one).

Two answers mention port 8080, and two mention all ports. Obviously we just need port 8080, so we can immediately eliminate those two questions that want all ports open. That gives us a 50/50 chance of getting this question right.

Of the remaining answers, one says "ingress" and the other "egress". We know that by default egress is permitted and ingress is not, so that makes "b" the only surviving choice.

upvoted 21 times

 **leogor**  2 weeks, 1 day ago

Selected Answer: B
B is correct obviously
upvoted 1 times

 **abirroy** 3 months, 1 week ago

Selected Answer: B
B is the correct answer
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right, We need to open firewall rules to allow port 8080 and It shouldn't be wide open... like /24 network.
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

go for B
upvoted 1 times

 **Jerickson** 9 months, 1 week ago

Selected Answer: B
B is correct
upvoted 2 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: B
B is correct
upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

B. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow TCP: 8080
upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - B
upvoted 1 times

 **sunilw** 1 year, 4 months ago

B is correct
upvoted 10 times

 **vmart** 1 year, 5 months ago

B is correct
upvoted 2 times

 **shankymore01** 1 year, 5 months ago

B is correct
upvoted 4 times

 **mcaromit** 1 year, 6 months ago

B is correct
upvoted 3 times

 **[Removed]** 1 year, 7 months ago

B is correct. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow TCP: 8080

upvoted 3 times

 **yuvি69** 1 year, 7 months ago

Correct answer is B

upvoted 2 times

 **sumanshu** 1 year, 7 months ago

Vote for 'B'

upvoted 1 times

 **sumanshu** 1 year, 7 months ago

Option D is egress - So eliminated

upvoted 1 times

Question #54

You are given a project with a single Virtual Private Cloud (VPC) and a single subnetwork in the us-central1 region. There is a Compute Engine instance hosting an application in this subnetwork. You need to deploy a new instance in the same project in the europe-west1 region. This new instance needs access to the application. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- B. 1. Create a VPC and a subnetwork in europe-west1. 2. Expose the application with an internal load balancer. 3. Create the new instance in the new subnetwork and use the load balancer's address as the endpoint.
- C. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Use Cloud VPN to connect the two subnetworks. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- D. 1. Create a VPC and a subnetwork in europe-west1. 2. Peer the 2 VPCs. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

Correct Answer: A*Community vote distribution*

A (100%)

 **Agents89** Highly Voted 2 years, 6 months ago

A is correct

upvoted 35 times

 **LaxmanTiwari** Highly Voted 5 months, 1 week ago

Selected Answer: A

JUST PASS THE EXAM THIS MORING , THIS ONE IS THERE AND I CHOOSE A

upvoted 18 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: A

A is correct

upvoted 1 times

 **gcpBeginner** 1 month, 1 week ago

if A is correct can someone explain what "use the first instance's private address as the endpoint" means? Does it mean to use IP from previous subnet? or does it mean use first IP from new subnet?

upvoted 2 times

 **Leutenant_Ololo** 1 month ago

A is correct. VPC allows you to spawn multiple subnets in different zones. Routing is handled automatically (because Routers are created automatically).

"use the first instance's private address as the endpoint" means that this new instance will be accessing the app via first instance's private IP (so there should be some routing rules created). Question says: "This new instance needs access to the application." ..

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 2 times

 **Bableves** 7 months ago

A cannot be good, I mean guys you are not good at networking, if you have two different subnets, you cannot use an IP from the other subnet, just randomly, you have to "give access" which means you have to connect to the two subnets somehow, it would be better with routing, but VPN does the job...

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

Cloud VPN securely connects your peer network to your Virtual Private Cloud (VPC) network through an IPsec VPN connection. Traffic traveling between the two networks is encrypted by one VPN gateway and then decrypted by the other VPN gateway. This action protects your data as it travels over the internet. You can also connect two instances of Cloud VPN to each other.

upvoted 1 times

 **Leutenant_Ololo** 1 month ago

I mean guys you are not good at networking <-- but we are glad you are proficient with it :-D

Routing between subnets in GCP is not the same as in "regular" networking.

<https://cloud.google.com/vpc/docs/routes#subnet-routes> : "When you add a subnet, Google Cloud creates a corresponding subnet route for the subnet's primary IP address range." RTFM..

upvoted 1 times

 **JelloMan** 6 months, 2 weeks ago

In GCP, VPC's are global - and subnets across different regions can be accessed using private IP's (no VPN setup required).

upvoted 16 times

 **AzureDP900** 4 months, 2 weeks ago

There is no need of setting VPN as you mentioned. AWS is different ..

upvoted 3 times

 **Vidyaji** 11 months, 1 week ago

Selected Answer: A

A is perfect

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - A

upvoted 2 times

 **Deeska** 1 year ago

Subnets are global. A is correct

upvoted 3 times

 **akshaychavan7** 5 months, 3 weeks ago

subnets are regional!

upvoted 1 times

 **piyu1515** 5 months, 2 weeks ago

VPC ARE GLOBA,SUBNETS ARE REGIONAL

upvoted 2 times

 **vmart** 1 year, 5 months ago

A is the correct option

upvoted 3 times

 **AD_0525** 1 year, 5 months ago

Could someone please explain if C can be an option as well.

upvoted 2 times

 **sanhoo** 1 year, 5 months ago

cloud VPN is not to connect two subnets in same VPN. It is usually to connect on-premises computer to cloud. we just need a separate subnet because of a separate region . The subnet can stay in same VPC that we they can talk to each other with private IP. Hence A is correct.

upvoted 5 times

 **Finger41** 1 year, 5 months ago

Cloud VPN is used when connecting on-premise infrastructure to the cloud. In this case creating a new subnet in a new region, and using the internal IP to communicate is sufficient enough to fulfill the requirements.

upvoted 8 times

 **arsh1916** 1 year, 5 months ago

A is correct

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

A is the correct option

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

A is correct. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct

upvoted 1 times

Question #55

Your projects incurred more costs than you expected last month. Your research reveals that a development GKE container emitted a huge number of logs, which resulted in higher costs. You want to disable the logs quickly using the minimum number of steps. What should you do?

- A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.
- B. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE Cluster Operations resource.
- C. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Logging.
- D. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Monitoring.

Correct Answer: A*Community vote distribution*

A (86%)	14%
---------	-----

 **Gini** Highly Voted 2 years, 6 months ago

The question mentioned that "GKE container emitted a huge number of logs", in my opinion A is correct.

upvoted 45 times

 **JackGlemins** Highly Voted 1 year, 9 months ago

I think A is right.

<https://cloud.google.com/logging/docs/api/v2/resource-list>

GKE Containers have more log than GKE Cluster Operations:

.-GKE Container:

cluster_name: An immutable name for the cluster the container is running in.
 namespace_id: Immutable ID of the cluster namespace the container is running in.
 instance_id: Immutable ID of the GCE instance the container is running in.
 pod_id: Immutable ID of the pod the container is running in.
 container_name: Immutable name of the container.
 zone: The GCE zone in which the instance is running.

VS

.-GKE Cluster Operations

project_id: The identifier of the GCP project associated with this resource, such as "my-project".

cluster_name: The name of the GKE Cluster.

location: The location in which the GKE Cluster is running.

upvoted 13 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: A

A can do it with fewest steps

upvoted 1 times

 **Untamables** 2 weeks, 4 days ago

Selected Answer: B

Currently B is correct. (Stackdriver Logging has been named Cloud Logging.)

<https://cloud.google.com/stackdriver/docs/solutions/gke/installing#migrating>

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: A

Go for A

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C,D are straight forward you can eliminate and decide between A & B .. A makes more sense .. logical thinking required to solve it quickly.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A right

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 2 times

 **haroldbenites** 5 months, 1 week ago

Beginning with GKE version 1.15.7, you can configure Cloud Operations for GKE to only capture system logs and not collect application logs.
<https://cloud.google.com/stackdriver/docs/solutions/gke/managing-logs>

upvoted 2 times

 **JelloMan** 6 months ago

Selected Answer: A

A is correct. Target GKE containers since that is what is emitting the most logs

upvoted 1 times

 **P00JA3808** 7 months, 3 weeks ago

Selected Answer: A

Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.

upvoted 1 times

 **Durgesh1997** 9 months, 1 week ago

Selected Answer: A

A is right answer

upvoted 1 times

 **ankit2021** 9 months, 2 weeks ago

Selected Answer: A

A is ans

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.

upvoted 2 times

 **Jaira1256** 11 months, 3 weeks ago

A is correct

upvoted 2 times

 **vmart** 1 year, 5 months ago

A is right

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

A is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

A is correct. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.

upvoted 3 times

 **yuvি69** 1 year, 7 months ago

Correct answer is A

upvoted 1 times

Question #56

You have a website hosted on App Engine standard environment. You want 1% of your users to see a new test version of the website. You want to minimize complexity. What should you do?

- A. Deploy the new version in the same application and use the --migrate option.
- B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.
- C. Create a new App Engine application in the same project. Deploy the new version in that application. Use the App Engine library to proxy 1% of the requests to the new version.
- D. Create a new App Engine application in the same project. Deploy the new version in that application. Configure your network load balancer to send 1% of the traffic to that new application.

Correct Answer: B*Community vote distribution*

B (100%)

 **yasu** Highly Voted 2 years, 7 months ago

I will prefer B as the answer.. why we need create new application?

upvoted 47 times

 **YAS007** 1 year, 3 months ago

more over, in app engine we cannot create "new application", we have to create a new Project to do that, an app engine projet has 1 application (which can have multiple versions and services)

upvoted 13 times

 **sanhoo** 1 year, 5 months ago

Agree B is correct. creating a new application in the same project for app engine is anyways not possible.

upvoted 8 times

 **Gini** Highly Voted 2 years, 7 months ago

I agree with yasu. And only one app engine can exist in one project. B is the best choice, simple and easy.

upvoted 17 times

 **leogor** Most Recent 2 weeks, 1 day ago

Selected Answer: B

B. deploy new version with --splits option

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **habros** 3 months ago

B! A very natural answer... Perfect for switching users over to new version. Imagine creating multiple projects to update App Engine deployments, isn't that logically unnecessary?

upvoted 1 times

 **Madj** 4 months, 2 weeks ago

Hint:

One app engine per project. So Option C,D eliminated. this hint will help in many similar questions.

Splitting traffic hint will help as well

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

B is right.

upvoted 1 times

 **mplibunao** 5 months ago

Selected Answer: B

b is my answer.

a: --migrate is for enabling gradual traffic migration as opposed to migrating traffic immediately

c & d: no need to create a project. You can split the traffic any time

upvoted 3 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **Durgesh1997** 9 months, 1 week ago

B is the correct answer as there is only one app engine can be created per project

upvoted 2 times

 **Priyankahere** 9 months, 3 weeks ago

This question was there , go with community answers.

upvoted 2 times

 **[Removed]** 10 months, 1 week ago

B,

--splits exists for such requirements

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.

upvoted 1 times

 **Jaira1256** 11 months, 3 weeks ago

Ans - B

upvoted 2 times

 **Rahul183** 1 year, 2 months ago

<https://cloud.google.com/appengine/docs/standard/python/splitting-traffic#gcloud>

Option-(B)

upvoted 3 times

 **vmart** 1 year, 5 months ago

B is correct

upvoted 2 times

 **arsh1916** 1 year, 5 months ago

B is correct

upvoted 1 times

Question #57

You have a web application deployed as a managed instance group. You have a new version of the application to gradually deploy. Your web application is currently receiving live web traffic. You want to ensure that the available capacity does not decrease during the deployment. What should you do?

- A. Perform a rolling-action start-update with maxSurge set to 0 and maxUnavailable set to 1.
- B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.
- C. Create a new managed instance group with an updated instance template. Add the group to the backend service for the load balancer. When all instances in the new managed instance group are healthy, delete the old managed instance group.
- D. Create a new instance template with the new application version. Update the existing managed instance group with the new instance template. Delete the instances in the managed instance group to allow the managed instance group to recreate the instance using the new instance template.

Correct Answer: C*Community vote distribution*

B (100%)

CarlS Highly Voted 2 years, 7 months ago

Correct option is B. We need to ensure the global capacity remains intact, for that reason we need to establish maxUnavailable to 0. On the other hand, we need to ensure new instances can be created. We do that by establishing the maxSurge to 1. Option C is more expensive and more difficult to set up and option D won't meet requirements since it won't keep global capacity intact.

upvoted 81 times

JavierCorrea Highly Voted 2 years, 2 months ago

I take my own previous comment back. It's definitely B.

upvoted 18 times

leogor Most Recent 1 week, 2 days ago

Selected Answer: B

B. maxUnavailable set to 0 is the key

upvoted 1 times

gcpBeginner 1 month, 1 week ago

it should be B. if we change template it cause outage and question said no outage allowed.

upvoted 1 times

ashit44244 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

king2658 1 month, 2 weeks ago

try searching on google, may be you can find it.

upvoted 1 times

RanjithK 4 months, 1 week ago

Answer is B

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

Correct option is B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0. This is also present on Tutorial Dojo practice questions.

upvoted 1 times

LaxmanTiwari 5 months, 1 week ago

WAS IN MY EXAM GO WITY THE MAJORITY

upvoted 3 times

haroldbenites 5 months, 1 week ago

Go for B

upvoted 1 times

Durgesh1997 9 months, 1 week ago

Selected Answer: B

max_surge to 1 and max_unavailable to 0
upvoted 3 times

 **Priyankahere** 9 months, 3 weeks ago

This was there in exam, I went with B.
upvoted 4 times

 **shputhan** 9 months, 3 weeks ago

Correct Answer is B
https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups#max_unavailable
upvoted 2 times

 **Raz0r** 9 months, 4 weeks ago

Selected Answer: B

B is right, that's exactly what the option is for
upvoted 1 times

 **rafsrod** 10 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 2 times

 **hellof20** 10 months, 3 weeks ago

Selected Answer: B

Option B is the right answer
upvoted 2 times

 **jaffarali** 11 months ago

Selected Answer: B

Option B is the right answer
upvoted 1 times

 **Vidyaji** 11 months, 1 week ago

Selected Answer: B

B is perfect
upvoted 1 times

Question #58

You are building an application that stores relational data from users. Users across the globe will use this application. Your CTO is concerned about the scaling requirements because the size of the user base is unknown. You need to implement a database solution that can scale with your user growth with minimum configuration changes. Which storage solution should you use?

- A. Cloud SQL
- B. Cloud Spanner
- C. Cloud Firestore
- D. Cloud Datastore

Correct Answer: B*Community vote distribution*

B (100%)

✉️  **Fidget_** Highly Voted 2 years, 3 months ago

B

Cloud SQL for small relational data, scaled manually
Cloud Spanner for relational data, scaled automatically
Cloud Firestore for app-based data(?)
Cloud Datastore for non-relational data
Correct me if i'm wrong

upvoted 78 times

✉️  **theBestStudent** 5 months ago

Just one detail: Cloud Firestore for non relational data (noSql)

upvoted 3 times

✉️  **karol_wu** Highly Voted 2 years, 7 months ago

in my opinion correct is B

upvoted 29 times

✉️  **leogor** Most Recent 1 week, 2 days ago

Selected Answer: B

B. Spanner for autoscale

upvoted 1 times

✉️  **learn_GCP** 1 month, 3 weeks ago

Selected Answer: B

B.

keywords: "global" and "relational"

upvoted 1 times

✉️  **gcp_world123** 2 months, 2 weeks ago

Correct answer is B

Cloud Firestore and Cloud Datastore - can easily be eliminated doesn't fall under Relational DB

upvoted 1 times

✉️  **abirroy** 3 months, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

✉️  **theBestStudent** 4 months, 1 week ago

B as many have mentioned already

upvoted 1 times

✉️  **AzureDP900** 4 months, 2 weeks ago

Global is keyword in this question, Cloud Spanner is the right option.. B

upvoted 5 times

✉️  **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

✉️  **TheScalator** 9 months ago

Selected Answer: B

B is correct

upvoted 2 times

  **k8sdr** 10 months, 2 weeks ago**Selected Answer: B**

B is correct Cloud spanner can scale across regions

upvoted 2 times

  **look1** 11 months, 1 week ago**Selected Answer: B**

B only

upvoted 1 times

  **vishnukumartr** 11 months, 3 weeks ago

B. Cloud Spanner

upvoted 1 times

  **Jaira1256** 11 months, 3 weeks ago

Ans - B

upvoted 1 times

  **player496** 1 year ago

its B. if there is a word 'global' or 'world', in the usecase then blindly go for spanner

upvoted 3 times

  **Litan** 1 year, 1 month ago

B Correct

upvoted 1 times

  **Bala13** 1 year, 5 months ago

for globalization and relational data...Cloud spanner is correct Option B

upvoted 1 times

Question #59

You are the organization and billing administrator for your company. The engineering team has the Project Creator role on the organization. You do not want the engineering team to be able to link projects to the billing account. Only the finance team should be able to link a project to a billing account, but they should not be able to make any other changes to projects. What should you do?

- A. Assign the finance team only the Billing Account User role on the billing account.
- B. Assign the engineering team only the Billing Account User role on the billing account.
- C. Assign the finance team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.
- D. Assign the engineering team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.

Correct Answer: D*Community vote distribution*

C (52%)	A (45%)
---------	---------

 **Bharathy**  2 years, 7 months ago

Option A is correct, as we don't want the engineering team to link projects to billing account and want only the Finance team. Billing Account User role will help to link projects to the billing account...

upvoted 65 times

 **mwoodm** 2 years, 2 months ago

Option A makes the most sense since Billing Account User can link projects to the billing account and the question reinforces principle of least privilege. Source: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 9 times

 **djgodzilla** 1 year, 5 months ago

wrong : you can't link project to billing accounts without Project billing manager.

C is Correct

upvoted 9 times

 **djgodzilla** 1 year, 5 months ago

my bad , it seems A,C are both technically correct

upvoted 1 times

 **Hasaaaan** 1 year, 5 months ago

Billing Account User also enables the user to make changes in resources.

upvoted 1 times

 **BobbyFlash** 1 year ago

I would also go with A. I would think they are trying to get a quick answer from you as "Billing Administrator": engineering team already has the project creator role; you just would want finance team to link (and only) link projects to billing accounts, nothing else. Maybe the key phrase here is "but they should not be able to make any other changes to projects" and that would include the action of unlinking projects.

upvoted 6 times

 **measme**  2 years, 5 months ago

for me is C:

https://cloud.google.com/billing/docs/how-to/modify-project#permissions_required_for_this_task_2

"Roles with adequate permissions to perform this task:

* Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account."

upvoted 47 times

 **obeythefist** 8 months, 2 weeks ago

The question states that the finance group should not be able to make changes to existing projects. Granting the finance team organizational level Billing Account Administrator will allow them to make changes to other projects. C cannot be correct.

upvoted 2 times

 **Robertolo** 4 weeks ago

Project Billing Manager does not allow to make any changes to projects. It's just about linking+unlinking projects to billing accounts

On the other hand, the single role "billing account user" does not grant any right to view projects. Even less likely to link them to any billing account. (see <https://cloud.google.com/iam/docs/job-functions/billing> "The Billing Account User role gives the service account the permissions to enable billing (associate projects with the organization's billing account for all projects in the organization) and thereby permit the service account to enable APIs that require billing to be enabled."). Thus A is not the correct answer.

The right answer is C, without any kind of doubt

upvoted 1 times

 **Frogfly** 6 days, 20 hours ago

Are you blind ? you posted link where its clearly stated in billing account user description: (associate projects with the organization's billing account for all projects in the organization)

So you literally posted link with clarification that answer A is correct.

answer C will give finance team additional permission to unlink billing account from projects and question clearly states that finance team should not be able to make any other changes to projects so C without any kind of doubt is wrong.

upvoted 2 times

 **Frogfly** Most Recent 6 days, 20 hours ago

A is the only correct answer

1. Only the finance team should be able to link a project to a billing account

2. but they should not be able to make any other changes to projects

A is correct because

1. Billing Account User definition from google

(roles/billing.user) Link projects to billing accounts.

https://cloud.google.com/billing/docs/how-to/billing-access#overview_of_billing_roles_in

2. Project billing manager got more permissions and question clearly states they should not be able to make any other changes

upvoted 1 times

 **hiromi** 1 week ago

Selected Answer: C

C is the ans

upvoted 1 times

 **klyaal** 1 week, 1 day ago

Selected Answer: C

Correct C

Billing Account User. This role has very restricted permissions, so you can grant it broadly. When granted in combination with Project Creator, the two roles allow a user to create new projects linked to the billing account on which the Billing Account User role is granted. Or, when granted in combination with the Project Billing Manager role, the two roles allow a user to link and unlink projects on the billing account on which the Billing Account User role is granted.

--- in combination with the Project Billing Manager role

Project Billing Manager. When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access.

<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

 **leogor** 1 week, 2 days ago

C is more correct in my opinion

upvoted 1 times

 **nonsense** 1 week, 5 days ago

Selected Answer: A

It's A - 100%

upvoted 2 times

 **Untamables** 2 weeks, 4 days ago

Selected Answer: C

The answer is obviously described in Google document.

<https://cloud.google.com/billing/docs/how-to/modify-project#required-permissions-change>

upvoted 1 times

 **Erie** 2 weeks, 5 days ago

Selected Answer: A

The finance team only needs to "Link" projects to billing accounts. They don't need to "Unlink", which if needed, then Project Billing Manager role would be required.

upvoted 1 times

 **PKookNN** 3 weeks, 1 day ago

Selected Answer: C

C is more correct. A is partial answer (so not correct) Per this statement " When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access." from

<https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

upvoted 1 times

 **GaneshSurwase** 3 weeks, 4 days ago

Option C is correct

upvoted 1 times

 **GaneshSurwase** 3 weeks, 4 days ago

with the billing manager role u can link/unlink the billing account to the project.

upvoted 1 times

 **PKookNN** 3 weeks, 5 days ago

Selected Answer: C

From <https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam> "When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access." So C is correct.

upvoted 1 times

 **PKookNN** 3 weeks, 5 days ago

Selected Answer: D

After reading this <https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam> - which said "When granted in combination with the Billing Account User role, the Project Billing Manager role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access." I came around to D as the correct answer but the wording on the answer is weird.

upvoted 1 times

 **Babillion** 4 weeks, 1 day ago

Selected Answer: A

Depending this phrase "they should not be able to make any other changes to projects" in question we can conclude that "Billing account user role " have the certain permission to link projects to billing account. whereas on the other hand, project billing manager have access to both link and unlink projects which is not recommended by google practices

Refer this link:

<https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

upvoted 1 times

 **sjagkoo7** 4 weeks, 1 day ago

this is the dumb answer given by examtopics. as question says engineering team do not need to link the project so how we can assign engineering team to billing account user role.very disappointg.

upvoted 1 times

 **Robertolo** 1 month ago

Selected Answer: C

tricky - anyway, C is the right answer

upvoted 1 times

 **Robertolo** 1 month ago

Selected Answer: D

tricky - anyway, D is the answer

upvoted 1 times

Question #60

You have an application running in Google Kubernetes Engine (GKE) with cluster autoscaling enabled. The application exposes a TCP endpoint. There are several replicas of this application. You have a Compute Engine instance in the same region, but in another Virtual Private Cloud (VPC), called gce-network, that has no overlapping IP ranges with the first VPC. This instance needs to connect to the application on GKE. You want to minimize effort. What should you do?

- A. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Set the service's externalTrafficPolicy to Cluster. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- B. 1. In GKE, create a Service of type NodePort that uses the application's Pods as backend. 2. Create a Compute Engine instance called proxy with 2 network interfaces, one in each VPC. 3. Use iptables on this instance to forward traffic from gce-network to the GKE nodes. 4. Configure the Compute Engine instance to use the address of proxy in gce-network as endpoint.
- C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- D. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add a Cloud Armor Security Policy to the load balancer that whitelists the internal IPs of the MIG's instances. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.

Correct Answer: A*Community vote distribution*

C (57%) A (43%)

 **someoneinthecloud** Highly Voted  2 years, 3 months ago

I believe it's A. It's never mentioned in the question that traffic cannot go through the Internet but it's mentioned that effort should be minimized. A requires a lot less effort than C to accomplish the same (no VPC peering, per example).

upvoted 46 times

 **ShakthiGCP** 1 year, 8 months ago

Ans: A . This sounds correct and avoids unnecessary steps in C. C is also correct but compared to it, A is much easier to achieve. Go over Kubernetes Loadbalancer concepts to get more details. Initially i was thinking C is the Answer. but after putting some time on K8's Network - changed my mind to A.

upvoted 9 times

 **AmitKM** 2 years, 2 months ago

Yeah, I feel the same. Nowhere does it say that the traffic has to be internal. But it does say "minimal effort" which I feel is option A.

upvoted 8 times

 **pgb54** 8 months, 2 weeks ago

Totally agree. I had the same thought and looked through the question for any indication that the traffic must be private.

upvoted 2 times

 **juancambb** Highly Voted  2 years, 5 months ago

i think C is better solution, the solution A pass traficrought public internet, also C by internal network and the "no overlap ips" in the statement suggest that.

upvoted 38 times

 **toni90ns** Most Recent  1 week, 6 days ago

Selected Answer: A

I have never done A in past, so I was not sure will it work. Just test it and it works. Question here is not saying what is the best option, it says to minimize the effort. So definitely it is easier to do steps from answer A.

upvoted 1 times

 **Robertolo** 4 weeks, 1 day ago

Selected Answer: C

C is the answer: <https://cloud.google.com/load-balancing/docs/choosing-load-balancer#external-internal>

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **ashit44244** 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 2 times

 **GaneshSurwase** 3 weeks, 4 days ago

search examtopics gcp ace question __ fill no. in blank space

upvoted 2 times

 **VietmanOfficiel** 2 months, 2 weeks ago

Selected Answer: C

[C]

"no overlapping IP's" so VPC peering will work. However one will need to configure firewall on both VPC's to allow internal traffic.

upvoted 1 times

 **joeMP** 2 months, 4 weeks ago

Selected Answer: A

A is th goof answer

upvoted 1 times

 **Prosecute** 3 months ago

Selected Answer: A

C. Keyword is "minimize effort".

If keyword was "securely" or "best practice" it would've been A.

upvoted 2 times

 **ryumada** 3 months ago

Vote for C, externalTrafficPolicy is used to support Internal Load Balancer. Also, You should see the comments from dark_3k03r and kimharsh. They also explain about externalTrafficPolicy.

<https://cloud.google.com/kubernetes-engine/docs/how-to/service-parameters#externalTrafficPolicy>

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: C

I will go with C

upvoted 1 times

 **jeffangel28** 3 months, 2 weeks ago

Selected Answer: A

someoneinthecloud is right

upvoted 1 times

 **amenur** 3 months, 4 weeks ago

In my opinion it is C.

I think that in answer A, what we are doing is deploying an internet facing LoadBalancer (because it is the default LoadBalancer type that is deployed if we do not specify an internal one in the annotations), that is, with the possibility of being called from outside our VPC, and forcing the compute engine instance to go outside the google network to be able to go back inside through the endpoint that we are publishing with the unicast ip of the balancer.

Answer C, however, performs a peering between the two VPC's (the statement makes sure that this option is feasible since it clearly specifies that there is no overlapping between the ip ranges of both vpc's), deploy the LoadBalancer as internal with the annotation, and configure the endpoint so that the compute engine instance can access the application internally, that is, without the need to have a public ip at any time and therefore, without the need to go outside the google network. The traffic, therefore, never crosses the public internet.

upvoted 2 times

 **dark_3k03r** 4 months, 2 weeks ago

Its not A. External traffic Policy, the name is a bit misleading. What it is referring to is the load balancing of the pods. Do you perform no load balancing of the pods and keep it internal to the node (local) or do you perform pod load balancing outside/externally from the local node (i.e. across the cluster). Has nothing to do with the routing over the internet or how traffic gets routed to it. It has all to do with how to route traffic once it gets there.

It's not B cause nodeport is a one for one mapping, so this wouldn't scale , not to mention the 2 network interfaces with proxy in two vpc statement...

It's not D cause there is no need for cloud armor security policy.

So that leaves C as the only viable source.

<https://medium.com/pablo-perez/k8s-externaltrafficpolicy-local-or-cluster-40b259a19404>

upvoted 4 times

 **kohsiangyu** 4 months, 3 weeks ago

Selected Answer: C

I will vote for C. Answer A did not state if there is a public IP for compute engine and the firewall policy is not set for GKE to allow connection from the internet.

upvoted 4 times

✉ **wolfie09** 4 months, 3 weeks ago

TCP is not encrypted so A is not a good idea

upvoted 3 times

✉ **Tirthankar17** 5 months ago

Except B (NodePort), all are valid answers. However, in A the traffic is public, so not secure. In D, it is mentioned MIG, which is not the case as per the Question. Lastly it is C that is the correct option remaining.

upvoted 1 times

Question #61

Your organization is a financial company that needs to store audit log files for 3 years. Your organization has hundreds of Google Cloud projects. You need to implement a cost-effective approach for log file retention. What should you do?

- A. Create an export to the sink that saves logs from Cloud Audit to BigQuery.
- B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.
- C. Write a custom script that uses logging API to copy the logs from Stackdriver logs to BigQuery.
- D. Export these logs to Cloud Pub/Sub and write a Cloud Dataflow pipeline to store logs to Cloud SQL.

Correct Answer: A

Reference:

<https://cloud.google.com/logging/docs/audit/>

Community vote distribution

B (92%)	8%
---------	----

✉  **yasu** Highly Voted 2 years, 7 months ago

Why not B? cost effective

upvoted 49 times

✉  **lxs** 1 year ago

BigQuery data after 90 days has the same cost for storage as Cloud Storage Nearline. Storing it in Cloud Storage adds more costs for data retrieval if the class is i.e archival

upvoted 1 times

✉  **uganeshku** 10 months, 2 weeks ago

B is correct because Coldline Storage is the perfect service to store audit logs from all the projects and is very cost-efficient as well. Coldline Storage is a very low-cost, highly durable storage service for storing infrequently accessed data.

upvoted 3 times

✉  **Gini** Highly Voted 2 years, 7 months ago

if it is all about cost, B is the best. However, speaking of "audit" you probably need to access the data once in a while, which Coldline storage might not be ideal for this case I guess? I would go for A in the exam though.

upvoted 19 times

✉  **pas77** 1 year, 3 months ago

The question is clearly saying cost effect. BQ is one of the most expensive services in GCP.

upvoted 5 times

✉  **Ale1973** 2 years, 2 months ago

Be strong!!! If B is the best, go for B!!!

upvoted 11 times

✉  **boof** 1 year, 1 month ago

I would play it safe and interpret the question literally, implying that they will only store the audit logs and not be accessing them a lot.

upvoted 2 times

✉  **leogor** Most Recent 1 week, 2 days ago

Selected Answer: B

Coldline Storage bucket

upvoted 1 times

✉  **dennydream** 1 week, 5 days ago

Selected Answer: B

Coldline would be the least expensive.

upvoted 1 times

✉  **sandipk91** 2 months, 1 week ago

Selected Answer: B

Option B because it talks about cost effective solution, I know BQ has the same cost as Coldline in GCS if data is kept for 90 days but in Cloud Storage we can save more by further moving the class to Archival which is cheaper than Coldline. SO DEFINATELY IT'S OPTION B

upvoted 1 times

✉  **joeMP** 2 months, 4 weeks ago

Selected Answer: A

My bad. A is good for bigquery can supports store many data

upvoted 1 times

 **joeMP** 2 months, 4 weeks ago

Selected Answer: B

Hundreds of projects means many logs. Bigquery is the good Storage.

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: B

Answer is B

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Since there is archive cold line is best in this scenario. Go with B as right answer.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **akshaychavan7** 5 months, 3 weeks ago

As per the Google recommended practice, audit logs are (by default) stored inside cloud storage. If cloud storage is a google recommended solution then it must be cost effective.

upvoted 1 times

 **rafsrod** 10 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 3 times

 **maahibhai** 10 months, 3 weeks ago

Selected Answer: B

b is the right option

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: B

B would be the right choice. The requirement doesn't say about analysis. Hence Coldline bucket storage would be cost effective.

upvoted 1 times

 **arapin** 11 months ago

Selected Answer: B

Logs are not structured data and relational databases are not a good fit. That excludes SQL.

Besides that, the most cost effective solution of the 4 is under B, the Coldline storage bucket.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.

upvoted 2 times

Question #62

Topic 1

You want to run a single caching HTTP reverse proxy on GCP for a latency-sensitive website. This specific reverse proxy consumes almost no CPU. You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes. You want to minimize cost. How should you run this reverse proxy?

- A. Create a Cloud Memorystore for Redis instance with 32-GB capacity.
- B. Run it on Compute Engine, and choose a custom instance type with 6 vCPUs and 32 GB of memory.
- C. Package it in a container image, and run it on Kubernetes Engine, using n1-standard-32 instances as nodes.
- D. Run it on Compute Engine, choose the instance type n1-standard-1, and add an SSD persistent disk of 32 GB.

Correct Answer: B*Community vote distribution*

A (75%)

B (25%)

 **jzh** Highly Voted  2 years, 2 months ago

Go to cloud console and create instance
select Memorystore with Basic tier, select us-central1 and us-central1-a, and capacity 32GB, the cost estimate is \$0.023/GB/hr

select VM instance with custom machine type with 6 vCPUs and 32 GB memory, the same region and zone as Memorystore setting, the cost estimate is \$0.239/hr

Option B will definitely cost more as it adds on CPU usage cost even it uses little in this scenario, but still charge you. So answer is A from real practice example.

upvoted 45 times

 **FenixRa73** 1 year, 10 months ago

$\$0.023 * 32 = \0.736

is it cheaper?

upvoted 4 times

 **SSPC** 2 years, 2 months ago

I agree with you

upvoted 1 times

 **mexblood1** 2 years, 2 months ago

Using pricing calculator matching 730 hrs per month for both.. Memorystore is 537.28 per month and vm (6 cpus 32 gb memory) is 174.41 per month. So vm is still cheaper even with 6 cpus.

upvoted 6 times

 **[Removed]** 2 years, 1 month ago

I agree its cheaper, but 2 drawbacks, 1 hit latency, 2 you need to install cache plain won't help even if check. So still think A.

upvoted 3 times

 **[Removed]** 2 years, 1 month ago

Typo correct my statements. There are 2 issues If you create a VM how every cheap there will be a hit with latency when communicating with it. Second, you will need to install/implement a caching system on that.

Whereas cloud memorystore for redis is designed for the sole purpose of HTTP caching which has very low latency compared to any other solution we are thinking of doing.

upvoted 3 times

 **Rothmansua** 1 year ago

and what about HTTP, how are you supporting that with Redis?

upvoted 2 times

 **obeythefist** 8 months, 1 week ago

A quick Bing search shows a number of solutions for caching HTTP services with Redis.

upvoted 2 times

 **smarty_arce** 4 months, 2 weeks ago

Who uses Bing at this present day and age?

upvoted 7 times

 **ESP_SAP** Highly Voted  2 years, 2 months ago

Correct Answer should be A:

The question mention "You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes"

What is Google Cloud Memorystore?

Overview. Cloud Memorystore for Redis is a fully managed Redis service for Google Cloud Platform. Applications running on Google Cloud Platform can achieve extreme performance by leveraging the highly scalable, highly available, and secure Redis service without the burden of managing complex Redis deployments.

upvoted 34 times

✉️ **ESP_SAP** 2 years, 2 months ago

Just to complement the answer:

We are looking for "latency-sensitive website"

What it's good for

Memorystore for Redis provides a fast, in-memory store for use cases that require fast, real-time processing of data. From simple caching use cases to real time analytics, Memorystore for Redis provides the performance you need.

Caching: Cache is an integral part of modern application architectures. Memorystore for Redis provides low latency access and high throughput for heavily accessed data, compared to accessing the data from a disk based backend store. Session management, frequently accessed queries, scripts, and pages are common examples of caching.

https://cloud.google.com/memorystore/docs/redis/redis-overview#what_its_good_for

upvoted 15 times

✉️ **[Removed]** 2 years, 2 months ago

I agree with your reasoning. Given that the question stresses that this is for a 'latency sensitive website', that's a clue that Redis is part of the answer. Even if spinning up a similarly sized VM were more cost effective, I can't find any documentation that this would provide sufficiently low latency as a memory cache. Yes, you want to keep costs low, but not if it causes your latency-sensitive website problems. Thus I agree that option A is the answer.

upvoted 2 times

✉️ **magistrum** 1 year, 10 months ago

Agree, don't think anything you create with the Compute Engine will meet the "You want to have a 30-GB in-memory cache" requirement...that's a very different technology

upvoted 2 times

✉️ **dennydream** [Most Recent] 1 week, 5 days ago

How do you figure the correct answer here? The votes overwhelmingly say one thing, but the correct answer is another.

upvoted 1 times

✉️ **PKookNN** 3 weeks, 1 day ago

Selected Answer: B

I changed my mind - cost effective is B (while A is easiest)

upvoted 1 times

✉️ **ashit44244** 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

✉️ **RanjithK** 4 months, 1 week ago

Selected Answer: A

Answer is A

upvoted 1 times

✉️ **AzureDP900** 4 months, 2 weeks ago

A is right ..Memorystore for Redis provides a fast.

upvoted 1 times

✉️ **kohsiangyu** 4 months, 3 weeks ago

Selected Answer: B

MemoryStore Pricing is \$0.023/GB/hr, for 32GB means \$0.736/hr compares to \$0.239/hr. The question states that it needs additional 2GB for its process, which mean if you choose A, you will need another vm with 2GB ram either.

upvoted 1 times

✉️ **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

✉️ **browneyes1985** 7 months, 1 week ago

Selected Answer: A

Memorystore offers various sizes to fit any budget. Pricing varies with settings—including how much capacity, how many replicas and which region you provision. Memorystore also offers per-second billing and instances and is easy to start and stop.

upvoted 1 times

✉️ **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

Memory store is for caching content
upvoted 2 times

✉ **RazOr** 9 months, 4 weeks ago

Selected Answer: A

A is right:
Is Redis a reverse proxy server?
nginx-redis-proxy is a reverse proxy based on nginx and redis to cache objects (web pages and more).
upvoted 2 times

✉ **[Removed]** 10 months, 1 week ago

I would go with A
upvoted 1 times

✉ **ahmadjw** 11 months, 3 weeks ago

Hello everybody, I purchased the exam from the certification-questions website ... 45% of questions are not the same as discussed in the comment. Also the certification-question help says that their dump are pretty sure. Would plz any one help where to get the right answer, from here or any other dump. And should I take the first comment as the right answer?

upvoted 4 times

✉ **vishnukumartr** 11 months, 3 weeks ago

A. Create a Cloud Memorystore for Redis instance with 32-GB capacity.
upvoted 2 times

✉ **AshisKumar** 12 months ago

A is correct
upvoted 1 times

✉ **gerhardbl** 1 year, 2 months ago

My answer is B. You guys are all forgetting that you cannot run a proxy just on Redis. The question even says: "need an additional 2 GB of memory for the rest of the processes". You cannot use Memorystore as memory for such other processes; you'll need a VM to run the proxy. 6 vCPUs says nothing about the compute power, it could be 6 vCPUs of a very modest CPU type. Cost does not play a role here, you simply cannot get this done with answer A.

upvoted 5 times

✉ **peter77** 1 year, 1 month ago

Redis behaves like a reverse proxy. And the whole point here is to have content cached so you keep the latency low. Having a caching server in a VM is an aberration and a terrible architectural choice.

It's A 100%.

upvoted 3 times

Question #63

You are hosting an application on bare-metal servers in your own data center. The application needs access to Cloud Storage. However, security policies prevent the servers hosting the application from having public IP addresses or access to the internet. You want to follow Google-recommended practices to provide the application with access to Cloud Storage. What should you do?

- A. 1. Use nslookup to get the IP address for storage.googleapis.com. 2. Negotiate with the security team to be able to give a public IP address to the servers. 3. Only allow egress traffic from those servers to the IP addresses for storage.googleapis.com.
- B. 1. Using Cloud VPN, create a VPN tunnel to a Virtual Private Cloud (VPC) in Google Cloud. 2. In this VPC, create a Compute Engine instance and install the Squid proxy server on this instance. 3. Configure your servers to use that instance as a proxy to access Cloud Storage.
- C. 1. Use Migrate for Compute Engine (formerly known as Velostrata) to migrate those servers to Compute Engine. 2. Create an internal load balancer (ILB) that uses storage.googleapis.com as backend. 3. Configure your new instances to use this ILB as proxy.
- D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in Google Cloud. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve *.googleapis.com as a CNAME to restricted.googleapis.com.

Correct Answer: C*Community vote distribution*

D (100%)

 **poogcp** Highly Voted 2 years, 5 months ago

D is the correct one as per Ref: <https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>
upvoted 49 times

 **sarva** Highly Voted 2 years, 5 months ago

D VPN/interconnect without exposing to public
upvoted 14 times

 **gcpBeginner** Most Recent 1 month, 1 week ago

correct answer is D. why will cx migrate it env. to GCP. easiest and faster approach is to have Cloud VPN setup and advertise route o cloud router
upvoted 1 times

 **gcpBeginner** 1 month, 1 week ago

in my above comment I meant advertise routes to local router used on-prem
upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today
upvoted 1 times

 **RAVI321** 2 months, 2 weeks ago

how you all know that everybody

mere to upar se nikal raha hai sab kuch pls help guys
upvoted 5 times

 **Souvik_79** 2 weeks, 4 days ago

Me too bro
upvoted 1 times

 **aakash4Chaudhary** 1 month, 1 week ago

tu akela nhi hai bro
upvoted 2 times

 **abirroy** 3 months ago

Selected Answer: D

Answer is D
upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: D

Answer is D
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is most appropriate, This is part of Tutorial Dojo practice questions.
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D
upvoted 1 times

 **obeythefist** 8 months, 1 week ago

What messy answers! I chose D and here is my reasoning per answer.

- A. It's bad practice to use nslookup to try find a permanent IP address because IPs can change. That's what DNS is for! Also, the security team aren't going to budge... this is just a silly answer.
- B. We're getting warmer. Any time a question mentions on-prem and cloud, Google wants you to think about Cloud VPN. This solution might even work, but installing Squid? This is a messy solution to a more simple problem.
- C. Talk about using a sledge hammer to swat a mosquito. I think this could work, but migrating servers to cloud to solve a simple networking problem?
- D. Once more Google's favorite Cloud VPN is in the answer. I'm not sure about the networking component of this question.

upvoted 10 times

 **obeythefist** 8 months, 1 week ago

Edit: Of course the reason D: is correct is because 199.36.153.4/30 is the network segment that you can direct traffic to if you want to use Google services "internally". So your on prem servers will resolve storage.googleapis.com to something in this 199.36.153.4/30 range. Then they will route using Cloud Router and your VPN tunnel into Google Cloud privately.

upvoted 4 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: D
D is the correct answer
upvoted 2 times

 **Durgesh1997** 9 months, 1 week ago

Selected Answer: D
D is practically good to implement so D is the answer
upvoted 2 times

 **[Removed]** 10 months, 1 week ago

D makes most sense
upvoted 3 times

 **hellof20** 10 months, 3 weeks ago

Selected Answer: D
D is Correct
upvoted 3 times

 **vishnukumartr** 11 months, 3 weeks ago

D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in Google Cloud. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve *.googleapis.com as a CNAME to restricted.googleapis.com.

upvoted 4 times

 **jitu12113** 1 year, 4 months ago

Can someone pls explain why D?
upvoted 4 times

 **arsh1916** 1 year, 5 months ago

D is Correct
upvoted 1 times

Question #64

You want to deploy an application on Cloud Run that processes messages from a Cloud Pub/Sub topic. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a Cloud Function that uses a Cloud Pub/Sub trigger on that topic. 2. Call your application on Cloud Run from the Cloud Function for every message.
- B. 1. Grant the Pub/Sub Subscriber role to the service account used by Cloud Run. 2. Create a Cloud Pub/Sub subscription for that topic. 3. Make your application pull messages from that subscription.
- C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.
- D. 1. Deploy your application on Cloud Run on GKE with the connectivity set to Internal. 2. Create a Cloud Pub/Sub subscription for that topic. 3. In the same Google Kubernetes Engine cluster as your application, deploy a container that takes the messages and sends them to your application.

Correct Answer: D*Community vote distribution*

C (100%)

 **Meix** Highly Voted 2 years, 5 months ago

C looks right for me as per <https://cloud.google.com/run/docs/tutorials/pubsub#integrating-pubsub>
upvoted 45 times

 **ChrisBelt5** 1 year, 3 months ago

great doc, its' C
upvoted 6 times

 **Bhagirathi** Highly Voted 1 year, 10 months ago

why c ?
explained>>
You can use Pub/Sub to push messages to the endpoint of your Cloud Run service, where the messages are subsequently delivered to containers as HTTP requests. You cannot use Pub/Sub pull subscriptions because Cloud Run only allocates CPU during the processing of a request.
upvoted 29 times

 **leogor** Most Recent 1 week, 2 days ago

C obviously
upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: C
Answer is C
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right .. Look for Meix explanation...
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for C
upvoted 1 times

 **gielda211** 7 months ago

Selected Answer: C
C. is the correct answer:
Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint
upvoted 3 times

 **gielda211** 7 months ago

C. is the correct answer:
Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint
upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: C

C is the correct answer since it follows Google recommended practices

upvoted 1 times

 **HansKloss611** 9 months, 2 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

 **emv** 9 months, 4 weeks ago

it is C

upvoted 1 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: C

Cloud run invoker role for this mentioned in the official doc as well

upvoted 1 times

 **sasonian** 10 months, 1 week ago

C it is

upvoted 1 times

 **RealEL40** 11 months ago

Answer - C:

<https://cloud.google.com/run/docs/tutorials/pubsub#integrating-pubsub>

upvoted 2 times

 **Vidyaji** 11 months, 1 week ago

Selected Answer: C

C is perfect

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.

upvoted 2 times

 **Rahul183** 1 year, 2 months ago

You can use Pub/Sub to push messages to the endpoint of your Cloud Run service, where the messages are subsequently delivered to containers as HTTP requests. You cannot use Pub/Sub pull subscriptions because Cloud Run only allocates CPU during the processing of a request.

upvoted 1 times

Question #65

You need to deploy an application, which is packaged in a container image, in a new project. The application exposes an HTTP endpoint and receives very few requests per day. You want to minimize costs. What should you do?

- A. Deploy the container on Cloud Run.
- B. Deploy the container on Cloud Run on GKE.
- C. Deploy the container on App Engine Flexible.
- D. Deploy the container on GKE with cluster autoscaling and horizontal pod autoscaling enabled.

Correct Answer: B*Community vote distribution*

A (100%)

 **Gurnoor** Highly Voted 2 years, 4 months ago

A should be cheapest as no infra needed.

upvoted 62 times

 **spudleymcdudley** 2 years, 4 months ago

Listen to this guy. Google says "Cloud Run abstracts away all infrastructure management by automatically scaling up and down from zero almost instantaneously—depending on traffic. Cloud Run only charges you for the exact resources you use."

upvoted 25 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer should be A:

Cloud Run takes any container images and pairs great with the container ecosystem: Cloud Build, Artifact Registry, Docker. ... No infrastructure to manage: once deployed, Cloud Run manages your services so you can sleep well. Fast autoscaling. Cloud Run automatically scales up or down from zero to N depending on traffic.

<https://cloud.google.com/run>

upvoted 24 times

 **Cornholio_LMC** Most Recent 1 month, 2 weeks ago

had this question today

upvoted 2 times

 **xharf** 3 months, 1 week ago

Don't forget that cloud run invoking rules

"Cloud Run redirects all HTTP requests to HTTPS but terminates TLS before they reach your web service."
in last case when i use cloud run with POST request and HTTP method, it will generate bug which redirect the request to GET https header :)

<https://cloud.google.com/run/docs/triggering/https-request#>

upvoted 1 times

 **orious** 3 months, 3 weeks ago

Selected Answer: A

What does It mean, "Cloud Run on GKE"?

B is not correct at all, there is no such thing – "Cloud Run on GKE". Or you deploy on "Cloud Run" Or on "GKE".

upvoted 4 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: A

Answer is A. CloudRun is the cheapest option of all

upvoted 3 times

 **AzureDP900** 4 months, 2 weeks ago

This is no brainier question.. If you know the GCP services well straight forward answer is Cloud Run to reduce the cost.. Option A is right!

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: A

A is correct as In B you will be paying for underlying infra

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **pfabio** 5 months, 1 week ago

Selected Answer: A
Cloud Run takes any container images and pairs great with the container ecosystem: Cloud Build, Artifact Registry, Docker.
upvoted 2 times

 **gitaexams** 5 months, 4 weeks ago

Selected Answer: A
A, because Cloud Run has free tier option and the question said: "...receives very few request per day.."
upvoted 4 times

 **rfmartinezv** 6 months ago

Selected Answer: A
A, because Cloud Run has free tier option and the question said: "...receives very few request per day.."
<https://cloud.google.com/hosting-options>
upvoted 2 times

 **subhas_cme** 6 months, 1 week ago

why not app engine Flex , as It will have free quota
upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A
CloudRun is the best option for this scenario
upvoted 1 times

 **Durgesh1997** 9 months, 1 week ago

Selected Answer: A
Cloud RUn as no infra needed to manage and autoscaling taken care by GCP
upvoted 2 times

 **sasonian** 10 months, 1 week ago

it is

https://www.splunk.com/en_us/blog/devops/gcp-serverless-comparison.html
upvoted 1 times

 **azure4life** 10 months, 3 weeks ago

Selected Answer: A
Cloud Run
upvoted 1 times

Question #66

Topic 1

Your company has an existing GCP organization with hundreds of projects and a billing account. Your company recently acquired another company that also has hundreds of projects and its own billing account. You would like to consolidate all GCP costs of both GCP organizations onto a single invoice. You would like to consolidate all costs as of tomorrow. What should you do?

- A. Link the acquired company's projects to your company's billing account.
- B. Configure the acquired company's billing account and your company's billing account to export the billing data into the same BigQuery dataset.
- C. Migrate the acquired company's projects into your company's GCP organization. Link the migrated projects to your company's billing account.
- D. Create a new GCP organization and a new billing account. Migrate the acquired company's projects and your company's projects into the new GCP organization and link the projects to the new billing account.

Correct Answer: D

Reference:

<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>*Community vote distribution*

A (67%)

B (33%)

 **GunjGupta** Highly Voted 2 years, 4 months ago

To me, A looks correct. projects are linked to another organization as well in the acquired company so migrating would need google cloud support. we can not do ourselves. however, we can link other company projects to an existing billing account to generate total cost.

<https://medium.com/google-cloud/google-cloud-platform-cross-org-billing-41c5db8fef6>

upvoted 45 times

 **uganeshku** 10 months, 2 weeks ago

A is correct because linking all projects of the acquired organization to the main organization's billing account will generate a single bill for all projects.

D is incorrect because there is no need to create a new organization for this.

upvoted 3 times

 **spudleymcdudley** 2 years, 4 months ago

Listen to this guy. It's 'A' as moving projects can take some time from Google. There's no need to create a new organisation and other options don't make any sense

upvoted 9 times

 **lxgywil** 1 year, 6 months ago

You're saying it as if "moving projects" was a viable option. What about B?

upvoted 2 times

 **ryumada** 3 months ago

I think B is not make sense. You don't want to do statistical analytic to the billing data. You want to consolidate all the costs as of tomorrow. So, the costs as of tomorrow should be billed in one billing account. That's what I've understand from the question.

upvoted 1 times

 **XRiddlerX** Highly Voted 2 years, 3 months ago

I could be missing something but where does it say in the question that the two orgs want to migrate projects? I believe the question and key points are "consolidate all GCP costs" and "consolidate all costs as of tomorrow". With that said, C and D would not be a 24 hour task and seems a bit cumbersome to perform for something simple as "creating a single invoice" AND that's a migration and not a consolidation of cost. With A, I can't find anywhere in GCP docs that this is a best practice, only a medium.com blog. IMHO, I won't go down this route because "Just because you can do something, doesn't mean you should." and I would consult GCP support for best practices on A before I do something like that.

That leaves B which is to export both detailed billing to BigQuery and create a invoice/report. This would be a temporary solution until you migrate Organizations. IMHO

I go with B.

upvoted 43 times

 **TAvenger** 1 year, 8 months ago

I am not sure that exporting some statistical data to BigQuery means anything for Google who creates the invoice. With "A" you are right, that is not the best practice, but the key word "for tomorrow" allows this custom approach. So the answer is "A"

upvoted 15 times

-  **zaxma** 7 months, 1 week ago
I will go with A in the exam as well, but just wondering, they are two different organisations, how can you link all projects from org2 to org1's billing account without the help of GCP support??
upvoted 1 times
-  **myuniqueusername** 1 year, 1 month ago
How can Bigquery consolidate into 1 invoice? It's for analytics, not invoicing. B is therefore incorrect, and A must be right.
upvoted 8 times
-  **RazOr** 9 months, 2 weeks ago
You can by querying the data and create output that groups costs. BigQuery is not analytics, it can create the datasets you need from given data.
upvoted 1 times
-  **obeythefist** 7 months, 4 weeks ago
That will show you the numbers, but that does not create an invoice.
upvoted 2 times
-  **ashrafh** 1 year, 2 months ago
I also vote B,
why?
agree with this technical explanation and my finance team not gonna pay some newly acquired company bill by tomorrow :)
upvoted 5 times
-  **Armne96X** 11 months, 1 week ago
Are you sure you can do all steps by tomorrow?
(You would like to consolidate all costs as of tomorrow)
upvoted 1 times
-  **Robertolo** Most Recent 4 weeks ago
Selected Answer: B
A. Can hundreds of projects be linked for tomorrow to the company's billing account? I highly doubt it.
B. Exporting to BigQuery seems a reasonable urgent solution to consolidate costs ("as of tomorrow").
C. Can hundreds of projects be migrated + linked for tomorrow? Same as with A, I doubt it.
D. Same as A and C: I doubt that we can create a new GCP organization, migrate all new and old projects and link them to the new billing account in just a single day. This is madness :-D
upvoted 1 times
-  **sjagkoo7** 4 weeks, 1 day ago
i will go with option D . As question states single invoice for all projects means all should be linked to one billing account. Re-organising the only option this case..
upvoted 1 times
-  **ashit44244** 1 month, 3 weeks ago
I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!
upvoted 1 times
-  **orious** 3 months, 2 weeks ago
Selected Answer: B
Cost consolidation is an Analytical Task (Account task), you rely on analytical tools, therefore you use BigQuery. So, there is no need for complexity.
-- "Consolidation Expenses means, for any period and with respect to any Person, the aggregate of all expenditures by such Person and its Subsidiaries on a consolidated basis during such period related to the consolidation of Station". (
Consolidation Expenses Definition | Law Insider <https://www.lawinsider.com/dictionary/consolidation-e...>)
upvoted 2 times
-  **AzureDP900** 4 months, 2 weeks ago
A is correct
upvoted 1 times
-  **S00999** 5 months ago
Selected Answer: B
"another company that also has hundreds of projects"
Reconfiguration = Too much work for the current day :)
Go B !
upvoted 2 times
-  **haroldbenites** 5 months, 1 week ago
Go for A
upvoted 1 times
-  **Sukon_Desknot** 7 months, 3 weeks ago
Selected Answer: A

The answer is most likely A, because you can move projects in 24 hours as stated here

https://cloud.google.com/resource-manager/docs/project-migration#oauth_consent_screen

upvoted 3 times

✉ **somenick** 7 months, 4 weeks ago

Selected Answer: A

The answer is definitely A. The keywords are "as of tomorrow" and "single invoice". The quickest(the only?) way of achieving this is by Linking. The proper way would be under C, but that's not possible in 24h or less

upvoted 5 times

✉ **pranavdhopey** 9 months, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

✉ **[Removed]** 10 months, 1 week ago

Selected Answer: A

The important thing here is that whichever option you choose should be easily doable in 24 hrs

upvoted 2 times

✉ **Naren080914** 10 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/resource-manager/docs/project-migration>

upvoted 2 times

✉ **arapin** 11 months ago

The answer is definitely A. The keywords are "as of tomorrow" and "single invoice". The quickest(the only?) way of achieving this is by Linking. The proper way would be under C, but that's not possible in 24h or less.

upvoted 4 times

✉ **jabrrJ68w02ond1** 11 months, 2 weeks ago

The question does not state any information that we need to migrate the two orgs and their projects into one big org. Exporting billing stuff to BigQuery indeed shows us the data consolidated, but we also need one single invoice. So I go with A. It's simple, does the job and I do not need to hassle around with moving projects / consult GCP support for a new org.

upvoted 3 times

✉ **Ridhanya** 11 months, 2 weeks ago

option a seems right

upvoted 1 times

Question #67

You built an application on Google Cloud that uses Cloud Spanner. Your support team needs to monitor the environment but should not have access to table data.

You need a streamlined solution to grant the correct permissions to your support team, and you want to follow Google-recommended practices. What should you do?

- A. Add the support team group to the roles/monitoring.viewer role
- B. Add the support team group to the roles/spanner.databaseUser role.
- C. Add the support team group to the roles/spanner.databaseReader role.
- D. Add the support team group to the roles/stackdriver.accounts.viewer role.

Correct Answer: B*Community vote distribution*

A (100%)

✉  **poogcp** Highly Voted 2 years, 5 months ago

its A, As you need to monitor only
upvoted 40 times

✉  **WindDriver** 1 year, 3 months ago

A, right, correct answer.
B and C are incorrect because allow to read data.
D also incorrect: Not for monitoring. roles/stackdriver.accounts.viewer Stackdriver Accounts Viewer:
Read-only access to get and list information about Stackdriver account structure (resourcemanager.projects.get,
resourcemanager.projects.list and stackdriver.projects.get)
upvoted 10 times

✉  **WindDriver** 1 year, 3 months ago

<https://cloud.google.com/iam/docs/understanding-roles>
upvoted 3 times

✉  **Gurnoor** Highly Voted 2 years, 4 months ago

A is correct as user should not have any access to data, so B and C cant be used in this scenario.
upvoted 17 times

✉  **Cornholio_LMC** Most Recent 1 month, 2 weeks ago

had this question today
upvoted 2 times

✉  **raghu09** 1 month, 3 weeks ago

Selected Answer: A
B is wrong because it grants write access also we only need monitoring access.
upvoted 1 times

✉  **RanjithK** 4 months, 1 week ago

Selected Answer: A
A. This is the only role that provides read-only access to get and list information about all monitoring data and configurations.
upvoted 2 times

✉  **AzureDP900** 4 months, 2 weeks ago

You only need to monitor so A is correct!
upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

roles/monitoring.viewer
Monitoring Viewer Grants read-only access to Monitoring in the Google Cloud console and API.
upvoted 1 times

✉  **wolfie09** 5 months ago

Selected Answer: A
A is correct
upvoted 2 times

✉  **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **Priyankahere** 9 months, 3 weeks ago

This was there in exam, go with community answers.

upvoted 7 times

 **NoniGeorge** 10 months ago

Selected Answer: A

You only need to monitor so A is correct!

upvoted 1 times

 **Behrin** 10 months, 4 weeks ago

Selected Answer: A

It's A

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: A

A is right option as monitoring is the requirement and reading data should not be allowed.

upvoted 1 times

 **jabrrJ68w02ond1** 11 months, 2 weeks ago

Selected Answer: A

B/C are out, they enable you to read data. D allows you to do stuff in Stackdriver, but it is not mentioned that any data is even being exported to that. So I go with A, which enables users to view monitoring of arbitrary resources (including Spanner-related Monitoring)

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Add the support team group to the roles/monitoring.viewer role

upvoted 1 times

 **TenshiD** 11 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **Rameshfmf** 1 year ago

All the options are given wrong. They should be roles/spanner.viewer. Since there is no option....I will go with Option B.

upvoted 1 times

 **sunilw** 1 year, 4 months ago

A is correct

https://cloud.google.com/monitoring/access-control#mon_roles_desc

Predefined roles

The following IAM roles are predefined by Cloud Monitoring. They grant permissions only for Monitoring.

Monitoring

The following roles grant general permissions for Monitoring:

Name

Title Description

roles/monitoring.viewer

Monitoring Viewer Gives you read-only access to Monitoring in the Google Cloud Console and API.

roles/monitoring.editor

Monitoring Editor Gives you read-write access to Monitoring in the Google Cloud Console and API, and lets you modify a metrics scope.

roles/monitoring.admin

Monitoring Admin Gives you full access to Monitoring in the Google Cloud Console and lets you modify a metrics scope.

The following role is used by service accounts for write-only access:

Name

Title Description

roles/monitoring.metricWriter

Monitoring Metric Writer Permits writing monitoring data to a metrics scope; doesn't permit access to Monitoring in the Google Cloud Console. For service accounts.

upvoted 6 times

Question #68

For analysis purposes, you need to send all the logs from all of your Compute Engine instances to a BigQuery dataset called platform-logs. You have already installed the Cloud Logging agent on all the instances. You want to minimize cost. What should you do?

- A. 1. Give the BigQuery Data Editor role on the platform-logs dataset to the service accounts used by your instances. 2. Update your instances' metadata to add the following value: logs-destination: bq://platform-logs.
- B. 1. In Cloud Logging, create a logs export with a Cloud Pub/Sub topic called logs as a sink. 2. Create a Cloud Function that is triggered by messages in the logs topic. 3. Configure that Cloud Function to drop logs that are not from Compute Engine and to insert Compute Engine logs in the platform-logs dataset.
- C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.
- D. 1. Create a Cloud Function that has the BigQuery User role on the platform-logs dataset. 2. Configure this Cloud Function to create a BigQuery Job that executes this query: `INSERT INTO dataset.platform-logs (timestamp, log) SELECT timestamp, log FROM compute.logs WHERE timestamp > DATE_SUB(CURRENT_DATE(), INTERVAL 1 DAY)` 3. Use Cloud Scheduler to trigger this Cloud Function once a day.

Correct Answer: C*Community vote distribution*

C (100%)

 **sumanshu** Highly Voted 1 year, 7 months ago
vote for "C"

https://cloud.google.com/logging/docs/export/configure_export_v2
upvoted 16 times

 **vmart** Highly Voted 1 year, 5 months ago
I vote for C
upvoted 7 times

 **Charumathi** Most Recent 1 month ago
C. is correct, Sinks control how Cloud Logging routes logs. Using sinks, you can route some or all of your logs to supported destinations.
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago
I will go with C..
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago
Go for C
upvoted 1 times

 **somenick** 7 months, 3 weeks ago
Selected Answer: C
Outdated question. It's now about Cloud Sink, but C is the closest option
upvoted 5 times

 **NoniGeorge** 10 months ago
I think it's C as all the other ones seem to get logs from everywhere not just Compute Engine!
upvoted 4 times

 **jaykumarjk99** 11 months, 2 weeks ago
Selected Answer: C
I vote for C
upvoted 4 times

 **vishnukumartr** 11 months, 3 weeks ago
C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.
upvoted 2 times

 **Nivesh93** 1 year, 2 months ago
c is correct as it uses filter so it basically reduces the cost of operation .
upvoted 2 times

 **Finger41** 1 year, 5 months ago

Its C.

upvoted 4 times

 **arsh1916** 1 year, 5 months ago

C should be correct

upvoted 3 times

 **mcaromit** 1 year, 6 months ago

C is correct as it would restrict the volume of data to be scanned by bq, thus reducing cost

upvoted 4 times

 **[Removed]** 1 year, 7 months ago

C is correct. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.

upvoted 5 times

 **GCP_Student1** 1 year, 7 months ago

C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.

upvoted 4 times

 **Hi2ALL** 1 year, 7 months ago

C is correct

upvoted 2 times

 **drizzydroo** 1 year, 7 months ago

C right? No need to establish pub/sub topic?

upvoted 2 times

Question #69

You are using Deployment Manager to create a Google Kubernetes Engine cluster. Using the same Deployment Manager deployment, you also want to create a DaemonSet in the kube-system namespace of the cluster. You want a solution that uses the fewest possible services. What should you do?

- A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.
- B. Use the Deployment Manager Runtime Configurator to create a new Config resource that contains the DaemonSet definition.
- C. With Deployment Manager, create a Compute Engine instance with a startup script that uses kubectl to create the DaemonSet.
- D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-access-for-kubectl>

Community vote distribution

A (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Adding an API as a type provider

This page describes how to add an API to Google Cloud Deployment Manager as a type provider. To learn more about types and type providers, read the Types overview documentation.

A type provider exposes all of the resources of a third-party API to Deployment Manager as base types that you can use in your configurations. These types must be directly served by a RESTful API that supports Create, Read, Update, and Delete (CRUD).

If you want to use an API that is not automatically provided by Google with Deployment Manager, you must add the API as a type provider.

<https://cloud.google.com/deployment-manager/docs/configuration/type-providers/creating-type-provider>

upvoted 59 times

 **magistrum** 1 year, 10 months ago

very good find, sounds like you hit the nail in the head

upvoted 6 times

 **kishoredeena** Highly Voted 2 years, 4 months ago

Option A is the right answer

upvoted 10 times

 **RanjithK** Most Recent 4 months, 1 week ago

Selected Answer: A

Answer is A.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

go with A as per ESP_SAP explanations..

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

A is the correct answe, the API need to be added as a type provider

upvoted 1 times

 **ArunTaneja** 9 months, 2 weeks ago

Selected Answer: A

A should be correct one

upvoted 1 times

 **thuvh** 10 months, 2 weeks ago

Selected Answer: A

<https://medium.com/google-cloud/cloud-deployment-manager-kubernetes-2dd9b8124223>

upvoted 1 times

 **MCMS** 11 months, 1 week ago

Selected Answer: A

Correct Answer is (A)

upvoted 2 times

 **PR0704** 11 months, 2 weeks ago

couldn't be more confusing

upvoted 4 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

dd the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

A is correct

upvoted 1 times

 **Linus11** 1 year, 7 months ago

A is correct.

upvoted 1 times

 **pca2b** 1 year, 7 months ago

A: in addition to ESP_SAP's comment, the following link explains exactly how to...

<https://github.com/GoogleCloudPlatform/deploymentmanager-samples/blob/master/examples/v2/gke/daemonsets/daemonset.jinja>

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

A is correct. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.

upvoted 3 times

 **sumanshu** 1 year, 7 months ago

vote for 'D'

upvoted 2 times

Question #70

You are building an application that will run in your data center. The application will use Google Cloud Platform (GCP) services like AutoML. You created a service account that has appropriate access to AutoML. You need to enable authentication to the APIs from your on-premises environment. What should you do?

- A. Use service account credentials in your on-premises application.
- B. Use gcloud to create a key file for the service account that has appropriate permissions.
- C. Set up direct interconnect between your data center and Google Cloud Platform to enable authentication for your on-premises applications.
- D. Go to the IAM & admin console, grant a user account permissions similar to the service account permissions, and use this user account for authentication from your data center.

Correct Answer: B

Reference:

<https://cloud.google.com/vision/automl/docs/before-you-begin>*Community vote distribution*

B (100%)

✉  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct answer should be (B):

To use a service account outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal.

<https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 41 times

✉  **Kopy** Most Recent 3 days, 13 hours ago

Selected Answer: B

Correct answer should be (B):

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

B is right

To use a service account from outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal. When you create a service account key, the public portion is stored on Google Cloud, while the private portion is available only to you. For more information about public/private key pairs, see Service account keys.

upvoted 3 times

✉  **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

✉  **NoniGeorge** 10 months ago

Selected Answer: B

Even thought A and B seem to be doing the same thing the best practice is to create a key so B is the right answer !

upvoted 1 times

✉  **vishnukumartr** 11 months, 3 weeks ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 3 times

✉  **shawnkkk** 11 months, 3 weeks ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 3 times

✉  **Vivekvkt123** 1 year, 1 month ago

Why not A? Aren't A and B getting the same key file?

upvoted 2 times

✉  **jabrrJ68w02ond1** 11 months, 3 weeks ago

A is not really telling you the steps to accomplish the task, it's only telling you the result of it (creating a SA with sufficient permissions and then use Console / gcloud to create a JSON token for it)

upvoted 1 times

 **sunilw** 1 year, 4 months ago

B is correct.

Creating service account keys

To use a service account from outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal. When you create a service account key, the public portion is stored on Google Cloud, while the private portion is available only to you. For more information about public/private key pairs, see Service account keys.

upvoted 3 times

 **arsh1916** 1 year, 5 months ago

B is correct

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

B is correct as a 1st step...the key file is to be referenced in the env variable GOOGLE_APPLICATION_CREDENTIALS which would then provide access to on-prem application using ADC library

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 2 times

 **nliaustemac** 1 year, 8 months ago

B. is the correct answer:

Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 2 times

 **lutoa** 1 year, 9 months ago

I think it is B. I have seen other variations on this question where the scenario specifies that the on-premise has no internet, in which case interconnect is required.

This question does NOT specify, and since it is using API connection (so HTTP) it must have internet so the JSON key should be fine in answer B.

upvoted 2 times

 **DucSiu** 1 year, 9 months ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 1 times

 **DucSiu** 1 year, 9 months ago

D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.

upvoted 1 times

Question #71

You are using Container Registry to centrally store your company's container images in a separate project. In another project, you want to create a Google Kubernetes Engine (GKE) cluster. You want to ensure that Kubernetes can download images from Container Registry. What should you do?

- A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.
- B. When you create the GKE cluster, choose the Allow full access to all Cloud APIs option under 'Access scopes'.
- C. Create a service account, and give it access to Cloud Storage. Create a P12 key for this service account and use it as an imagePullSecrets in Kubernetes.
- D. Configure the ACLs on each image in Cloud Storage to give read-only access to the default Compute Engine service account.

Correct Answer: C*Community vote distribution*

A (83%)

R (17%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer (A):

IAM permissions

IAM permissions determine who can access resources. All users, service accounts, and other identities that interact with Container Registry must have the appropriate Cloud Storage permissions.

By default, Google Cloud use default service accounts to interact with resources within the same project. For example, the Cloud Build service account can both push and pull images when Container Registry is in the same project.

You must configure or modify permissions yourself if:

You are using a service account in one project to access Container Registry in a different project

You are using a default service account with read-only access to storage, but you want to both pull and push images

You are using a custom service account to interact with Container Registry

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 55 times

 **[Removed]** 2 years, 1 month ago

A is correct, practical implementation in video <https://www.youtube.com/watch?v=R16z7Sjrkxs>

upvoted 8 times

 **XRiddlerX** Highly Voted 2 years, 2 months ago

A is correct...

Container Registry uses Cloud Storage buckets as the underlying storage for container images. You control access to your images by granting appropriate Cloud Storage permissions to a user, group, service account, or other identity.

If the service account needs to access Container Registry in another project, you must grant the required permissions in the project with Container Registry.

Reference:

<https://cloud.google.com/container-registry/docs/access-control#permissions>

upvoted 16 times

 **GaneshSurwase** Most Recent 3 weeks, 4 days ago

CORRET ANS is A

upvoted 1 times

 **ravip12345** 1 month, 4 weeks ago

Answer A is correct.

Storage Object Viewer (roles/storage.objectViewer) -- Grant the role on the registry storage bucket.

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Storage Object viewer is enough, A is right.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A.

the option A is specific with the role that will be used. In GCP , the recommendations is using the specific permission. The others options not are specific and are not correct .

upvoted 1 times

✉ **somenick** 7 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/container-registry/docs/access-control>

upvoted 2 times

✉ **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

A is the correct answe because it follows google recommended practices to grant permissions to service accounts, also the object viewer is the appropriate role for the GKE to pull the image

upvoted 2 times

✉ **jdx000** 9 months ago

Selected Answer: A

A is the most obvious

upvoted 1 times

✉ **ArunTaneja** 9 months, 2 weeks ago

A should be correct

https://cloud.google.com/container-registry/docs/pushing-and-pulling#pulling_images_from_a_registry

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

upvoted 2 times

✉ **NoniGeorge** 10 months ago

It's A as you need Storage Object Viewer IAM role in order to have access to the images!

upvoted 1 times

✉ **DiegoCG** 10 months ago

Selected Answer: R

ES A Y PÚNTO

upvoted 1 times

✉ **Pret** 11 months, 1 week ago

Option A:

Pulling images from a registry

Pulling an image requires the Storage Object Viewer for the registry storage bucket, or a role with the same permissions.

To pull from Container Registry, use the command:

`docker pull HOSTNAME/PROJECT-ID/IMAGE:TAG`

<https://cloud.google.com/container-registry/docs/pushing-and-pulling>

upvoted 1 times

✉ **PradeepPen** 11 months, 1 week ago

C is correct answer: imagePullSecrets is use for pulling images from the private registry.

upvoted 1 times

✉ **vishnukumartr** 11 months, 3 weeks ago

A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.

upvoted 1 times

✉ **shawnkkk** 11 months, 3 weeks ago

A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.

upvoted 1 times

✉ **Zesn** 1 year ago

How A is correct??

Storage Object Viewer IAM role can only view or list the object. Here the ask is to pull the image. Must be C

upvoted 1 times

✉ **Zesn** 1 year ago

sorry Ignore this comment.

upvoted 1 times

Question #72

You deployed a new application inside your Google Kubernetes Engine cluster using the YAML file specified below.

```
apiVersion: apps/v1           apiVersion: v1
kind: Deployment               kind: Service
metadata:
  name: myapp-deployment     metadata:
  spec:
    selector:
      matchLabels:
        app: myapp
    replicas: 2
    template:
      metadata:
        labels:
          app: myapp
      spec:
        containers:
          - name: myapp
            image: myapp:1.1
          ports:
            - containerPort: 80
spec:
  selector:
    app: myapp
```

You check the status of the deployed pods and notice that one of them is still in PENDING status:

```
kubectl get pods -l app=myapp
NAME                  READY   STATUS    RESTART   AGE
myapp-deployment-58ddbbb995-1p86m   0/1     Pending   0          9m
myapp-deployment-58ddbbb995-qjpkd   1/1     Running   0          9m
```

You want to find out why the pod is stuck in pending status. What should you do?

- A. Review details of the myapp-service Service object and check for error messages.
- B. Review details of the myapp-deployment Deployment object and check for error messages.
- C. Review details of myapp-deployment-58ddbbb995-1p86m Pod and check for warning messages.
- D. View logs of the container in myapp-deployment-58ddbbb995-1p86m pod and check for warning messages.

Correct Answer: C

Reference:

<https://cloud.google.com/run/docs/gke/troubleshooting>

Community vote distribution

C (75%)

B (25%)

✉ **spudleymcdudley** Highly Voted 2 years, 4 months ago

It's C - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#debugging-pods>
upvoted 25 times

✉ **someoneinthecloud** Highly Voted 2 years, 3 months ago

Answer is C - You can't view logs of a pod that isn't deployed, so D is incorrect.
C allows you to check the pod deployment messages and look for errors

upvoted 19 times

✉ **sidharthwader** 1 year, 4 months ago

What u said is incorrect you can view pod's log even in pending state.
kubectl logs <pod-name> -n <namespace>
upvoted 3 times

✉ **Charumathi** Most Recent 1 month ago

C is correct,
Debugging Pods
The first step in debugging a Pod is taking a look at it. Check the current state of the Pod and recent events with the following command:

kubectl describe pods \${POD_NAME}
upvoted 1 times

Letahrgicbeagle 1 month, 4 weeks ago

Selected Answer: C

Definitely

upvoted 1 times

Dheeraj1986 2 months, 3 weeks ago

Selected Answer: B

I guess it's B. its deployment that creates the pod and it has the information why it is not able to create. it shows the information if you describe the deployment (kubectl describe deployment)

upvoted 1 times

abirroy 3 months ago

Selected Answer: C

Answer is C

upvoted 1 times

RanjithK 4 months, 1 week ago

Selected Answer: C

Answer is C

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

C is right

upvoted 1 times

haroldbenites 5 months, 1 week ago

Go for C

upvoted 1 times

sharmaishu 8 months, 3 weeks ago

C is the correct answer: The first step in debugging a Pod is taking a look at it. Check the current state of the Pod and recent events with the following command:

kubectl describe pods \${POD_NAME}

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

upvoted 1 times

sharmaishu 8 months, 3 weeks ago

C is the correct answer.

If a Pod is stuck in Pending it means that it can not be scheduled onto a node. Generally this is because there are insufficient resources of one type or another that prevent scheduling. Look at the output of the kubectl describe ... command above. There should be messages from the scheduler about why it can not schedule your Pod.

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

upvoted 1 times

vishnukumartr 11 months, 3 weeks ago

C. Review details of myapp-deployment-58dbbbb995-lp86m Pod and check for warning messages.

upvoted 1 times

shawnkkk 11 months, 3 weeks ago

C. Review details of myapp-deployment-58dbbbb995-lp86m Pod and check for warning messages.

upvoted 2 times

AshisKumar 11 months, 3 weeks ago

C is correct . As the pod is not running it can't able to send logs .

so D is wrong

upvoted 2 times

latislal 1 year, 4 months ago

Why not D.

upvoted 1 times

associatecloudexamuser 1 year, 4 months ago

You can't view logs of a pod that isn't deployed

upvoted 1 times

mcaromit 1 year, 6 months ago

C is correct

upvoted 1 times

[Removed] 1 year, 7 months ago

C is correct. Review details of myapp-deployment-58dbbbb995-lp86m Pod and check for warning messages.

upvoted 2 times

Question #73

You are setting up a Windows VM on Compute Engine and want to make sure you can log in to the VM via RDP. What should you do?

- A. After the VM has been created, use your Google Account credentials to log in into the VM.
- B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.
- C. When creating the VM, add metadata to the instance using 'windows-password' as the key and a password as the value.
- D. After the VM has been created, download the JSON private key for the default Compute Engine service account. Use the credentials in the JSON file to log in to the VM.

Correct Answer: D

Community vote distribution

B (100%)

 **John_Iam**  2 years, 5 months ago

Correct Answer is B.

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

<https://cloud.google.com/sdk/gcloud/reference/beta/compute/reset-windows-password>

upvoted 51 times

 **voler** 2 years, 4 months ago

Yes! "If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

upvoted 2 times

 **ragu123** 2 years, 2 months ago

Correct answer is B.

gcloud beta compute reset-windows-password allows a user to reset and retrieve a password for a Windows virtual machine instance. If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned.

upvoted 3 times

 **dan80** 2 years, 5 months ago

did you even look at the link you provide ? it clearly say gcloud beta compute reset-windows-password my-instance and not gcloud compute reset-windows-password. D is correct - <https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 3 times

 **dan80** 2 years, 5 months ago

nobody talk on reset the password but how to access the Windows - best way - Service Account

upvoted 1 times

 **lxgywil** 1 year, 6 months ago

Oh yes? Then what about this link (for non-beta command)?

<https://cloud.google.com/sdk/gcloud/reference/compute/reset-windows-password>

"If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

The answer is obviously B. Just test it and it'll become very clear

upvoted 5 times

 **ale_brd_**  1 month, 3 weeks ago

Selected Answer: B

After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is perfect for this question..

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

<https://cloud.google.com/sdk/gcloud/reference/compute/reset-windows-password>

upvoted 2 times

 **Himadhar1997** 5 months, 2 weeks ago

B is correct

after VM creating , you can reset the password

upvoted 1 times

 **Harbeeb** 6 months ago

Selected Answer: B
After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM
upvoted 2 times

 **dinesh198728** 7 months, 4 weeks ago

Selected Answer: B
b is correct
upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: B
B is the correct answer, the command is the correct one
upvoted 1 times

 **ttttttttt** 9 months, 1 week ago

Test. Yyyyy
upvoted 1 times

 **xianpxian** 10 months ago

Yes, It's B
upvoted 1 times

 **Zufair** 10 months, 2 weeks ago

Selected Answer: B
Correct Answer is B. Tested it.
upvoted 1 times

 **Zufair** 10 months, 2 weeks ago

Correct Answer is B .
Tested it .
upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.
upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.
upvoted 1 times

 **dtncl** 1 year ago

Answer is B.

Steps are clearly stated here:
<https://cloud.google.com/compute/docs/instances/windows/generating-credentials#gcloud>
upvoted 1 times

 **PriyaGCP** 1 year, 4 months ago

Answer B
<https://cloud.google.com/compute/docs/quickstart-windows>
upvoted 3 times

 **alperenhasan** 1 year, 5 months ago

folks B is correct. you can try on console google cloud.
upvoted 2 times

Question #74

You want to configure an SSH connection to a single Compute Engine instance for users in the dev1 group. This instance is the only resource in this particular Google Cloud Platform project that the dev1 users should be able to connect to. What should you do?

- A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.
- B. Set metadata to enable-oslogin=true for the instance. Set the service account to no service account for that instance. Direct them to use the Cloud Shell to ssh to that instance.
- C. Enable block project wide keys for the instance. Generate an SSH key for each user in the dev1 group. Distribute the keys to dev1 users and direct them to use their third-party tools to connect.
- D. Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

Correct Answer: D

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>*Community vote distribution*

A (100%)

 **poogcp** Highly Voted 2 years, 5 months ago

A correct one

upvoted 43 times

 **nithinpb180** 2 years, 4 months ago

Agree with that

upvoted 3 times

 **spudleymcdudley** 2 years, 4 months ago

For further evidence... <https://cloud.google.com/compute/docs/instances/managing-instance-access>

upvoted 6 times

 **student002** Highly Voted 2 years ago

Pure from logic thinking: A can't be right. If the group get access to that instance with enable-oslogin=true, then they could have access to every instance that has enable-oslogin=true. Or do I miss something?

upvoted 10 times

 **akshaychavan7** 5 months, 3 weeks ago

Note the sentence "Set metadata to enable-oslogin=true for the instance." This means the metadata for oslogin has been set to that particular instance only, and not for all.

upvoted 3 times

 **magistrum** 1 year, 10 months ago

I'm convinced with this logic

upvoted 1 times

 **bgallet** 11 months ago

clearly, question say "the only ressource they need to access in this project"
as you said, all ressources will be available if we set the role

upvoted 2 times

 **Charumathi** Most Recent 1 month ago

A is the correct answer

Granting OS Login IAM roles

After you enable OS Login on one or more instances in your project, those VMs accept connections only from user accounts that have the necessary IAM roles in your project or organization.

roles/compute.osLogin, which doesn't grant administrator permissions

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: A

Answer is A

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is correct..

upvoted 1 times

 **LaxmanTiwari** 5 months, 2 weeks ago

Selected Answer: A

For further evidence... <https://cloud.google.com/compute/docs/instances/managing-instance-access>

upvoted 1 times

 **pfabio** 6 months ago

Selected Answer: A

A is correct and recommended option.

D is incorrect because block project-wide restrict access to this instance, evidence: <https://cloud.google.com/compute/docs/connect/restrict-ssh-keys>

upvoted 1 times

 **juliusali** 6 months, 2 weeks ago

D. Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

upvoted 1 times

 **somenick** 7 months, 3 weeks ago

Selected Answer: A

A is the correct answer. It is recommended by Google at here <https://cloud.google.com/compute/docs/instances/connecting-advanced#linux-and-macos>

upvoted 2 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

A is the correct answer

upvoted 1 times

 **ElenaL** 10 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.

upvoted 3 times

 **shawnkkk** 11 months, 3 weeks ago

Selected Answer: A

A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.

upvoted 3 times

 **Icemen** 1 year ago

A is the correct answer

upvoted 1 times

 **jxzhjq** 1 year, 2 months ago

A is correct

<https://cloud.google.com/compute/docs/instances/managing-instance-access>

upvoted 1 times

 **Ruyal** 1 year, 3 months ago

A is correct .

You can grant roles/compute.osLogin instance access roles at the instance level by using the gcloud compute instances add-iam-policy-binding command.

<https://cloud.google.com/compute/docs/instances/managing-instance-access#grant-iam-roles>

upvoted 4 times

Question #75

You need to produce a list of the enabled Google Cloud Platform APIs for a GCP project using the gcloud command line in the Cloud Shell. The project name is my-project. What should you do?

- A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.
- B. Run gcloud init to set the current project to my-project, and then run gcloud services list --available.
- C. Run gcloud info to view the account value, and then run gcloud services list --account <Account>.
- D. Run gcloud projects describe <project ID> to verify the project value, and then run gcloud services list --available.

Correct Answer: A*Community vote distribution*

A (100%)

 **dan80** Highly Voted 2 years, 5 months ago

A is the correct answer, log to gcloud and run the commands, doesn't make sense to run cloud init and gcloud services list --available gives you the full services that are available.

upvoted 48 times

 **raffiq** 2 years, 4 months ago

Yes, Answer A correct. it shows only enabled services of API

upvoted 6 times

 **Ixgywil** Highly Voted 1 year, 6 months ago

"A" is correct.

For those, who have doubts:

`gcloud services list --available` returns not only the enabled services in the project but also services that CAN be enabled. Therefore, option B is incorrect.

<https://cloud.google.com/sdk/gcloud/reference/services/list#--available>

upvoted 10 times

 **squishy_fishy** 1 year, 5 months ago

Best answer!

upvoted 1 times

 **Charumathi** Most Recent 1 month ago

A is correct answer,

Run the following command to list the enabled APIs and services in your current project:

gcloud services list

whereas, Run the following command to list the APIs and services available to you in your current project:

gcloud services list --available

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is the correct answer

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

A is the correct answer

upvoted 2 times

 **sharmaishu** 8 months, 3 weeks ago

A is the correct answer.

<https://cloud.google.com/sdk/gcloud/reference/services/list#--available>

--available

Return the services available to the project to enable. This list will include any services that the project has already enabled.

To list the services the current project has enabled for consumption, run:

gcloud services list --enabled

To list the services the current project can enable for consumption, run:

gcloud services list --available

upvoted 1 times

 **[Removed]** 10 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 1 times

 **Litan** 1 year, 1 month ago

A Correct

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

A is correct

upvoted 2 times

 **Rabbitfur** 1 year, 6 months ago

A is correct answer --project option available in gcloud command

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

A is correct. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 2 times

 **Vic1043** 1 year, 8 months ago

B is the right answer

Run the following command to list the APIs and services available to you in your current project:

gcloud services list --available

<https://cloud.google.com/service-usage/docs/list-services#gcloud>

upvoted 1 times

 **Abalagu** 1 year, 5 months ago

B is wrong here. The question asks about enabled APIs, not available APIs. The `gcloud list` command returns the enabled APIs by default.

upvoted 1 times

 **sumanshu** 1 year, 7 months ago

B wrong as there is no need to run init command

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 2 times

Question #76

You are building a new version of an application hosted in an App Engine environment. You want to test the new version with 1% of users before you completely switch your application over to the new version. What should you do?

- A. Deploy a new version of your application in Google Kubernetes Engine instead of App Engine and then use GCP Console to split traffic.
- B. Deploy a new version of your application in a Compute Engine instance instead of App Engine and then use GCP Console to split traffic.
- C. Deploy a new version as a separate app in App Engine. Then configure App Engine using GCP Console to split traffic between the two apps.
- D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

Correct Answer: D*Community vote distribution*

D (100%)

 **SIX** Highly Voted 2 years, 5 months ago

Correct answer is D

upvoted 52 times

 **mohdafiuddin** Highly Voted 1 year, 10 months ago

Splitting the question to the key requirements

1. new version of an application hosted in an App Engine environment.
2. test the new version with 1% of users

App engine supports versioning and traffic splitting so no need to involve anything else
(source - <https://cloud.google.com/appengine#all-features>)

A.'Google Kubernetes Engine'.... - No need to involve GKE. Not the right option

B.'Compute Engine instance'.... - No need to involve Compute Engine.

C.'Separate app in App Engine'....- No need to deploy as a separate app. versioning is supported already. Not the right option.

D. This is the right answer.

upvoted 24 times

 **akshaychavan7** 5 months, 3 weeks ago

Just to add, for option C you cannot have two applications deployed inside an app engine project. In order to do so, you need to create the application inside a new project.

So, we just eliminate option C.

upvoted 2 times

 **gcpBeginner** Most Recent 1 month, 1 week ago

without any doubt answer is D

upvoted 1 times

 **learn_GCP** 1 month, 2 weeks ago

Selected Answer: D

D is correct

A. no need to use GKE

B. Compute engine not-relevant.

C. cant deploy 2 applications in one project.

D. --splits & --split-by flags are available for gcloud app deploy.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is correct

upvoted 1 times

 **mplibunao** 5 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 3 times

 **jabrJ68w02nd1** 11 months, 3 weeks ago

D is correct. Practical tip: When deploying your service with gcloud, you can also choose to split traffic without visiting the Console via setting some flags.

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

Correct answer is D

upvoted 1 times

 **arsh1916** 1 year, 5 months ago

Ans: D

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

D is the only correct option

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

D is correct. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 3 times

 **GCP_Student1** 1 year, 8 months ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 3 times

 **lutoa** 1 year, 9 months ago

D. Remember, you can only have a single App Engine app per project.

And the question has nothing to do with GKE. Traffic splitting over versions is the obvious answer (D)

upvoted 4 times

 **DucSiu** 1 year, 9 months ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 2 times

 **victory108** 1 year, 9 months ago

D - Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 1 times

Question #77

You need to provide a cost estimate for a Kubernetes cluster using the GCP pricing calculator for Kubernetes. Your workload requires high IOPs, and you will also be using disk snapshots. You start by entering the number of nodes, average hours, and average days. What should you do next?

- A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.
- B. Fill in local SSD. Add estimated cost for cluster management.
- C. Select Add GPUs. Fill in persistent disk storage and snapshot storage.
- D. Select Add GPUs. Add estimated cost for cluster management.

Correct Answer: C

Reference:

<https://cloud.google.com/products/calculator#tab=container>*Community vote distribution*

A (100%)

✉  **dan80** Highly Voted 2 years, 5 months ago

This one is Tricky, local SSD is require for High IOPS - <https://cloud.google.com/compute/docs/disks/local-ssd> , but it say using disk snapshots. A is correct.

upvoted 53 times

✉  **poogcp** Highly Voted 2 years, 5 months ago

A is correct .

upvoted 16 times

✉  **Racinely** Most Recent 3 weeks, 2 days ago

The question is about storage then A is the correct answer.

upvoted 1 times

✉  **PKookNN** 3 weeks, 5 days ago

Selected Answer: A

for high IOPS needs local ssd and snapshot wording is clear enough so A is correct

upvoted 1 times

✉  **gcpBeginner** 1 month, 1 week ago

we should not use GPU in this case. SSD is correct solution for storage . correct answer is A

upvoted 1 times

✉  **sandipk91** 2 months, 1 week ago

Selected Answer: A

Option A is correct

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A is correct, Local SSD provides high performance.

upvoted 1 times

✉  **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

✉  **LaxmanTiwari** 5 months, 2 weeks ago

Selected Answer: A

This one is Tricky, local SSD is require for High IOPS - <https://cloud.google.com/compute/docs/disks/local-ssd> , but it say using disk snapshots. A is correct.

upvoted 3 times

✉  **gitaexams** 5 months, 4 weeks ago

Selected Answer: A

A is the correct answer since Local SSD provide best IOPS performance

upvoted 1 times

✉  **rsuresh27** 6 months, 2 weeks ago

Selected Answer: A

A is correct. For high IOPS use SSD. And there is no need to include cluster management fees as stated in B because it is already included in the cost.

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

A is the correct answer since Local SSD provide best IOPS performance

upvoted 1 times

 **ArunTaneja** 9 months, 2 weeks ago

Selected Answer: A

A should be correct one

upvoted 2 times

 **NoniGeorge** 10 months ago

Not sure what the question is implying like are we making the cluster now or do we just want an estimate for a new cluster! If new yes A is correct as we do not need a gpu for this kind of job but if the cluster already exist we need to input all the date from that cluster hence if the cluster has gpu you would have to add that too so in that kind of scenario C would be more suited. I'm gonna assumed that we need an estimate for a cluster that we are going to create now so I'll go with A.

upvoted 2 times

 **rafsrod** 10 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **maahibhai** 10 months, 3 weeks ago

Selected Answer: A

A is correct .

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: A

Local SSD is sufficient for high IOPS. Hence A is right choice.

upvoted 1 times

Question #78

You are using Google Kubernetes Engine with autoscaling enabled to host a new application. You want to expose this new application to the public, using HTTPS on a public IP address. What should you do?

- A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.
- B. Create a Kubernetes Service of type ClusterIP for your application. Configure the public DNS name of your application using the IP of this Service.
- C. Create a Kubernetes Service of type NodePort to expose the application on port 443 of each node of the Kubernetes cluster. Configure the public DNS name of your application with the IP of every node of the cluster to achieve load-balancing.
- D. Create a HAProxy pod in the cluster to load-balance the traffic to all the pods of the application. Forward the public traffic to HAProxy with an iptable rule. Configure the DNS name of your application using the public IP of the node HAProxy is running on.

Correct Answer: A

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>*Community vote distribution*

A (100%)

  **dan80** Highly Voted 2 years, 5 months ago

A is correct.

upvoted 32 times

  **magistrum** 1 year, 10 months ago

Saw this which provides good context <https://medium.com/google-cloud/kubernetes-nodeport-vs-loadbalancer-vs-ingress-when-should-i-use-what-922f010849e0>

upvoted 16 times

  **nitinz** 1 year, 8 months ago

you nailed it.

upvoted 2 times

  **arsav** Highly Voted 1 year, 8 months ago

HAProxy is HTTP only, doesn't support HTTPS, so you can reject option D

<https://www.haproxy.org/#desc>

Cluster IP - is an internal IP, you cannot expose public externally. reject option B

out of option A and C

C, port 443 is https but public DNS is not going to give you a load balancing

A is the right choice,

kubernetes ingress exposes HTTPS

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

and cloud load balancer is the right choice which will help to expose the app to public

upvoted 29 times

  **NoniGeorge** 10 months ago

Pretty sure that option D works more from on premise than cloud because with cloud you pretty much don't have to configure your ip tables !

upvoted 1 times

  **gcpBeginner** Most Recent 1 month, 1 week ago

Correct answer is A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 1 times

  **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

  **NOOGLeR** 4 months ago

A is the correct

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

This is A without any second thought.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **BalajiMBala** 6 months, 3 weeks ago

option A is an correct answer

upvoted 1 times

 **cysteine** 7 months ago

Selected Answer: A

A is correct. Because, you want to access Global HTTP(s) Load Balancing with GKE services, you must use Ingress Object.

<https://cloud.google.com/kubernetes-engine/docs/concepts/service>

upvoted 3 times

 **dev_sami** 7 months ago

Selected Answer: A

option B - WRONG --> cluster IP is an internal IP, so we cannt expose publically.

option C - WRONG-->port 443 is HTTPS but public DNS is not provide load balancing.

option D - WRONG -->HAProxy is HTTP only not HTTPS

so

A is right option.

upvoted 4 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

A is correct

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

A is correct. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 5 times

 **yubi69** 1 year, 7 months ago

correct option is A.

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct

upvoted 2 times

 **JackGlemins** 1 year, 8 months ago

A is right: <https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types>

upvoted 2 times

Question #79

You need to enable traffic between multiple groups of Compute Engine instances that are currently running two different GCP projects. Each group of Compute Engine instances is running in its own VPC. What should you do?

- A. Verify that both projects are in a GCP Organization. Create a new VPC and add all instances.
- B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
- C. Verify that you are the Project Administrator of both projects. Create two new VPCs and add all instances.
- D. Verify that you are the Project Administrator of both projects. Create a new VPC and add all instances.

Correct Answer: B*Community vote distribution*

B (100%)

- ✉ **dan80** Highly Voted 2 years, 5 months ago
B - <https://cloud.google.com/vpc/docs/shared-vpc>
upvoted 30 times
- ✉ **glam** Highly Voted 2 years ago
B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 9 times
- ✉ **darcal95** Most Recent 1 month, 3 weeks ago
ok, is B, but that means that the VMs in the "other project" have to change their ip?
upvoted 1 times
- ✉ **AzureDP900** 4 months, 2 weeks ago
It is shared VPC concept. Go with B
upvoted 1 times
- ✉ **AzureDP900** 4 months, 2 weeks ago
Shared VPC allows an organization to connect resources from multiple projects to a common Virtual Private Cloud (VPC) network, so that they can communicate with each other securely and efficiently using internal IPs from that network. When you use Shared VPC, you designate a project as a host project and attach one or more other service projects to it. The VPC networks in the host project are called Shared VPC networks. Eligible resources from service projects can use subnets in the Shared VPC network
upvoted 2 times
- ✉ **mplibunao** 5 months ago
Selected Answer: B
shared vpc is the answer
upvoted 1 times
- ✉ **haroldbenites** 5 months, 1 week ago
Go for B
upvoted 1 times
- ✉ **JohnnyStudy** 7 months ago
B. You use "Shared VPC Network" to share a network across several projects in your GCP organization. You designate a project as a host project, and attach one or more other service projects to it. Only works within the same organization. Only works across projects.
upvoted 2 times
- ✉ **vishnukumartr** 11 months, 3 weeks ago
B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 1 times
- ✉ **shawnkkk** 11 months, 3 weeks ago
B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
upvoted 1 times
- ✉ **Litan** 1 year, 1 month ago
B is correct

upvoted 1 times

 **esandips** 1 year, 1 month ago

B looks incorrect

"For example, an existing instance in a service project cannot be reconfigured to use a Shared VPC network, but a new instance can be created to use available subnets in a Shared VPC network."

<https://cloud.google.com/vpc/docs/shared-vpc>

upvoted 1 times

 **Swapsss** 1 year, 2 months ago

answer is B

upvoted 1 times

 **vmart** 1 year, 5 months ago

b is correct

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

B is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.

upvoted 4 times

 **chithra1990** 1 year, 8 months ago

So when the question talks about communication between 2 VPC's in same organization - go for Shared VPC If diff organization, we can consider VPC Peering

upvoted 8 times

 **GCP_Student1** 1 year, 8 months ago

B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.

upvoted 3 times

Question #80

You want to add a new auditor to a Google Cloud Platform project. The auditor should be allowed to read, but not modify, all project items. How should you configure the auditor's permissions?

- A. Create a custom role with view-only project permissions. Add the user's account to the custom role.
- B. Create a custom role with view-only service permissions. Add the user's account to the custom role.
- C. Select the built-in IAM project Viewer role. Add the user's account to this role.
- D. Select the built-in IAM service Viewer role. Add the user's account to this role.

Correct Answer: C

Reference:

<https://cloud.google.com/resource-manager/docs/access-control-project>

Community vote distribution

C (67%)

A (33%)

✉  **cloudenthu01** Highly Voted 2 years, 4 months ago

C is correct
roles/Viewer role provides access to all resources under the projects but do not alter the state of these resources
upvoted 36 times

✉  **mav3r1ck** 3 months ago

It should be A.

https://cloud.google.com/iam/docs/faq#when_would_i_use_basic_roles

When would I use basic roles?

You can use basic roles in development and test environments, where it might be appropriate for some principals to have wide-ranging permissions. Avoid basic roles in production environments.

upvoted 2 times

✉  **mav3r1ck** 3 months ago

Principle of least privilege

upvoted 1 times

✉  **glam** Highly Voted 2 years ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.

upvoted 13 times

✉  **PKookNN** Most Recent 23 hours, 42 minutes ago

A is more correct (best practice is to avoid basic roles)

upvoted 1 times

✉  **Erie** 2 weeks, 5 days ago

Selected Answer: A

Avoid assigning basic roles.

upvoted 1 times

✉  **ale_brd_** 1 month, 3 weeks ago

Selected Answer: C

roles/Viewer role provides access to all resources under the projects but do not alter the state of these resources

upvoted 1 times

✉  **Frozen_fire** 2 months, 3 weeks ago

Answer - A

Avoid using basic roles except when absolutely necessary. These roles are very powerful, and include a large number of permissions across all Google Cloud services.

<https://cloud.google.com/resource-manager/docs/access-control-project>

upvoted 1 times

✉  **patashish** 4 months, 1 week ago

A is correct answer.

upvoted 1 times

✉  **mplibunao** 5 months ago

Selected Answer: C

C is the answer

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for C

upvoted 1 times

 **P0OJA3808** 9 months, 1 week ago

Project Viewer provides read-only access to all services and resources in a project. Project viewer roles are good for security auditors who need to see everything but build nothing. so, c is correct. https://ondemand.roitraining.com/members/course_play/4246/preview_course_lesson/26798

upvoted 2 times

 **kvkarthikeyan** 11 months, 1 week ago

A is Right Answer.

Caution: Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.
<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 3 times

 **MariGK** 11 months, 1 week ago

IAM role is not a basic role.

upvoted 5 times

 **kvkarthikeyan** 11 months, 1 week ago

Caution: Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

I think Answer is A.

upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.

upvoted 1 times

 **zaxxon** 1 year ago

it should be A: Caution: Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.

upvoted 3 times

 **A_ksh_a_Y** 1 year, 5 months ago

there is nothing like a project Viewer role ? its only a projectBroswer role.

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

C is correct

upvoted 2 times

Question #81

You are operating a Google Kubernetes Engine (GKE) cluster for your company where different teams can run non-production workloads. Your Machine Learning (ML) team needs access to Nvidia Tesla P100 GPUs to train their models. You want to minimize effort and cost. What should you do?

- A. Ask your ML team to add the `accelerator: gpu` annotation to their pod specification.
- B. Recreate all the nodes of the GKE cluster to enable GPUs on all of them.
- C. Create your own Kubernetes cluster on top of Compute Engine with nodes that have GPUs. Dedicate this cluster to your ML team.
- D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

Correct Answer: B*Community vote distribution*

D (100%)

 **John_lam** Highly Voted 2 years, 5 months ago

D is the correct answer.

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

upvoted 45 times

 **tablet444** 2 years, 3 months ago

the documentation states "Limitations

Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.

GPU nodes cannot be live migrated during maintenance events."

upvoted 8 times

 **nightflyer** 1 year, 11 months ago

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 13 times

 **glam** Highly Voted 2 years ago

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

upvoted 14 times

 **raghu09** Most Recent 1 month, 3 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

 **iadarsh** 1 month, 3 weeks ago

Selected Answer: D

D is correct

Because if you create entirely new node pool then its not cost efficient and also the pods which not require that much high GPU is get scheduled into it. So instead of that add a new node pool with GPU and in the pod YAML file mention the node affinity to get scheduled into the GPU enabled node pool.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

By looking at all answers first 3 can be eliminated without any second thought. D is correct.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D

upvoted 1 times

 **LaxmanTiwari** 5 months, 2 weeks ago

Selected Answer: D

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 2 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: D

D is the correct answer
upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: D

D would be the right option when there is possibility to add GPUs without recreating the nodes.
upvoted 3 times

 **ME_MYSELF** 11 months ago

I have a doubt. Will these questions repeat in ace exam or not??Please answer if u have written the exam...it will be very helpful
upvoted 5 times

 **ahsangh** 10 months, 1 week ago

from past experience yes, haven't written this one yet, will revert after writing in a few days.
upvoted 3 times

 **pnVino27** 11 months ago

Selected Answer: D

ANSWER D
upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.
upvoted 1 times

 **Litan** 1 year, 1 month ago

yes D right
upvoted 1 times

 **shankymre01** 1 year, 5 months ago

D is a very wrong candidate but I would say B. because we also need to care take of cost .
if we create a new node pool with GPU and don't delete exiting node pool it's increase pricing.
i would go with B to recreate node pool with GPU as it's test env so I don't need to care about downtime.
upvoted 3 times

 **boof** 1 year, 1 month ago

You are operating a Google Kubernetes Engine (GKE) cluster for your company where DIFFERENT TEAMS can run non-production workloads.

Note the different teams part. Recreating the entire GKE cluster when only the ML team needs access to the nvidia-tesla GPUs would be very disruptive and waste a lot of resources as it's redundant for the other teams. D would be the best pick in this case since it only marginally adds to the cluster instead of recreating the whole thing.

upvoted 4 times

 **mcaromit** 1 year, 6 months ago

D is correct
upvoted 1 times

 **[Removed]** 1 year, 7 months ago

D is correct. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.
upvoted 2 times

 **pondai** 1 year, 7 months ago

D & B can add gpu to this cluster ,but "minimize effort" so I pick D
upvoted 1 times

Question #82

Your VMs are running in a subnet that has a subnet mask of 255.255.255.240. The current subnet has no more free IP addresses and you require 10 IP addresses for new VMs. The existing and new VMs should all be able to reach each other without additional routes. What should you do?

- A. Use gcloud to expand the IP range of the current subnet.
- B. Delete the subnet, and recreate it using a wider range of IP addresses.
- C. Create a new project. Use Shared VPC to share the current network with the new project.
- D. Create a new subnet with the same starting IP but a wider range to overwrite the current subnet.

Correct Answer: C*Community vote distribution*

A (100%)

 **JustLearning** Highly Voted 2 years, 5 months ago

A: Expand the existing subnet.
<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>
 upvoted 47 times

 **glam** Highly Voted 2 years ago

A. Use gcloud to expand the IP range of the current subnet.
 upvoted 11 times

 **dennydream** Most Recent 1 week, 4 days ago

Why in the world would C be the answer? A is def the answer.
 upvoted 1 times

 **Charumathi** 1 month ago

A. is correct,
 gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

gcloud compute networks subnets expand-ip-range NAME --prefix-length=PREFIX_LENGTH [--region=REGION] [GCLOUD_WIDE_FLAG ...]
 upvoted 1 times

 **gcpBeginner** 1 month, 1 week ago

correct answer is A. Use gcloud to expand the IP range of the current subnet.
 upvoted 1 times

 **sandipk91** 2 months, 1 week ago

Selected Answer: A
 Option A is the answer
 upvoted 1 times

 **Shlok27_gcloud** 1 month, 3 weeks ago

Hello can anyone send me the questions to my mail : avinash.shlok0@gmail.com I have exam in next week.
 upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right answer
 upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A
 upvoted 1 times

 **Harish_AP** 6 months, 4 weeks ago

255.255.255.240 = /28 which has 8 Ips , if expand it will go to /29 which has 4 Ips so A is wrong and C is correct
 upvoted 1 times

 **mplibunao** 5 months ago

When you expand the ip range it should go down to /27 or lower not increase. You can't even reduce the IP range to /29
 upvoted 1 times

 **Ery** 7 months, 1 week ago

Selected Answer: A<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 2 times

  **luciorifa** 8 months, 2 weeks ago**Selected Answer: A**

A is correct, it's better and easier if you expand the IP range

upvoted 2 times

  **ayushrockzzz** 9 months, 1 week agoAA
AA
AA
AA
AA
AA
AA
AAAAAAAAAAAAAAAA
upvoted 7 times  **jaffarali** 10 months, 3 weeks ago**Selected Answer: A**

A is the right option

upvoted 1 times

  **pnVino27** 11 months ago

Answer A

upvoted 1 times

  **look1** 11 months, 1 week ago**Selected Answer: A**

Expand

upvoted 1 times

  **jabrrJ68w02ond1** 11 months, 1 week ago**Selected Answer: A**

A. According to the subnet mask, there's still some bits left for sizing it up.

upvoted 3 times

  **sid0127** 11 months, 2 weeks ago**Selected Answer: A**

AAAAAAAAAAAAAA

upvoted 1 times

Question #83

Topic 1

Your organization uses G Suite for communication and collaboration. All users in your organization have a G Suite account. You want to grant some G Suite users access to your Cloud Platform project. What should you do?

- A. Enable Cloud Identity in the GCP Console for your domain.
- B. Grant them the required IAM roles using their G Suite email address.
- C. Create a CSV sheet with all users' email addresses. Use the gcloud command line tool to convert them into Google Cloud Platform accounts.
- D. In the G Suite console, add the users to a special group called cloud-console-users@yourdomain.com. Rely on the default behavior of the Cloud Platform to grant users access if they are members of this group.

Correct Answer: B

Reference:

<https://cloud.google.com/resource-manager/docs/creating-managing-organization>*Community vote distribution*

B (100%)

 **austin1** Highly Voted 2 years, 5 months ago

B is correct

upvoted 27 times

 **Ciumela** Highly Voted 2 years, 4 months ago

B is correct: To actively adopt the Organization resource, the G Suite or Cloud Identity super admins need to assign the Organization Administrator Cloud IAM role to a user or group

upvoted 18 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

B is right

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **Majkl93** 8 months, 3 weeks ago

Selected Answer: B

B as per the comments

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

B is Correct

upvoted 1 times

 **bubblegumbeach** 1 year ago

B is correct

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

B is correct

upvoted 2 times

 **ChristianK** 1 year, 7 months ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

B is correct. Grant them the required IAM roles using their G Suite email address.

upvoted 3 times

 **GCP_Student1** 1 year, 8 months ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 1 times

 **victory108** 1 year, 9 months ago

B - Grant them the required IAM roles using their G Suite email address.

upvoted 1 times

 **swatititame** 1 year, 11 months ago

• A. Use gcloud to expand the IP range of the current subnet.

upvoted 3 times

 **TAvenger** 1 year, 8 months ago

this is for previous question :)

upvoted 4 times

 **nwk** 2 years ago

Should be B

Default behavior does not grant access to the "your GCP Project"

Default behavior allow only create billing account and project - When the organization is created, all users in your domain are automatically granted Project Creator and Billing Account Creator IAM roles at the organization level. This enables users in your domain to continue creating projects with no disruption.

upvoted 3 times

 **glam** 2 years ago

B. Grant them the required IAM roles using their G Suite email address.

upvoted 6 times

 **[Removed]** 2 years, 1 month ago

D is better with best practice.

upvoted 1 times

 **[Removed]** 2 years, 1 month ago

Take back B is correct. There are a lot of uncertain words in the option.

upvoted 1 times

Question #84

You have a Google Cloud Platform account with access to both production and development projects. You need to create an automated process to list all compute instances in development and production projects on a daily basis. What should you do?

- A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.
- B. Create two configurations using gsutil config. Write a script that sets configurations as active, individually. For each configuration, use gsutil compute instances list to get a list of compute resources.
- C. Go to Cloud Shell and export this information to Cloud Storage on a daily basis.
- D. Go to GCP Console and export this information to Cloud SQL on a daily basis.

Correct Answer: A

Community vote distribution

A (100%)

✉  **cludenthu01** Highly Voted 2 years, 4 months ago

A is correct

upvoted 33 times

✉  **glam** Highly Voted 2 years ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 13 times

✉  **110kc3** Most Recent 1 month, 3 weeks ago

Selected Answer: A

"gcloud" can create and manage Google Cloud resources while "gsutil" cannot do so. "gsutil" can manipulate buckets, bucket's objects and bucket ACLs on GCS(Google Cloud Storage) while "gcloud" cannot do so

upvoted 1 times

✉  **Shlok27_gcloud** 1 month, 2 weeks ago

Which one is correct in all cases , suggested one or community one . Im confused totally for all questions

upvoted 1 times

✉  **habros** 3 months ago

A. Gsutil is used for cloud storage bucket

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A is right, This is part of Tutorial Dojo practice questions

upvoted 2 times

✉  **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

✉  **cysteine** 7 months ago

Selected Answer: A

Obviously A is correct

upvoted 1 times

✉  **RegisFTM** 10 months, 3 weeks ago

This looks to be a multiple-choice question. The answer A is correct, and the C completes the task... Does it make sense?

upvoted 1 times

✉  **Zufair** 10 months, 2 weeks ago

C is not automated task

A is correct

upvoted 2 times

✉  **vishnukumartr** 11 months, 3 weeks ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

A is Correct

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 1 times

 **mcaromit** 1 year, 6 months ago

A is correct

upvoted 3 times

 **[Removed]** 1 year, 7 months ago

A is correct. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 4 times

 **GCP_Student1** 1 year, 8 months ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 4 times

 **GCP_Student1** 1 year, 9 months ago

A is the correct answer

upvoted 2 times

 **DucSiu** 1 year, 9 months ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.

upvoted 2 times

 **EABDAJA** 1 year, 9 months ago

Answer A

upvoted 1 times

Question #85

You have a large 5-TB AVRO file stored in a Cloud Storage bucket. Your analysts are proficient only in SQL and need access to the data stored in this file. You want to find a cost-effective way to complete their request as soon as possible. What should you do?

- A. Load data in Cloud Datastore and run a SQL query against it.
- B. Create a BigQuery table and load data in BigQuery. Run a SQL query on this table and drop this table after you complete your request.
- C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.
- D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries.

Correct Answer: C

Community vote distribution

C (100%)

✉  **mohdafiuddin** Highly Voted 1 year, 10 months ago

Breaking down the question into key points -

1. 5-TB AVRO file stored in a Cloud Storage bucket.
2. Analysts are proficient only in SQL
3. cost-effective way to complete their request as soon as possible

A.Load data in Cloud Datastore... (Not Correct because Cloud Datastore is not a good option to run SQL Queries)

B. ...Load data in BigQuery.... (Not Cost Effective because loading the data which is already present in the bucket into BigQuery again is expensive)

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

(This is the right answer as it meets all the requirements from the question)

D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries.

(Too roundabout and indirect. Not the right option)

upvoted 106 times

✉  **pondai** 1 year, 7 months ago

listem this guy

upvoted 12 times

✉  **Ciumela** Highly Voted 2 years, 4 months ago

C is correct: <https://cloud.google.com/bigquery/external-data-sources>

upvoted 21 times

✉  **Untamables** Most Recent 2 weeks, 3 days ago

Selected Answer: C

Similar to Athena

upvoted 1 times

✉  **Charumathi** 1 month ago

C. is correct,

An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Amazon S3
Azure Storage
Cloud Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive

upvoted 1 times

✉  **fifi1907** 1 month, 4 weeks ago

answer is c

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

mohdafiuddin explanation is very detailed .. C is right

upvoted 1 times

✉ **haroldbenites** 5 months, 1 week ago

Go for C

upvoted 1 times

✉ **JelloMan** 7 months ago

Selected Answer: C

@mohdafiuddin covered it exactly but heres more info:

<https://cloud.google.com/bigquery/docs/loading-data-cloud-storage-avro>

upvoted 2 times

✉ **luciorifa** 8 months, 2 weeks ago

Selected Answer: C

C is the correct answer

upvoted 2 times

✉ **ericyev** 11 months, 1 week ago

C is good. <https://cloud.google.com/bigquery/external-data-cloud-storage>

upvoted 1 times

✉ **vishnukumartr** 11 months, 3 weeks ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 1 times

✉ **alaahakim** 11 months, 3 weeks ago

Selected Answer: C

The Correct Ans: C

upvoted 1 times

✉ **shawnkkk** 11 months, 3 weeks ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 1 times

✉ **mcaromit** 1 year, 6 months ago

C is the most cost effective option

upvoted 3 times

✉ **mistryminded** 1 year, 6 months ago

C seems to be correct https://cloud.google.com/bigquery/docs/loading-data-cloud-storage-avro#avro_schemas

upvoted 2 times

✉ **[Removed]** 1 year, 7 months ago

C is correct. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 2 times

✉ **drizzyd़roo** 1 year, 7 months ago

has anyone actually tried this? I don't know how to have a table point to a bucket without actually loading the data into the table.

upvoted 3 times

Question #86

You need to verify that a Google Cloud Platform service account was created at a particular time. What should you do?

- A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.
- B. Filter the Activity log to view the Configuration category. Filter the Resource type to Google Project.
- C. Filter the Activity log to view the Data Access category. Filter the Resource type to Service Account.
- D. Filter the Activity log to view the Data Access category. Filter the Resource type to Google Project.

Correct Answer: D

Community vote distribution

A (100%)

 **John_Jam** Highly Voted 2 years, 5 months ago

Correct Answer is A.

Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

upvoted 66 times

 **mlantonis** 2 years, 5 months ago

I agree A

upvoted 8 times

 **shafiqeee1** Highly Voted 2 years, 3 months ago

A - I reproduced in my project.

upvoted 23 times

 **sausalito** Most Recent 1 day, 14 hours ago

Selected Answer: A

Answer is A

upvoted 1 times

 **anolive** 6 days, 1 hour ago

Selected Answer: A

for sure, i simulated here.

upvoted 1 times

 **alleinallein** 1 month, 4 weeks ago

Selected Answer: A

A key word is configuration

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: A

Answer is A. Tried and tested

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Answer is A

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right ..

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **Sac3433** 6 months, 1 week ago

I think now C seems more probably, as activity logs are deprecated. There are 4 types of audit logs - System, data access, System, and Policy denial. So need to select and check SYSTEM LOGS to view service account creation.

upvoted 1 times

 **Sac3433** 6 months, 1 week ago

Sorry for the errors...The 4 types of logs are: Admin Activity audit logs, Data Access audit logs, System Event audit logs, and Policy Denied audit logs. <https://developers.google.com/cloud-search/docs/guides/audit-logging-manual>

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: A

Correct answer is A because you don't need data access logs and configuration counts as creation

upvoted 3 times

 **NoniGeorge** 10 months ago

Selected Answer: A

A and I love how the solution provided by the site is literally the worst one of the bunch :))))

upvoted 7 times

 **maahibhai** 10 months, 3 weeks ago

Selected Answer: A

A is the right option

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: A

A is the right option

upvoted 1 times

 **pnVino27** 11 months ago

Selected Answer: A

Answer A

upvoted 1 times

 **VidyaJi** 11 months, 1 week ago

Selected Answer: A

A is perfect

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

upvoted 1 times

Question #87

You deployed an LDAP server on Compute Engine that is reachable via TLS through port 636 using UDP. You want to make sure it is reachable by clients over that port. What should you do?

- A. Add the network tag allow-udp-636 to the VM instance running the LDAP server.
- B. Create a route called allow-udp-636 and set the next hop to be the VM instance running the LDAP server.
- C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.
- D. Add a network tag of your choice to the instance running the LDAP server. Create a firewall rule to allow egress on UDP port 636 for that network tag.

Correct Answer: C*Community vote distribution*

C (100%)

 **kishoredeena** Highly Voted 2 years, 4 months ago

Option C is the right one

upvoted 31 times

 **cloudenthalu01** Highly Voted 2 years, 4 months ago

C is correct

You tag the instances ,then create ingress firewall rules to allow udp on desired port for target-tags name applied to instances

upvoted 20 times

 **GA2022** Most Recent 2 months, 3 weeks ago

In this site I can see only 92 questions , can someone have access to 93-195 questions ?

Please mail me gajare.amar2022@gmail.com

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

C is right.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for C

upvoted 1 times

 **Fayi** 9 months, 2 weeks ago

You are developing a new web application that will be deployed on Google Cloud Platform. As part of your release cycle, you want to test updates to your application on a small portion of real user traffic. The majority of the users should still be directed towards a stable version of your application. What should you do?

- A. Deploy me application on App Engine For each update, create a new version of the same service Configure traffic splitting to send a small percentage of traffic to the new version
- B. Deploy the application on App Engine For each update, create a new service Configure traffic splitting to send a small percentage of traffic to the new service.
- C. Deploy the application on Kubernetes Engine For a new release, update the deployment to use the new version
- D. Deploy the application on Kubernetes Engine For a now release, create a new deployment for the new version Update the service e to use the now deployment.

upvoted 3 times

 **shax910** 9 months, 1 week ago

A create new version

upvoted 1 times

 **akshaychavan7** 5 months, 2 weeks ago

A without any doubt.

upvoted 1 times

 **Teyo** 9 months, 2 weeks ago

Selected Answer: C

C is correct as you can use tags and then set firewall rules for instances with such tag.

upvoted 2 times

 **Vidyaji** 11 months, 1 week ago

C IS PERFECT

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.
upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

C is Correct
upvoted 2 times

 **devopsbatch** 1 year, 5 months ago

C A tag is simply a character string added to a tags field in a resource, such as Compute Engine virtual machine (VM) instances or instance templates. A tag is not a separate resource, so you cannot create it separately. All resources with that string are considered to have that tag. Tags enable you to make firewall rules and routes applicable to specific VM instances.

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

C is correct
upvoted 1 times

 **kopper2019** 1 year, 6 months ago

C, D is for egress so for outgoing traffic
upvoted 2 times

 **[Removed]** 1 year, 7 months ago

C is correct. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.
upvoted 2 times

 **bart54** 1 year, 7 months ago

Defo C as the LDAP server will sit internally in the network protected by a Firewall, so an ingress rule will allow traffic to be routed internally to the LDAP Server.
upvoted 1 times

 **pentium2000** 1 year, 8 months ago

C, LDAP can be reached by customers means ingress traffic.
upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

C is correct
upvoted 1 times

Question #88

You need to set a budget alert for use of Compute Engine services on one of the three Google Cloud Platform projects that you manage. All three projects are linked to a single billing account. What should you do?

- A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.
- B. Verify that you are the project billing administrator. Select the associated billing account and create a budget and a custom alert.
- C. Verify that you are the project administrator. Select the associated billing account and create a budget for the appropriate project.
- D. Verify that you are project administrator. Select the associated billing account and create a budget and a custom alert.

Correct Answer: B*Community vote distribution*

A (100%)

 **kishoredeena** Highly Voted 2 years, 5 months ago

I think the answer is A, You can rely on default alert. No need for custom alert

upvoted 41 times

 **alan9999** 2 years, 1 month ago

Right its not asking to set custom alert

upvoted 1 times

 **Eshkrkrkr** 1 year, 12 months ago

One point - there is no such role as Project Billing Administrator - it should be Project Billing Manager but he can't create budgets, the only one who can - Billing Account Administrator. Nor Project Administrator exists. Very tricky question, maybe the option a wrong, hope smb will catch it on exam and pass some light on real variants.

<https://cloud.google.com/iam/docs/understanding-roles#billing-roles>

upvoted 5 times

 **Wachy** 1 year, 10 months ago

Eshkrkrkr read the question calmly. The role there is Billing Administrator. Not Project Billing Administrator.

It's more like: "Verify you are the project; billing administrator"

upvoted 11 times

 **ryumada** 3 months, 1 week ago

more like: "Verify you are the project's billing administrator"

upvoted 1 times

 **BobbyFlash** 1 year ago

I agree. If I'm not wrong, project admin doesn't have billing permissions so C and D discarded. Between A and B, option B looks like it works but we would be creating a budget and alert receiving info about billing as a whole; so A delimits billing for the project you want to get info from.

upvoted 6 times

 **Ciumela** Highly Voted 2 years, 4 months ago

A is correct, as you can set a default alert also on a single project: <https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 14 times

 **mwoodm** 2 years, 1 month ago

Agreed. Per the link included: "To create a budget for your Cloud Billing account, you must be a Billing Account Administrator on the Cloud Billing account." So that eliminates C & D. Then no need for custom alert, eliminating B. The answer is A.

upvoted 3 times

 **nickyshil** 3 months ago

why nobody is talking about "set a budget alert for use of Compute Engine services" only.. why not custom alert ?how default alert ?

upvoted 2 times

 **Charumathi** Most Recent 1 month ago

A. is correct,

You can define the scope of the budget. For example, you can scope the budget to apply to the spend in an entire Cloud Billing account, or narrow the scope to one or more projects, and/or one or more services, and/or other budget filters applicable to your Cloud Billing account.

upvoted 1 times

 **Adeem1981** 3 months, 3 weeks ago

Selected Answer: A

Billing Administrator is in between organization and projects, So he can chose and select appropriate project.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

You need billing admin not a project admin .. A is right

upvoted 1 times

 **Aksher** 5 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A

upvoted 1 times

 **oaP12** 6 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **browneyes1985** 7 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

 **pnVino27** 11 months ago

Selected Answer: A

Answer A

upvoted 2 times

 **pnVino27** 11 months ago

Answer A

upvoted 1 times

 **Vidyaji** 11 months, 1 week ago

A is perfect

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

The Correct Ans : A

upvoted 1 times

 **dhr21** 1 year, 1 month ago

It must be A, because when you are going to create alert under billing -> budget and alert then there will be one option to select project. You can select any project or all, depending on your needs. So, I think A is justifiable because it's telling to select appropriate project.

upvoted 3 times

 **Swapsss** 1 year, 2 months ago

option A.

upvoted 1 times

 **Chotebhaishab** 1 year, 2 months ago

A is the correct option

upvoted 1 times

Question #89

You are migrating a production-critical on-premises application that requires 96 vCPUs to perform its task. You want to make sure the application runs in a similar environment on GCP. What should you do?

- A. When creating the VM, use machine type n1-standard-96.
- B. When creating the VM, use Intel Skylake as the CPU platform.
- C. Create the VM using Compute Engine default settings. Use gcloud to modify the running instance to have 96 vCPUs.
- D. Start the VM using Compute Engine default settings, and adjust as you go based on Rightsizing Recommendations.

Correct Answer: B*Community vote distribution*

A (91%) 9%

 **dan80** Highly Voted 2 years, 5 months ago

A is correct - <https://cloud.google.com/compute/docs/machine-types>
upvoted 42 times

 **glam** Highly Voted 2 years ago

A. When creating the VM, use machine type n1-standard-96.
upvoted 12 times

 **Charumathi** Most Recent 1 month ago

A. is correct,
N1
* CPU types --> Skylake, Broadwell, Haswell, Sandy Bridge, and Ivy Bridge
* Architecture --> x86
* vCPUs 1 to 96
upvoted 1 times

 **Sivaprakash674** 3 months, 3 weeks ago

Go with A, It is the Correct one .
upvoted 1 times

 **Adeem1981** 3 months, 3 weeks ago

Selected Answer: B
SKYLAKE IS 96VCPU. B is correct.
upvoted 1 times

 **vividso** 2 months, 2 weeks ago

<https://cloud.google.com/compute/docs/machine-types>
N1 machine series have up to 96 vCPUs, 6.5 GB of memory per vCPU, and are available on Intel Sandy Bridge, Ivy Bridge, Haswell, Broadwell, and Skylake CPU platforms.
upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: A
Go with A
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

<https://cloud.google.com/compute/docs/general-purpose-machines> - A is correct
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for A
<https://cloud.google.com/compute/all-pricing?hl=es>
upvoted 2 times

 **LaxmanTiwari** 5 months, 2 weeks ago

Selected Answer: A
Why not update to A .. doesn't look like a rocket science to me.
upvoted 1 times

 **Lowballed** 9 months, 2 weeks ago

A is correct.

upvoted 2 times

 **DhrubaD** 10 months ago

The question says 96vCPU required, it never said 360GB (or nearest) memory is required. n1-standard-96 comes with 96vCPU and 360GB mem, but Skylake comes with n1-standard-96, n1-highmem-96 and n1-highcpu-96, all three are for 96vCPU but different memory. So I believe Correct answer should be: B

upvoted 1 times

 **pgb54** 8 months, 2 weeks ago

I think the reason A is better is that the question specifically calls out 96 cores as a requirement, so picking Skylake by itself (option B) doesn't necessarily meet that requirement.

upvoted 2 times

 **andresbp** 10 months, 3 weeks ago

Selected Answer: A

N1 offer up 96 CPU's, then A is correct

upvoted 3 times

 **pnVino27** 11 months ago

Selected Answer: A

ANSWER A

upvoted 2 times

 **Vidyaji** 11 months, 1 week ago

Selected Answer: A

A perfect

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

A. When creating the VM, use machine type n1-standard-96.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

A is Correct

upvoted 2 times

 **Litan** 1 year, 1 month ago

A is correct

upvoted 2 times

Question #90

You want to configure a solution for archiving data in a Cloud Storage bucket. The solution must be cost-effective. Data with multiple versions should be archived after 30 days. Previous versions are accessed once a month for reporting. This archive data is also occasionally updated at month-end. What should you do?

- A. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Coldline Storage.
- B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.
- C. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Coldline Storage.
- D. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Nearline Storage.

Correct Answer: B

Reference:

<https://cloud.google.com/storage/docs/managing-lifecycles>*Community vote distribution*

B (100%)

 **neelesh88** Highly Voted 2 years, 4 months ago

B is correct

upvoted 31 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer (B):

NumberOfNewerVersions

The NumberOfNewerVersions condition is typically only used in conjunction with Object Versioning. If the value of this condition is set to N, an object version satisfies the condition when there are at least N versions (including the live version) newer than it. For a live object version, the number of newer versions is considered to be 0. For the most recent noncurrent version, the number of newer versions is 1 (or 0 if there is no live object version), and so on.

Important: When specifying this condition in a .json configuration file, you must use numNewerVersions instead of NumberOfNewerVersions.

<https://cloud.google.com/storage/docs/lifecycle#numberofnewerversions>

upvoted 21 times

 **kadc** Most Recent 1 month, 3 weeks ago

Selected Answer: B

B should be correct:

Nearline has min storage of 30 days, while Coldline has 90 days.

Since "archive data is also occasionally updated at month-end", updating object before min storage period is allowed but causes early deletion fees as if the object was stored for the min duration, so using Coldline will always charge for 90 days and not likely to save cost.

<https://cloud.google.com/storage/pricing#early-delete>

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right and straight forward.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

regional storage after 30 days to Nearline Storage option is trick you :)

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

 **Vidyaji** 11 months, 1 week ago

B is perfect

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 1 times

 **Chotebhaisahab** 1 year, 2 months ago

Agree B is the correct

upvoted 2 times

 **Rahul183** 1 year, 2 months ago

B- <https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

 **jcols** 1 year, 5 months ago

B is correct. C and D don't mention newer object versions. Then, it's Coldline vs. Nearline, the best option is Nearline because:

- * "Previous versions are accessed once a month for reporting"
- * "Nearline Storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline Storage is a great choice." -- <https://cloud.google.com/storage/docs/storage-classes#nearline>

upvoted 2 times

 **mcaromit** 1 year, 6 months ago

B is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

B is correct

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 2 times

 **arsav** 1 year, 8 months ago

"Previous versions are accessed once a month for reporting"

only option C and D says, that you are storing the data in "regional" storage and then trying to move from either "coldline" or "nearline"

so, i will go with option D. to hold the data which is accessed once a month in "regional" storage and then after 30 days i archive it and move to "nearline"

also note, question says you occasionally update the archive data. so near-line would be right choice.

upvoted 3 times

Question #91

Topic 1

Your company's infrastructure is on-premises, but all machines are running at maximum capacity. You want to burst to Google Cloud. The workloads on Google Cloud must be able to directly communicate to the workloads on-premises using a private IP range. What should you do?

- A. In Google Cloud, configure the VPC as a host for Shared VPC.
- B. In Google Cloud, configure the VPC for VPC Network Peering.
- C. Create bastion hosts both in your on-premises environment and on Google Cloud. Configure both as proxy servers using their public IP addresses.
- D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

Correct Answer: D*Community vote distribution*

D (100%)

 **SIX** Highly Voted 2 years, 5 months ago

I believe D is the right answer

upvoted 53 times

 **dan80** 2 years, 5 months ago

B is correct - <https://cloud.google.com/solutions/best-practices-vpc-design>. this answer also on all machines are running at maximum capacity.

upvoted 1 times

 **JustLearning** 2 years, 5 months ago

vpc network peering does not connect to on-prem. Cloud VPN is the correct solution.

<https://cloud.google.com/vpn/docs/concepts/overview>

upvoted 22 times

 **mlantonis** 2 years, 5 months ago

You need VPN, so D is the correct. VPC network peering is between VPCs.

upvoted 14 times

 **xharf** 3 months, 1 week ago

"Google Cloud VPC Network Peering allows internal IP address connectivity across two Virtual Private Cloud (VPC) networks regardless of whether they belong to the same project or the same organization."
<https://cloud.google.com/vpc/docs/vpc-peering>

while

"Cloud Interconnect provides low latency, high availability connections that enable you to reliably transfer data between your on-premises and Google Cloud Virtual Private Cloud (VPC) networks."

<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview>

and

"HA VPN is a high-availability (HA) Cloud VPN solution that lets you securely connect your on-premises network to your VPC network through an IPsec VPN connection in a single region."

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

so, cloud vpn is the best answer for the question requirement

upvoted 1 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (D):

Access internal IPs directly

Your VPC network's internal (RFC 1918) IP addresses are directly accessible from your on-premises network with peering, no NAT device or VPN tunnel required.

Hybrid made easy

Today's business climate demands flexibility. Connecting your on-premises resources to your cloud resources seamlessly, with minimum latency or interruption, is a business-critical requirement. The speed and reliability of Cloud Interconnect lets you extend your organization's data center network into Google Cloud, simply and easily, while options such as Cloud VPN provide flexibility for all your workloads. This unlocks the potential of hybrid app development and all the benefits the cloud has to offer.

In the graphic below: What GCP Connection is right for you? shows clearly what is the method for extend your on premise network (IP Private communication).

What GCP Connection is right for you?

<https://cloud.google.com/hybrid-connectivity>

upvoted 28 times

 **VaneA** Most Recent ⓘ 1 month, 1 week ago

Selected Answer: D

It is the answer

upvoted 1 times

 **RanjithK** 4 months, 1 week ago

Selected Answer: D

Go with D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is right answer

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Cloud VPN is way to establish connection between on prem to cloud . D is correct.

upvoted 2 times

 **Uqqasha** 10 months ago

Selected Answer: D

D is the right answer.

upvoted 2 times

 **ARVII** 10 months, 3 weeks ago

On-premise -> GCP

There are 2 ways

1. Cloud VPN

2. Interconnect

Since we have VPN as an option, others is not recommended

upvoted 2 times

 **Vidyaji** 11 months, 1 week ago

D is perfect

upvoted 1 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Ans : D

upvoted 1 times

 **Chotebhaisahab** 1 year, 2 months ago

yes VPN is needed. D is correct

upvoted 2 times

 **nenoAZ** 1 year, 2 months ago

D is correct. 100%.

upvoted 1 times

 **mcaromit** 1 year, 5 months ago

D is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

Not A and B, because for network peering, shared VPC you need to either be in the same project or same organization. Here we are in on-premise and cloud.

Maybe C, but why do you need an extra server on both GCP and On-premise for a proxy?. We already have solutions for proxy inside GCP.
D is correct because, we generally connect GCP to on-premise using VPN

upvoted 4 times

 **[Removed]** 1 year, 7 months ago

D is correct. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

upvoted 2 times

 **Hi2ALL** 1 year, 8 months ago

D is correct since VPC peering only works on GCP cloud network

upvoted 1 times

Question #92

You want to select and configure a solution for storing and archiving data on Google Cloud Platform. You need to support compliance objectives for data from one geographic location. This data is archived after 30 days and needs to be accessed annually. What should you do?

- A. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.
- B. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- C. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

Correct Answer: D

Community vote distribution

D (100%)

 **SIX**  2 years, 5 months ago

D

Google Cloud Coldline is a new cold-tier storage for archival data with access frequency of less than once per year. Unlike other cold storage options, Nearline has no delays prior to data access, so now it is the leading solution among competitors.

upvoted 42 times

 **dan80** 2 years, 5 months ago

C is correct - This data is archived after 30 days - Nearline Storage 30 days , Coldline Storage 90 days
<https://cloud.google.com/storage/docs/storage-classes>

upvoted 11 times

 **lxgywil** 1 year, 6 months ago

The answer is D.

The main thing here is how often the data is retrieved. The question is saying that data needs to be accessed annually - i.e. once a year. Therefore, you should choose Coldline Storage, as it implies less frequent access than Nearline. (Archival Storage would fit even better but there's no such option)

From the link you provided:

"Nearline Storage is ideal for data you plan to read or modify on average once per month or less."

and

"Coldline Storage is ideal for data you plan to read or modify at most once a quarter. "

upvoted 12 times

 **mlantonis** 2 years, 5 months ago

dan80 is right

upvoted 1 times

 **nithinpb180** 2 years, 4 months ago

That is the minimum storage duration. I would go with D. Coldline storage is more suitable for infrequent data access.

upvoted 7 times

 **JustLearning** 2 years, 5 months ago

D is correct. Coldline is a better choice.

upvoted 9 times

 **ESP_SAP**  2 years, 2 months ago

Correct Answer is (D):

<https://cloud.google.com/storage/docs/storage-classes>

Nearline Storage

Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline Storage is a better choice than Standard Storage in scenarios where slightly lower availability, a 30-day minimum storage duration, and costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Nearline Storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline Storage is a great choice.

Nearline Storage is also appropriate for data backup, long-tail multimedia content, and data archiving. Note, however, that for data accessed less frequently than once a quarter, Coldline Storage or Archive Storage are more cost-effective, as they offer lower storage costs.
<https://cloud.google.com/storage/docs/storage-classes#nearline>

upvoted 13 times

 **ESP_SAP** 2 years, 2 months ago

CORRECTION.

Correct Answer is (D):

The Real description is about Coldline storage Class:

Coldline Storage

Coldline Storage is a very-low-cost, highly durable storage service for storing infrequently accessed data. Coldline Storage is a better choice than Standard Storage or Nearline Storage in scenarios where slightly lower availability, a 90-day minimum storage duration, and higher costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Coldline Storage is ideal for data you plan to read or modify at most once a quarter. Note, however, that for data being kept entirely for backup or archiving purposes, Archive Storage is more cost-effective, as it offers the lowest storage costs.

<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 10 times

 **biren111** Most Recent 3 weeks, 1 day ago

As in question it is asking for "one geographic location". So multi region options A & B is eliminated. And Between C & D "D is correct" as data will be accessed once a year.

upvoted 1 times

 **gcpBeginner** 1 month, 1 week ago

its C because archive data is 30 days and Nearline storage support that. <https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is for sure.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D

"... and needs to be accessed annually"

Cold line is the better choice.

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: D

D is the right answer.

upvoted 1 times

 **Uqqasha** 10 months ago

Selected Answer: D

D is the right answer.

upvoted 1 times

 **Vidyaji** 11 months, 1 week ago

D is perfect

upvoted 2 times

 **vishnukumartr** 11 months, 3 weeks ago

D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Ans: D

upvoted 1 times

 **shawnkkk** 11 months, 3 weeks ago

D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

upvoted 1 times

 **ankatsu2010** 1 year, 1 month ago

"One geographic location" and "needs to be accessed annually"

-> "Archive Storage" class is most appropriate, but in this question, "D" is the best choice.

upvoted 2 times

 **vikram___** 1 year, 2 months ago

Ans D , even though coldline storage is for data being accessed in a quarter. As archive storage is not mentioned the next feasible option is coldline hence , option D

<https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

 **Chotebhaisahab** 1 year, 2 months ago

Needed regional storage for one geographic location cold line for annual access of data. D is correct.

upvoted 1 times

 **Israel** 1 year, 3 months ago

D is the right answer

upvoted 1 times

Question #93

Your company uses BigQuery for data warehousing. Over time, many different business units in your company have created 1000+ datasets across hundreds of projects. Your CIO wants you to examine all datasets to find tables that contain an employee_ssn column. You want to minimize effort in performing this task.

What should you do?

- A. Go to Data Catalog and search for employee_ssn in the search box.
- B. Write a shell script that uses the bq command line tool to loop through all the projects in your organization.
- C. Write a script that loops through all the projects in your organization and runs a query on INFORMATION_SCHEMA.COLUMNS view to find the employee_ssn column.
- D. Write a Cloud Dataflow job that loops through all the projects in your organization and runs a query on INFORMATION_SCHEMA.COLUMNS view to find employee_ssn column.

Correct Answer: D

Community vote distribution

A (100%)

 **poogcp** Highly Voted 2 years, 5 months ago

Its A.

upvoted 34 times

 **filco72** Highly Voted 2 years, 3 months ago

Correct is A.

I tested on my account following this procedure: <https://cloud.google.com/bigquery/docs/quickstarts/quickstart-web-ui?authuser=4>
I created a data set and through Data Catalog I easily and effortlessly searched for the column name "gender"

upvoted 25 times

 **CELx_12** Most Recent 4 days, 6 hours ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my cesexylia@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 2 times

 **Unicorn_12** 5 days, 19 hours ago

I am taking the exam soon, please someone share the full questions on roms.jiw@gmail.com

upvoted 1 times

 **Rajehs** 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my rajeshkannafeb6@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 1 times

 **rixson** 4 weeks, 1 day ago

I'm taking the exam this week. I can't access all questions. Can someone send me the full version please. ricsoncoc@gmail.com. Thanks

upvoted 1 times

 **Charumathi** 1 month ago

A is the correct answer, Data Catalog can be used to search the column with keyword:value pair,

Filter your search by adding a keyword:value to your search terms in the search box:

Keyword Description

name: Match data asset name

***column: Match column name or nested column name

description: Match table description

upvoted 3 times

 **CELx_12** 4 days, 6 hours ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my cesexylia@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 1 times

 **mdvp** 1 month, 3 weeks ago

ITS A COLUMN WE ARE SEARCHING FOR AND I DONT SEE DATA CATALOGUE CAN SEARCH FOR COLUMN

upvoted 1 times



👤 **ashit44244** 1 month, 3 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 92 questions only, if anyone has the entire questions please share them with my reachme.ashit@gmail.com address. Thanks in advance!

upvoted 1 times

✉️ 🚩 **Cornholio_LMC** 1 month, 2 weeks ago

but replied to question 93

upvoted 6 times

✉️ 🚩 **JackyWorld** 2 months, 1 week ago

Selected Answer: A

Definitely, A is Right.

upvoted 1 times

✉️ 🚩 **patashish** 4 months, 1 week ago

D is right

upvoted 1 times

✉️ 🚩 **AzureDP900** 4 months, 2 weeks ago

D is right

upvoted 2 times

✉️ 🚩 **AzureDP900** 4 months, 2 weeks ago

I misread the question, A is right answer.

upvoted 1 times

✉️ 🚩 **haroldbenites** 5 months, 1 week ago

Go for A.

upvoted 1 times

✉️ 🚩 **hiranfilho** 6 months, 1 week ago

Selected Answer: A

use data Catalog for search the column

upvoted 1 times

✉️ 🚩 **Mr_Tiwariji** 7 months, 2 weeks ago

Selected Answer: A

IT IS THE SIMPLEST WAY

upvoted 2 times

✉️ 🚩 **[Removed]** 7 months, 4 weeks ago

Selected Answer: A

<https://cloud.google.com/data-catalog/docs/how-to/search>

upvoted 1 times

✉️ 🚩 **icecoolcat** 9 months, 1 week ago

Its A.

upvoted 1 times

Question #94

You create a Deployment with 2 replicas in a Google Kubernetes Engine cluster that has a single preemptible node pool. After a few minutes, you use kubectl to examine the status of your Pod and observe that one of them is still in Pending status:

```
$ kubectl get pods -l app=myapp
NAME           READY   STATUS    RESTART   AGE
myapp-deployment-58ddbbb995-1p86m   0/1     Pending   0          9m
myapp-deployment-58ddbbb995-qjpkqg  1/1     Running   0          9m
```

What is the most likely cause?

- A. The pending Pod's resource requests are too large to fit on a single node of the cluster.
- B. Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.
- C. The node pool is configured with a service account that does not have permission to pull the container image used by the pending Pod.
- D. The pending Pod was originally scheduled on a node that has been preempted between the creation of the Deployment and your verification of the Pods' status. It is currently being rescheduled on a new node.

Correct Answer: B*Community vote distribution*

B (58%) D (42%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Reasons for a Pod Status Pending:

Troubleshooting Reason #1: Not enough CPU

Troubleshooting Reason #2: Not enough memory

Troubleshooting Reason #3: Not enough CPU and memory

<https://managedkube.com/kubernetes/k8sbot/troubleshooting/pending/pod/2019/02/22/pending-pod.html>

upvoted 42 times

 **SSPC** 2 years, 2 months ago

I agree with you. The correct answer is B

upvoted 5 times

 **[Removed]** 2 years, 1 month ago

D gives you the reason why the resource could not be available as it was preempted

upvoted 3 times

 **Finger41** 1 year, 5 months ago

Its in a deployment, the pod will be recreated. There is insufficient resources in the node, not because its preemptible but because there is no memory/cpu.....

upvoted 2 times

 **tavva_prudhvi** 1 year, 7 months ago

If it was preempted, then it has to be restarted right? then it will show its failing not pending, check the articles mentioned by ESP_SAP

upvoted 3 times

 **Linus11** 1 year, 6 months ago

The real crux of this question is the mention about "Pre-emptible Node pool". That needs to take into consideration while determining the answer. If we choose B, then the importance of "Pre-emptible node pool" is not there. Whether the node pool is pre-emptible or not, resource scarcity can lead to pending pods.

When we consider the mention of "Pre-emptible Node Pool", then the answer is obviously D. if a pre-emptible Node gets pre-empted there will be a delay in cluster to sync it.

Answer is D.

upvoted 12 times

 **alexgrig** 1 year, 1 month ago

Questions says "Single Node" at that case the second pod can't be in running state.

upvoted 3 times

 **MidhunJose** 9 months, 3 weeks ago

It says a single node pool, not a single node. Meaning there can be multiple nodes, right?

upvoted 7 times

 **brvinod** 9 months ago

A node can have multiple pods. So that is not a problem.

upvoted 1 times

 **Finger41** 1 year, 5 months ago

This is to throw you off, when there is insufficient resources for a Pod to stand up, then the status will equal pending :
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#:~:text=If%20a%20Pod%20is%20stuck,be%20scheduled%20onto%20a%20node.&text=You%20don't%20have%20enough,new%20nodes%20to%20your%20cluster.>

upvoted 2 times

 **brvinod** 9 months ago

Pre-emptible would have been an issue if the cluster had more than one node. The question clearly states that it is a single node cluster. That means if that single VM was pre-empted, neither of the pods should have been running. Since one pod is running, that means that (the only) VM is running. So, the reason the second pod is still pending because the VM is not having enough resources to run both the pods. Hence B.

upvoted 8 times

 **mplibunao** 5 months ago

Actually the question stated "single preemptible node pool" and not "single node" so it's possible that there are multiple nodes and one of the node on which the pod was scheduled on was preempted

upvoted 3 times

 **cloudenthu01** Highly Voted  2 years, 4 months ago

D is correct as the node on which pod was scheduled to run was preempted & now this pod is scheduled to run on different preemptible node from the node-pool

upvoted 16 times

 **myuniqueusername** 1 year, 1 month ago

Incorrect. There is a single preemptible instance, if it was preempted then both pods would show as 'Pending'. B is correct.

upvoted 4 times

 **obeythefist** 7 months, 3 weeks ago

> There is a single preemptible instance

Where does it say that? It doesn't. Don't make things up. There's a single pre-emptible node pool. A single pool is not the same as a single node.

upvoted 6 times

 **nonsense** Most Recent  2 weeks, 1 day ago

Selected Answer: D

Pre-emptible Node pool keyword.

upvoted 1 times

 **darcal95** 1 month, 3 weeks ago

B is correct, but why not A?

upvoted 2 times

 **sandipk91** 2 months ago

Selected Answer: D

Option D , I have practically experienced this before on live envs whenever a node gets tainted and pods are supposed to be rescheduled it goes to pending state

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 2 times

 **Passerofexams** 4 months, 1 week ago

After doing a bit of research within the Google Docs, I agree.

There is specific logic that can and should be put in place to prevent the scenario mentioned in “D” from happening.

<https://cloud.google.com/blog/products/containers-kubernetes/cutting-costs-with-google-kubernetes-engine-using-the-cluster-autoscaler-and-preemptible-vms>

upvoted 1 times

 **theBestStudent** 1 month ago

if you meant of using <https://github.com/estafette/estafette-gke-preemptible-killer>, then that is an out of the box solution. So can not be assumed as part of the possible components in the answers.

upvoted 1 times

 **lamboz** 4 months, 3 weeks ago

I think:

- Not enough CPU & RAM --> status: OMMKilled
- Cannot pull image or syntax errors --> status: Crash...
- Cannot mount volumes --> status: ContainerCreating or Evicted

- Cannot scheduled because Node not ready --> status: Pending

If a preemptible node not ready, pod cannot start, so pod's status is Pending.
So, the answer is D

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: B

Correct answer is B. If the node was pre-empted the other pod would have been terminated

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for B

upvoted 1 times

 **salehm** 6 months, 1 week ago

Correct Answer is (D):

B options is not correct because if the node resources are not enough kubernetes scheduler will reschedule the pod deployment to an another cluster node and pod would not enter pending status except if all cluster nodes resources are fully consumed.

Option D makes more since

upvoted 1 times

 **hiranfilho** 6 months, 1 week ago

Selected Answer: B

<https://managedkube.com/kubernetes/k8sbot/troubleshooting/pending/pod/2019/02/22/pending-pod.html>

upvoted 1 times

 **zaxma** 7 months, 1 week ago

Selected Answer: D

I vote 'D'

B says: Too many 'Pods' are already running in the cluster, and there are not enough resources left to schedule the pending Pod.

There are only 2 pods and setup as 2 pods, where is the 'too many' coming from??

upvoted 3 times

 **wjtb** 5 months, 1 week ago

look at the command, the pods are filtered on label: app=myapp, this doesn't mean all the pods are listed (only the pods with that label)

upvoted 2 times

 **Mr_Tiwariji** 7 months, 2 weeks ago

Selected Answer: D

D is correct as the node on which pod was scheduled to run was preempted & now this pod is scheduled to run on different preemptible node from the node-pool

upvoted 1 times

 **ryzior** 7 months, 4 weeks ago

Selected Answer: B

age is important

both pods have been scheduled the same time (so if the pending one got pending due to restart of a preempted one it should be rescheduled to another node and the age would be shorter right?)

in this case it is pending because there is no such node which fits the pod capacity requirement so either the pod (assuming there are only 2 nodes) is being restarted/preempted and the other running node is unable to host the pod due to low capacity

to be honest this is again both options are correct but I'm rather with option B looking at the result not the cause of the issue (pending because unable to schedule not because anything got restarted)

upvoted 3 times

 **special1** 8 months ago

Selected Answer: B

Correct Answer is (B)

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

B is the correct answer, just because the node is preemptible doesn't mean that it was removed while there's a pod already running.

upvoted 1 times

 **raimangsxr** 8 months, 2 weeks ago

B is correct

upvoted 1 times

Question #95

You want to find out when users were added to Cloud Spanner Identity Access Management (IAM) roles on your Google Cloud Platform (GCP) project. What should you do in the GCP Console?

- A. Open the Cloud Spanner console to review configurations.
- B. Open the IAM & admin console to review IAM policies for Cloud Spanner roles.
- C. Go to the Stackdriver Monitoring console and review information for Cloud Spanner.
- D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.

Correct Answer: B*Community vote distribution*

D (100%)

 **samvegas** Highly Voted 2 years, 3 months ago

Answer = D, I have simple rule; if metrics then Monitoring, if Auditing then Logging.
upvoted 58 times

 **Meix** Highly Voted 2 years, 5 months ago

I think the answer is D
upvoted 44 times

 **Anand2608** 1 year, 12 months ago

As per the Cloud Audit logs documentation.
upvoted 3 times

 **klyaal** Most Recent 1 week, 4 days ago

I'm taking the exam next week. I can't access all questions. Can someone please send me the full version to klacewicz@gmail.com
Thanks a lot!
upvoted 1 times

 **Rajehs** 1 week, 5 days ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my rajeshkannafeb6@gmail.com address. Thanks in advance. I would be forever grateful.
upvoted 1 times

 **Charumathi** 1 month ago

D is correct, Activity logs captures the time when the users were given the IAM roles for Cloud Spanner
upvoted 1 times

 **Priyanka109** 1 month ago

You need to see when not what so D.
upvoted 1 times

 **TinkalS** 1 month ago

ANS- D
upvoted 1 times

 **king2658** 1 month, 2 weeks ago

Answer should be D
upvoted 1 times

 **ATANGA** 3 months ago

WhatsApp : +1(956)-520-4006 to obtain PMP, CISM, CCNA, CEH, PRINCE2, CISCO, ISTQB, PRINCE2, AWS/Azure/Sales force/ITIL

Foundation/EC-COUNCIL...

Get Certified with 100% pass guarantee. Pay after exam.

all CISCO, ISACA & EC-COUNCIL certifications

For the Below certificates we offer 100% pass guarantee:

1. AWS Certification
2. Sales force
3. Scrum Master
4. Oracle Certification: OCA, OCP
5. Cisco Certification: CCNA, CCNP, CCIE
6. ITIL Foundation & Intermediate
7. Prince 2 Foundation and Practitioner
8. VMWARE Certification
9. Check Point Certification
10. EC-COUNCIL Certification (CEH V-9, CCISO, CND)

- 11. Cloud Certification
- 12. IBM Certification
- 13. HP Certification
- 14. Citrix Certification
- 15. Juniper certification
- 16. Azure
- 17. Skype 70-333/34
- 18. PMI (PMP/CAPM/ACP/PBA)
- 19. ISTQB
- 20. SAP
- 21. ISACA (CISA, CISM, CRISC, CGEIT, COBIT)

PAYMENT ONLY AFTER CERTIFICATION AND RESULT CONFIRMATION.

WhatsApp : +1(956)-520-4006

upvoted 1 times

 **ikychughtai** 3 months, 1 week ago

Selected Answer: D

bow before my divine intellect

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

There is similar practice question in Tutorials Dojo .. D is correct. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.

upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D

upvoted 1 times

 **PrithvirajLohar619** 5 months, 1 week ago

Answer D: use Logging

upvoted 2 times

 **hiranfilho** 6 months, 1 week ago

Selected Answer: D

We use Logging to know who when where what

<https://cloud.google.com/monitoring/audit-logging>

upvoted 2 times

 **Priyankahere** 9 months, 3 weeks ago

This was there in exam, go with community answers.

upvoted 5 times

 **dishum** 10 months, 1 week ago

Answer is B B B B

The question is "What should you do" ? i think it is about the responsibility of checking security, policies and best practices, whether rightly assigned the roles.

The question is not about the timestamp, my view.

upvoted 1 times

 **Rakesh_SAP** 9 months, 3 weeks ago

The question is - "When users were added" - to check "When" need to check logs

upvoted 1 times

 **pnVino27** 11 months ago

Selected Answer: D

Answer is D

upvoted 2 times

Question #96

Your company implemented BigQuery as an enterprise data warehouse. Users from multiple business units run queries on this data warehouse. However, you notice that query costs for BigQuery are very high, and you need to control costs. Which two methods should you use? (Choose two.)

- A. Split the users from business units to multiple projects.
- B. Apply a user- or project-level custom query quota for BigQuery data warehouse.
- C. Create separate copies of your BigQuery data warehouse for each business unit.
- D. Split your BigQuery data warehouse into multiple data warehouses for each business unit.
- E. Change your BigQuery query model from on-demand to flat rate. Apply the appropriate number of slots to each Project.

Correct Answer: BE*Community vote distribution*

BE (58%)

D (42%)

 **CPBach** Highly Voted 2 years, 4 months ago

I'd say B and E. So either you do B or E to reduce costs.
upvoted 40 times

 **Jignesh_Gamda** Highly Voted 2 years ago

B & E
Refer below link - first of all you can define quotas on project or user level and 2nd one is you can change from on demand to flat rate model and define the parameters based on your requirement ---

<https://cloud.google.com/bigquery/docs/custom-quotas>
https://cloud.google.com/bigquery/pricing#flat_rate_pricing
upvoted 19 times

 **Ashishyadav420** Most Recent 1 day, 16 hours ago

I am able to access 96 questions only, if anyone has the entire questions please share on ashish.uk2022@gmail.com
upvoted 1 times

 **jyotimishra1** 1 day, 21 hours ago

I am preparing for GCP Associate engineer exam.
As I was able to access 96 questions only, if anyone has the entire questions please share them to my jyotimishra.jec@gmail.com address.
upvoted 1 times

 **rjav3** 3 days, 17 hours ago

Can someone with contributor access, send me the remaining questions to this email : rjav333@gmail.com Thanks..
upvoted 1 times

 **fai06** 2 weeks ago

I am able to access 96 questions only. If anyone has remaining questions. Can you please send it to fairojm@gmail.com
upvoted 1 times

 **Rajehs** 2 weeks ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my rajeshkannafeb6@gmail.com address. Thanks in advance. I would be forever grateful.
upvoted 1 times

 **Shailender24** 2 weeks ago

If anyone has contributor access , please email the full list of questions to shailender.ravindranath@gmail.com. Not able to access beyond 96 questions
upvoted 1 times

 **Abc1235555** 2 weeks, 1 day ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my etcdelhi01@gmail.com address. Thanks in advance. I would be forever grateful.
upvoted 1 times

 **Untamables** 2 weeks, 3 days ago

Selected Answer: BE
<https://cloud.google.com/bigquery/docs/custom-quotas>
<https://cloud.google.com/blog/products/data-analytics/cost-optimization-best-practices-for-bigquery>
upvoted 1 times

 **gcpcloud** 3 weeks, 2 days ago

Hello everyone! I am preparing for the ACE examination and as I was able to access 96 questions only, if anyone has all questions and can send them at vinayk2455@gmail.com I would highly appreciate it. Thanks in advance!

upvoted 1 times

 **alexandra29** 3 weeks, 3 days ago

¡Hola a todos! Me estoy preparando para el examen de ingeniero asociado de Google Cloud y solo pude acceder a 96 preguntas. Si alguien tiene todas las preguntas, ¿puede compartirlas en alexacolombia2013@gmail.com? Realmente lo aprecio. ¡Gracias por adelantado!

upvoted 1 times

 **Robertolo** 3 weeks, 5 days ago

Selected Answer: BE

Quotas and flat rate are ways to cut spending.

Splitting the data warehouse does not reduce the costs. Instead, it makes cross-business unit queries rather difficult

upvoted 1 times

 **Akai007** 4 weeks ago

Hello everyone! I am preparing for the Google Cloud Associate Engineer examination and I was able to access 96 questions only. If anyone has all the questions, can you please share them at detconan915@gmail.com. Really appreciate it. Thanks in advance!

upvoted 1 times

 **itsphase** 1 month ago

Hello everyone! I am preparing for the ACE examination and as I was able to access 96 questions only, if anyone has all questions and can send them at aceexaminationquestions@gmail.com I would highly appreciate it. Thanks in advance!

upvoted 1 times

 **Neha_1** 1 month ago

I am preparing for the GCP-ACE exam, I was able to access 96 questions only, if anyone has the entire questions please share them with my neharika1596@gmail.com address. Thanks in advance!

upvoted 1 times

 **Abc1235555** 2 weeks, 1 day ago

If you received the docs can you share them with my etcdelhi01@gmail.com address. Thanks in advance. I would be forever grateful.

upvoted 1 times

 **TinkalS** 1 month ago

Ans - D,E

upvoted 1 times

Question #97

You are building a product on top of Google Kubernetes Engine (GKE). You have a single GKE cluster. For each of your customers, a Pod is running in that cluster, and your customers can run arbitrary code inside their Pod. You want to maximize the isolation between your customers' Pods. What should you do?

- A. Use Binary Authorization and whitelist only the container images used by your customers' Pods.
- B. Use the Container Analysis API to detect vulnerabilities in the containers used by your customers' Pods.
- C. Create a GKE node pool with a sandbox type configured to gvisor. Add the parameter runtimeClassName: gvisor to the specification of your customers' Pods.
- D. Use the cos_containerd image for your GKE nodes. Add a nodeSelector with the value cloud.google.com/gke-os-distribution: cos_containerd to the specification of your customers' Pods.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/sandbox/>*Community vote distribution*

C (100%)

akshaychavan7 Highly Voted 5 months, 2 weeks ago

Let me be honest, I did not have any clue to answer this question. However, I spotted the keyword, 'isolation', from the question and a keyword, 'sandbox' from the answers and guessed the answer which turned out to be correct.

So, yes it is C!

upvoted 7 times

mattcl Most Recent 2 days, 12 hours ago

GKE Sandbox <https://cloud.google.com/kubernetes-engine/docs/concepts/sandbox-pods>

upvoted 1 times

theBestStudent 3 months, 1 week ago**Selected Answer: C**

As it has been mentioned already: <https://cloud.google.com/kubernetes-engine/docs/how-to/sandbox-pods?hl=en>

https://cloud.google.com/kubernetes-engine/docs/how-to/sandbox-pods?hl=en#working_with

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

gVisor is the way to isolate. Those who already preparing for CKS can answer this question without even thinking further. C is right
upvoted 2 times

haroldbenites 5 months, 1 week ago

Go for C

upvoted 1 times

Sac3433 6 months, 1 week ago

Correct answer is C: You can enable GKE Sandbox on your cluster to isolate untrusted workloads in sandboxes on the node. GKE Sandbox is built using gVisor, an open source project: https://cloud.google.com/kubernetes-engine/docs/concepts/security-overview?hl=en#protecting_nodes_from_untrusted_workloads

upvoted 3 times

PAUGURU 6 months, 1 week ago**Selected Answer: C**

https://cloud.google.com/kubernetes-engine/docs/concepts/security-overview?hl=en#protecting_nodes_from_untrusted_workloads

upvoted 2 times

Question #98

Your customer has implemented a solution that uses Cloud Spanner and notices some read latency-related performance issues on one table. This table is accessed only by their users using a primary key. The table schema is shown below.

```
CREATE TABLE Persons (
    person_id INT64 NOT NULL,      // sequential number based on number of registration
    account_creation_date DATE,   // system date
    birthdate DATE,               // customer birthdate
    firstname STRING (255),       // first name
    lastname STRING (255),        // last name
    profile_picture BYTES (255)  // profile picture
) PRIMARY KEY (person_id)
```

You want to resolve the issue. What should you do?

- A. Remove the profile_picture field from the table.
- B. Add a secondary index on the person_id column.
- C. Change the primary key to not have monotonically increasing values.

- D. Create a secondary index using the following Data Definition Language (DDL):

```
CREATE INDEX person_id_ix
ON Persons (
    person_id,
    firstname,
    lastname
) STORING (
    profile_picture
)
```

Correct Answer: D

Community vote distribution

C (54%)

D (46%)

 **BenKenGo6**  2 months, 1 week ago

Selected Answer: D

Create a secondary index using the following Data Definition.
If we watch the next video, he talks about a change to monotonically when we insert rows.
Finally when we talk about read and we have a performance issues, we must create a index.
<https://www.youtube.com/watch?v=r6uj0HMNQNQ>
upvoted 5 times

 **mattcl**  2 days, 12 hours ago

C <https://cloud.google.com/spanner/docs/schema-design#primary-key-prevent-hotspots>
upvoted 1 times

 **Sozan** 2 days, 19 hours ago

Selected Answer: C

C is the right answer. Why? "This table is accessed only by their users using a primary key." So adding additional indexes on firstname and lastname won't help.
upvoted 1 times

 **Aninina** 1 week, 1 day ago

Selected Answer: C

you should be careful when choosing a primary key to not accidentally create hotspots in your database. One cause of hotspots is having a column whose value monotonically increases as the first key part, because this results in all inserts occurring at the end of your key space. This pattern is undesirable because Spanner divides data among servers by key ranges, which means all your inserts will be directed at a single server that will end up doing all the work.
upvoted 1 times

 **Erie** 2 weeks, 5 days ago

Selected Answer: C

C: this is to avoid having hotspots. If the PK is monotonic, then there is a higher chance of requests being routed to the same spanner server and thus overloading it.
D: Primary key is an index already. The question states that the users only accesses this table by PK.
upvoted 1 times

 **leliodesouza** 4 weeks, 1 day ago

Selected Answer: D

The answer is D.

upvoted 1 times

  **theBestStudent** 4 weeks, 1 day ago**Selected Answer: C**

Answer C: Schema design best practice #1: Do not choose a column whose value monotonically increases or decreases as the first key part for a high write rate table.

<https://cloud.google.com/spanner/docs/schema-design#postgresql>

upvoted 2 times

  **alleinallein** 1 month, 4 weeks ago**Selected Answer: C**

IMO C

<https://cloud.google.com/spanner/docs/schema-design>

upvoted 2 times

  **shreykul** 2 months agoAnswer should be c. <https://cloud.google.com/spanner/docs/schema-design>

upvoted 1 times

Question #99

Your finance team wants to view the billing report for your projects. You want to make sure that the finance team does not get additional permissions to the project. What should you do?

- A. Add the group for the finance team to roles/billing user role.
- B. Add the group for the finance team to roles/billing admin role.
- C. Add the group for the finance team to roles/billing viewer role.
- D. Add the group for the finance team to roles/billing project/Manager role.

Correct Answer: C*Community vote distribution*

C (100%)

- ✉ **miniso8153** Highly Voted 2 years, 5 months ago
- C
"Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account."
- <https://cloud.google.com/billing/docs/how-to/billing-access>
upvoted 46 times
- ✉ **dans80** Highly Voted 2 years, 5 months ago
- Answer is C - Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account.
upvoted 20 times
- ✉ **krsourav** 1 year, 9 months ago
- Hey, look at this:)
upvoted 3 times
- ✉ **AzureDP900** Most Recent 4 months, 2 weeks ago
- It is no brainer question, C is right
upvoted 1 times
- ✉ **haroldbenites** 5 months, 1 week ago
- Go for C
upvoted 1 times
- ✉ **ElenaL** 10 months ago
- Selected Answer: C**
C - the only role appropriate answer to view and not change anything in the project is the billing viewer role.
upvoted 2 times
- ✉ **Vidyaji** 11 months, 1 week ago
- C is perfect
upvoted 1 times
- ✉ **vishnukumartr** 11 months, 3 weeks ago
- C. Add the group for the finance team to roles/billing viewer role.
upvoted 1 times
- ✉ **Chotebhaisahab** 1 year, 2 months ago
- agree c is the correct option. needs only billing viewer role.
upvoted 1 times
- ✉ **Katheri7689** 1 year, 5 months ago
- C is the correct answer
upvoted 2 times
- ✉ **mcaromit** 1 year, 5 months ago
- C is correct
upvoted 1 times
- ✉ **[Removed]** 1 year, 7 months ago
- C is correct. Add the group for the finance team to roles/billing viewer role.
upvoted 1 times

 **cloud_guru** 1 year, 8 months ago

Obvious choice here is C-- viewer to make sure they dont have additional access.

Source: this is my job

upvoted 1 times

 **Vic1043** 1 year, 8 months ago

ANS - C

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times

 **swatititame** 1 year, 11 months ago

• C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times

 **glam** 2 years ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 4 times

 **GopinathM** 2 years, 1 month ago

C is correct

upvoted 2 times

Question #100

Your organization has strict requirements to control access to Google Cloud projects. You need to enable your Site Reliability Engineers (SREs) to approve requests from the Google Cloud support team when an SRE opens a support case. You want to follow Google-recommended practices. What should you do?

- A. Add your SREs to roles/iam.roleAdmin role.
- B. Add your SREs to roles/accessapproval.approver role.
- C. Add your SREs to a group and then add this group to roles/iam.roleAdmin.role.
- D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.

Correct Answer: B*Community vote distribution*

D (100%)

 **reinocd21** Highly Voted 2 years, 5 months ago

D. Add your SREs to a group and then add this group to roles/accessapproval approver role.
-Google recommendation.

upvoted 49 times

 **Priyankahere** Highly Voted 9 months, 3 weeks ago

This was there in exam, go with community answers.
upvoted 12 times

 **Stella_1999** 9 months, 3 weeks ago

how was your exam? is this website qts useful?
upvoted 2 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

D is right ..
upvoted 1 times

 **haroldbenites** 5 months, 1 week ago

Go for D
upvoted 1 times

 **hiranfilho** 6 months, 1 week ago

Selected Answer: D
Answers C and D are correct, but it doesn't say if the SRE already has a group and as it is Google's recommendation to make a group to add users and privileges to the group, the right one is D
upvoted 2 times

 **WTY** 6 months, 2 weeks ago

It mentioned more than one SRE, so adding the user to group is most suitable approach, Answer is D.
upvoted 2 times

 **ErenYeager** 7 months ago

Passed my exams today. Not because just because of the questions I practiced here, but because of you guys, your knowledge and experience and breakdown of questions. Too bad this site can't go legit. It such an wholesome resource.

Some final words... KEEP MOVING FORWARD UNTIL ALL THE QUESTIONS ARE DESTROYED TATAKAE!!!!!!

upvoted 10 times

 **obeythefist** 8 months, 1 week ago

I've seen about 5 questions which are like this, always asking how to grant access and "follow Google best practice", and every time it's just making sure you know to use a group to control access to resources for users, and not adding users directly to objects.

Remember that keyword, "Google best practice" means "make sure you use a group"

upvoted 10 times

 **Mhkgupta** 8 months, 3 weeks ago

D. Add your SREs to a group and then add this group to roles/accessapproval approver role.
upvoted 3 times

 **pnVino27** 11 months ago

Selected Answer: D
D is Correct

upvoted 3 times

 **Vidyaji** 11 months, 1 week ago

D is perfect

upvoted 3 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: D

D Correct

upvoted 4 times

 **shawnkkk** 11 months, 3 weeks ago

D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.

upvoted 2 times

 **TenshiD** 11 months, 3 weeks ago

Selected Answer: D

D is correct

upvoted 3 times

 **Litan** 1 year, 1 month ago

D correct

upvoted 3 times

 **shrutisri** 1 year, 1 month ago

Option D

upvoted 3 times

 **akshay_jadhav** 1 year, 1 month ago

D is correct option as per google recommended practice

upvoted 4 times

Question #101

You need to host an application on a Compute Engine instance in a project shared with other teams. You want to prevent the other teams from accidentally causing downtime on that application. Which feature should you use?

- A. Use a Shielded VM.
- B. Use a Preemptible VM.
- C. Use a sole-tenant node.
- D. Enable deletion protection on the instance.

Correct Answer: D*Community vote distribution*

D (70%)

C (30%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (D):

Preventing Accidental VM Deletion

This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an `Instance` resource. To learn more about VM instances, read the Instances documentation.

As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.

By setting the `deletionProtection` flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the `deletionProtection` flag, the request fails. Only a user that has been granted a role with `compute.instances.create` permission can reset the flag to allow the resource to be deleted.

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 43 times

 **professor** Highly Voted 2 years, 4 months ago

Agree with D

You can enable Termination protection

upvoted 13 times

 **patrickakpgcp** Most Recent 2 months, 2 weeks ago

Selected Answer: D

D seems the most logical choice

upvoted 1 times

 **Adeem1981** 3 months, 3 weeks ago

Selected Answer: C

<https://cloud.google.com/compute/docs/nodes/sole-tenant-nodes> look for "host" application and "Downtime". answer is C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Please see ESP_SAP detailed explanation. This is part Tutorials Dojo practice tests.

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: D

D is the correct answer.

upvoted 2 times

 **haroldbenites** 5 months ago

Go for D

upvoted 1 times

 **nhadi82** 5 months, 3 weeks ago

I would go for D as it prevents users from deleting instances and make the downtime

upvoted 1 times

 **hiranfilho** 6 months, 1 week ago

Selected Answer: D

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 1 times

Bableves 7 months ago

Selected Answer: C

I mean read the question, how could be D the correct answer? The solution is even in the link you provided... What causes downtime? VM deletion only? Gosh what going on here...

Specifications

Deletion prevention does not prevent the following actions:

Terminating an instance from within the VM (such as running the shutdown command)

Stopping an instance

Resetting an instance

Suspending an instance

Instances being removed due to fraud and abuse after being detected by Google

Instances being removed due to project termination

upvoted 1 times

Passerofexams 4 months, 1 week ago

Sole-tenancy more or less refers to not sharing the host hardware with other GCP customers.

I believe it is out-of-scope in regard to this particular scenario.

<https://cloud.google.com/compute/docs/nodes/sole-tenant-nodes>

upvoted 2 times

Priyankahere 9 months, 3 weeks ago

This was there in exam, go with community answers.

upvoted 3 times

sqwert 9 months, 4 weeks ago

Correct Answer is D!

upvoted 1 times

Vidyaji 11 months, 1 week ago

Selected Answer: D

D is perfect

upvoted 2 times

FernandoJ 11 months, 1 week ago

Selected Answer: C

If we are talking about downtimes, there are multiple ways to cause downtime in an instance not just deleting it. That option in this list that can protect from most of them is: isolation, answer C.

upvoted 1 times

Ridhanya 11 months, 2 weeks ago

Ans is D because C is applicable project wise while in this case everyone are using the same project. so it doesn't sense.

Sole-tenancy lets you have exclusive access to a sole-tenant node, which is a physical Compute Engine server that is dedicated to hosting only your project's VMs. Use sole-tenant nodes to keep your VMs physically separated from VMs in other projects, or to group your VMs together on the same host hardware

option A is not applicable too: Shielded VMs are virtual machines (VMs) on Google Cloud hardened by a set of security controls that help defend against rootkits and bootkits. Using Shielded VMs helps protect enterprise workloads from threats like remote attacks, privilege escalation, and malicious insiders.

option B is not at all an option. so left with D

upvoted 7 times

jabrrJ68w02ond1 11 months, 3 weeks ago

Although I could tell countless ways to cause downtime even if all four were right, D is the most reasonable answer.

upvoted 1 times

alaahakim 11 months, 3 weeks ago

Selected Answer: D

Correct_Ans: D

upvoted 1 times

Question #102

Your organization needs to grant users access to query datasets in BigQuery but prevent them from accidentally deleting the datasets. You want a solution that follows Google-recommended practices. What should you do?

- A. Add users to roles/bigquery user role only, instead of roles/bigquery dataOwner.
- B. Add users to roles/bigquery dataEditor role only, instead of roles/bigquery dataOwner.
- C. Create a custom role by removing delete permissions, and add users to that role only.
- D. Create a custom role by removing delete permissions. Add users to the group, and then add the group to the custom role.

Correct Answer: B*Community vote distribution*

D (65%)	A (30%)	5%
---------	---------	----

 **someoneinthecloud** Highly Voted  2 years, 3 months ago

I believe the key part is the "following Google Best Practices" phrase.

- A - Works, but doesn't follow GCP best practices
- B - Doesn't work as the role grants permission to delete datasets
- C - Works, but is more complicated than A and doesn't follow Google best practices
- D - Correct, more complicated than A, but it follows Google Best Practices.

upvoted 64 times

 **kyo** 1 year, 3 months ago

I don't think A works properly.

roles/bigquery.user has bigquery.datasets.create. And the documentation states:

> Additional, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

If bigquery.user creates a new dataset, it's likely that bigquery.user will get permission to delete that dataset. This means that bigquery.user may have permission to delete data.

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 9 times

 **brvinod** 9 months ago

A bigquery.user will get a "data owner" role on the datasets he creates. That means he can delete those data sets he created. In that sense A fails to that extent.

upvoted 2 times

 **Bossam** 3 months, 1 week ago

See the question carefully "accidentally deleting the datasets" it is saying not to delete "the" datasets which means original dataset which existed before his creation .So answer is A.

upvoted 1 times

 **YuvrajK** 1 year, 4 months ago

I think A is the Answer and it follow GCP best practices.

<https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>

We do have the role - BigQuery User which does the below permissions

When applied to a project, this role also provides the ability to run jobs, including queries, within the project.

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

upvoted 6 times

 **TAvenger** 1 year, 8 months ago

Read description carefully "prevent from accidentally deleting the datasets". Not tables, datasets! option B does not allow to delete datasets either.

Check dataset permissions in the roles/bigquery.dataEditor:

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

bigquery.datasets.updateTag

You CANNOT delete dataset with option "B"

upvoted 11 times

 **Bableves** 7 months ago

Neither with A.

upvoted 1 times

 **afooh** 3 months, 1 week ago

But it means you will have to add the users one by one which doesn't follow Google best practices...
upvoted 2 times

 **ESP_SAP** Highly Voted  2 years, 2 months ago

Correct Answer is (D):

The proper answer regarding to bigquery roles is the listed in the options, the proper role that resolve this requirement is:
roles/bigquery.dataViewer
https://cloud.google.com/bigquery/docs/access-control#custom_roles

on the other hand, the question explicitly is asking to use the GCP best practices on IAM :
GCP Best Practices explain clearly these rules:

Policy management

- Set organization-level IAM policies to grant access to all projects in your organization.
- Grant roles to a Google group instead of individual users when possible. It is easier to add members to and remove members from a Google group instead of updating an IAM policy to add or remove users.
- If you need to grant multiple roles to allow a particular task, create a Google group, grant the roles to that group, and then add users to that group.

https://cloud.google.com/iam/docs/using-iam-securely#policy_management

upvoted 45 times

 **JackGlemins** 1 year, 8 months ago

Other best practice is use predefine roles over custom roles. Maybe A is correct

upvoted 7 times

 **JackGlemins** 1 year, 8 months ago

I correct myself: <https://cloud.google.com/iam/docs/understanding-custom-roles>

Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.

upvoted 3 times

 **prashuG** 2 years, 2 months ago

Answer is A: roles/bigquery.user is a BigQuery User role which when applied to a project provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project.

Ref: <https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>

upvoted 5 times

 **j1c4b** 1 year, 7 months ago

you can create data set with bigquery.user role because it has bigquery.datasets.create permissions. And if a user has bigquery.datasets.create permissions, when that user creates a dataset, they are granted bigquery.dataOwner access to it. So A is NOT a choice

upvoted 5 times

 **hiromi** Most Recent  1 week ago

Selected Answer: D

D is more complete than A

upvoted 1 times

 **PKookNN** 3 weeks ago

Selected Answer: B

I went back to read the description of IAM for BQ pertaining to datasets:

role/data editor has bigquery.datasets.create, bigquery.datasets.get, bigquery.datasets.getIamPolicy, bigquery.datasets.updateTag

But NO delete for datasets (the role can delete tables in datasets)

role/data owner can "When applied to a dataset, this role provides permissions to: Read, update, and delete the dataset."

Very confusing but B is correct in this case.

upvoted 1 times

 **JagatSingh** 1 month, 3 weeks ago

Selected Answer: D

As per Link: <https://cloud.google.com/bigquery/docs/access-control#bigquery.dataEditor> Correct answer is D

A - Works, but doesn't follow GCP best practices

B - Doesn't work as the role grants permission to delete datasets

C - Works, but is more complicated than A and doesn't follow Google best practices

D - Correct, more complicated than A, but it follows Google Best Practices

upvoted 1 times

 **JackyWorld** 2 months, 1 week ago

Selected Answer: D

D is the right one...

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: D

if your company has 10.000 employees that need to have a specific action for BigQuery, good luck adding the role manually to each of them. OBVIOUSLY adding them to a group and then adding the group to the respective custom role is a common sense solution here. And a GCP good practice.

upvoted 1 times

 **sam2022hi** 3 months, 1 week ago

Thoughts according to docs (<https://cloud.google.com/bigquery/docs/access-control>) which list out all permissions:

Definitely not B) As (roles/bigquery.dataEditor) role has access to bigquery.tables.delete

Not C) Not best practice

Not D) This is a possibility but no need to go to extra effort of creating a custom role (not best practice) when A, according to the docs, not not have access to the permission of deleting a table.

Tutorial Dojo has a similar question and they also say a custom role could also add deletion permissions in an update and this is a problem.

So the answer should be A) for simplicity.

upvoted 1 times

 **mrNicky** 3 months, 2 weeks ago

Best practices when we talk about IAM / rôles is:

Create group => assign role to the group => add users

ANSWER D

upvoted 2 times

 **Passerofexams** 4 months, 1 week ago

From Google - 'Controlling access to datasets'

"Note: For a user to be able to query the tables in a dataset, it is not sufficient for the user to have access to the dataset. A user must also have permission to run a query job in a project. If you want to give a user permission to run a query from your project, give the user the bigquery.jobs.create permission for the project. You can do this by assigning the user the roles/bigquery.jobUser role for your project. For more information, see Access control examples."

upvoted 1 times

 **ahmadjubair33** 4 months, 2 weeks ago

The q is saying about dataset not dataset table so answer will be B

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

D is correct, group need to add it to role

upvoted 1 times

 **haroldbenites** 5 months ago

Go for D

Best practice is create a group and then assign role to that group .

upvoted 1 times

 **rsuresh27** 6 months, 2 weeks ago

Correct answer is D. Best practice is to use Google groups to manage multiple people and organizations. A does not do this, so A cannot be correct. Thus, D is the best option out of the remaining choices.

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: D

Ans: D

Google Best Practices = Users -> Group -> IAM Role

upvoted 2 times

 **dinesh198728** 7 months, 4 weeks ago

Selected Answer: D

cutome role

upvoted 1 times

 **Akash7** 8 months ago

D is correct. Best practice is to use groups rather then assigning the role directly.

upvoted 1 times

Question #103

You have a developer laptop with the Cloud SDK installed on Ubuntu. The Cloud SDK was installed from the Google Cloud Ubuntu package repository. You want to test your application locally on your laptop with Cloud Datastore. What should you do?

- A. Export Cloud Datastore data using gcloud datastore export.
- B. Create a Cloud Datastore index using gcloud datastore indexes create.
- C. Install the google-cloud-sdk-datastore-emulator component using the apt get install command.
- D. Install the cloud-datastore-emulator component using the gcloud components install command.

Correct Answer: D*Community vote distribution*

D (54%)

C (46%)

 **someoneinthecloud** Highly Voted 2 years, 3 months ago

I believe answer is C

<https://cloud.google.com/sdk/docs/downloads-apt-get>

The question is not about the datastore command itself but from where we should run the update command on the Ubuntu to install the component.

upvoted 34 times

 **XRiddlerX** 2 years, 3 months ago

I agree with this comment. The answer is C.

If you installed the SDK from the Ubuntu repo and try to do the following:

\$ gcloud components install cloud-datastore-emulator

You will receive this message:

ERROR: (gcloud.components.install)

You cannot perform this action because the Cloud SDK component manager

is disabled for this installation. You can run the following command

to achieve the same result for this installation:

sudo apt-get install google-cloud-sdk-datastore-emulator

upvoted 36 times

 **myuniqueusername** 1 year, 1 month ago

absolutely insane if that question comes up during the associate exam, who on earth would know that off the top of their heads?

upvoted 17 times

 **Ale1973** 2 years, 2 months ago

WOW!!! Today I have learned a new and interesting thing thanks to you...

upvoted 6 times

 **stepkurniawan** 2 years, 2 months ago

it says that in your Ubuntu, you have Cloud SDK installed already. So it should be able to run the command in D

upvoted 4 times

 **Ale1973** 2 years, 2 months ago

Yes, but it says that "The Cloud SDK was installed from the Google Cloud Ubuntu package repository", then to install datastore emulator you should use the command in Option C.

upvoted 9 times

 **professor** Highly Voted 2 years, 4 months ago

Ans is D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 21 times

 **Eshkrkrkr** 1 year, 12 months ago

Wrong! The answer is C! When you install SDK using apt Cloud SDK Component Manager is disabled and you need to install extra packages again using apt.

https://cloud.google.com/sdk/docs/components#managing_cloud_sdk_components

Note: These instructions will not work if you have installed Cloud SDK using a package manager such as APT or yum because Cloud SDK Component Manager is disabled when using that method of installation.

upvoted 24 times

 **SWObaby** 1 year, 11 months ago

I believe the answer is C...

It is a tricky question!! The question states, "The Cloud SDK was installed from the Google Cloud Ubuntu package repository." For those, who aren't that familiar with Debian/Ubuntu, D seems like an attractive answer. It works as a way to install Datastore...but it does NOT fit the context of the question.

I recommend looking back to G Cloud SDK installation (Debian/Ubuntu): <https://cloud.google.com/sdk/docs/install#deb>

Read the "Installation Steps" in the documentation. In Step 3, "sudo apt-get update && sudo apt-get install google-cloud-sdk". Then, Step 4 is additionally adding other components, such as "sudo apt-get install google-cloud-sdk-datastore-emulator".

Proving C the correct answer.

upvoted 21 times

 **ShakthiGCP** 1 year, 8 months ago

Go With 'C' ... just tried creating a ubuntu server and verified these.. Dont worry about any other options.

<https://cloud.google.com/sdk/docs/quickstart#deb> check this link .

upvoted 4 times

 **mattcl** Most Recent 1 day, 9 hours ago

it is D : https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing_the_emulator

upvoted 1 times

 **nonsense** 1 week, 1 day ago

Selected Answer: C

it's C

upvoted 1 times

 **FaisalJ** 1 week, 2 days ago

Selected Answer: D

It is D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 1 times

 **Untamables** 2 weeks, 3 days ago

Selected Answer: D

Absolutely D

upvoted 1 times

 **Erie** 2 weeks, 5 days ago

Selected Answer: C

If "gcloud" was not downloaded directly from Google's website, and was installed through a third-party package manager, gcloud components install will not work.

See this Stackoverflow post: <https://stackoverflow.com/questions/42697026/install-google-cloud-components-error-from-gcloud-command>

upvoted 1 times

 **PKookNN** 3 weeks ago

D is correct

upvoted 1 times

 **GaneshSurwase** 3 weeks, 3 days ago

The ans is C

upvoted 1 times

 **anolive** 3 weeks, 6 days ago

I believe is D gcloud components install cloud-datastore-emulator

upvoted 1 times

 **AwesomeGCP** 4 weeks, 1 day ago

C. Install the google-cloud-sdk-datastore-emulator component using the apt get install command. END OF THE DISCUSSION!!!

upvoted 1 times

 **ESP_SAP** 2 months ago

Correct Answer is (D):

The Datastore emulator provides local emulation of the production Datastore environment. You can use the emulator to develop and test your application locally. In addition, the emulator can help you generate indexes for your production Firestore in Datastore mode instance and delete unneeded indexes. This page guides you through installing the emulator, starting the emulator, and setting environment variables to connect your application to the emulator.

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 3 times

 **Shara1997** 2 months ago

Selected Answer: D

gcloud components install cloud-datastore-emulator option d

upvoted 1 times

 **akg001** 2 months, 3 weeks ago

Correct ans : D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>

upvoted 2 times

 **abirroy** 3 months ago

Selected Answer: D

Install the cloud-datastore-emulator component using the gcloud components install command

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: C

<https://cloud.google.com/sdk/docs/downloads-apt-get>

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: D

<https://cloud.google.com/sdk/docs/downloads-apt-get>

upvoted 1 times

Question #104

Your company set up a complex organizational structure on Google Cloud. The structure includes hundreds of folders and projects. Only a few team members should be able to view the hierarchical structure. You need to assign minimum permissions to these team members, and you want to follow Google-recommended practices. What should you do?

- A. Add the users to roles/browser role.
- B. Add the users to roles/iam.roleViewer role.
- C. Add the users to a group, and add this group to roles/browser.
- D. Add the users to a group, and add this group to roles/iam.roleViewer role.

Correct Answer: C*Community vote distribution*

C (100%)

 **SIX** Highly Voted 2 years, 5 months ago

C is the better answer.

upvoted 30 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (C):

We need to apply the GCP Best practices.

roles/browser Browser Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 25 times

 **Charumathi** Most Recent 1 month ago

C is the correct answer,

ID

roles/browser

Role launch stage

General Availability

Description

Access to browse GCP resources.

6 assigned permissions

resourcemanager.folders.get
resourcemanager.folders.list
resourcemanager.organizations.get
resourcemanager.projects.get
resourcemanager.projects.getIamPolicy
resourcemanager.projects.list

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: C

<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 1 times

 **ikychugtai** 3 months, 1 week ago

C is the better answer

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: C

C is correct

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Correct Answer is (C):

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: C

C all the way
upvoted 1 times

 **haroldbenites** 5 months ago

Go for C
upvoted 1 times

 **avik001** 7 months, 2 weeks ago

C is the better ans
upvoted 1 times

 **casokan** 9 months ago

Correct Answer: C
upvoted 1 times

 **pililili_bed** 9 months, 2 weeks ago

C is correct
upvoted 2 times

 **jaravu** 9 months, 3 weeks ago

C is right
upvoted 2 times

 **vishnukumartr** 11 months ago

C. Add the users to a group, and add this group to roles/browser.
upvoted 2 times

 **jabrrJ68w02ond1** 11 months, 3 weeks ago

A is correct but not following GCP best practices. GCP wants you to put users into groups whenever it's possible. So the answer is C.
upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Correct Ans is : C
upvoted 2 times

 **shawnkkk** 11 months, 3 weeks ago

C. Add the users to a group, and add this group to roles/browser.
upvoted 1 times

Question #105

Your company has a single sign-on (SSO) identity provider that supports Security Assertion Markup Language (SAML) integration with service providers. Your company has users in Cloud Identity. You would like users to authenticate using your company's SSO provider. What should you do?

- A. In Cloud Identity, set up SSO with Google as an identity provider to access custom SAML apps.
- B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
- C. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Mobile & Desktop Apps.
- D. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Web Server Applications.

Correct Answer: A*Community vote distribution*

B (100%)

 **poogcp** Highly Voted 2 years, 5 months ago

For me its B option

upvoted 26 times

 **PhilipAWS** Highly Voted 1 year, 10 months ago

Only option B make sense to me as per - https://support.google.com/cloudidentity/answer/6262987?hl=en&ref_topic=7558767

upvoted 22 times

 **nitinz** 1 year, 8 months ago

you nailed it. B is correct.

upvoted 3 times

 **hiromi** Most Recent 6 days, 21 hours ago

Selected Answer: B
For me b is the ans

upvoted 1 times

 **leliodesouza** 4 weeks ago

Selected Answer: B
B is the correct answer

upvoted 1 times

 **AwesomeGCP** 1 month ago

Selected Answer: B
B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.

upvoted 1 times

 **Angel_99** 2 months, 3 weeks ago

Selected Answer: B
Only option B make sense to me

upvoted 1 times

 **habros** 3 months ago

Same and OAuth were two different standards! Google Account is an IdP on its own, whereas question calls for external IdP. Definitely B.
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Option B
upvoted 1 times

 **haroldbenites** 5 months ago

Go for B
upvoted 1 times

 **Terzlightyear** 6 months, 2 weeks ago

Selected Answer: B
B is the right one
upvoted 2 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: B

B is the correct answer
upvoted 2 times

 **Priyankahere** 9 months, 3 weeks ago

This was there in exam, go with community answers.
upvoted 5 times

 **azure4life** 10 months, 3 weeks ago

Selected Answer: B
Option B
upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: B
B would be the right choice
upvoted 1 times

 **jackwillis** 1 year, 1 month ago

B. Why would you setup google as identity provider when you have already has SSO
upvoted 2 times

 **vamgcp** 1 year, 1 month ago

Option B - In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
upvoted 1 times

 **sudav** 1 year, 2 months ago

B - This is the only possible option. You configure applications (service providers) to accept SAML assertions from the company's existing identity provider and users in Cloud Identity can sign in to various applications through the third-party single sign-on (SSO) identity provider. It is important to note that user authentication occurs in the third-party IdP so the absence of a Gmail login is not an issue for signing in.
Ref: <https://cloud.google.com/identity/solutions/enable-sso>
If you have a third-party IdP, you can still configure SSO for third-party apps in the Cloud Identity catalog. User authentication occurs in the third-party IdP, and Cloud Identity manages the cloud apps.
To use Cloud Identity for SSO, your users need Cloud Identity accounts. They sign in through your third-party IdP or using a password on their Cloud Identity accounts.

upvoted 3 times

Question #106

Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. You need to assign this person the minimum role for projects. What should you do?

- A. Add the user to roles/iam.roleAdmin role.
- B. Add the user to roles/iam.securityAdmin role.
- C. Add the user to roles/iam.serviceAccountUser role.
- D. Add the user to roles/iam.serviceAccountAdmin role.

Correct Answer: D*Community vote distribution*

D (100%)

 **PhilipAWS** Highly Voted 1 year, 10 months ago

Whoever say C is right answer, please read the question 1000000000 times if not understand - "Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. " Dedicated person who creates and manages all service... Now read below;

To allow a user to manage service accounts, grant one of the following roles:

Service Account User (roles/iam.serviceAccountUser): Includes permissions to list service accounts, get details about a service account, and impersonate a service account.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

Now look in which role mentioned "CREATE"?

Obviously - roles/iam.serviceAccountAdmin..... So Answer is????

1M% - D only

upvoted 42 times

 **Jhelum** 9 months, 3 weeks ago

Calm down Jamal, don't pull out the knife...

upvoted 10 times

 **SIX** Highly Voted 2 years, 5 months ago

The right answer is D.

upvoted 39 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

dedicated person who creates and manages all service accounts is key word makes me select D as right answer.

upvoted 1 times

 **haroldbenites** 5 months ago

Go for D

upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

D is the correct answer

upvoted 1 times

 **Priyankahere** 9 months, 3 weeks ago

This was there in exam, go with community answers.

upvoted 2 times

 **kped21** 10 months, 1 week ago

D - ServiceAccountUser has no privs on ServiceAccounts so C is wrong.

upvoted 1 times

 **jaffarali** 11 months ago

D is the right option

upvoted 1 times

 **sharmamohitkr700** 11 months, 1 week ago

Selected Answer: D

To allow a user to manage service accounts, grant one of the following roles:

Service Account User (roles/iam.serviceAccountUser): Includes permissions to list service accounts, get details about a service account, and impersonate a service account.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

upvoted 3 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: D

D is Correct

upvoted 2 times

 **TenshiD** 11 months, 3 weeks ago

Selected Answer: D

D is correct.

upvoted 1 times

 **maggieli** 1 year ago

D is correct.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

upvoted 1 times

 **mcaromit** 1 year, 5 months ago

D is correct

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

D is correct. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 3 times

 **yuvি69** 1 year, 7 months ago

option D is correct. because it is the only one in the 4 options that can create and manage service accounts

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

D. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

D is correct

upvoted 1 times

Question #107

You are building an archival solution for your data warehouse and have selected Cloud Storage to archive your data. Your users need to be able to access this archived data once a quarter for some regulatory requirements. You want to select a cost-efficient option. Which storage option should you use?

- A. Cold Storage
- B. Nearline Storage
- C. Regional Storage
- D. Multi-Regional Storage

Correct Answer: A

Nearline, Coldline, and Archive offer ultra low-cost, highly-durable, highly available archival storage. For data accessed less than once a year, Archive is a cost-effective storage option for long-term preservation of data. Coldline is also ideal for cold storage data your business expects to touch less than once a quarter.

For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year. All storage classes are available across all GCP regions and provide unparalleled sub-second access speeds with a consistent API.

Reference:

<https://cloud.google.com/storage/archival>

Community vote distribution

A (56%)	B (44%)
---------	---------

✉ **Teegongkia** Highly Voted 2 years, 3 months ago

Took ACE last week and the exact question came out. I go with B as i felt A is a trick answer. There is no Cold Storage in GCP.
upvoted 21 times

✉ **ssankar** 2 years, 2 months ago

Hello Teegongkia ,
is the questions are still valid ??
Thanks
upvoted 1 times

✉ **BenAji** 1 year, 6 months ago

Cold data tiering refers to the storage of less frequently, or sporadically accessed data in low cost media such as HDFS (Hadoop Distributed File System) and cloud storage options including Amazon Web Services (AWS), Google Cloud Platform (GCP), and Azure Data Lake Storage (ADLS) that are managed separately from the SAP HANA database, but still accessible at any time. blogs.sap.com/2018/12/03/what-is-sap-hana-cold-data-tiering/

upvoted 1 times

✉ **droogie** Highly Voted 2 years, 4 months ago

This one is confusing. First, there's no 'Cold' storage. It's Coldline.
Nearline Storage is ideal for data you plan to read or modify on average once per month or less. Coldline Storage is ideal for data you plan to read or modify at most once a quarter.
<https://cloud.google.com/storage/docs/storage-classes>

So with the misspelling of 'Cold' and these guys accessing it every 90 days, I'm leaning towards Nearline

upvoted 14 times

✉ **DickDastardly** 1 year, 8 months ago

"Cold" is not a typo. I took the exam today and the answers appeared exactly as listed here.
upvoted 9 times

✉ **Eshkrkrkr** 1 year, 12 months ago

It's a typo. Google wouldn't force to consume knowledge that is a non-best practice from Google. Answer is A.
upvoted 4 times

✉ **Ixgywil** 1 year, 6 months ago

For Google, these exams are just another business.
upvoted 6 times

✉ **sarahf** 1 year, 10 months ago

At the page for data archiving (<https://cloud.google.com/storage/archival>) the first paragraph says: "Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter."
So there is such thing as Cold storage according to Google.

Also at (<https://cloud.google.com/storage/docs/storage-classes#archive>) they talk about Cold storage: "Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive Storage, yet still be available if you need it."

upvoted 10 times

 **sanhoo** 1 year, 5 months ago

Thanks for bringing this up. Really helpfull.

upvoted 1 times

 **sanhoo** 1 year, 5 months ago

This line indicates that cold storage term is used for - archival / coldline

"With low latency and a consistent API across Cloud Storage, Archive and Coldline introduce cold storage you can actually use"

<https://cloud.google.com/storage/archival>

upvoted 2 times

 **TAvenger** 1 year, 8 months ago

I believe the question is old, when Regional and Multi-Regional were also storage classes of the GCS.

Before changes: (Multi-Region, Regional, Nearline, Coldline)

After recent changes we have

- Storage Classes (Standard, Nearline, Coldline, Archive)

- Storage Locations (Regional, Dual-region, Multi-Region)

It's tricky for exam because we don't to answer according to old version or new version.

For the latest version, costs for 1Gb for storing (3 month) + retrieval

Nearline: $0.01 * 3 + 0.01 = 0.04$

Coldline: $0.004 * 3 + 0.02 = 0.032$

Coldline is more cost effective.

If "Cold" means Coldline (not Archive) the answer is A

If "Cold" means Archive the answer is B

I hope that "Cold" means Coldline. I would try with A

upvoted 2 times

 **ri_unhou119** 1 year, 5 months ago

A:

Google Cloud doc:

<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 1 times

 **obeythefist** 8 months, 2 weeks ago

Yes, but the question says the data will be accessed once per quarter, Google's documentation tells us that Coldline is most suitable for data accessed less than once per quarter. This direct part of the question tells us how we must answer.

upvoted 2 times

 **gneknurd** Most Recent 1 week, 4 days ago

This question on the test I took today and passed.. I chose B cause the other answers to questions I received mentioned coldline. So google was specific on the answer.

upvoted 1 times

 **Untamables** 2 weeks, 3 days ago

Selected Answer: A

I believe 'Cold Storage' is a typo.

upvoted 1 times

 **nonsense** 2 weeks, 6 days ago

Selected Answer: B

cold not exist

upvoted 1 times

 **AwesomeGCP** 4 weeks ago

Selected Answer: A

Nearline Storage: Low cost. Good for data that can be stored for at least 30 days, including data backup and long-tail multimedia content so ofcourse its a wrong option.

But A. Cold Storage is also not correct as it means that With low latency and a consistent API across Cloud Storage, both Archive and Coldline tiers deliver cold storage you can actually use. When it comes to archival and business continuity, both Archive and Coldline change what the industry can expect from COLD STORAGE in the cloud.

Correct answer should be provided as Coldline Storage NOT Cold Storage.

WRONG AMBIGUOUS QUESTION!!!

upvoted 1 times

 **ale_brd_** 1 month, 2 weeks ago

Selected Answer: A

Cold is a typo, Coldline.

Option A is correct: Coldline storage is a very-low-cost, highly durable storage service for storing infrequently accessed data. Coldline storage is a better choice than Standard storage or Nearline storage in scenarios where slightly lower availability, a 90-day minimum storage duration, and higher costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Coldline storage is ideal for data you plan to read or modify at most once a quarter. Note, however, that for data being kept entirely for backup or archiving purposes, Archive storage is more cost-effective, as it offers the lowest storage costs.

upvoted 1 times

✉️ **CyberG** 1 month, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

✉️ **Krishna171802** 3 months ago

what is the exact answer, its too confusing

upvoted 1 times

✉️ **ryumada** 3 months ago

Selected Answer: A

See sarahf answer for the explanation. I agree with her. See also on this link:

<https://cloud.google.com/storage>

In this link, on "Archival storage you can actually use". They explain on the first sentence: "With low latency and a consistent API across Cloud Storage, the Archive and Coldline tiers deliver !!cold storage!! you can actually use." Even google says their Coldline storage as Cold Storage too. So, This make Cold Storage makes sense as the answer for this question.

upvoted 1 times

✉️ **theBestStudent** 3 months ago

Selected Answer: B

If explicitly the question would have said coldLine storage, all good. But ColdLine is not an option listed here. B, Nearline is the only valid option for this specific question.

upvoted 1 times

✉️ **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: B

Answer is B, Nearline storage

Here is why?

Let's assume for a min that this is typo or it is an alias for Archival storage in gcp, google documentation clearly states that "Archive storage is the best choice for data that you plan to access less than once a year" and now if you look at the question it says you need to be able to access this data once per quarter (which is 3 times in a year). Hence, the only option that makes sense is Nearline storage.

<https://cloud.google.com/storage/docs/storage-classes#archive>

upvoted 1 times

✉️ **zolthar_z** 3 months ago

Wrong. Did you read the link? "Nearline storage is ideal for data you plan to read or modify on average once per month or less."

upvoted 2 times

✉️ **zolthar_z** 3 months, 2 weeks ago

Selected Answer: A

Coldline storage is ideal for data you plan to read or modify at most once a quarter.<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 1 times

✉️ **abirroy** 3 months, 2 weeks ago

I think A should have been coldline not cold storage. Must be a typo

upvoted 1 times

✉️ **JasmineS1975** 3 months, 3 weeks ago

Selected Answer: B

Google doc says if accessed less than once per quarter, go to Coldline. If they're accessing once per quarter, I'd go with B - Nearline. Also, there's no such thing as "Cold" storage. <https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

✉️ **JasmineS1975** 4 months ago

Google doc says if accessed less than once per quarter, go to Coldline. If they're accessing once per quarter, I'd go with B - Nearline. Also, there's no such thing as "Cold" storage. <https://cloud.google.com/storage/docs/storage-classes>

upvoted 1 times

✉️ **Bumbah** 4 months ago

Selected Answer: B

Once a quarter access, so nearline.

upvoted 1 times

Question #108

A team of data scientists infrequently needs to use a Google Kubernetes Engine (GKE) cluster that you manage. They require GPUs for some long-running, non-restartable jobs. You want to minimize cost. What should you do?

- A. Enable node auto-provisioning on the GKE cluster.
- B. Create a VerticalPodAutoscaler for those workloads.
- C. Create a node pool with preemptible VMs and GPUs attached to those VMs.
- D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

Correct Answer: C

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

Community vote distribution

A (67%)	D (33%)
---------	---------

✉  **Polok** Highly Voted 2 years, 5 months ago

If you need something for long-running, non-restartable jobs you don't use preemptible VMs

Think answer is D.

upvoted 54 times

✉  **[Removed]** Highly Voted 1 year, 7 months ago

Incorrect options are

- B. VerticalPodAutoscaler scales PODS based on the app you deploy.
For handle infrequently GPU access, you need infrequently GPU nodes
VerticalAutoscaler Pod deployed on a non GPU node it useless,
[We can't have the node always have GPU for infrequent requests]
- C. Preemptible VMs can't last long
- D. For infrequent access, you don't want to have a permanent homogenous cluster.

The correct option is "A"

auto-provisioning = Attaches and deletes node pools to cluster based on the requirements.
Hence creating a GPU node pool, and auto-scaling would be better

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 14 times

✉  **dttnc1** 1 year ago

I do agree A is the answer. Since this is for infrequent needs, autoscaling in letter D is not cost effective as it will always run min. of 1 instance.
If we need to infrequently use a cluster, the nodes should be able to adjust based on the current need.

"With node auto-provisioning, new node pools are created and deleted automatically." <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 5 times

✉  **kyo** 1 year, 3 months ago

I think using NAP is the correct answer.

→ Node Auto Provisioning (NAP a.k.a., Nodepool Auto Provisioning)

There is an introduction of NAP described below on the blog.

> The above recommendations optimize for cost. NAP, for instance, reduces costs by taking down nodes during underutilized periods.

<https://cloud.google.com/blog/products/containers-kubernetes/best-practices-for-creating-a-highly-available-gke-cluster>

upvoted 2 times

✉  **JCH760310** 10 months, 3 weeks ago

they "require GPUs" - so after checking in Udemy practice tests there is similar question there. And the D answer seems to be the best fit for our scenario here.

"This option is the most optimal solution for the requirement. Rather than recreating all nodes, you create a new node pool with GPU enabled. You then modify the pod specification to target particular GPU types by adding node selector to your workload's Pod specification. You still have a single cluster, so you pay Kubernetes cluster management fee for just one cluster, thus minimizing the cost." Still better option than creating new GKE cluster with GPUs.

Ref: <https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

Ref: <https://cloud.google.com/kubernetes-engine/pricing>

upvoted 2 times

 **kimharsh** 11 months, 3 weeks ago

A is not correct because you can't add a GPU node to an existing GKE cluster

Limitations

Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.

GPU nodes cannot be live migrated during maintenance events.

GPUs are only supported with general-purpose N1 machine types.

GPUs are not supported in Windows Server node pools

REF: <https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#limitations>

So the answer should be D

upvoted 6 times

 **rachee** 11 months, 1 week ago

Your reference says existing "node pools" not GKE cluster. Auto-provisioning creates new "node pools":
<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 2 times

 **Ridhanya** 11 months, 1 week ago

but node pools are homogenous, so how can we be sure that option A will create a GPU node pool

upvoted 1 times

 **wjtb** 5 months, 1 week ago

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>
 Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.

GPU requests

Pending Pods' node affinities and label selectors.

Pending Pods' node taints and tolerations.

upvoted 3 times

 **mattcl** Most Recent 1 day, 9 hours ago

it is A) <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 1 times

 **Untamables** 2 weeks, 3 days ago

Selected Answer: A

Vote A

upvoted 1 times

 **diasporabro** 3 weeks, 1 day ago

I'll go with A.

Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.

GPU requests

Pending Pods' node affinities and label selectors.

Pending Pods' node taints and

upvoted 1 times

 **anolive** 3 weeks, 6 days ago

Selected Answer: A

Node auto-provisioning automatically manages a set of node pools on the user's behalf. Without node auto-provisioning, GKE starts new nodes only from user-created node pools. With node auto-provisioning, new node pools are created and deleted automatically.

upvoted 1 times

 **AwesomeGCP** 4 weeks ago

D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

upvoted 1 times

 **Charumathi** 1 month ago

A is the correct answer,

Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.

GPU requests

Pending Pods' node affinities and label selectors.

Pending Pods' node taints and tolerations.

GPU requests:

GPU availability in GKE

In GKE Autopilot and Standard, you can attach GPU hardware to nodes in your clusters, and then allocate GPU resources to containerized

workloads running on those nodes. You can use these accelerators to perform resource-intensive tasks, such as the following:

Machine learning (ML) inference and training
Large-scale data processing

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: A

Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.
GPU requests
Pending Pods' node affinities and label selectors.
Pending Pods' node taints and tolerations.

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning#operation>

upvoted 1 times

 **jeffangel28** 3 months, 2 weeks ago

Selected Answer: A

Based on this, you can request GPU with node auto-provisioning. it's A

https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning?hl=es_419#operation

upvoted 2 times

 **RanjithK** 4 months, 1 week ago

Go with D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is my answer .. long-running, non-restartable jobs is key to answer this question

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: D

D is the correct answer

upvoted 1 times

 **haroldbenites** 5 months ago

Go for D

The key word is "GPU"

upvoted 1 times

 **peugeotdude** 5 months, 2 weeks ago

Selected Answer: A

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 4 times

 **LaxmanTiwari** 5 months, 2 weeks ago

Selected Answer: D

Creating an autoscaling GPU node pool

To take the best, most cost-effective advantage of GPUs on GKE, and to take advantage of cluster autoscaling, we recommend creating separate GPU node pools in your clusters.

When you add a GPU node pool to an existing cluster that already runs a non-GPU node pool, GKE automatically taints the GPU nodes with the following node taint:

Key: nvidia.com/gpu

Effect: NoSchedule

Note: If a GPU node pool is added to a cluster where all the existing node pools are GPU node pools, or if you are creating a new cluster with a GPU attached default pool, the above taint will not be added to the GPU nodes. The taint will also not be added to the existing GPU nodes retrospectively when a non-GPU node pool is added afterwards.

Additionally, GKE automatically applies the corresponding tolerations to Pods requesting GPUs by running the ExtendedResourceToleration admission controller.

This causes only Pods requesting GPUs to be scheduled on GPU nodes, which enables more efficient autoscaling: your GPU nodes can quickly scale down if there are not enough Pods requesting GPUs.

upvoted 1 times

 **nhadi82** 5 months, 3 weeks ago

I would go for D as it is non-restartable jobs which is not perfect with preemptive node

upvoted 1 times

Question #109

Topic 1

Your organization has user identities in Active Directory. Your organization wants to use Active Directory as their source of truth for identities. Your organization wants to have full control over the Google accounts used by employees for all Google services, including your Google Cloud Platform (GCP) organization. What should you do?

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.
- B. Use the cloud Identity APIs and write a script to synchronize users to Cloud Identity.
- C. Export users from Active Directory as a CSV and import them to Cloud Identity via the Admin Console.
- D. Ask each employee to create a Google account using self signup. Require that each employee use their company email address and password.

Correct Answer: A

Reference:

<https://cloud.google.com/solutions/federating-gcp-with-active-directory-introduction>

Community vote distribution

A (100%)

 **professor** Highly Voted 2 years, 4 months ago

Ans is A

<https://tools.google.com/dlpage/dirsync/>

upvoted 26 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer (A):

Directory Sync

Google Cloud Directory Sync enables administrators to synchronize users, groups and other data from an Active Directory/LDAP service to their Google Cloud domain directory

<https://tools.google.com/dlpage/dirsync/>

upvoted 18 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

A is right, this is part of Tutorials Dojo practice test

upvoted 1 times

 **haroldbenites** 5 months ago

Go for A

upvoted 1 times

 **crisyeb** 8 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **LeonNip** 10 months ago

Selected Answer: A

A is right

upvoted 3 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

A is Correct

upvoted 4 times

 **mcaromit** 1 year, 5 months ago

A is correct

upvoted 3 times

 **[Removed]** 1 year, 7 months ago

A is correct. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.
upvoted 3 times

 **devscorpio2001** 1 year, 11 months ago

This is A , you can use Google Cloud Sync
upvoted 1 times

 **swatititame** 1 year, 11 months ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.
upvoted 1 times

 **RockAJ** 2 years, 1 month ago

A also for me!
upvoted 1 times

 **hiteshrup** 2 years, 2 months ago

Straight Answer: A

Google has Cloud Directory Sync service to link LDAP in GCP

upvoted 2 times

 **SSPC** 2 years, 2 months ago

A is correct for me
upvoted 2 times

 **Verve** 2 years, 3 months ago

A it is
upvoted 2 times

Question #110

You have successfully created a development environment in a project for an application. This application uses Compute Engine and Cloud SQL. Now you need to create a production environment for this application. The security team has forbidden the existence of network routes between these 2 environments and has asked you to follow Google-recommended practices. What should you do?

- A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.
- B. Create a new production subnet in the existing VPC and a new production Cloud SQL instance in your existing project, and deploy your application using those resources.
- C. Create a new project, modify your existing VPC to be a Shared VPC, share that VPC with your new project, and replicate the setup you have in the development environment in that new project in the Shared VPC.
- D. Ask the security team to grant you the Project Editor role in an existing production project used by another division of your company. Once they grant you that role, replicate the setup you have in the development environment in that project.

Correct Answer: A*Community vote distribution*

A (83%) C (17%)

 **poogcp** Highly Voted 2 years, 5 months ago

Correct answer is A.

upvoted 32 times

 **pYWORLD** 1 year, 3 months ago

Correct answer!

upvoted 4 times

 **[Removed]** Highly Voted 2 years, 2 months ago

A - correct. Best practice is to create a new project for each environment, such as production and testing. There are no routes between VPCs in these projects by default, so that satisfies the requirement by the security team.

B. Nope. not best practice and allows communication.

C. While this is best practice to create a new project for a different environment, it explicitly breaks the security team's rule of having no path between environments by nature of the shared VPC. The shared VPC allows entities in both VPCs to communicate as if they were in the same VPC. That's definitely wrong.

D. One - not best practice to replicate in the setup in that project. Two - why do they suddenly need the project editor rule? Just a bad answer. Wrong.

upvoted 26 times

 **diasporabro** Most Recent 3 weeks, 1 day ago

Selected Answer: A

Satisfies requirements by the security team

upvoted 1 times

 **anolive** 3 weeks, 5 days ago

Selected Answer: A

make sense

upvoted 1 times

 **alexandercamachop** 3 months ago

Selected Answer: A

A is definitely the answer.

upvoted 1 times

 **theBestStudent** 3 months ago

<https://cloud.google.com/architecture/framework/system-design/resource-management#decouple>

upvoted 1 times

 **andreherwanto** 3 months, 3 weeks ago

Selected Answer: A

Correct answer is A.

upvoted 1 times

 **csrazdan** 4 months ago

Selected Answer: C

Technically we should create a new VPC if the network is not shared. Creating resources in a new project even within a new subnet will not separate unless firewall rules are not explicitly denying the traffic. The best answer is to create a shared VPC where DEV and PROD are service projects.

My Answer is: C

upvoted 1 times

 **theBestStudent** 3 months ago

Why you want to share environments? they should be isolated. Therefore Answer should be A.

upvoted 3 times

 **haroldbenites** 5 months ago

Go for A

I thought that the correct answer was the C , but the question did not say to communicate both environments.

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: A

Ans: A

Agreed!

upvoted 1 times

 **JieHeng** 1 year, 4 months ago

Should be A

it's a best practice "to have one project per application per environment." - <https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#project-structure>

upvoted 5 times

 **arsh1916** 1 year, 5 months ago

A answer

upvoted 1 times

 **mcaromit** 1 year, 5 months ago

A is correct

upvoted 1 times

 **EJJ** 1 year, 7 months ago

i will go with A

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

A is correct. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

C is correct

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.

upvoted 2 times

Question #111

Your management has asked an external auditor to review all the resources in a specific project. The security team has enabled the Organization Policy called Domain Restricted Sharing on the organization node by specifying only your Cloud Identity domain. You want the auditor to only be able to view, but not modify, the resources in that project. What should you do?

- A. Ask the auditor for their Google account, and give them the Viewer role on the project.
- B. Ask the auditor for their Google account, and give them the Security Reviewer role on the project.
- C. Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.
- D. Create a temporary account for the auditor in Cloud Identity, and give that account the Security Reviewer role on the project.

Correct Answer: C

Community vote distribution

C (100%)

 **dan80** Highly Voted 2 years, 5 months ago

C - https://cloud.google.com/iam/docs/roles-audit-logging#scenario_external_auditors
upvoted 41 times

 **spudleymcdudley** 2 years, 4 months ago

This guy is right!
upvoted 6 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (C):

roles/viewer Read access to all resources. Get and list access for all resources.

Using primitive roles

The following table lists the primitive roles that you can grant to access a project, the description of what the role does, and the permissions bundled within that role. Avoid using primitive roles except when absolutely necessary. These roles are very powerful, and include a large number of permissions across all Google Cloud services. For more details on when you should use primitive roles, see the Identity and Access Management FAQ.

IAM predefined roles are much more granular, and allow you to carefully manage the set of permissions that your users have access to. See Understanding Roles for a list of roles that can be granted at the project level. Creating custom roles can further increase the control you have over user permissions.

https://cloud.google.com/resource-manager/docs/access-control-proj#using_primitive_roles

upvoted 18 times

 **alexandercamachop** Most Recent 3 months ago

Selected Answer: C
C - https://cloud.google.com/iam/docs/roles-audit-logging#scenario_external_auditors
upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: C
Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right
upvoted 1 times

 **haroldbenites** 5 months ago

Go for C
upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: C
Ans: C
upvoted 1 times

 **Cthakker** 1 year, 3 months ago

C is right

upvoted 2 times

 **nana1995** 1 year, 5 months ago

Why is not D?

upvoted 1 times

 **happyBoo** 1 year, 5 months ago

roles/iam.securityReviewer will give access to view IAM roles in addition, which is not required.

upvoted 6 times

 **mcaromit** 1 year, 5 months ago

C is correct

upvoted 1 times

 **meh009** 1 year, 6 months ago

Easy - C

Domain Restricted Sharing is on (restrict the set of identities that are allowed to be used in Identity and Access Management policies) which takes out A/B of the equation.

Easy choice between C/D

upvoted 2 times

 **[Removed]** 1 year, 7 months ago

C is correct. Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.

upvoted 1 times

 **yuv169** 1 year, 7 months ago

option C is correct. because viewer role grants viewing all the resources in the project whereas securityreviewer role only grants permission to list the resources

<https://cloud.google.com/iam/docs/understanding-roles#iam-roles>

upvoted 1 times

 **SSunny** 1 year, 8 months ago

C- <https://cloud.google.com/iam/docs/understanding-custom-roles>

The IAM Security Reviewer role (roles/iam.securityReviewer) enables the ability to view custom roles but not administer them.

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

C - Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

Answer C

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

Answer C

upvoted 1 times

Question #112

You have a workload running on Compute Engine that is critical to your business. You want to ensure that the data on the boot disk of this workload is backed up regularly. You need to be able to restore a backup as quickly as possible in case of disaster. You also want older backups to be cleaned automatically to save on cost. You want to follow Google-recommended practices. What should you do?

- A. Create a Cloud Function to create an instance template.
- B. Create a snapshot schedule for the disk using the desired interval.
- C. Create a cron job to create a new disk from the disk using gcloud.
- D. Create a Cloud Task to create an image and export it to Cloud Storage.

Correct Answer: B*Community vote distribution*

B (100%)

 **DarioFama23** Highly Voted 2 years, 4 months ago

B is correct for this question

upvoted 20 times

 **stepkurniawan** 2 years, 2 months ago

Question: One cannot delete the old disk when using snapshot, right?

upvoted 3 times

 **Ale1973** 2 years, 2 months ago

Snapshots and disks are independent objects on GCP, you could create a snapshot from disk and then delete the disk, the snapshot will stay in place. Actually, you could use this snapshot to create a new disk, assign to another VM, mount it, and use it (all the information that the original disk had at the time of the snapshot will still be there).

upvoted 6 times

 **Ridhanya** 11 months, 1 week ago

In snapshot schedule, there is autodelete and you can specify the days after which auto delete can happen

upvoted 5 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer (B):

Best practices for persistent disk snapshots

You can create persistent disk snapshots at any time, but you can create snapshots more quickly and with greater reliability if you use the following best practices.

Creating frequent snapshots efficiently

Use snapshots to manage your data efficiently.

Create a snapshot of your data on a regular schedule to minimize data loss due to unexpected failure.

Improve performance by eliminating excessive snapshot downloads and by creating an image and reusing it.

Set your snapshot schedule to off-peak hours to reduce snapshot time.

Snapshot frequency limits

Creating snapshots from persistent disks

You can snapshot your disks at most once every 10 minutes. If you want to issue a burst of requests to snapshot your disks, you can issue at most 6 requests in 60 minutes.

If the limit is exceeded, the operation fails and returns the following error:

<https://cloud.google.com/compute/docs/disks/snapshot-best-practices>

upvoted 18 times

 **abirroy** Most Recent 3 months ago

Selected Answer: B

Create a snapshot schedule for the disk using the desired interval.

upvoted 1 times

 **csrazdan** 4 months ago

Selected Answer: B

Snapshot is a better option because they are incremental and you can configure them to consolidate and delete snapshots that are not required for recovery. Image can also provide this functionality but the image is full backup which is inefficient in cases where the content of the file system is changing frequently.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 2 times

 **haroldbenites** 5 months ago

Go for B

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: B

Ans: B

upvoted 2 times

 **ryzior** 8 months, 1 week ago

say no more:

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>

"Use snapshot schedules as a best practice to back up your Compute Engine workloads."

upvoted 4 times

 **alaahakim** 11 months, 3 weeks ago

The right Ans is : B

upvoted 1 times

 **mcaromit** 1 year, 5 months ago

B is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Create a snapshot schedule for the disk using the desired interval.

upvoted 1 times

 **cloud_guru** 1 year, 7 months ago

B is correct due to snapshots being the perfect solution to the problem question

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B. Create a snapshot schedule for the disk using the desired interval.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

B is Correct

upvoted 2 times

 **Vikash211982** 1 year, 9 months ago

B is correct

upvoted 1 times

 **Bhagirathi** 1 year, 10 months ago

B 200%

upvoted 2 times

 **hiteshrup** 2 years, 1 month ago

Answer B is correct. Snapshot can be best option available from listed option

- Option A is just creating instance template and not taking boot disk backup
- Option C is to create a disk, but to run this cron job, you need another gcloud resource which is not cost effective and introducing another single failure point.
- Option D is not a definite not a good cost effective solution.

upvoted 3 times

Question #113

You need to assign a Cloud Identity and Access Management (Cloud IAM) role to an external auditor. The auditor needs to have permissions to review your Google Cloud Platform (GCP) Audit Logs and also to review your Data Access logs. What should you do?

- A. Assign the auditor the IAM role roles/logging.privateLogViewer. Perform the export of logs to Cloud Storage.
- B. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
- C. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Perform the export of logs to Cloud Storage.
- D. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Direct the auditor to also review the logs for changes to Cloud IAM policy.

Correct Answer: C*Community vote distribution*

B (100%)

 **ESP_SAP**  2 years, 2 months ago

Correct Answer is (B):

Background

Google Cloud provides Cloud Audit Logs, which is an integral part of Cloud Logging. It consists of two log streams for each project: Admin Activity and Data Access.

Admin Activity logs contain log entries for API calls or other administrative actions that modify the configuration or metadata of resources. Admin Activity logs are always enabled. There is no charge for your Admin Activity audit logs.

Data Access logs record API calls that create, modify, or read user-provided data. Data Access audit logs are disabled by default because they can be large.

logging.viewer: The logging.viewer role gives the security admin team the ability to view the Admin Activity logs.

logging.privateLogViewer : The logging.privateLogViewer role gives the ability to view the Data Access logs.

upvoted 48 times

 **ESP_SAP** 2 years, 2 months ago

Correct Answer is (B): (Continuation).

Scenario: External auditors

In this scenario, audit logs for an organization are aggregated and exported to a central sink location. A third-party auditor is granted access several times a year to review the organization's audit logs. The auditor is not authorized to view PII data in the Admin Activity logs.

During normal access, the auditors' Google group is only granted access to view the historic logs stored in BigQuery. If any anomalies are discovered,

the group is granted permission to view the actual Cloud Logging Admin Activity logs via the dashboard's elevated access mode. At the end of each audit period, the group's access is then revoked.

Data is redacted using Cloud DLP before being made accessible for viewing via the dashboard application.

upvoted 20 times

 **ESP_SAP** 2 years, 2 months ago

Correct Answer is (B): (Continuation).

The table below explains IAM logging roles that an Organization Administrator can grant to the service account used by the dashboard, as well as the resource level at which the role is granted:

logging.viewer Organization Dashboard service account The logging.viewer role permits the service account to read the Admin Activity logs in Cloud Logging.

bigquery.dataViewer BigQuery dataset Dashboard service account The bigquery.dataViewer role permits the service account used by the dashboard application to read the exported Admin Activity logs.

upvoted 17 times

 **DarioFama23**  2 years, 4 months ago

for me B is the correct answer..

upvoted 17 times

 **Eshkrkrkr** 1 year, 12 months ago

Yes, B is correct because:

- 1) Question doesn't ask us to export and store logs for any long period of time.

2) Custom role with only logging.privateLogEntries.list permission won't let the auditor to access Log Explorer at all (https://cloud.google.com/logging/docs/access-control#console_permissions - Minimal read-only access: logging.logEntries.list)
upvoted 5 times

 **AzureDP900** [Most Recent] 4 months, 2 weeks ago

B is right. Similar practice question in tutorials dojo
upvoted 1 times

 **Rutu_98** 5 months, 2 weeks ago

Selected Answer: B

B is correct ans
upvoted 1 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: B

B is the correct answer
upvoted 3 times

 **lazyabhi606** 10 months, 3 weeks ago

Selected Answer: B

Correct Answer is (B)
upvoted 1 times

 **maggieli** 1 year ago

Correct Answer is B.
upvoted 1 times

 **ankatsu2010** 1 year, 1 month ago

A is the correct answer. Exporting logging data to Cloud Storage is ideal, and 'Cloud IAM Policy' is not mentioned in this question.
upvoted 2 times

 **mcaromit** 1 year, 5 months ago

B is correct
upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
upvoted 1 times

 **yubi69** 1 year, 7 months ago

answer is B
upvoted 2 times

 **Sathy22** 1 year, 7 months ago

why here cloud storage is mentioned ? they are mentioning only access and why this is coming in the middle
upvoted 3 times

 **EABDAJA** 1 year, 8 months ago

B is correct
upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B - Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
upvoted 2 times

 **victory108** 1 year, 9 months ago

B - Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
upvoted 1 times

 **Morgan91** 1 year, 11 months ago

A is correct
roles/logging.privateLogViewer (Private Logs Viewer) includes roles/logging.viewer, plus the ability to read Access Transparency logs and Data Access audit logs. This role applies only to the _Required and _Default buckets.
upvoted 2 times

 **Bhagirathi** 1 year, 11 months ago

it is B .
No other choices to make.
upvoted 1 times

Question #114

You are managing several Google Cloud Platform (GCP) projects and need access to all logs for the past 60 days. You want to be able to explore and quickly analyze the log contents. You want to follow Google-recommended practices to obtain the combined logs for all projects. What should you do?

- A. Navigate to Stackdriver Logging and select resource.labels.project_id="*"
- B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.
- C. Create a Stackdriver Logging Export with a Sink destination to Cloud Storage. Create a lifecycle rule to delete objects after 60 days.
- D. Configure a Cloud Scheduler job to read from Stackdriver and store the logs in BigQuery. Configure the table expiration to 60 days.

Correct Answer: B

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

Community vote distribution

B (100%)

 **Verve** Highly Voted 2 years, 3 months ago

Its B.

upvoted 23 times

 **[Removed]** Highly Voted 2 years, 1 month ago

The question is to view log past 60 days. B, c, D talks about deleting an object or truncation of table data

upvoted 10 times

 **[Removed]** 2 years, 1 month ago

Answer should be A

upvoted 2 times

 **[Removed]** 2 years, 1 month ago

Also A specifically talks about aggregation

upvoted 3 times

 **[Removed]** 2 years, 1 month ago

Also by default, you have a lot of flexibility when viewing logging in stack driver , to filter and query.

upvoted 2 times

 **xian2900** 2 years, 1 month ago

what about minimum retention is 30 days ? is it true ?

upvoted 3 times

 **[Removed]** 2 years, 1 month ago

Ur correct so minimally is 30 for data access logs <https://cloud.google.com/logging/quotas>
then B is the way to go.

upvoted 3 times

 **sai_learner** Most Recent 3 months, 4 weeks ago

All options are wrong , they are talking about deletion after 60 days, but questions asks us to analyse logs of past 60 days

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

B is right for sure

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: B

B is the correct answer.

upvoted 2 times

 **dtncl** 1 year ago

I believe B is the answer.

All that matters in this scenario is the logs for the past 60 days.

We can use BigQuery to analyze contents so C is incorrect. We need to configure a BQ as the sink for the logs export so we can query and analyze log data in the future. Therefore D is incorrect.

<https://cloud.google.com/logging/docs/audit/best-practices#export-best-practices>

Since we only care about the logs within 60 days, we can set the expiration time to 60 to retain only the logs within that time frame. Once data is beyond 60 days old, it wouldn't be included in future analyzations.

https://cloud.google.com/bigquery/docs/managing-tables#updating_a_tables_expiration_time

upvoted 6 times

 **ryzior** 8 months, 1 week ago

I think here we have the case described in details:

<https://cloud.google.com/architecture/exporting-stackdriver-logging-for-security-and-access-analytics>

upvoted 1 times

 **ankatsu2010** 1 year, 1 month ago

D should be the correct answer. To 'quickly analyze', you need to use BQ, next, you always need access to the logs 'for past 60days'. This means you have to export logs on a daily basis. You don't want to do this job manually right?

upvoted 1 times

 **ankatsu2010** 1 year, 1 month ago

My apologies, B is correct... 'Sink' can route logging data to BQ automatically.

upvoted 3 times

 **AD_0525** 1 year, 4 months ago

B is the correct one, option A does not give you the flexibility to analyze.

upvoted 3 times

 **mcaromit** 1 year, 5 months ago

B is correct as analysis of the log contents is a key requirement

upvoted 2 times

 **tifo16** 1 year, 6 months ago

Firstly i thought that A was correct but when considering "default" retention period which is 30 days, i go for B.

<https://cloud.google.com/blog/products/it-ops/best-practices-for-working-with-google-cloud-audit-logging>

Log entries are held in Stackdriver Logging for a limited time known as the retention period. After that, the entries are deleted. To keep log entries longer, you need to export them outside of Stackdriver Logging by configuring log sinks.

upvoted 4 times

 **Crad** 1 year, 6 months ago

I think it's C.

If you set expiration date to 60days on a table then that table will be deleted after 60 days of creation.

Let's say you created some logs on 31st day - they'll be gone within 29 days and not 60.

upvoted 1 times

 **Crad** 1 year, 6 months ago

The output is in JSON so it's still possible to analyze the logs using BigQuery easily as an external source of data.

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

B is correct. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 1 times

 **GCP_Student1** 1 year, 7 months ago

B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 4 times

 **cloud_guru** 1 year, 7 months ago

The answer is B because that's the simplest solution. In order to view AND analyze logs quickly as question asks, we need it to be in BigQuery

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

B is correct

upvoted 1 times

 **JKRowlings** 1 year, 10 months ago

Ans is B. <https://cloud.google.com/solutions/exporting-stackdriver-logging-for-security-and-access-analytics>

upvoted 2 times

 **Bhagirathi** 1 year, 11 months ago

It is B. that best fits.

upvoted 1 times

Question #115

You need to reduce GCP service costs for a division of your company using the fewest possible steps. You need to turn off all configured services in an existing GCP project. What should you do?

- A. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.
- B. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.
- C. 1. Verify that you are assigned the Organizational Administrator IAM role for this project. 2. Locate the project in the GCP console, enter the project ID and then click Shut down.
- D. 1. Verify that you are assigned the Organizational Administrators IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.

Correct Answer: C*Community vote distribution*

A (100%)

✉  **DarioFama23** Highly Voted 2 years, 4 months ago

for me is A the correct answer
upvoted 41 times

✉  **shafiqeee1** Highly Voted 2 years, 3 months ago

A - I reproduced in my project
upvoted 15 times

✉  **patashish** Most Recent 4 months ago

A is right
Hint : You need to turn off all configured services in an ***existing GCP project***.
So C and D out from selection
upvoted 1 times

✉  **RanjithK** 4 months, 1 week ago

Selected Answer: A
Tried and tested
upvoted 1 times

✉  **keep_it_on** 4 months, 1 week ago

correct ans is A
upvoted 2 times

✉  **taiyi078** 4 months, 2 weeks ago

<https://cloud.google.com/run/docs/tutorials/gcloud>
Clean up

In the dialog, type the project ID, and then click Shut down to delete the project.
upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A is right
upvoted 2 times

✉  **Rutu_98** 5 months, 2 weeks ago

Selected Answer: A
It's A
Because we have to follow the least required permission.
Also here given is that they want to reduce the services for a GCP project so the Project Owner role would be sufficient.
upvoted 3 times

✉  **bigbenben** 6 months ago

Selected Answer: A
answer is A.
upvoted 1 times

 **rikku33** 6 months, 2 weeks ago

Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.

upvoted 1 times

 **crisyeb** 7 months, 4 weeks ago

Selected Answer: A

A right

upvoted 1 times

 **dishum** 9 months, 3 weeks ago

Answer is C

The question says ' reduce GCP service costs for a division of your company' means the authority person is going to take few more actions on other projects to reduce costs.

So the right permission to have is at the organisational level.

upvoted 3 times

 **rachee** 11 months, 1 week ago

A - <https://cloud.google.com/resource-manager/docs/creating-managing-projects>

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

The Right Ans is : A

upvoted 2 times

 **maggieli** 1 year ago

Correct Answer is A.

upvoted 1 times

 **Rothmansua** 1 year ago

But.. there is no "Project Owner" role. There is a basic "Owner" role.

There is also no "Organizational Administrator" role, there is "Organization Administrator" role instead. And that role doesn't have project modifying permissions.

upvoted 3 times

 **sudav** 1 year, 2 months ago

A - The primitive Project Owner role provides permissionst to delete project

https://cloud.google.com/iam/docs/understanding-roles#primitive_roles

You can shut down projects using the Cloud Console. When you shut down a project, this immediately happens: All billing and traffic serving stops, You lose access to the project, The owners of the project will be notified and can stop the deletion within 30 days, The project will be scheduled to be deleted after 30 days. However, some resources may be deleted much earlier.

upvoted 3 times

Question #116

You are configuring service accounts for an application that spans multiple projects. Virtual machines (VMs) running in the web-applications project need access to BigQuery datasets in `crm-databases-proj`. You want to follow Google-recommended practices to give access to the service account in the web-applications project. What should you do?

- A. Give `project owner` for web-applications appropriate roles to `crm-databases-proj`.
- B. Give `project owner` role to `crm-databases-proj` and the web-applications project.
- C. Give `project owner` role to `crm-databases-proj` and `bigrquery.dataViewer` role to web-applications.
- D. Give `bigrquery.dataViewer` role to `crm-databases-proj` and appropriate roles to web-applications.

Correct Answer: C

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

Community vote distribution

D (59%)	C (41%)
---------	---------

✉  **ezat** [Highly Voted] 2 years, 4 months ago

D cuz u just need read for DB at the other project
upvoted 27 times

✉  **DarioFama23** 2 years, 4 months ago

U re right, D is the correct answee
upvoted 4 times

✉  **tavva_prudhvi** 1 year, 7 months ago

See the option correctly, as the web app needs access to the big query datasets we have to give access to the web app the data viewer role to only read the datasets! Hence, C
upvoted 6 times

✉  **AnasAloan84** 1 year, 10 months ago

Question didn't specify if the required access is Read only or more, its saying "access" which could be write permissions as well. I will go with C
upvoted 2 times

✉  **AnasAloan84** 1 year, 10 months ago

Ur right, it D. why to give "project owner" as stated on C. correct answer is D
upvoted 5 times

✉  **GCPACE2020** 1 year, 4 months ago

but why giving `bigrquery.dataViewer` to `crm-databases-proj`. we should give for web-application.
upvoted 7 times

✉  **S_marquez** 7 months, 3 weeks ago

You can technically give `bigrquery.dataviewer` to `crm-databases-proj` service account then create a Key and use that key on the VMs, there for making it correct to use D as answer but is way to dumb I would prefer C BUUUUUUT WHY would I give Project Owner to `crm-databases-proj`? they really do not evaluate your knowledge
upvoted 3 times

✉  **dtncl** 1 year ago

It is D because you're right, the question doesn't specify any specific kind of access, however, we need to follow the principle of least-privilege. Hence, we can only assume that read-only access is needed.

`bigrquery.dataViewer` should be assigned to the group of analysts in the `crm-databases-proj` project.

https://cloud.google.com/bigquery/docs/access-control-examples#read_access_to_data_in_a_different_project
upvoted 9 times

✉  **DarioFama23** [Highly Voted] 2 years, 4 months ago

C is correct..
upvoted 11 times

✉  **BigQuery** 7 months, 4 weeks ago

THAT SO DUM
upvoted 3 times

✉  **GCPACE2020** 1 year, 4 months ago

But why giving project owner role to `crm-databases-proj` ?

upvoted 2 times

 **Untamables** (Most Recent) 2 weeks, 3 days ago

Selected Answer: C

All options are hard to understand.

Only C mentions bigquery.dataViewer for web-applications

upvoted 1 times

 **AwesomeGCP** 3 weeks, 6 days ago

Selected Answer: D

No doubt about this Question!!!

Answer is simple and its D. Give bigquery.dataViewer role to crm-databases-proj and appropriate roles to web-applications.

END OF THE DISCUSSION NOW!!!

upvoted 1 times

 **ale_brd_** 1 month, 2 weeks ago

Selected Answer: D

Option D is correct: as per google best practices you should not give primitive roles to a service account, this already exclude as options the answer A to C that give too many rights to the service account as Project owners.

upvoted 1 times

 **theBestStudent** 3 months ago

Selected Answer: D

There are some stuff to assume based on the question:

1- You are configuring Service accounts for projects that span multiple projects, therefore, assume your service account in project X is at least "accepted" in project Y.

2- That being said, make sure the project Y allows SA in project X every time it hits project Y to have the action of having access to big query enabled. Therefore you have to configure that role in project Y (project Y = crm-databases-proj)

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: C

Give project owner role to crm-databases-proj and bigquery.dataViewer role to web-applications.

upvoted 1 times

 **sonuricky** 3 months, 1 week ago

C is the right answer

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is right

upvoted 1 times

 **zaxma** 7 months, 1 week ago

Selected Answer: D

In my opinion, both C and D are wrong.

C - you do not need to give the owner role, it has no meaning and over power

D - the viewer is grant to the wrong project

Then C still works, but give extra permission

D does not work, but no damage to the existing security

God help me here. I will choose D, just based on instinct

upvoted 3 times

 **rsuresh27** 6 months, 2 weeks ago

This is the correct way of thinking about it. You never want to give project owner role to a service account. Even though D does not grant the role to the correct project, it is much safer from security standpoint to do with D.

upvoted 3 times

 **bobthebuilder55110** 3 months, 1 week ago

This does not makes sense, If it is not giving permissions to correct project then simply it will NOT work and what is the point of having it if it will not work. Yes you are giving extra permissions but at the least you are getting the job done in option D you are securing but no Performing the Job, what good would that be ?

upvoted 2 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: D

Ans: D

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: C

Ans: C

Because providing bigquery.dataViewer role to crm-databases-proj doesn't full fill the requirement. basically, the web application needs this access. So its 'C'

upvoted 3 times

 **Rukman** 7 months, 3 weeks ago

I am really sorry folks, the ans should be 'D'

Basically the webapplication [VM(service-account)] need access to read bigquery data in crm-databases-proj project. so, we need to grant the service account the bigquery role on crm-databases-proj. so D is the right answer.

upvoted 2 times

 **Rukman** 7 months, 3 weeks ago

also google doesn't recommend to provide primitive role for service account!

upvoted 2 times

 **bobthebuilder55110** 3 months, 1 week ago

Where do you see the mention of service account in the option D, it say "to srm-databases-proj" NOT service account of that project?

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: D

D would be the perfect choice.

upvoted 1 times

 **FernandoJ** 11 months, 2 weeks ago

Selected Answer: C

<https://cloud.google.com/bigquery/docs/access-control?hl=en#bigquery.dataViewer>

upvoted 2 times

 **FernandoJ** 11 months, 2 weeks ago

The correct option is "D"

upvoted 3 times

 **GreenTick** 11 months, 3 weeks ago

why not A ? give the web app's project owner the required permission to run biq query on crm-database-proj. The scope of the question is about cross project access. It do not ask how crm resources access the big query,

upvoted 1 times

 **jabrrJ68w02ond1** 11 months, 3 weeks ago

The project Owner is mentioned nowhere in the question itself. It only says that the VM instances need 'some' access to the data in other projects. Following least privilege, you let the VM instances access the data read-only.

upvoted 1 times

 **abbottWang** 11 months, 3 weeks ago

Selected Answer: D

you just need read DB permission

upvoted 2 times

 **Mandy** 1 year ago

Correct answer must be D .. Give bigquery.dataViewer role to crm-databases-proj and appropriate roles to web-applications.

upvoted 1 times

Question #117

An employee was terminated, but their access to Google Cloud was not removed until 2 weeks later. You need to find out if this employee accessed any sensitive customer information after their termination. What should you do?

- A. View System Event Logs in Cloud Logging. Search for the user's email as the principal.
- B. View System Event Logs in Cloud Logging. Search for the service account associated with the user.
- C. View Data Access audit logs in Cloud Logging. Search for the user's email as the principal.
- D. View the Admin Activity log in Cloud Logging. Search for the service account associated with the user.

Correct Answer: B*Community vote distribution*

C (100%)

 **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: C

View Data Access audit logs in Cloud Logging. Search for the user's email as the principal

upvoted 2 times

 **Rutu_98** 5 months, 2 weeks ago

Selected Answer: C

ANSWER IS C

As we want to find out whether the user has accessed the data or not , so Data Acess Logs would be correct option to view that

upvoted 2 times

 **akshaychavan7** 5 months, 2 weeks ago

I will go with option C.

upvoted 1 times

 **bigbenben** 6 months ago

Selected Answer: C

C. <https://cloud.google.com/logging/docs/audit#data-access>

upvoted 2 times

 **rsuresh27** 6 months, 2 weeks ago

Selected Answer: C

C is the correct answer. We are trying to find out if any sensitive data was accessed. Data access logs are the only logs that show this. C is the only option that mentions data access logs.

upvoted 4 times

 **rtnk22** 6 months, 2 weeks ago

Shouldn't the correct option be A here?

What does service account have to do here.

upvoted 2 times

 **amrith501** 6 months, 2 weeks ago

It should be C User activity generally come under audit logs

upvoted 3 times

Question #118

Topic 1

You need to create a custom IAM role for use with a GCP service. All permissions in the role must be suitable for production use. You also want to clearly share with your organization the status of the custom role. This will be the first version of the custom role. What should you do?

- A. Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- B. Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to BETA while testing the role permissions.
- C. Use permissions in your role that use the 'testing' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- D. Use permissions in your role that use the 'testing' support level for role permissions. Set the role stage to BETA while testing the role permissions.

Correct Answer: C

Community vote distribution

A (100%)

  **raksteer**  2 years, 3 months ago

You need a custom role with permissions supported in prod and you want to publish the status of the role.

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

You can't use TESTING as it is not good for prod. And you need first version which should be ALPHA. Answer should be A.

upvoted 70 times

  **passmepls** 2 years, 2 months ago

good job

upvoted 2 times

  **BigQuery** 7 months, 4 weeks ago

WAY TO GO. VERY CLEAR EXP INDEED

upvoted 1 times

  **ESP_SAP**  2 years, 2 months ago

Correct Answer is (A):

Testing and deploying

Custom roles include a launch stage, which is stored in the stage property for the role. The launch stage is informational; it helps you keep track of whether each role is ready for widespread use.

Each custom role can have one of the following launch stages:

Launch stages

ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

BETA The role has been tested on a limited basis, or it includes permissions for Google Cloud services or features that are not generally available.

GA The role has been widely tested, and all of its permissions are for Google Cloud services or features that are generally available.

upvoted 29 times

  **ESP_SAP** 2 years, 2 months ago

Correct Answer is (A): Continuation

Support levels for permissions in custom roles

You can include many, but not all, Identity and Access Management (IAM) permissions in custom roles. Each permission has one of the following support levels:

Support level Description

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

The first version of the Custom Role is ALPHA then suitable to productions all permissions in "Supported"...

upvoted 13 times

 **GCP_Student1** 1 year, 8 months ago

ESP_SAP

There is a discrepancy between your first post and the second post. Compare these two sentences;

1st POST - ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

2nd POST - SUPPORTED The permission is fully supported in custom roles.

Are you still going to go with A ?

upvoted 2 times

 **learn_GCP** 1 month, 2 weeks ago

Here ALPHA is for Google cloud feature, only informational. given to identify whether the feature is fully available as a service. and SUPPORTED -- is for a custom role which is supported by Google cloud, meaning any support is provided by Google cloud

upvoted 1 times

 **theBestStudent** Most Recent 3 months ago

Selected Answer: A

It must be suitable for production so Supported permissions only. Plus, it is your first version of the custom role, so you need to check if all is good, then ALPHA.

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: A

Use permissions in your role that use the 'supported' support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **Razor** 9 months, 2 weeks ago

Selected Answer: A

A is the only right solution.

upvoted 2 times

 **sid0127** 11 months, 2 weeks ago

Selected Answer: A

A is the correct answer

upvoted 2 times

 **suday** 1 year, 2 months ago

A -

SUPPORTED -The permission is fully supported in custom roles.

role stage, the stage transitions from ALPHA → BETA → GA

The only option that satisfies “ALPHA” stage with “SUPPORTED” support level is

Use permissions in your role that use the SUPPORTED support level for role permissions. Set the role stage to ALPHA while testing the role permissions

upvoted 3 times

 **mcaromit** 1 year, 5 months ago

A is correct

upvoted 1 times

 **[Removed]** 1 year, 7 months ago

A is correct. Use permissions in your role that use the “supported” support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

 **yubi69** 1 year, 7 months ago

answer is A. because it contains SUPPORTED which we must see when creating custom roles and as it first version we must set it to ALPHA
upvoted 1 times

 **cloud_guru** 1 year, 7 months ago

Looks like A to be because the role is the first iteration and still under testing, we need to select 'Alpha' for the custom role. Additionally, we shouldn't use 'testing' permissions in the role after they said not to

upvoted 1 times

 **SSunny** 1 year, 8 months ago

A:

Supporting Points:

- a. The status of the custom rule is to be presented to the organization. And since this is yet in testing state, the custom rule state will be Alpha.
 - b. The Support level, must be SUPPORTED so as to be used in production.
- If the support level is at TESTING, then in that case - the custom role must not be used as its results may change and unexpected behaviour may be expected.

SUPPORTED The permission is fully supported in custom roles.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

NOT_SUPPORTED The permission is not supported in custom roles.

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

upvoted 3 times

 **EABDAJA** 1 year, 8 months ago

A is correct.

upvoted 1 times

 **JackGlemins** 1 year, 8 months ago

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

SUPPORTED or TESTING is by Google, and refer to the permissions (Not the custom role)

Hence C & D Is wrong: Not recommended for production use.

TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

Now A & B is tricky because is the stages of the Custom role: https://cloud.google.com/iam/docs/understanding-custom-roles#testing_and_deploying

The Key word is: "This will be the first version of the custom role"

To me A is the correct:

ALPHA: The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

BETA: The role has been tested on a limited basis.

It is the first version it cant be tested yet. I think B is wrong.

upvoted 2 times

 **victory108** 1 year, 9 months ago

A - Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

 **DucSiu** 1 year, 9 months ago

A. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.

upvoted 1 times

Question #119

Your company has a large quantity of unstructured data in different file formats. You want to perform ETL transformations on the data. You need to make the data accessible on Google Cloud so it can be processed by a Dataflow job. What should you do?

- A. Upload the data to BigQuery using the bq command line tool.
- B. Upload the data to Cloud Storage using the gsutil command line tool.
- C. Upload the data into Cloud SQL using the import function in the console.
- D. Upload the data into Cloud Spanner using the import function in the console.

Correct Answer: B

Reference:

<https://cloud.google.com/solutions/performing-etl-from-relational-database-into-bigquery>*Community vote distribution*

B (100%)

✉  **DarioFama23** [Highly Voted] 2 years, 4 months ago

B looks correct. Key work unstructured data
upvoted 22 times

✉  **obeythefist** 8 months, 1 week ago

Also "different" file formats, this further supports B as the correct choice.
upvoted 1 times

✉  **Untamables** [Most Recent] 2 weeks, 3 days ago

Selected Answer: B
Cloud Storage as a datalake
upvoted 1 times

✉  **VietmanOfficiel** 2 months, 2 weeks ago

Answer B :
"large quantity" : Cloud Storage or BigQuery
"files" a file is nothing but an Object

So Cloud Storage is the better option.

upvoted 1 times

✉  **abirroy** 3 months ago

Selected Answer: B
Upload the data to Cloud Storage using the gsutil command line tool.
upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

B looks correct
upvoted 1 times

✉  **pfabio** 5 months, 3 weeks ago

Selected Answer: B
For unstructured data use cloud storage. Use Big Query for analytics, data warehouse with structured data
upvoted 3 times

✉  **PAUGURU** 6 months, 2 weeks ago

Selected Answer: B
B looks correct
upvoted 2 times

✉  **Akash7** 7 months ago

Answer is B, Cloud Storage for unstructured data
upvoted 1 times

✉  **Rukman** 7 months, 3 weeks ago

Selected Answer: B
Ans: B

We can upload unstructured data to Cloud storage not to bigquery.

upvoted 2 times

✉️ **associatecloudexamuser** 1 year, 3 months ago

B is correct. Since GCS is a datalake

upvoted 2 times

✉️ **mcaromit** 1 year, 5 months ago

B is correct

upvoted 1 times

✉️ **[Removed]** 1 year, 7 months ago

B is correct. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 2 times

✉️ **cloud_guru** 1 year, 7 months ago

B because it's unstructured file data and needs to be ETL transformed

upvoted 2 times

✉️ **GCP_Student1** 1 year, 8 months ago

B. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 3 times

✉️ **Bhagirathi** 1 year, 10 months ago

B 2000%

upvoted 2 times

✉️ **swatititame** 1 year, 11 months ago

• B. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 1 times

✉️ **BabaMat** 2 years ago

Answer is B

upvoted 1 times

Question #120

You need to manage multiple Google Cloud projects in the fewest steps possible. You want to configure the Google Cloud SDK command line interface (CLI) so that you can easily manage multiple projects. What should you do?

- A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- B. 1. Create a configuration for each project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project
- C. 1. Use the default configuration for one project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- D. 1. Use the default configuration for one project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project.

Correct Answer: D*Community vote distribution*

A (100%)

S  **SSunny**  1 year, 8 months ago

A

Cloud SDK comes with a default configuration. To create multiple configurations, use gcloud config configurations create, and gcloud config configurations activate to switch between them.

<https://cloud.google.com/sdk/gcloud/reference/config/set>

upvoted 21 times

S  **GCP_Student1**  1 year, 8 months ago

A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 7 times

S  **VietmanOfficial**  2 months, 2 weeks ago**Selected Answer: A**

1. Generate your configurations with "gcloud config configurations create <config_id> ..." then activate the one you need according to the project you are working on with "gcloud config activate <config_id>"

upvoted 1 times

S  **abirroy** 3 months ago**Selected Answer: A**

Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 1 times

S  **AzureDP900** 4 months, 2 weeks ago

A is right.

gcloud config set project PROJECT_ID

<https://cloud.google.com/sdk/gcloud/reference/config/set>

upvoted 1 times

S  **Rutu_98** 5 months, 2 weeks ago**Selected Answer: A**

A is correct Answer

upvoted 1 times

S  **pfabio** 5 months, 3 weeks ago**Selected Answer: A**

You have to create a config for each project and activate to use it

upvoted 1 times

S  **luciorifa** 8 months, 2 weeks ago**Selected Answer: A**

A is the correct answer

upvoted 1 times

S  **Raju44** 10 months, 2 weeks ago

why we need to create config, when project is created so its config .

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: A

A is the right option. Multiple configurations. Activate to switch between configurations.

upvoted 1 times

 **Ridhanya** 11 months, 1 week ago

A, C seems hectic, B and D surely are eliminated

upvoted 1 times

 **rachee** 11 months, 1 week ago

A - <https://cloud.google.com/sdk/docs/configurations>

upvoted 1 times

 **sid0127** 11 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Correct Ans: A

upvoted 1 times

 **arsh1916** 1 year, 5 months ago

A 100%

upvoted 2 times

 **mcaromit** 1 year, 5 months ago

A is correct

upvoted 2 times

 **ri_unhou119** 1 year, 6 months ago

Why not C?

I think C is correct.

https://cloud.google.com/sdk/docs/configurations?hl=ja#default_configuration

upvoted 1 times

 **bubblegumbeach** 1 year ago

C doesnt tell anything about creating another project

upvoted 1 times

 **csrazdan** 4 months ago

You cannot activate a configuration unless it is created first. Option C gives an impression that configurations are already created and they need to just be activated.

upvoted 1 times

Question #121

Your managed instance group raised an alert stating that new instance creation has failed to create new instances. You need to maintain the number of running instances specified by the template to be able to process expected application traffic. What should you do?

- A. Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.
- B. Create an instance template that contains valid syntax that will be used by the instance group. Verify that the instance name and persistent disk name values are not the same in the template.
- C. Verify that the instance template being used by the instance group contains valid syntax. Delete any persistent disks with the same name as instance names. Set the disks.autoDelete property to true in the instance template.
- D. Delete the current instance template and replace it with a new instance template. Verify that the instance name and persistent disk name values are not the same in the template. Set the disks.autoDelete property to true in the instance template.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>

Community vote distribution

A (62%)	C (38%)
---------	---------

✉️  **TAvenger** Highly Voted 1 year, 8 months ago

Really tricky question.

Ideal scenario would be

1. create new template, while creating ensure that in the new template disks.autoDelete=true, 3. delete existing persistent disks, 4. make rolling update ...

In order to switch to new template we need "Rolling update". Unfortunately, it is not mentioned.

With current options

C - not correct, we cannot update existing template

D - not correct, we cannot delete existing template when it is in use (just checked in GCP) (We need rolling update)

B - will not solve our problem without Rolling update

A - This is the only option (I know that it can be temporary) that will work without Rolling update according to <https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

upvoted 40 times

✉️  **ShakthiGCP** 1 year, 8 months ago

Well reasoned. I'm also going with A.

upvoted 2 times

✉️  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (C):

Your instance template has set the disks.autoDelete option to false for boot persistent disks so that when a VM has been deleted (for example, because of autohealing), the persistent disk was not deleted. When the managed instance group attempted to recreate the VM with the same name, it ran into the same issue where a persistent disk already exists with the same name. Delete the existing persistent disk to resolve the immediate problem and update the instance template to set the disks.autoDelete to true if you would like boot persistent disks to be deleted alongside the instance

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#troubleshooting>

upvoted 38 times

✉️  **magistrum** 1 year, 10 months ago

Can't update instance templates, see below

upvoted 17 times

✉️  **JackGlemins** 1 year, 8 months ago

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

upvoted 11 times

✉️  **Untamables** Most Recent 2 weeks, 2 days ago

Selected Answer: C

Vote C

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

<https://cloud.google.com/sdk/gcloud/reference/compute/instances/set-disk-auto-delete>

upvoted 1 times

✉️  **sneha1607** 2 months, 1 week ago

C is the correct answer.

1. Ensure you don't have any persistent disks with the same name as the VM instance.
2. Ensure the disk autodelete property is turned on (disks.autoDelete set to true).
3. Ensure instance template syntax is valid

upvoted 1 times

 **sneha1607** 2 months, 1 week ago

As described in this article, "My managed instance group keeps failing to create a VM. What's going on?"

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#troubleshooting>

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: A

Answer : A

As many have mentioned here C and D is straight NO, C talks about updating the instance template and D talks about deleting an instance template BUT as per google documentation you cannot delete an instance template nor update it if it is in use. So A is the only one the makes sense here.

https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates

upvoted 4 times

 **gscharly** 3 months, 2 weeks ago

Selected Answer: A

Correct answer is A

You don't have the option to edit an existing instance.

upvoted 1 times

 **rdumitru** 3 months, 3 weeks ago

Selected Answer: A

Correct answer is A

You don't have the option to edit an existing instance.

upvoted 1 times

 **Bumbah** 4 months ago

Selected Answer: A

You cannot update a template, need to create a new one.

upvoted 1 times

 **patashish** 4 months, 1 week ago

A is correct answer. [Instance template resources are immutable in GCP.]

How to update instance templates

Instance templates are designed to create instances with identical configurations. So you cannot update an existing instance template or change an instance template after you create it.

If you need to make changes to the configuration, create a new instance template. You can create a template based on an existing instance template, or based on an existing instance. You can also override instance template fields when creating a VM instance from an instance template.

Ref : <https://cloud.google.com/compute/docs/instance-templates#:~:text=So%20you%20cannot%20update%20an,base%20on%20an%20existing%20instance>.

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: C

C is the correct answer because your instance exists you don't have to recreate it

upvoted 2 times

 **deadlydeb** 4 months, 4 weeks ago

d IS RIGHT

upvoted 1 times

 **pfabio** 5 months, 3 weeks ago

Selected Answer: A

Incorrect:

B - incompleted

C - Instance template cannot be modified

D - you cannot delete a instance template while its in use

upvoted 1 times

 **somenick** 7 months, 3 weeks ago

Selected Answer: C

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

upvoted 2 times

 **ryzior** 8 months, 1 week ago

I tried to reproduce the issue and I managed to prepare a template which creates an instance with the same name the boot disk has - and there is no issue with starting it up - maybe this is new functionality or something but it seems the question is about the template lifecycle not the vm issues. since you are not able to modify the templat but only to "create similar" the only option is to create a new one.

You can't set the disk name but only the device name , I was able to mimic this by setting the vm name during vm name creation (not by mig but manually from the template).

I think this is an old question...

upvoted 1 times

 **AminT** 1 year, 3 months ago

D is correct answer.

upvoted 1 times

 **Blaaa** 1 year, 4 months ago

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>

upvoted 2 times

 **Umesh09** 1 year, 5 months ago

D looks better option C is close but an existing instance template cannot be modified hence my answer would be D

upvoted 2 times

Question #122

Your company is moving from an on-premises environment to Google Cloud. You have multiple development teams that use Cassandra environments as backend databases. They all need a development environment that is isolated from other Cassandra instances. You want to move to Google Cloud quickly and with minimal support effort. What should you do?

- A. 1. Build an instruction guide to install Cassandra on Google Cloud. 2. Make the instruction guide accessible to your developers.
- B. 1. Advise your developers to go to Cloud Marketplace. 2. Ask the developers to launch a Cassandra image for their development work.
- C. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Use the snapshot to create instances for your developers.
- D. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Upload the snapshot to Cloud Storage and make it accessible to your developers. 3. Build instructions to create a Compute Engine instance from the snapshot so that developers can do it themselves.

Correct Answer: D*Community vote distribution*

B (100%)

 **DarioFama23** Highly Voted  2 years, 4 months ago
B is correct for me.. launch a solution from marketplace
upvoted 26 times

 **ESP_SAP** Highly Voted  2 years, 2 months ago
Correct Answer is (B):

<https://medium.com/google-cloud/how-to-deploy-cassandra-and-connect-on-google-cloud-platform-with-a-few-clicks-11ee3d7001d1>
upvoted 20 times

 **xtian2900** 2 years, 1 month ago
thanks, i always look for your insight
upvoted 3 times

 **nightflyer** 1 year, 10 months ago
But we are moving from on premises to gcp
upvoted 1 times

 **Untamables** Most Recent  2 weeks, 2 days ago

Selected Answer: B
B
<https://cloud.google.com/blog/products/databases/open-source-cassandra-now-managed-on-google-cloud>
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago
B absolutely correct, there is no need of manual installs.
upvoted 1 times

 **pfabio** 5 months, 3 weeks ago

Selected Answer: B
B is correct: You want to move to Google Cloud quickly and with minimal support effort.
<https://cloud.google.com/marketplace>
upvoted 1 times

 **obeythefist** 7 months, 4 weeks ago
I suspect every time an answer features Google Marketplace, that is the correct answer. In these kinds of exams the purpose is often just to check general product knowledge.
upvoted 3 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: B
B is the correct answer
upvoted 2 times

 **Raju44** 10 months, 2 weeks ago
Minimum support when you have steps laydown to your developer to do task on their own. So its D.
upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: B

B is the right option.

upvoted 2 times

 **sid0127** 11 months, 2 weeks ago

Selected Answer: B

B. its easier to launch from the marketplace

upvoted 3 times

 **jabrrJ68w02ond1** 11 months, 3 weeks ago

For me it's B, on upfront it offers the lowest effort to launch Marketplace solutions.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: B

B is Correct

upvoted 1 times

 **sudav** 1 year, 2 months ago

B - You can deploy Cassandra as a Service, called Astra, on the Google Cloud Marketplace. Not only do you get a unified bill for all GCP services, your Developers can now create Cassandra clusters on Google Cloud in minutes and build applications with Cassandra as a database as a service without the operational overhead of managing Cassandra

upvoted 4 times

 **arsh1916** 1 year, 5 months ago

B is correct

upvoted 1 times

 **mcaromit** 1 year, 5 months ago

B is correct

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

B is correct

upvoted 2 times

 **GCP_Student1** 1 year, 8 months ago

B. 1. Advise your developers to go to Cloud Marketplace. 2. Ask the developers to launch a Cassandra image for their development work.

upvoted 2 times

Question #123

Topic 1

You have a Compute Engine instance hosting a production application. You want to receive an email if the instance consumes more than 90% of its CPU resources for more than 15 minutes. You want to use Google services. What should you do?

- A. 1. Create a consumer Gmail account. 2. Write a script that monitors the CPU usage. 3. When the CPU usage exceeds the threshold, have that script send an email using the Gmail account and smtp.gmail.com on port 25 as SMTP server.
- B. 1. Create a Cloud Monitoring Workspace and associate your Google Cloud Platform (GCP) project with it. 2. Create a Cloud Monitoring Alerting Policy that uses the threshold as a trigger condition. 3. Configure your email address in the notification channel.
- C. 1. Create a Cloud Monitoring Workspace and associate your GCP project with it. 2. Write a script that monitors the CPU usage and sends it as a custom metric to Cloud Monitoring. 3. Create an uptime check for the instance in Cloud Monitoring.
- D. 1. In Cloud Logging, create a logs-based metric to extract the CPU usage by using this regular expression: CPU Usage: ([0-9]{1,3})%. 2. In Cloud Monitoring, create an Alerting Policy based on this metric. 3. Configure your email address in the notification channel.

Correct Answer: D

Community vote distribution

B (100%)

 **anolive** 3 weeks, 5 days ago

Selected Answer: B

is correct

upvoted 1 times

 **ale_brd_** 1 month, 2 weeks ago

Selected Answer: B

answer is B, but I would write it this way as stackdriver is deprecated and Operation Suite uses scopes now.

1. Create a Cloud Monitoring metric scope and associate your Google Cloud Platform (GCP) project with it.

2. Create a Cloud Monitoring Alerting Policy that uses the threshold as a trigger condition.

3. Configure your email address in the notification channel.

upvoted 2 times

 **Akso** 2 months ago

Selected Answer: B

The dump has been changed some days ago. This answer was the old and best version:

1. Create a Stackdriver Workspace, and associate your Google Cloud Platform (GCP) project with it.

2. Create an Alerting Policy in Stackdriver that uses the threshold as a trigger condition.

3. Configure your email address in the notification channel.

upvoted 4 times

 **thimai** 1 month, 3 weeks ago

And the answer is still the same?

upvoted 1 times

Question #124

You have an application that uses Cloud Spanner as a backend database. The application has a very predictable traffic pattern. You want to automatically scale up or down the number of Spanner nodes depending on traffic. What should you do?

- A. Create a cron job that runs on a scheduled basis to review Cloud Monitoring metrics, and then resize the Spanner instance accordingly.
- B. Create a Cloud Monitoring alerting policy to send an alert to oncall SRE emails when Cloud Spanner CPU exceeds the threshold. SREs would scale resources up or down accordingly.
- C. Create a Cloud Monitoring alerting policy to send an alert to Google Cloud Support email when Cloud Spanner CPU exceeds your threshold. Google support would scale resources up or down accordingly.
- D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

Correct Answer: D*Community vote distribution*

D (100%)

 **GCP_Student1** Highly Voted 1 year, 8 months ago

D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

upvoted 24 times

 **erikamrqz** Highly Voted 1 year, 1 month ago

D 205%

upvoted 5 times

 **bubblegumbeach** 1 year ago

seriously? 205%?

upvoted 3 times

 **theBestStudent** Most Recent 2 months, 4 weeks ago

Selected Answer: D

Without knowing that much, you can discard easily B,C as they don't make any sense. Automation should be a key in this answer. Also you should discard "A" as with a CronJob you won't spann on time as it will be a fixed time checking. So the only one that is left is D, as just creating an alert and sending it to "something else" (in this case a webhook) in an automated way, should be the common sense way of handling this.

upvoted 2 times

 **abirroy** 3 months, 2 weeks ago

Selected Answer: D

Correct answer is D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is correct, It is part of Tutorials Dojo practice test

upvoted 1 times

 **rsuresh27** 6 months, 3 weeks ago

Answer is D. The keyword to look for is "automatically". A, B, C all have steps that are not automatic. Thus, only D is left.

upvoted 3 times

 **alaahakim** 11 months, 3 weeks ago

D is the answer

upvoted 3 times

 **Ricky_92** 12 months ago

<https://cloud.google.com/architecture/autoscaling-cloud-spanner>

upvoted 4 times

 **kukabura** 1 year ago

D is the answer

upvoted 3 times

 **ankatsu2010** 1 year ago

It's a tricky question. The answer is 'A'.

Question says traffic pattern is predictable. This means you can schedule the scalability.

You can achieve this easily by using Cloud Spanner's API.

<https://cloud.google.com/spanner/docs/reference/rest/v1/projects.instances.patch>

upvoted 3 times

✉️  **BobbyFlash** 1 year ago

Really tricky indeed...But if you think about it and you don't stick religiously to "traffic pattern is predictable" sentence (that sentence is to trick you), you would think that despite the pattern is very predictable, it could vary at some point in time. With a cron job you can schedule a job using a fixed starting point and a a fixed ending, can't you? Well, what would happen if the app traffic suddenly spikes before your cron job starts running? Or after?. At this point, I go for D, making scalability automatic thanks to Cloud Functions.

upvoted 5 times

✉️  **jcols** 1 year, 5 months ago

Option D. There's an official repository that does something similar to provide autoscaling to Cloud Spanner.
<https://github.com/cloudspannerecosystem/autoscaler>

upvoted 2 times

✉️  **zaxxon** 1 year ago

Which uses a Cloud Scheduler so answer A

upvoted 1 times

✉️  **ryumada** 3 months ago

"...and then resize the Spanner instance accordingly." I think the last sentence says that the resize action is done manually. If so, then D is still the right answer.

upvoted 1 times

✉️  **ryumada** 3 months ago

For the D option, "Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly." The listen and resize actions is done by Cloud Function.

upvoted 1 times

✉️  **arsh1916** 1 year, 5 months ago

D no doubt

upvoted 2 times

✉️  **mcaromit** 1 year, 5 months ago

D is perfect

upvoted 2 times

✉️  **kopper2019** 1 year, 6 months ago

D the way to go

upvoted 1 times

✉️  **pondai** 1 year, 7 months ago

D

A cron job not what we want.

B Call SRE to do this not automatic

C Google don't help you do this.

upvoted 2 times

✉️  **SSunny** 1 year, 8 months ago

D

Excerpt from - <https://cloud.google.com/spanner/docs/instances>

Note: You can scale the number of nodes in your instance based on the Cloud Monitoring metrics on CPU or storage utilization in conjunction with Cloud Functions.

upvoted 3 times

Question #125

Your company publishes large files on an Apache web server that runs on a Compute Engine instance. The Apache web server is not the only application running in the project. You want to receive an email when the egress network costs for the server exceed 100 dollars for the current month as measured by Google Cloud.

What should you do?

- A. Set up a budget alert on the project with an amount of 100 dollars, a threshold of 100%, and notification type of email.
- B. Set up a budget alert on the billing account with an amount of 100 dollars, a threshold of 100%, and notification type of email.
- C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.
- D. Use the Cloud Logging Agent to export the Apache web server logs to Cloud Logging. Create a Cloud Function that uses BigQuery to parse the HTTP response log data in Cloud Logging for the current month and sends an email if the size of all HTTP responses, multiplied by current Google Cloud egress prices, totals over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.

Correct Answer: A

Community vote distribution

C (78%)	11%	11%
---------	-----	-----

 **GCP_Student1** Highly Voted 1 year, 8 months ago

C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.

upvoted 22 times

 **MyName_** Highly Voted 1 year, 7 months ago

[C]
I think the keyword here is "as measured by Google Cloud". In Answer D you calculate the price yourself, in C you use the billing provided by GCP. Thus I think the Answer is C.

upvoted 10 times

 **Charumathi** Most Recent 1 month ago

Selected Answer: C

C is the correct answer,
You export the bill to BigQuery and filter for the Egress cost for the particular application, and send an email if the cost is over 100 dollars, to send an email you need to use cloud function to monitor and trigger based on the conditions.

upvoted 1 times

 **sylva91** 1 month, 3 weeks ago

Selected Answer: C

it can only be the C because "it's not the only app running"

upvoted 1 times

 **VietmanOfficial** 2 months, 2 weeks ago

Selected Answer: D

[D]
Complex but works for resume network egres cost of ONLY the Apache instance

upvoted 1 times

 **theBestStudent** 2 months, 4 weeks ago

Selected Answer: C

I would say this is the approach without knowing that much:

A & B discarded as they are not resource oriented. We need to charge for the apache server (so focusing on the VM where it is hosted) in order to charge this server.

D can't be as you are not charged in this case for that Response payload received.

upvoted 1 times

 **ryumada** 3 months ago

Selected Answer: C

Answer C and D are correct to me. You can calculate the current Google Cloud egress prices using Cloud billing catalog API which can be used in D option.

ref: <https://cloud.google.com/blog/topics/cost-management/introducing-cloud-billing-catalog-api-gcp-pricing-in-real-time>

But, if you choose option C, then you will get the usage cost directly from cloud billing data. You should add a label to the Apache web server in order to select its cost.

D has the more complex step, using more services too which is Cloud Logging to store the logging data of the VM and Cloud Billing Catalog API. The using of more services makes the D option to have more expensive costs.

upvoted 1 times

 **andreherwanto** 3 months, 3 weeks ago

Selected Answer: C

C is best options

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is best options

upvoted 1 times

 **PAUGURU** 6 months, 1 week ago

Selected Answer: B

C cannot be the right answer, you export the data to big Query only once. What is the use of a cloud function running every hour on the same data? It doesn't say "you export the data to BigQuery hourly".

upvoted 1 times

 **PAUGURU** 6 months, 1 week ago

I was wrong, when you configure an export to Bigquery billing data are updated periodically, so answer is actually C.
<https://cloud.google.com/billing/docs/how-to/export-data-bigquery-tables?hl=en>

upvoted 3 times

 **Leo_li_leo** 7 months, 1 week ago

I feel the answer should be B as projects can be specified and also services.

upvoted 2 times

 **Pitchino80** 9 months, 2 weeks ago

Selected Answer: C

It's C

upvoted 2 times

 **Ridhanya** 11 months ago

C. Even I thought its D initially, but D is a lot of steps while C has minimum steps

upvoted 1 times

 **obeythefist** 8 months, 1 week ago

C and D are not correct, because running a project for a week when it has no load is not an accurate way to estimate costs.

upvoted 1 times

 **JackalHuang** 11 months, 3 weeks ago

"C" is not completely correct, but I think it should be the one closest to the answer.
Because Cloud Scheduler not support Cloud Function, It must integrate the Pub/Sub service.

upvoted 1 times

 **Xumbegnows** 1 year, 4 months ago

How about B?

You can define the scope of the budget. For example, you can scope the budget to apply to the spend in an entire Cloud Billing account, or narrow the scope to one or more projects, and/or one or more products, and/or other budget filters applicable to your Cloud Billing account.

<https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 1 times

 **YAS007** 1 year, 3 months ago

the billing account might contains other projets

upvoted 2 times

 **obeythefist** 8 months, 1 week ago

The difference between "A" and "B" is that "A" uses a Google tool made specifically for price estimation, "B" does not, it's not as good of an answer.

Questions like this are intended to probe candidates knowledge of Google Cloud tools and their usage.

upvoted 1 times

 **obeythefist** 8 months, 1 week ago

Disregard my answer here :)

upvoted 1 times

 **arsh1916** 1 year, 5 months ago

C, remember you need to filter.

upvoted 2 times

 **mcaromit** 1 year, 5 months ago

C is best option given that the web-server is not the only application in the project emitting traffic

upvoted 1 times

Question #126

You have designed a solution on Google Cloud that uses multiple Google Cloud products. Your company has asked you to estimate the costs of the solution. You need to provide estimates for the monthly total cost. What should you do?

- A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.
- B. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Create a Google Sheet that summarizes the expected monthly costs for each product.
- C. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Navigate to the Billing Report page in the Cloud Console. Multiply the 1 week cost to determine the monthly costs.
- D. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Use Cloud Monitoring to determine the provisioned and used resource amounts. Multiply the 1 week cost to determine the monthly costs.

Correct Answer: C*Community vote distribution*

A (100%)

 **GCP_user** Highly Voted  1 year, 7 months ago

Yes I agree with A. Makes more sense.

upvoted 26 times

 **BobbyFlash** 1 year ago

Why not B?? Even though answer A makes sense, they are also stating to provide estimates for "monthly total costs". One would think that it is not only necessary to get estimates from every resource, but also consolidate them to inform the monthly total cost required.

upvoted 1 times

 **BobbyFlash** 1 year ago

Ignore the comment folks. Sentence B is missing the little thing about the pricing calculator. I go with A.

upvoted 5 times

 **GCP_Student1** Highly Voted  1 year, 8 months ago

A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.

upvoted 6 times

 **Untamables** Most Recent  2 weeks, 2 days ago

Selected Answer: A

Vote A

<https://cloud.google.com/free/docs/estimate-costs-google-cloud-platform>

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: A

A is the correct answer, use the pricing calculator to estimate the pricing for a month and download the estimate to csv file, or you can share the URL of the pricing calculator or email the estimate to the respective people in the company.

upvoted 1 times

 **VietmanOfficial** 2 months, 2 weeks ago

Selected Answer: A

[A]

When "estimate" you need to read "price calculator"

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

I agree with A

upvoted 1 times

 **dishum** 9 months, 3 weeks ago

Answer is C

The question says 'cost of solution of the design' means, how it is going to work in practical, means the traffic load, no of users, uploads, downloads, transactions etc etc.

In this case, the more nearer option is to run for a week, then calculate.

Answer A is not correct, becoz A is applicable when the cost of the product is to be determined.

upvoted 3 times

✉ **SunnyDey** 1 day, 16 hours ago

Question has asked to estimate the costs of the solution i.e. the initial setup cost not the running cost. Therefore, C and D discarded and B is of no use without pricing calculator. Question is testing whether you are aware of thes calculator service from Google.

upvoted 1 times

✉ **sid0127** 11 months, 2 weeks ago

Selected Answer: A

answer is a

upvoted 2 times

✉ **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

The Right Ans: A

upvoted 3 times

✉ **TenshiD** 11 months, 3 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

✉ **GCPdump** 1 year, 1 month ago

As per GCP best practices, A makes more sense

upvoted 2 times

✉ **jcloud965** 1 year, 3 months ago

Correct is A. It is the main purpose of Pricing calculator. You have to check pricing details on products pages to understand what charges apply and how to size it.

B. It works bur it is more complicated than A.

C. It works but you need to wait one week and you pay for this. Moreoever, you won't be in production so somes charges may lack like net egress charges

D. No, because Cloud Monitoring can't be used for billing purpose

upvoted 4 times

✉ **tanito83** 1 year, 5 months ago

The correct answer is A. Please, modify it.

upvoted 2 times

✉ **arsh1916** 1 year, 5 months ago

A, is best option

upvoted 1 times

✉ **mcaromit** 1 year, 5 months ago

A is correct

upvoted 1 times

✉ **lio123** 1 year, 8 months ago

it is should be A

upvoted 1 times

✉ **hasnaa** 1 year, 8 months ago

For me the correct answer is A

upvoted 2 times

Question #127

Topic 1

You have an application that receives SSL-encrypted TCP traffic on port 443. Clients for this application are located all over the world. You want to minimize latency for the clients. Which load balancing option should you use?

- A. HTTPS Load Balancer
- B. Network Load Balancer
- C. SSL Proxy Load Balancer
- D. Internal TCP/UDP Load Balancer. Add a firewall rule allowing ingress traffic from 0.0.0.0/0 on the target instances.

Correct Answer: C

Reference:

<https://cloud.google.com/load-balancing/docs/ssl>*Community vote distribution*

C (100%)

Bituz Highly Voted 2 years, 3 months ago

SSL Proxy Load Balancing support for the following ports: 25, 43, 110, 143, 195, 443, 465, 587, 700, 993, 995, 1883, 3389, 5222, 5432, 5671, 5672, 5900, 5901, 6379, 8085, 8099, 9092, 9200, and 9300. When you use Google-managed SSL certificates with SSL Proxy Load Balancing, the frontend port for traffic must be 443 to enable the Google-managed SSL certificates to be provisioned and renewed.

upvoted 25 times

DarioFama23 Highly Voted 2 years, 4 months ago

C is correct

upvoted 19 times

Charumathi Most Recent 1 month ago**Selected Answer: C**

C is the correct Answer,

SSL proxy load balancer with Traffic Type - TCP with SSL offload, for global IPv4, IPv6, external ports for load balancing - 25,43,110,143,195,443,465,587,700,993,995,1883,5222

upvoted 1 times

abirroy 3 months ago**Selected Answer: C**

SSL Proxy Load Balancer

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

SSL proxy load balancer - C is right

upvoted 1 times

JelloMan 5 months, 2 weeks ago**Selected Answer: C**

Network LB = Regional

SSL Proxy = Global <----That alone tells you the answer is C

upvoted 4 times

pfabio 5 months, 3 weeks ago**Selected Answer: C**

A appear to be correct. But TCP was mentioned on question.

More information about:

Layer 4 and Layer 7 load balancing. Use Layer 4-based load balancing to direct traffic based on data from network and transport layer protocols such as TCP, UDP, ESP, ICMP, and ICMPv6. Use Layer 7-based load balancing to add request routing decisions based on attributes, such as the HTTP header and the uniform resource identifier.

<https://cloud.google.com/load-balancing/docs/load-balancing-overview>

upvoted 1 times

rikku33 6 months, 2 weeks ago

C

SSL Proxy Load Balancing can be configured as a global load balancing service.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

backspace_19 7 months, 1 week ago**Selected Answer: C**

Refer: <https://cloud.google.com/load-balancing/images/choose-lb.svg>

upvoted 1 times

✉ **emv** 9 months, 3 weeks ago

SSL Proxy Load Balancing is intended for non-HTTP(S) traffic !!!

upvoted 3 times

✉ **fishnoodlesoup** 10 months, 3 weeks ago

For me "A" is the correct answer.

We need a Global LB for 443, which is HTTPS.

SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing.

Source: <https://cloud.google.com/load-balancing/docs/ssl>

upvoted 9 times

✉ **mahesh0049** 9 months, 3 weeks ago

Answer should be SSL proxy LB because it handles tcp traffic on port 443. where https LB not handles tcp traffic.

upvoted 4 times

✉ **Ridhanya** 11 months ago

C seems right

upvoted 1 times

✉ **Dogbrasil** 1 year ago

C is correct

The type of traffic that you need your load balancer to handle is another factor in determining which load balancer to use:

HTTP and HTTPS traffic: External HTTP(S) Load Balancing / Regional external HTTP(S) / load balancer Internal HTTP(S) Load Balancing

TCP traffic: TCP Proxy Load Balancing / External TCP/UDP Network Load Balancing / Internal TCP/UDP Load Balancing /

SSL traffic: SSL Proxy Load Balancing

UDP traffic: External TCP/UDP Network Load Balancing / Internal TCP/UDP Load Balancing

upvoted 7 times

✉ **ankatsu2010** 1 year ago

"A" is the answer. Think simple, 443/tcp is HTTPS.

upvoted 6 times

✉ **zaxxon** 1 year ago

SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing.

upvoted 3 times

✉ **sudav** 1 year, 2 months ago

B - Google Cloud external TCP/UDP Network Load Balancing (after this referred to as Network Load Balancing) is a regional, non-proxied load balancer. The network load balancer supports any and all ports. You can use Network Load Balancing to load balance TCP and UDP traffic. Because the load balancer is a pass-through load balancer, your backends terminate the load-balanced TCP connection or UDP packets themselves. For example, you might run an HTTPS web server on your backends (which is our scenario) and use a Network Load Balancing to route requests to it, terminating TLS on your backends themselves.

Ref: <https://cloud.google.com/load-balancing/docs/network>

upvoted 4 times

✉ **JelloMan** 6 months, 3 weeks ago

You just mentioned a keyword here. Network Load balancing = regional. Question says clients are located all over the world which means global. Just from that you can eliminate B and choose C

upvoted 2 times

✉ **gerhardbl** 1 year, 2 months ago

Minimize latency for global users means SSL offloading close to those users, while sending the traffic as much through the Google network as possible as opposed to over the internet. This implies SSL Proxy. If you think worldwide users implies "geographic control over where TLS is terminated" - then your understanding of load balancing is flawed. Control over where TLS is terminated with a Network Load Balancer means the load balancer forwards the raw TCP as is and you'll be terminating SSL on your own backends, which increases latency to users (unless you have backends in each Region close to users but that is unlikely and mentioned nowhere in the question).

upvoted 4 times

✉ **aa_desh** 1 year, 5 months ago

Answer is:B

he SSL proxy load balancer terminates TLS in locations that are distributed globally, so as to minimize latency between clients and the load balancer. If you require geographic control over where TLS is terminated, you should use Network Load Balancing instead, and terminate TLS on backends that are located in regions appropriate to your needs.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times



👤 **theBestStudent** 2 months, 4 weeks ago

As per this table Network load balancer is Regional: <https://cloud.google.com/load-balancing/docs/choosing-load-balancer#lb-summary>. But clients are all over the world

upvoted 1 times

Question #128

You have an application on a general-purpose Compute Engine instance that is experiencing excessive disk read throttling on its Zonal SSD Persistent Disk. The application primarily reads large files from disk. The disk size is currently 350 GB. You want to provide the maximum amount of throughput while minimizing costs.

What should you do?

- A. Increase the size of the disk to 1 TB.
- B. Increase the allocated CPU to the instance.
- C. Migrate to use a Local SSD on the instance.
- D. Migrate to use a Regional SSD on the instance.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/disks/performance>

Community vote distribution

C (75%)	13%	13%
---------	-----	-----

✉️  **DarioFama23** [Highly Voted] 2 years, 4 months ago

C is correct, local SSD has more IOPS

upvoted 33 times

✉️  **TAvenger** 1 year, 8 months ago

Agree. This is also cheaper than having 350 Gb persistent SSD:

Here are calculations (taken from GCP when creating instance)

350 Gb SSD Persistent disk: 59.50\$/month, read IOPS: 10 500 with n1-standard-1

1000 Gb SSD Persistent disk: 170.00\$/month, read IOPS: 15 000 with n1-standard-1

375 Gb Local SSD (NVMe): 30.00\$/month, read IOPS: 170 000 with n1-standard-1

upvoted 11 times

✉️  **RegisFTM** 10 months, 3 weeks ago

trick question... locak-ssd is not persistent. increasing the size of of the disk will also increase the iops. A is correct imho.

upvoted 4 times

✉️  **shax910** 9 months, 2 weeks ago

Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted.

upvoted 1 times

✉️  **obeythefist** 8 months, 1 week ago

I once thought that A was the correct response because of the persistence problem, but reading the question carefully, we must choose C over A.

The question does not stipulate that the local files must be persistent, and this is the only reason why you would choose "A" over "C".

Also, the question has an important key word: Minimising costs.

1TB of zonal persistent disk costs a huge amount more than 350GB of local disk.

We should choose C.

upvoted 5 times

✉️  **ace_2021_1** 1 month, 2 weeks ago

they have been using persistent disk....there was a reason during the initial design.....persistent disk was a req....we should find a solution without affecting that...increasing the size would be the best solution for this scenario.

upvoted 2 times

✉️  **pfabio** 5 months, 3 weeks ago

Very Nice, agree C is correct

upvoted 1 times

✉️  **ESP_SAP** [Highly Voted] 2 years, 2 months ago

Correct Answer is (C):

Performance

Standard persistent disks are efficient and economical for handling sequential read/write operations, but they aren't optimized to handle high

rates of random input/output operations per second (IOPS). If your apps require high rates of random IOPS, use SSD persistent disks. SSD persistent disks are designed for single-digit millisecond latencies. Observed latency is application specific.

<https://cloud.google.com/compute/docs/disks#performance>

upvoted 16 times

✉ **hogtrough** 9 months, 1 week ago

A local SSD is not the same as an SSD persistent disk.

"Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted."

<https://cloud.google.com/compute/docs/disks#localssds>

The answer is C.

upvoted 2 times

✉ **hogtrough** 9 months, 1 week ago

Sorry, A not C.

upvoted 1 times

✉ **AwesomeGCP** Most Recent 3 weeks, 6 days ago

Selected Answer: C

C. Migrate to use a Local SSD on the instance

upvoted 1 times

✉ **Charumathi** 1 month ago

Selected Answer: C

C is the correct Answer,

Local SSDs

Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The data that you store on a local SSD persists only until the instance is stopped or deleted. Each local SSD is 375 GB in size, but you can attach a maximum of 24 local SSD partitions for a total of 9 TB per instance.

Performance

Local SSDs are designed to offer very high IOPS and low latency. Unlike persistent disks, you must manage the striping on local SSDs yourself. Combine multiple local SSD partitions into a single logical volume to achieve the best local SSD performance per instance, or format local SSD partitions individually.

Local SSD performance depends on which interface you select. Local SSDs are available through both SCSI and NVMe interfaces.

upvoted 1 times

✉ **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 1 times

✉ **gcp_world123** 2 months, 3 weeks ago

C is correct

Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The performance gains from local SSDs require certain trade-offs in availability, durability, and flexibility. Because of these trade-offs, Local SSD storage isn't automatically replicated and all data on the local SSD might be lost if the instance terminates for any reason.

Ref: <https://cloud.google.com/compute/docs/disks#localssds>

Ref: https://cloud.google.com/compute/docs/disks/performance#type_comparison

upvoted 1 times

✉ **theBestStudent** 2 months, 4 weeks ago

Selected Answer: B

C can't B for the reasons explained here in the answers (local disk info will be totally deleted if you restart or delete your vm). Then A or B. Google recommends to increase the size of the disk as performance of the disk is linear to the size of the same. But Im wondering if 1TB is too much (it is almost 2 times more the original size of the disk, and that sounds like too much, plus if with 350GB the files already fit there, then increasing the size of the disk "a little bit more" should have been a better approach to test how the performance increases. Then Google Also recommends to add more CPU to get a better IOPS, it might be cheaper than option A if the right machine with more CPU is chosen.

Answer: B

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

Local SSD .. C is right

upvoted 1 times

✉ **Tirthankar17** 5 months ago

Selected Answer: C

C is correct.

upvoted 1 times

✉ **Rutu_98** 5 months, 2 weeks ago

Selected Answer: C

Go with C

As Local SSDs have high IOPS

upvoted 1 times

✉ **Akash7** 7 months ago

C for me.

<https://cloud.google.com/compute/docs/disks#performance>

upvoted 1 times

✉ **Soso_zozo** 7 months ago

Correct Answer is B

From both tables:

https://cloud.google.com/compute/docs/disks/performance#performance_by_disk_size

<https://cloud.google.com/compute/docs/disks/performance#machine-type-disk-limits>

The general purpose vm limit is less than ssd persistent disk size limit of 250-500

upvoted 2 times

✉ **s_a_t_y_a_m** 8 months, 2 weeks ago

Selected Answer: C

Increasing Disk size would increase cost while question is asking for minimising the cost

upvoted 2 times

✉ **Mixer5** 9 months, 2 weeks ago

Selected Answer: A

We have to assume that persistent type of disk was used on purpose so we can't switch to non-persistent solution. Larger disk will- in itself-increase IOPS as well so A is an answer.

upvoted 1 times

✉ **theBestStudent** 2 months, 4 weeks ago

But "A" is expensive.

upvoted 1 times

✉ **emv** 9 months, 3 weeks ago

When you need more disk space or better performance, resize your disks and possibly add more vCPUs to add more storage space, throughput, and IOPS. Persistent disk performance is based on the total persistent disk capacity attached to an instance and the number of vCPUs that the instance has.

* Persistent disk IOPS and throughput performance depends on disk size, instance vCPU count, and I/O block size, among other factors.

** Persistent disks can achieve greater throughput performance on instances with more vCPUs. Read Network egress caps on write throughput.

upvoted 1 times

✉ **rachee** 11 months, 1 week ago

C cannot be right. The data only persists until VM is stopped or deleted. <https://cloud.google.com/compute/docs/disks>.

A is better answer to increase disk size and "possibly" CPU - <https://cloud.google.com/compute/docs/disks>.

upvoted 1 times

✉ **RazOr** 9 months, 2 weeks ago

Not fully true, see: https://cloud.google.com/compute/docs/disks/local-ssd#data_persistence

upvoted 1 times

✉ **sudav** 1 year, 2 months ago

C - Local SSDs are physically attached to the server that hosts your VM instance. Local SSDs have higher throughput and lower latency than standard persistent disks or SSD persistent disks. The performance gains from local SSDs require certain trade-offs in availability, durability, and flexibility. Because of these trade-offs, Local SSD storage isn't automatically replicated and all data on the local SSD might be lost if the instance terminates for any reason.

Ref: <https://cloud.google.com/compute/docs/disks/localssds>

upvoted 3 times

Question #129

Your Dataproc cluster runs in a single Virtual Private Cloud (VPC) network in a single subnet with range 172.16.20.128/25. There are no private IP addresses available in the VPC network. You want to add new VMs to communicate with your cluster using the minimum number of steps. What should you do?

- A. Modify the existing subnet range to 172.16.20.0/24.
- B. Create a new Secondary IP Range in the VPC and configure the VMs to use that range.
- C. Create a new VPC network for the VMs. Enable VPC Peering between the VMs' VPC network and the Dataproc cluster VPC network.
- D. Create a new VPC network for the VMs with a subnet of 172.32.0.0/16. Enable VPC network Peering between the Dataproc VPC network and the VMs VPC network. Configure a custom Route exchange.

Correct Answer: A*Community vote distribution*

A (52%) C (48%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correction.

Correct Answers is (A):

gcloud compute networks subnets expand-ip-range

NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

upvoted 34 times

 **ccieman2016** 2 months, 2 weeks ago

I think, you can't expand ip range subnet, if there isn't space in VPC. I read this question a lot, VPC CIDR like with 172.16.20.128/25 and there's only one subnet 172.16.20.128/25 inside this VPC, so you can't expand nothing. for me, there's Letter C and D works, but letter D is necessary extra work. LETTER C is right.

upvoted 1 times

 **francisco_guerra** Highly Voted 2 years, 3 months ago

I think is A,

upvoted 33 times

 **Priyanka109** 1 month ago

No it can't be as you can't modify ip address but can expand. There is no ip in the existing vpc so you have to create a new vpc and connect it using peering.

upvoted 2 times

 **passmepls** 2 years, 2 months ago

thank you

upvoted 4 times

 **mattcl** Most Recent 9 hours, 54 minutes ago

Why not B? If you create a new subnet you can communicate your VM with the cluster.

upvoted 1 times

 **hiromi** 5 days, 20 hours ago

Selected Answer: A

Modify the existing subnet range to 172.16.20.0/24

upvoted 1 times

 **PKookNN** 3 weeks ago

Selected Answer: C

after re-reading the question and spotted "There are no private IP addresses available in the VPC network." so the correct answer is C (and not A).

upvoted 1 times

 **CyberG** 3 weeks, 1 day ago

Selected Answer: C

Correct is C, Reference: <https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 1 times

 **anolive** 3 weeks, 4 days ago

Selected Answer: C

A subnet has a single primary IP address range and, optionally, one or more secondary IP address ranges. For each subnet IP address range, Google Cloud creates a subnet route. When you use VPC Network Peering, Google Cloud always exchanges the subnet routes that don't use privately reused public IP addresses between the two peered networks. If firewall rules in each network permit communication, VM instances in one network can communicate with instances in the peered network.

Reference: <https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 1 times

 **AwesomeGCP** 3 weeks, 6 days ago

Selected Answer: A

A. Modify the existing subnet range to 172.16.20.0/24.

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: A

A is correct answer, we can increase the CIDR range with minimal steps, below is the calculation for both /25 and /24 CIDR for your understanding.

/25:

CIDR to IP Range

Result

CIDR Range 172.16.20.128/25

Netmask 255.255.255.128

Wildcard Bits 0.0.0.127

First IP 172.16.20.128

First IP (Decimal) 2886734976

Last IP 172.16.20.255

Last IP (Decimal) 2886735103

Total Host 128

CIDR

172.16.20.128/25

/24:

CIDR to IP Range

Result

CIDR Range 172.16.20.128/24

Netmask 255.255.255.0

Wildcard Bits 0.0.0.255

First IP 172.16.20.0

First IP (Decimal) 2886734848

Last IP 172.16.20.255

Last IP (Decimal) 2886735103

Total Host 256

CIDR

172.16.20.128/24

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

Selected Answer: A

A. is correct

upvoted 1 times

 **ale_brd_** 1 month, 2 weeks ago

Selected Answer: A

Correct answer is A:

let's analyze the key phrases...

"runs in a SINGLE Virtual Private Cloud (VPC) network in a single subnet with range 172.16.20.128/25." >>> it means that only one subnet has been created for the VPC.

>> consequently "There are no private IP addresses available in the VPC network."

So we can assume that we can expand the only subnet created, as the question does not tell us there are in the same VPC thousands of other subnetworks and VMs that occupied all the possible available IP ranges.

upvoted 3 times

 **ryumada** 3 months ago

Selected Answer: A

Vote for A. See ESP_SAP's comment.

upvoted 1 times

 **patashish** 4 months, 1 week ago

A is correct.

Hint from question - Minimum Step

We can expand subnet range which will fit in VPC

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: A

A is right, GCP allows you to expand subnet range without any downtime
upvoted 3 times

 **AzureDP900** 4 months, 2 weeks ago

A is right, GCP allows you to expand subnet range without any downtime
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

changing my answer to C... This is exact scenario in Tutorial Dojo .. point here is no available IP's is key.

There are no private IP addresses available in the VPC network.

upvoted 2 times

 **csrazdan** 4 months ago

A is the correct answer /24 means subnets will have 64 IP addresses and /25 means the subnet is extended and will now have 128 IP addresses.

upvoted 1 times

 **csrazdan** 4 months ago

Correction /24 = 128 IP Addresses and /25 = 256 IP Addresses.

upvoted 1 times

 **Bossko** 3 months, 2 weeks ago

Correction /24 = 256 IP Addresses and /25 = 128 IP Addresses.

upvoted 3 times

 **S00999** 4 months, 3 weeks ago

Selected Answer: A

Correct is A

VPC networks do not have any IP address ranges associated with them (we may probably say that VPC is "something" under network OSI layer 3).

Each VPC network consists of one or more IP range partitions called subnets and IP ranges are defined for the subnets.

Subnet IP range can be expanded by modifying its subnet mask.

(...if the expanded range is not overlapping an other existing subnet)

upvoted 1 times

 **SergijF** 5 months, 1 week ago

Selected Answer: C

No IPs in the VPC, not in the subnet.

upvoted 3 times

Question #130

You manage an App Engine Service that aggregates and visualizes data from BigQuery. The application is deployed with the default App Engine Service account.

The data that needs to be visualized resides in a different project managed by another team. You do not have access to this project, but you want your application to be able to read data from the BigQuery dataset. What should you do?

- A. Ask the other team to grant your default App Engine Service account the role of BigQuery Job User.
- B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer.
- C. In Cloud IAM of your project, ensure that the default App Engine service account has the role of BigQuery Data Viewer.
- D. In Cloud IAM of your project, grant a newly created service account from the other team the role of BigQuery Job User in your project.

Correct Answer: B

The Owner, Editor, and Viewer primitive roles include the BigQuery Admin (roles/bigquery.dataOwner), BigQuery Data Editor (roles/bigquery.dataEditor), and

BigQuery Data Viewer (roles/bigquery.dataViewer) roles, respectively. This means the Owner, Editor, and Viewer primitive roles have BigQuery access as defined for the respective BigQuery roles.

Reference:

<https://cloud.google.com/bigquery/docs/access-control>

Community vote distribution

B (54%)

A (46%)

 **Hjameel** Highly Voted 2 years, 2 months ago

I think B is the answer

upvoted 28 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Sorry, I copied/pasted the wrong statement.

This is the proper explanation regarding to Big Query Data Viewer Role.

The resource that you need to get access is in the other project.

[roles/bigquery.dataViewer](#) BigQuery Data Viewer

When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.

This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.

Read data and metadata from the dataset's tables.

When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 20 times

 **[Removed]** 2 years, 1 month ago

A is correct, data viewer role does not allow you to execute query, that can be done if you are the user role.

upvoted 8 times

 **[Removed]** 2 years, 1 month ago

The question states you want to aggregate and visualize data, that is run aggregate SQL on data before visualizing.

upvoted 2 times

 **Untamables** Most Recent 2 weeks, 2 days ago

Selected Answer: A

Applications generally get data from BigQuery with querying. It needs BigQuery Job User role to query.

<https://cloud.google.com/bigquery/docs/access-control#bigquery>

I think using the default App Engine service account is recommended.

<https://cloud.google.com/appengine/docs/legacy/standard/python/service-account>

upvoted 1 times

 **AwesomeGCP** 3 weeks, 6 days ago

Selected Answer: B

you want your application to be able to read data from the BigQuery dataset so it's simple B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: B

B is the correct answer,

As the question states you need only read access to the BigQuery datasets residing in other project, we need to request the other project team to give data viewer role to access the Bigquery datasets.

upvoted 1 times

 **ATANGA** 3 months ago

WhatsApp : +1(956)-520-4006 to obtain PMP, CISM, CCNA, CEH, PRINCE2, CISCO, ISTQB, PRINCE2, AWS/Azure/Sale force/ITIL Foundation/EC- COUNCIL...

Get Certified with 100% pass guarantee. Pay after exam.

all CISCO, ISACA & EC- COUNCIL certifications

For the Below certificates we offer 100% pass guarantee:

1. AWS Certification
2. Sales force
3. Scrum Master
4. Oracle Certification: OCA, OCP
5. Cisco Certification: CCNA, CCNP, CCIE
6. ITIL Foundation & Intermediate
7. Prince 2 Foundation and Practitioner
8. VMWARE Certification
9. Check Point Certification
10. EC-COUNCIL Certification (CEH V-9, CCISO, CND)
11. Cloud Certification
12. IBM Certification
13. HP Certification
14. Citrix Certification
15. Juniper certification
16. Azure
17. Skype 70-333/34
18. PMI (PMP/CAPM/ACP/PBA)
19. ISTQB
20. SAP
21. ISACA (CISA, CISM, CRISC, CGEIT, COBIT)

PAYMENT ONLY AFTER CERTIFICATION AND RESULT CONFIRMATION.

WhatsApp : +1(956)-520-4006

upvoted 1 times

 **ikychughtai** 3 months, 1 week ago

Selected Answer: A

The question states that you need to be able to export the data to the app engine application so that you can visualise the data there. You do not need to run queries directly in BigQuery on the other team's project, and hence B is the correct answer, you only need DataViewer permissions by principle of least privilege

upvoted 3 times

 **ikychughtai** 3 months, 1 week ago

I meant to select answer B haha

upvoted 1 times

 **pspandher** 4 months, 1 week ago

Read and Query are both different terminologies, Google it, Read means only fetch and Query means Insert, update, delete and remove so if read is mentioned the data viewer role is the BEST.

upvoted 3 times

 **AzureDP900** 4 months, 2 weeks ago

A is right, This is part of Tutorial Dojo practice questions.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

roles/bigquery.user – When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset

upvoted 1 times

 **sneha1607** 2 months, 1 week ago

its job user in the option not User. so the correct answer is option B

upvoted 1 times

 **backspace_19** 7 months, 1 week ago

Selected Answer: B

The question says- "you want your application to be able to READ data from the BigQuery dataset", nothing about the jobs.

upvoted 1 times

 **akshaychavan7** 5 months, 2 weeks ago

How would you perform the data aggregation then?

upvoted 1 times

 **somenick** 7 months, 3 weeks ago

Selected Answer: B

roles/bigquery.dataViewer BigQuery Data Viewer

When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.

This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.

Read data and metadata from the dataset's tables.

When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 2 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: A

I vote for 'A'

Firstly, How would you get the data without querying from bigquery?

If your application wants to visualize the data from bq then it should have permission to run the queries. Does it make sense?

upvoted 4 times

 **Deax** 9 months, 1 week ago

Selected Answer: B

I'd go with B. The question states that the app needs to 'read data from the big query dataset' and nothing else. In that case dataViewer is the least privilege role needed. From the docs:

roles/bigquery.dataViewer

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.

Read data and metadata from the dataset's tables.

upvoted 5 times

 **RazOr** 9 months, 2 weeks ago

Selected Answer: B

B is right imo, because of the docs:

BigQuery Data Viewer

(roles/bigquery.dataViewer)

When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.

This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.

Read data and metadata from the dataset's tables.

upvoted 3 times

 **shanj910** 9 months, 2 weeks ago

Selected Answer: A

BigQuery Job User

(roles/bigquery.jobUser)

Provides permissions to run jobs, including queries, within the project.

Lowest-level resources where you can grant this role:

Project

upvoted 3 times

 **Ridhanya** 11 months ago

it is B

BigQuery Data Viewer

(roles/bigquery.dataViewer)

When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.

This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.

Read data and metadata from the dataset's tables.

BigQuery Job User
(roles/bigquery.jobUser)

Provides permissions to run jobs, including queries, within the project.

so job user can get the data while data viewer can
upvoted 2 times

 **ankatsu2010** 1 year ago

I think A is correct. Data Viewer role can only 'view' the table data. To run the query or copy, export data to AppEngine app, the service account needs Job User role.

upvoted 6 times

Question #131

You need to create a copy of a custom Compute Engine virtual machine (VM) to facilitate an expected increase in application traffic due to a business acquisition.

What should you do?

- A. Create a Compute Engine snapshot of your base VM. Create your images from that snapshot.
- B. Create a Compute Engine snapshot of your base VM. Create your instances from that snapshot.
- C. Create a custom Compute Engine image from a snapshot. Create your images from that image.
- D. Create a custom Compute Engine image from a snapshot. Create your instances from that image.

Correct Answer: D

A custom image belongs only to your project. To create an instance with a custom image, you must first have a custom image.

Reference:

<https://cloud.google.com/compute/docs/instances/create-start-instance>

Community vote distribution

B (63%)

D (37%)

✉️  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (D):

Preparing your instance for an image

You can create an image from a disk even while it is attached to a running VM instance. However, your image will be more reliable if you put the instance in a state that is easier for the image to capture. Use one of the following processes to prepare your boot disk for the image:

Stop the instance so that it can shut down and stop writing any data to the persistent disk.

If you can't stop your instance before you create the image, minimize the amount of writes to the disk and sync your file system.

Pause apps or operating system processes that write data to that persistent disk.

Run an app flush to disk if necessary. For example, MySQL has a FLUSH statement. Other apps might have similar processes.

Stop your apps from writing to your persistent disk.

Run sudo sync.

After you prepare the instance, create the image.

https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images#create_instance_for_image

upvoted 28 times

✉️  **pca2b** Highly Voted 1 year, 7 months ago

B:

we just need to make 'a copy' of the VM, B works well for that.

not D: Had the question mentioned more copies, we would need to go the way of images...templates etc. D will work but not needed here.

upvoted 7 times

✉️  **djgodzilla** 1 year, 4 months ago

custom images are better a fit if its for a new business workload you just acquired

upvoted 2 times

✉️  **wolfie09** 5 months ago

What about the answer that says create your instanceS ??

upvoted 1 times

✉️  **anolive** Most Recent 2 weeks, 1 day ago

Selected Answer: B

I not sure about d, there is no mention that already hás a snaoshot crrated.

upvoted 1 times

✉️  **jeanfmc** 2 weeks, 5 days ago

Selected Answer: D

agree with D

upvoted 1 times

✉️  **AwesomeGCP** 3 weeks, 6 days ago

Selected Answer: D

D. Create a custom Compute Engine image from a snapshot. Create your instances from that image.

Preparing your instance for an image

You can create an image from a disk even while it is attached to a running VM instance. However, your image will be more reliable if you put the instance in a state that is easier for the image to capture. Use one of the following processes to prepare your boot disk for the image:

Stop the instance so that it can shut down and stop writing any data to the persistent disk.

If you can't stop your instance before you create the image, minimize the amount of writes to the disk and sync your file system.

Pause apps or operating system processes that write data to that persistent disk.

Run an app flush to disk if necessary. For example, MySQL has a FLUSH statement. Other apps might have similar processes.

Stop your apps from writing to your persistent disk.

Run sudo sync.

After you prepare the instance, create the image.

https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images#prepare_instance_for_image

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: D

D is the correct answer,

you need to create an image before you create a similar instance, this will automate the creation of new instances from the image, when there is an increase in traffic. This is the concept used in MIGs for scaling up and down the instances.

upvoted 1 times

 **Angel_99** 2 months, 3 weeks ago

Selected Answer: D

snapshot --> image -----> creating our instance from image that we got from snapshot

upvoted 3 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: D

D is Correct here, look at ESP_SAP comments for more details.

upvoted 1 times

 **koko2314** 3 months, 3 weeks ago

I think D is wrong because there is no option available in console to create an image from a Snapshot as suggested in Option D.

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

You can create an image from snapshot here, [https://cloud.google.com/compute/docs/images/create-custom#:~:text=You%20can%20create%20custom%20images,virtual%20machine%20\(VM\)%20instances](https://cloud.google.com/compute/docs/images/create-custom#:~:text=You%20can%20create%20custom%20images,virtual%20machine%20(VM)%20instances).

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 1 times

 **pfabio** 5 months, 3 weeks ago

Selected Answer: B

Create copy not copies

upvoted 2 times

 **rsuresh27** 6 months, 2 weeks ago

Selected Answer: B

B is correct. You are only creating a single custom VM as stated in the problem. This can be done best with a snapshot. However, if you are creating multiple VMs it would be best to use an image. But in this case we are only creating a single copy. Thus, B is correct.

upvoted 4 times

 **Letahrgicbeagle** 1 month, 3 weeks ago

What you said is correct however in D, it says to create the image from snapshot which will bring us to the same point as B

upvoted 1 times

 **csrazdan** 4 months ago

The Key is "copy of a custom Compute Engine virtual machine (VM)" - Snapshot is going to copy only the data. It is not going to copy any bootstrap command which you might have used when this instance was created. For Example, install the HPPT server. So D is the correct answer.

upvoted 1 times

 **somenick** 7 months, 3 weeks ago

B&D are both possible but to facilitate an increase in application traffic, likely multiple VM instances need to be created. You should use an image when creating multiple copies of VMs.

Hence D.

upvoted 2 times

 **ryzior** 8 months, 1 week ago

Selected Answer: B

You can create an instance from both an image and a snapshot. here B is the fastest path I suppose. Other options require you to create an image from a snapshot , while you can create an instance from a snapshot directly.

So I opt for B

upvoted 3 times

 **RazOr** 9 months, 2 weeks ago

Selected Answer: B

B! Because you can create instances from a snapshot and creating an image is the wrong way for this situation. You only need a custom image when:

What is custom image in GCP?

A custom image is a boot disk image that you own and control access to. Use custom images for the following tasks: Import a boot disk image to Compute Engine from your on-premises environment, or import virtual disks from VMs that are running on your local workstation or on another cloud platform.

upvoted 2 times

 **mahesh0049** 9 months, 3 weeks ago

option D

snapshot --> image -----> creating our instance from image that we got from snapshot

upvoted 1 times

 **emv** 9 months, 3 weeks ago

it's D <https://cloud.google.com/compute/docs/instances/copy-vm-between-projects?hl=en>

upvoted 1 times

Question #132

You have deployed an application on a single Compute Engine instance. The application writes logs to disk. Users start reporting errors with the application. You want to diagnose the problem. What should you do?

- A. Navigate to Cloud Logging and view the application logs.
- B. Connect to the instance's serial console and read the application logs.
- C. Configure a Health Check on the instance and set a Low Healthy Threshold value.
- D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging.

Correct Answer: D*Community vote distribution*

D (58%)	B (33%)	8%
---------	---------	----

 **hiteshrup**  2 years, 1 month ago

Answer: D

App logs can't be visible to Cloud Logging until we install Cloud Logging Agent on GCE
upvoted 23 times

 **ashrafh** 1 year, 2 months ago

Hi all
check this document and decide :)
<https://cloud.google.com/logging/docs/agent/logging/installation>
upvoted 2 times

 **hiteshrup** 2 years, 1 month ago

Continuation of reasoning.

If Problem statement is not having this statement "The application writes logs to disk", then we might assume that application is writing logs on Cloud Logging with google-fluentd agent API library. However, problem statement is clearly mentioned that logs are writing down on disk, we need agent installed on GCE to fetch those logs from disk to Cloud Logging. If that is not desirable, then option B is left
upvoted 4 times

 **hiteshrup** 2 years, 1 month ago

(Correction) Answer is A after rethinking and doing some research by focusing words "App Engine", which has by default enabled Request Logs which has App logs on each request and those logs are enabled for Cloud Logging ..
https://cloud.google.com/appengine/docs/standard/python/logs#request_logs_vs_application_logs
upvoted 3 times

 **abirroy** 3 months ago

For compute engine you need to install logging agent. Answer is D

upvoted 1 times

 **rezavage** 2 years, 1 month ago

Cloud logging without agent only works for App engine as you stated . but the question is about the compute engine which has to be equipped first with Logging Agent in order to write logs into Cloud Logging. so based your assumption the correct answer is "D"
upvoted 3 times

 **ESP_SAP**  2 years, 2 months ago

Correct Answer is (D):

In its default configuration, the Logging agent streams logs from common third-party applications and system software to Logging; review the list of default logs. You can configure the agent to stream additional logs; go to Configuring the Logging agent for details on agent configuration and operation.

It is a best practice to run the Logging agent on all your VM instances. The agent runs under both Linux and Windows. To install the Logging agent, go to Installing the agent.

<https://cloud.google.com/logging/docs/agent>

upvoted 15 times

 **ESP_SAP** 2 years, 2 months ago

CORRECTION.

Correct Answer is (A):

Activity logging is enabled by default for all Compute Engine projects.

You can see your project's activity logs through the Logs Viewer in the Google Cloud Console:

In the Cloud Console, go to the Logging page.
Go to the Logging page

When in the Logs Viewer, select and filter your resource type from the first drop-down list.
From the All logs drop-down list, select compute.googleapis.com/activity_log to see Compute Engine activity logs.
https://cloud.google.com/compute/docs/logging/activity-logs#viewing_logs

Besides:

Activity logs are provided as part of the Cloud Logging service. For more information about Logging in general, read the Cloud Logging documentation.

<https://cloud.google.com/compute/docs/logging/activity-logs>

upvoted 9 times

 **ESP_SAP** 2 years, 2 months ago

Additional information about VM Image for AWS EC2:

The Logging agent streams logs from your VM instances and from selected third-party software packages to Cloud Logging. It is a best practice to run the Logging agent on all your VM instances.

The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances. The agent runs under both Linux and Windows.

If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page.

upvoted 2 times

 **magistrum** 1 year, 10 months ago

This points to D then

upvoted 3 times

 **babusartop17** 1 year, 2 months ago

I feel sorry for the woman in your life.

upvoted 14 times

 **DamonSalvatore** 1 year, 1 month ago

Haha! That was funny

upvoted 1 times

 **mexblood1** 2 years, 2 months ago

Activity Logs do not include 2nd party application logs. Activity logs are more related to operations and changes in the infrastructure. This question is tricky, I think it's either D or B, because if it's only an application on a single instance, you can connect to the instance and read the application logs directly and you save the cost of logging agent.

upvoted 9 times

 **mexblood1** 2 years, 2 months ago

Maybe I was assuming serial console is the same than system console, technically I guess they're not the same, hence I guess D will be my chosen answer.

upvoted 1 times

 **sapguru** 2 years, 2 months ago

Cloud logging enabled by default for compute engine

upvoted 1 times

 **csrazdan** 4 months ago

Do you mean the logging agent is installed by default? It depends on the OS you decide. For example, it is installed in Ubuntu but not on RedHat or Windows. Besides installing of the agent is not enough. You have to configure and let the agent know where your application is writing the logs on the disk so that it can monitor and stream the log to cloud monitoring. D is the correct answer

upvoted 1 times

 **AwesomeGCP** Most Recent 3 weeks, 6 days ago

Selected Answer: D

D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: D

D is the correct answer,

Install fluentD logging agent to compute engine and view the logs from cloud logging.

upvoted 1 times

 **theBestStudent** 2 months, 3 weeks ago

Selected Answer: D

The assumption is errors are related to THE APPLICATION, and NOT to the VM instance itself (like booting, starting up, crashing, etc).

That being said, Cloud Logging is not by default installed in the VMS, and we need to install the agent first in order to see the application logs in Cloud Logging.

Of course another option would have been to just go to the file where the logs are located in the VM, but that is not an option in this question.

CORRECT ANSWER: D.

upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: D

For compute engine you need to install logging agent. Answer is D

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: B

Answer should be B, Read the question it clearly mentions that " The application writes logs to disk." What is the point of writing the log to disk if you don't want to use them. D does not makes sense since why Install and agent if you are already writing logs + the errors have already happened, How will the agent go back in time and give an error messages back ?

upvoted 2 times

 **S00999** 4 months, 3 weeks ago

Selected Answer: D

Answer: D

Cloud Logging knows nothing about applications installed on the system without an agent collecting logs.

Using the serial console is not a best-practice and is impractical on a large scale.

Exam questions always invite responses in favor of the tools offered by the publisher.

upvoted 3 times

 **SergijF** 5 months, 1 week ago

Selected Answer: D

D is correct. Not B. You can't read anything from the disc with the serial console. You can only see the system events, not the logs of a running application that are already written to the disc. With the agent you are able to access your logs

upvoted 1 times

 **PAUGURU** 6 months, 1 week ago

Selected Answer: B

The question asks for debugging of events that happened in the past, so how can installing the cloud logging agent help debug, maybe the problem won't repeat in the future.

To debug you have to read logs stored on disk so only answer is B.

If it asked for a solution to centralize future logs the answer would be D.

upvoted 3 times

 **rsuresh27** 6 months, 3 weeks ago

The answer is B. The logs are being written to the disk itself and thus we need to use the serial port to view the output. Cloud Logging is not needed because, as the question says, the application logs are being written to the disk itself.

upvoted 1 times

 **BigQuery** 7 months, 4 weeks ago

Selected Answer: B

Activity Logs do not include 2nd party application logs. Activity logs are more related to operations and changes in the infrastructure. This question is tricky, I think it's either D or B, because if it's only an application on a single instance, you can connect to the instance and read the application logs directly and you save the cost of logging agent.

upvoted 2 times

 **ryzior** 8 months, 1 week ago

Selected Answer: D

I did a simple test.

First of all the monitoring logging agent is not installed on a GCE VM by default, so I could only see some activity logs in the Logging page (reached from the instance screen).

After OpsAgent installation I was able to see a lot more logs in Cloud Logging but the trace starts the same time the agent's been installed. So if you need historical error details, you won't get them this way I suppose (or there is some more tuning needed).

I'd vote for serial connectivity and check the logs from the console (maybe they put serial to make some confusion because it is disabled by default and you need explicitly enable serial ports in your vm , which does not need a restart or anything - checked myself) but if the errors persist in the application you will see them too after Log agent installation, so I vote finally for D - Agent install and Cloud Logging.

upvoted 4 times

 **injarapu** 8 months, 1 week ago

Selected Answer: D

The Logging agent streams logs from your VM instances and from selected third-party software packages to Cloud Logging. It is a best practice to run the Logging agent on all your VM instances.

The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances. The agent runs under both Linux and Windows.

Source: <https://cloud.google.com/logging/docs/agent/logging/installation>

upvoted 1 times

 **vdc_cdv** 8 months, 1 week ago

Selected Answer: B

This is a talk about application log, so logs that the app/service hosted on CE is writing to disk. Activity logs are related to the infra not the apps. Question is tricky.

upvoted 1 times

 **vruizrob** 8 months, 3 weeks ago

Answer: A; If the application already writes logs to disk, the cloud logging is already configured

upvoted 2 times

 **Eminenza22** 9 months ago

It is a best practice to run the Logging agent on all your VM instances. If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page. Answer: A

upvoted 1 times

 **torontowhale** 8 months, 3 weeks ago

It's compute an engine, not app engine. - The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances. The agent runs under both Linux and Windows. - <https://cloud.google.com/logging/docs/agent/logging/installation#:~:text=The%20VM%20images%20for%20Compute,under%20both%20Linux%20and%20Windows>.

upvoted 1 times

Question #133

An application generates daily reports in a Compute Engine virtual machine (VM). The VM is in the project corp-iot-insights. Your team operates only in the project corp-aggregate-reports and needs a copy of the daily exports in the bucket corp-aggregate-reports-storage. You want to configure access so that the daily reports from the VM are available in the bucket corp-aggregate-reports-storage and use as few steps as possible while following Google-recommended practices. What should you do?

- A. Move both projects under the same folder.
- B. Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.
- C. Create a Shared VPC network between both projects. Grant the VM Service Account the role Storage Object Creator on corp-iot-insights.
- D. Make corp-aggregate-reports-storage public and create a folder with a pseudo-randomized suffix name. Share the folder with the IoT team.

Correct Answer: A

Reference:

<https://cloud.google.com/billing/docs/onboarding-checklist>*Community vote distribution*

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Predefined roles

The following table describes Identity and Access Management (IAM) roles that are associated with Cloud Storage and lists the permissions that are contained in each role. Unless otherwise noted, these roles can be applied either to entire projects or specific buckets.

Storage Object Creator (roles/storage.objectCreator) Allows users to create objects. Does not give permission to view, delete, or overwrite objects.

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

upvoted 30 times

 **ESP_SAP** 2 years, 2 months ago

Basically, you are giving the permissions to the VM Service Account to create a copy of the daily report on the bucket that the other team has access.

upvoted 4 times

 **francisco_guerra** Highly Voted 2 years, 3 months ago

i think is B

upvoted 13 times

 **francisco_guerra** 2 years, 3 months ago

Object creator cant see object so i think is D

upvoted 1 times

 **Ixgywil** 1 year, 6 months ago

VM doesn't need to see the obects - just to create them. It's B:

The VM is located in project "corp-iot-insights" - give its SA the Storage Object Creator role for bucket "corp-aggregate-reports-storage" that is located in project "corp-aggregate-reports", where your team operates.

upvoted 4 times

 **StefiJohnson** Most Recent 1 month, 3 weeks ago

Correct Answer is (B)

upvoted 1 times

 **theBestStudent** 2 months, 3 weeks ago

Selected Answer: B

If that is the default service Account of the Compute Instance, then we should do nothing. As the role is already included. Either way, we should do nothing as the role is already covered. Also we shouldn't modify Compute instance Service account. But again, I will assume it is not the default.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 1 times

 **deadlydeb** 4 months, 3 weeks ago

Selected Answer: B

b it is

upvoted 1 times

 **[Removed]** 9 months, 2 weeks ago

Selected Answer: B

You should be able to add a service account to another project:

Create the first service account in project A in the Cloud Console. Activate it using gcloud auth activate-service-account.

In the Cloud Console, navigate to project B. Find the "IAM & admin" > "IAM" page. Click the "Add" button. In the "New members" field paste the name of the service account (it should look like a strange email address) and give it the appropriate role.

Run gcloud commands with --project set to project B. They should succeed (I just manually verified that this will work).

Automatic creation of service accounts is something that we're hesitant to do until we can work through all of the security ramifications.

<https://stackoverflow.com/a/35558464>

upvoted 4 times

 **Mixxer5** 9 months, 2 weeks ago

Selected Answer: B

It's B since bucket names are globally unique so it's enough to refer to them when you've proper role assigned

upvoted 2 times

 **arsh1916** 1 year, 5 months ago

B, assign access is less step

upvoted 1 times

 **pradr** 1 year, 7 months ago

From stackoverflow: Bucket names are globally unique, so your app will refer to an existing bucket in another project in the same way that it refers to buckets in its own project. Hence the shared VPC is not required to access the bucket. Just the IAM role.

upvoted 2 times

 **EABDAJA** 1 year, 8 months ago

B is correct

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 2 times

 **victory108** 1 year, 9 months ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 1 times

 **Bhagirathi** 1 year, 11 months ago

why not A ?

upvoted 1 times

 **magistrum** 1 year, 10 months ago

lots of steps vs B

upvoted 1 times

 **RockAJ** 2 years, 1 month ago

For me is B

upvoted 1 times

 **shan316** 2 years, 1 month ago

Why the answer can't be C since the communication is between two different projects. If there is no VPC peering then there won't be any communication at the first place. What you guys think?

upvoted 1 times

 **Eshkrkrkr** 1 year, 12 months ago

We think that you are digging in the wrong direction. So you've shared a network, what next? How does C helps to solve the question problem?

upvoted 2 times

 **Examan1** 1 year, 9 months ago

The VM is in the project corp-iot-insights. Your team operates only in the project corp-aggregate-reports and needs a copy of the daily exports in the bucket corp-aggregate-reports-storage. Hence you need Storage Object Creator on the bucket?

upvoted 1 times

 **Ale1973** 2 years, 2 months ago

You are right, but to solve the scenario only is required that VM could write the report on the bucket.

upvoted 1 times

Question #134

You built an application on your development laptop that uses Google Cloud services. Your application uses Application Default Credentials for authentication and works fine on your development laptop. You want to migrate this application to a Compute Engine virtual machine (VM) and set up authentication using Google-recommended practices and minimal changes. What should you do?

- A. Assign appropriate access for Google services to the service account used by the Compute Engine VM.
- B. Create a service account with appropriate access for Google services, and configure the application to use this account.
- C. Store credentials for service accounts with appropriate access for Google services in a config file, and deploy this config file with your application.
- D. Store credentials for your user account with appropriate access for Google services in a config file, and deploy this config file with your application.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

Community vote distribution

B (63%)

A (38%)

ESP_SAP Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Best practices

In general, Google recommends that each instance that needs to call a Google API should run as a service account with the minimum permissions necessary for that instance to do its job. In practice, this means you should configure service accounts for your instances with the following process:

Create a new service account rather than using the Compute Engine default service account.

Grant IAM roles to that service account for only the resources that it needs.

Configure the instance to run as that service account.

Grant the instance the <https://www.googleapis.com/auth/cloud-platform> scope to allow full access to all Google Cloud APIs, so that the IAM permissions of the instance are completely determined by the IAM roles of the service account.

Avoid granting more access than necessary and regularly check your service account permissions to make sure they are up-to-date.

https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices

upvoted 45 times

Ridhanya 11 months ago

you just gave justification for option A which is right

upvoted 1 times

ryumada 3 months ago

Maybe for the option A you are modifying the default service account because it's not explain which service account used by the VM, is it the default one or the new one?

The best practice is to Create a new service account rather than using the Compute Engine default service account.

B still has the bigger prove here as the answer.

upvoted 2 times

ryumada 3 months ago

You should read lxgywil comment. His comment explains how authentication works to access Google Services in your application.

a relevant link also:

https://cloud.google.com/storage/docs/reference/libraries#setting_up_authentication

upvoted 1 times

cRobert 1 year, 10 months ago

From your quote:

Configure the "instance" to run as that service account.

From answer B:

and configure the "application" to use this account.

You don't add service accounts to applications, ans A

upvoted 17 times

lxgywil 1 year, 6 months ago

When you use a GCP service within your app (code), you have to use its client libraries. When you instantiate a client with client libraries you can pass it a Service Account key, which will define on behalf of which SA the client will be acting. That's how you can configure your app to use a particular service account.

E.g. https://cloud.google.com/storage/docs/reference/libraries#using_the_client_library

upvoted 3 times

 **magistrum** 1 year, 10 months ago

wording is the clue :)

upvoted 1 times

 **TAvenger** 1 year, 8 months ago

It's dirty play with words... All understand that we need custom SA, grant required permissions and attach this SA to the VM... Why Google does this?

upvoted 4 times

 **filco72**  2 years, 2 months ago

I would choose: A. Assign appropriate access for Google services to the service account used by the Compute Engine VM. as there is no need to create a new service account.

upvoted 20 times

 **Hjameel** 2 years, 2 months ago

I agree, there is no need to create a new service account

upvoted 9 times

 **Untamables**  2 weeks, 2 days ago

Selected Answer: A

A service account is able to be attached to VMs, not applications in VM.

upvoted 1 times

 **PKookNN** 3 weeks ago

Selected Answer: B

B since I am focusing on recommended practices which is to create a new service account and not to use the default one.

upvoted 1 times

 **theBestStudent** 3 weeks, 3 days ago

Selected Answer: A

- Assuming its not default Service answer is A.

-Assuming the GCE already have a Default Service Account associated, then B "sounds good". My problem here is it says "to configure the application with that Service account".

You don't configure the "application" with the service account. What you configure with your Service account is the GCE Instance. That would kill option B and would leave as "the most correct answer" option A.

upvoted 1 times

 **Lethargicbeagle** 1 month, 4 weeks ago

Selected Answer: B

Never edit the default service account as this may solve the purpose here but may pose a risk for future VMs to be created

upvoted 1 times

 **theBestStudent** 2 months, 3 weeks ago

Selected Answer: B

B, and never ever modify a default service account of a Compute Instance.

upvoted 1 times

 **theBestStudent** 2 months, 3 weeks ago

DO NOT modify (default) Service account of the VM! big no no. Not recommended by google.

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: A

Answer should be A,

Since it says you already have a service account and you just need to ensure you have proper access and that service account is assigned to the VM, Now I have tested this, whenever you use Google Application Credentials, they use the service account which is used to configure the instance by default. In the local environment there must be an environment variable setup with that service account creds.

upvoted 1 times

 **zolthar_z** 3 months, 1 week ago

Selected Answer: B

Answer is B, Service Account can be use in application too <https://cloud.google.com/iam/docs/service-accounts>

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: B

B is correct.

upvoted 1 times

 **mrNicky** 9 months, 3 weeks ago

B I THINK IS B. When you edit existing service account for example an instance, we editing compute engine default service account, of course there is a minimal step but we have to follow best google pratice, and best google pratice is to create ALWAYS a service account for the ressource (instance..) when we have different autorizations compared to compute engine default service account.

upvoted 1 times

 **akshaym87** 9 months, 3 weeks ago

Options A & B both talk about the service accounts.

Option A talks about the VM service account.

Option B talks about a service account specially for the application. Therefore we need to assign the permissions for this new application user. And this account is different from the VM default service account.

The question speaks about the application usage. Hence i think option B is correct.

upvoted 3 times

 **Ridhanya** 11 months ago

" application uses Application Default Credentials for authentication and works fine on your development laptop" Also the application was already using gcp services from the development laptop, so in that case, using minimal changes, B seems right

upvoted 2 times

 **kimharsh** 11 months, 2 weeks ago

I would say A, but this answer doesn't guarantee that the application will use the VM service account to access Google Services
I think the answer is B , it's clearly stated that the application will use the service account

upvoted 1 times

 **ExpensiveMouse** 1 year ago

I would go with B here although the question does not state that it is using the default service accounts and you could therefore assume one has previously been made which would make you consider A.

The question also does not state a service account other than default has been created or is in use, so I would assume worst case here and assume as part of the migration a service account will need to be created, as we cannot edit the permissions of default service accounts B is most suitable and best practice

upvoted 2 times

 **liyux21** 1 year ago

A is correct! A does not state that it uses compute engine default service account, just say assigne IAM role to service account used by instance. When your application runs on an instance, Application Default Credentials can authenticate using the service account that has been enabled on the instance. So important is, the service account is used by VM instance.

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#clientlib>

upvoted 4 times

Question #135

You need to create a Compute Engine instance in a new project that doesn't exist yet. What should you do?

- A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.
- B. Enable the Compute Engine API in the Cloud Console, use the Cloud SDK to create the instance, and then use the --project flag to specify a new project.
- C. Using the Cloud SDK, create the new instance, and use the --project flag to specify the new project. Answer yes when prompted by Cloud SDK to enable the Compute Engine API.
- D. Enable the Compute Engine API in the Cloud Console. Go to the Compute Engine section of the Console to create a new instance, and look for the Create In A New Project option in the creation form.

Correct Answer: B*Community vote distribution*

A (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Quickstart: Creating a New Instance Using the Command Line

Before you begin

1. In the Cloud Console, on the project selector page, select or create a Cloud project.
2. Make sure that billing is enabled for your Google Cloud project. Learn how to confirm billing is enabled for your project.

To use the gcloud command-line tool for this quickstart, you must first install and initialize the Cloud SDK:

1. Download and install the Cloud SDK using the instructions given on [Installing Google Cloud SDK](#).

2. Initialize the SDK using the instructions given on [Initializing Cloud SDK](#).

To use gcloud in Cloud Shell for this quickstart, first activate Cloud Shell using the instructions given on [Starting Cloud Shell](#).

<https://cloud.google.com/ai-platform/deep-learning-vm/docs/quickstart-cli#before-you-begin>

upvoted 33 times

 **filco72** Highly Voted 2 years, 2 months ago

I would choose A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

as first I need to create a project. Instance creation cannot automatically create a project.

upvoted 11 times

 **dr1ka** Most Recent 10 months, 2 weeks ago

Selected Answer: A

Vote A

upvoted 2 times

 **jaffarali** 11 months ago

Selected Answer: A

A is the correct option

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A

Correct Ans is : A

upvoted 2 times

 **arsh1916** 1 year, 5 months ago

A is correct

upvoted 2 times

 **kopper2019** 1 year, 6 months ago

A the way to go

upvoted 1 times

 **pondai** 1 year, 7 months ago

Vote A

upvoted 1 times

 **Devgela** 1 year, 7 months ago

A for me

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct

upvoted 1 times

 **neerajgoyal** 1 year, 8 months ago

A is the right answer

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

upvoted 2 times

 **JackGlemins** 1 year, 8 months ago

I think is A: <https://cloud.google.com/sdk/gcloud/reference/projects/create>

upvoted 1 times

 **lutoa** 1 year, 8 months ago

yep has to be A, new project must be created first

upvoted 1 times

 **RockAJ** 2 years, 1 month ago

A for me

upvoted 1 times

 **hiteshrup** 2 years, 1 month ago

Answer: A

Project creation is always a first step as we need to link billing account before creating any new resources. So Answer is A

upvoted 5 times

 **stepkurniawan** 2 years, 2 months ago

2 things that you need to create VM:

1. a project
2. Compute API

So answer A is correct

upvoted 2 times

Question #136

Your company runs one batch process in an on-premises server that takes around 30 hours to complete. The task runs monthly, can be performed offline, and must be restarted if interrupted. You want to migrate this workload to the cloud while minimizing cost. What should you do?

- A. Migrate the workload to a Compute Engine Preemptible VM.
- B. Migrate the workload to a Google Kubernetes Engine cluster with Preemptible nodes.
- C. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.
- D. Create an Instance Template with Preemptible VMs On. Create a Managed Instance Group from the template and adjust Target CPU Utilization. Migrate the workload.

Correct Answer: B

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool>

Community vote distribution

C (100%)

✉  **juliandm** Highly Voted  2 years, 2 months ago

i understand preemptible as a no-go because of "must be restarted if interrupted" here meaning "starting from scratch" . So C seems right
upvoted 41 times

✉  **jcloud965** 1 year, 3 months ago

I agree, C.
you won't run 30 hours job on preemptible instances that can be stopped at any time and can't run more than 24 hours.
If the job could be splitted, then preemptible VM is an option.
upvoted 5 times

✉  **Vador** 1 year, 1 month ago

Preemptible seems fine on batch jobs for at least 24hours, not the case in here
upvoted 1 times

✉  **dttnc1** 1 year ago

I agree with C. You can't risk running a processes that take 30 hours on a preemptible VM (Compute Engine always stops preemptible instances after they run for 24 hours). They are good for "short-lived" batch jobs. The scenario is NOT fault tolerant as the whole process restarts if interrupted.

<https://cloud.google.com/compute/docs/instances/preemptible>

upvoted 4 times

✉  **stepkurniawan** Highly Voted  2 years, 2 months ago

Preemptible will be perfect for a batch job that takes less than 24 hours. But it's not in this case.
upvoted 14 times

✉  **Linus11** 1 year, 6 months ago

What if it is a managed group of Pre emptible instances like in D. If one instance stops, another instance will take over.

I choose D.

upvoted 4 times

✉  **sanhoo** 1 year, 5 months ago

is there an option to specify Pre emptible instances while creating template? I couldn't find that. If so then D can't be true
upvoted 1 times

✉  **djgodzilla** 1 year, 4 months ago

Yes under management> Availability policy > preemptibility ON/OFF
upvoted 2 times

✉  **Charumathi** Most Recent  1 month ago

Selected Answer: C

C is the correct answer,

Install the workload in a compute engine VM, start and stop the instance as needed, because as per the question the VM runs for 30 hours, process can be performed offline and should not be interrupted, if interrupted we need to restart the batch process again. Preemptible VMs are cheaper, but they will not be available beyond 24hrs, and if the process gets interrupted the preemptible VM will restart.

upvoted 1 times

✉  **KapilDhamija** 2 months, 4 weeks ago

Selected Answer: C

C is the correct answer
upvoted 1 times

 **ryumada** 3 months ago

Selected Answer: C
The preemptible instance in GKE is same as Compute Engine Instance. They have same behavior that will be last for 24 hours.

Also, see the key here "...and must be restarted if interrupted.". That means the job will start from the scratch again if the preemptible instance terminated. So, you will just waste your preemptible instances because the job will never be finished.

<https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#overview>

upvoted 1 times

 **koko2314** 3 months, 2 weeks ago

Option D will achieve the goal here. If a preemptible VM goes down in 24 hours a new one will be built for running the batch process through the instance group configuration.

upvoted 1 times

 **ryumada** 3 months ago

"...and must be restarted if interrupted."

The job will start again from a scratch, then run again for another 24 hours in a new preemptible instance, the job will be terminated again after 24 hours, then start again from a scratch for another 24 hours in a new preemptible instance. This makes the process to be an infinite-loop process with wasted resources.

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

A preemptible VM is an instance that you can create and run at a much lower price than normal instances. However, Compute Engine might stop (preempt) these instances if it requires access to those resources for other tasks. Preemptible instances are excess Compute Engine capacity, so their availability varies with usage.

If your apps are fault-tolerant and can withstand possible instance preemptions then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances.

Hence, the correct answer is: Build an instance template configured to launch a Preemptible VM. Provision a managed instance group (MIG) from the template you just created. Adjust the Target CPU Utilization setting.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is correct, since job needs to run 30 hours.

upvoted 1 times

 **pfabio** 5 months, 3 weeks ago

Selected Answer: C

A - D are not possible because: For example, preemptible VMs can only run for up to 24 hours at a time (<https://cloud.google.com/compute/docs/instances/preemptible>)

I agree with C

upvoted 1 times

 **rsuresh27** 6 months, 2 weeks ago

The answer is C. Many people are saying D but this is incorrect. Preemptible instances are finite Compute Engine resources, so they might not always be available. If there is a long running job that must be restarted, it does not make sense to use preemptible VMs as if there are no resources available, we will not be able to even run or restart the job. Thus, C is correct.

upvoted 2 times

 **dinesh198728** 7 months, 3 weeks ago

Selected Answer: C

auto restart not in Prem VM

upvoted 1 times

 **ArunTaneja** 9 months, 2 weeks ago

Selected Answer: C

It should be C Because a Preemptible VM can't run more than 24 hours

upvoted 1 times

 **shputhan** 9 months, 3 weeks ago

Answer is C

- Since it is run on-premises server, the expectation is to run it on Compute Engine.
- Preemptible VMs are not an option as it runs for 24 hours only continuously. The batch job takes 30 hours and expects it to be restarted in case of interruptions. (Note it is restart & not resume)

upvoted 1 times

 **emv** 9 months, 3 weeks ago

If your apps are fault-tolerant and can withstand possible instance preemptions, then preemptible instances can reduce your Compute Engine costs significantly. For example, batch processing jobs can run on preemptible instances.

upvoted 1 times

 **movingsteady** 10 months ago

Its D...guys the the preemptible VM is restarted if its in a MIG

upvoted 1 times

 **Ridhanya** 11 months ago

D seems right

upvoted 2 times

 **alex90fin** 11 months, 2 weeks ago

C

VM preemptible are stoped after 24h mean task will never been completed wiht option A

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

I Choose - C

upvoted 2 times

Question #137

You are developing a new application and are looking for a Jenkins installation to build and deploy your source code. You want to automate the installation as quickly and easily as possible. What should you do?

- A. Deploy Jenkins through the Google Cloud Marketplace.
- B. Create a new Compute Engine instance. Run the Jenkins executable.
- C. Create a new Kubernetes Engine cluster. Create a deployment for the Jenkins image.
- D. Create an instance template with the Jenkins executable. Create a managed instance group with this template.

Correct Answer: C

Reference:

<https://cloud.google.com/solutions/jenkins-on-kubernetes-engine>*Community vote distribution*

A (100%)

filco72 Highly Voted 2 years, 2 months ago

I would choose A. Deploy Jenkins through the Google Cloud Marketplace.
as this is a well known opportunity on the GCP Marketplace

upvoted 26 times

ESP_SAP Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Installing Jenkins

In this section, you use Cloud Marketplace to provision a Jenkins instance. You customize this instance to use the agent image you created in the previous section.

Go to the Cloud Marketplace solution for Jenkins.

Click Launch on Compute Engine.

Change the Machine Type field to 4 vCPUs 15 GB Memory, n1-standard-4.

Machine type selection for Jenkins deployment.

Click Deploy and wait for your Jenkins instance to finish being provisioned. When it is finished, you will see:

Jenkins has been deployed.

https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine#installing_jenkins

upvoted 14 times

Charumathi Most Recent 1 month ago**Selected Answer: A**

A is correct answer,
To quickly deploy Jenkins, deploy it through google cloud marketplace.

upvoted 1 times

KapilDhamija 2 months, 4 weeks ago**Selected Answer: A**

remember as quickly as possible, also Google encourage things to be performed in minimal steps so A is the quickest and easiest choice

upvoted 1 times

kiwi123 2 months, 4 weeks ago

Go for A, the easiest

upvoted 1 times

pspandher 4 months, 1 week ago

This is Repeat Question.

upvoted 2 times

AzureDP900 4 months, 2 weeks ago

Cloud Market Place is fastest and best .. A is right

upvoted 1 times

pfabio 5 months, 3 weeks ago

Selected Answer: A

A - ... as quickly and easily as possible
upvoted 1 times

 **Razor** 9 months, 2 weeks ago

Selected Answer: A
A is right the rest is nonsense
upvoted 1 times

 **Raju44** 10 months, 2 weeks ago

As we need to automate installation of Jenkins as pre-requisite (upgrade etc for future), option C coz with Kubernetes you can automate installation using charts (easy to change few parameters)
upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: A
A is right. Using Cloud Marketplace
upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: A
I agree with A
upvoted 1 times

 **maggieli** 1 year ago

A. The faster and best solution.
upvoted 1 times

 **jackwillis** 1 year, 1 month ago

its A.
upvoted 1 times

 **vvkds** 1 year, 2 months ago

A is the shortest path.
upvoted 1 times

 **vinaybattula** 1 year, 5 months ago

Option A is the quickest way.
upvoted 1 times

 **arsh1916** 1 year, 5 months ago

A is fastest way
upvoted 1 times

Question #138

You have downloaded and installed the gcloud command line interface (CLI) and have authenticated with your Google Account. Most of your Compute Engine instances in your project run in the europe-west1-d zone. You want to avoid having to specify this zone with each CLI command when managing these instances.

What should you do?

- A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.
- B. In the Settings page for Compute Engine under Default location, set the zone to europe-west1-d.
- C. In the CLI installation directory, create a file called default.conf containing zone=europe-west1-d.
- D. Create a Metadata entry on the Compute Engine page with key compute/zone and value europe-west1-d.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/gcloud-compute>

Community vote distribution

A (90%)	10%
---------	-----

✉  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Change your default zone and region in the metadata server

Note: This only applies to the default configuration.

You can change the default zone and region in your metadata server by making a request to the metadata server. For example:

```
gcloud compute project-info add-metadata \
--metadata google-compute-default-region=europe-west1,google-compute-default-zone=europe-west1-b
```

The gcloud command-line tool only picks up on new default zone and region changes after you rerun the gcloud init command. After updating your default metadata, run gcloud init to reinitialize your default configuration.

https://cloud.google.com/compute/docs/gcloud-compute#change_your_default_zone_and_region_in_the_metadata_server
upvoted 25 times

✉  **xtian2900** 2 years, 1 month ago

does your comment imply that the answer is D ? i'm confused

upvoted 1 times

✉  **mahesh0049** 9 months, 3 weeks ago

every thing is correct in your explanation but instead of using gcloud compute command they used gcloud config.

upvoted 4 times

✉  **bobthebuilder55110** 3 months, 1 week ago

You can use the gcloud config set command here, https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client

upvoted 2 times

✉  **dtncl** 1 year ago

I agree the answer is A.

gcloud config - view and edit Cloud SDK properties

zone

Default zone to use when working with zonal Compute Engine resources.

<https://cloud.google.com/sdk/gcloud/reference/config>

upvoted 3 times

✉  **Examan1** 1 year, 8 months ago

Using gcloud config you can set the zone in your active configuration only. This setting does not apply to other gcloud configurations and does not become the default for the project.

Ref: <https://cloud.google.com/sdk/gcloud/reference/config/set>

So I believe correct answer is B as per <https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>

In the Cloud Console, go to the Settings page.

From the Zone drop-down menu, select a default zone.

upvoted 3 times

 **tavva_prudhvi** 1 year, 7 months ago

bro, it mentioned going into the console settings, not the compute engine settings!
To change your default region or zone:

In the Cloud Console, go to the Settings page.

Go to the Settings page

From the Region drop-down menu, select a default region.

From the Zone drop-down menu, select a default zone.

upvoted 1 times

 **jcloud965** 1 year, 3 months ago

This setting in the Cloud Console won't be taken into account for gcloud on your active config

upvoted 1 times

 **SSPC** Highly Voted  2 years, 2 months ago

I would go with the answer A

upvoted 10 times

 **Angel_99** Most Recent  2 months, 3 weeks ago

Selected Answer: A

It is clearly mentioned it is to be done via CLI not console

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: A

A, it is clearly mentioned it is to be done via CLI not console.

upvoted 1 times

 **pfabio** 5 months, 1 week ago

Selected Answer: A

Just run this command on CLI
gcloud config set compute/zone ZONE

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: A

Ans: A

"You want to avoid having to specify this zone with each CLI command"
its about CLI not console!

upvoted 2 times

 **injarapu** 8 months, 1 week ago

Selected Answer: B

<https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>

upvoted 1 times

 **oracle111** 8 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/compute/docs/gcloud-compute>

gcloud config set compute/zone ZONE

upvoted 3 times

 **ElenaL** 10 months ago

Selected Answer: A

in GCP qwiklabs, this is how they ask you set the zone primarily, so I go with A

upvoted 1 times

 **Raju44** 10 months, 2 weeks ago

Answer A. There is no file like default.config. In my labs i set default values Zone/Region/project etc default values stored in the following location

cat ~/.config/gcloud/configurations/config_default.

We can also set zone default using " gcloud config set compute/zone <zona_name>"

upvoted 1 times

✉ **jaffarali** 11 months ago

A is the right option
upvoted 1 times

✉ **Krishna0014** 11 months ago

Answer is A: Here we are talking about local CLI client and setting a default zone. It has nothing to do in cloud console any directory. simply can be done with below one command.

As per : https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client

"Set default region and zone in your local client

The default region and zone set in your local client override the default region and zone set in the metadata server."

To set the default region and zone in your local client, run the following commands:

gcloud config set compute/region REGION
gcloud config set compute/zone ZONE

upvoted 1 times

✉ **wh1t4k3r** 11 months ago

I would say A
<https://cloud.google.com/kubernetes-engine/docs/how-to/enable-oslogin?hl=nl#gcloud-config>
upvoted 1 times

✉ **rachee** 11 months, 1 week ago

A - https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client
upvoted 1 times

✉ **alaahakim** 11 months, 3 weeks ago

I Pick A
upvoted 1 times

✉ **Rameshfm** 1 year ago

I hope answer is D - Please refer the link : <https://cloud.google.com/compute/docs/gcloud-compute>
upvoted 1 times

Question #139

The core business of your company is to rent out construction equipment at large scale. All the equipment that is being rented out has been equipped with multiple sensors that send event information every few seconds. These signals can vary from engine status, distance traveled, fuel level, and more. Customers are billed based on the consumption monitored by these sensors. You expect high throughput up to thousands of events per hour per device and need to retrieve consistent data based on the time of the event. Storing and retrieving individual signals should be atomic. What should you do?

- A. Create a file in Cloud Storage per device and append new data to that file.
- B. Create a file in Cloud Filestore per device and append new data to that file.
- C. Ingest the data into Datastore. Store data in an entity group based on the device.
- D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.

Correct Answer: D*Community vote distribution*

D (88%)

13%

 **hiteshrup** Highly Voted 2 years, 1 month ago

Answer: D

Keyword need to look for

- "High Throughput",
- "Consistent",
- "Property based data insert/fetch like engine status, distance traveled, fuel level, and more." which can be designed in column,
- "Large Scale Customer Base + Each Customer has multiple sensor which send event in seconds" This will go for pera bytes situation,
- Export data based on the time of the event.
- Atomic

- o BigTable will fit all requirement.

- o DataStore is not fully Atomic

- o CloudStorage is not a option where we can export data based on time of event. We need another solution to do that

- o FireStore can be used with MobileSDK.

So go with Option D: Big Table

upvoted 34 times

 **[Removed]** 2 years, 1 month ago

Its File store not firestore. But still, the argument is correct here as file store will not be automatic like cloud store that no SQL.

If it would be Firestore I would have gone with it, then big data for throughout v/s cost.

upvoted 1 times

 **Hjameel** Highly Voted 2 years, 2 months ago

D is the best answer , Cloud Bigtable

upvoted 8 times

 **har_riy** 1 year, 8 months ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 5 times

 **Cornholio_LMC** Most Recent 1 month, 2 weeks ago

had this question today

upvoted 1 times

 **tomis2** 4 months ago

Selected Answer: D

Timeseries + IoT = Bigtable

upvoted 3 times

 **AzureDP900** 4 months, 2 weeks ago

This is related to IoT , it is no sql means BigTable. D is right

upvoted 1 times

 **pfabio** 5 months, 3 weeks ago

Selected Answer: D

D - Ideal for use cases such as personalization, ad tech, fintech, digital media, and IoT
<https://cloud.google.com/bigtable>

upvoted 2 times

 **JelloMan** 6 months ago

Selected Answer: D

Big table for the exact reasons as @hitrshrup mentioned. Time related data flowing in and out will point you directly to BigTable without reading anything else

upvoted 1 times

 **wael_tn** 6 months, 2 weeks ago

Selected Answer: D

When we want to store data based on timestamp => Cloud Bigtable

upvoted 1 times

 **Akash7** 7 months ago

Answer is D, large streaming data = bigtable

upvoted 1 times

 **zaxma** 7 months, 1 week ago

Selected Answer: C

I vote c with DataStore (maybe even better with firestore includes pubsub ability)

Looks like the internet agrees with me.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Answer: D

upvoted 2 times

 **CompteBidon91** 1 year, 4 months ago

D,

Keyword : IOT, Database, large scale = Big Table

upvoted 3 times

 **vinaybattula** 1 year, 5 months ago

D is correct. If Data is huge, unstructured and related time Bigtable is best option

upvoted 1 times

 **arsh1916** 1 year, 5 months ago

D is ans 100%

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.

upvoted 3 times

 **har_riy** 1 year, 8 months ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 3 times

 **har_riy** 1 year, 8 months ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 1 times

Question #140

Topic 1

You are asked to set up application performance monitoring on Google Cloud projects A, B, and C as a single pane of glass. You want to monitor CPU, memory, and disk. What should you do?

- A. Enable API and then share charts from project A, B, and C.
- B. Enable API and then give the metrics.reader role to projects A, B, and C.
- C. Enable API and then use default dashboards to view all projects in sequence.
- D. Enable API, create a workspace under project A, and then add projects B and C.

Correct Answer: C

Community vote distribution

-  **jlcclaude** Highly Voted 2 years, 2 months ago
D. workspaces is made for monitoring multiple projects.
upvoted 33 times
-  **Hjameel** Highly Voted 2 years, 2 months ago
D , Workspace to monitor multiple projects.
upvoted 10 times
-  **Khoka** 1 year, 12 months ago
<https://cloud.google.com/monitoring/workspaces>
upvoted 1 times
-  **Charumathi** Most Recent 1 month ago
Selected Answer: D
D is the correct answer,
Keep Project A as host project in workspace and Project B and C as Service Project, and monitor the metrics of the Project A for a centralized view.
upvoted 1 times
-  **ale_brd_** 1 month, 2 weeks ago
Selected Answer: D
Stackdriver workspaces are deprecated, now in the monitoring page of the Project you want, you need to select the "Scopes". Anyway he closest answer is D.

Scopes allow you to monitor multiple projects.

<https://cloud.google.com/monitoring/settings/multiple-projects>
upvoted 1 times
-  **KapilDhamija** 2 months, 4 weeks ago
Selected Answer: D
D should be the correct answer
upvoted 1 times
-  **AzureDP900** 4 months, 2 weeks ago
D is correct
upvoted 1 times
-  **Rutu_98** 5 months, 2 weeks ago
Selected Answer: D
D is correct
upvoted 1 times
-  **JelloMan** 6 months ago
D. One project must be the host (A) in this case and all others can be linked (B&C)
upvoted 2 times
-  **Ken701** 7 months, 3 weeks ago
Selected Answer: D
You can add/link multiple project to a single workspace for monitoring
upvoted 1 times

✉  **gkb** 9 months ago

Wanted to check, if somebody has appeared ACE exam recently and if yes, what is the percentage of questions that come from this site? Thanks for answering.

upvoted 5 times

✉  **Eminenza22** 9 months ago

D is correct

Workspace can monitor multiple projects but a Google Cloud project can be monitored by exactly 1 Workspace.

upvoted 2 times

✉  **HansKloss611** 9 months ago

Selected Answer: D

D is the correct

upvoted 2 times

✉  **alex90fin** 11 months, 2 weeks ago

workspace is mandatory for monitoring

upvoted 1 times

✉  **alaahakim** 11 months, 3 weeks ago

i Vote D

upvoted 2 times

✉  **maggieli** 1 year ago

I vote D. Group dashboards are limited to 40 charts. Charts on a group dashboard display data only from the members of the group.
<https://cloud.google.com/monitoring/groups>

upvoted 1 times

✉  **blan_ak** 1 year, 2 months ago

D should be the answer

upvoted 1 times

✉  **vinaybattula** 1 year, 5 months ago

Option D is correct. For monitoring first thing we should create workspace for one project and then can add more projects to the same workspace for monitoring.

upvoted 1 times

Question #141

You created several resources in multiple Google Cloud projects. All projects are linked to different billing accounts. To better estimate future charges, you want to have a single visual representation of all costs incurred. You want to include new cost data as soon as possible. What should you do?

- A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.
- B. Visit the Cost Table page to get a CSV export and visualize it using Data Studio.
- C. Fill all resources in the Pricing Calculator to get an estimate of the monthly cost.
- D. Use the Reports view in the Cloud Billing Console to view the desired cost information.

Correct Answer: A

Reference:

<https://cloud.google.com/billing/docs/how-to/visualize-data>

Community vote distribution

A (75%)	D (25%)
---------	---------

✉  **Zakaullah** Highly Voted 2 years, 2 months ago

Highly recommended website for exam prep.just passed the exam.

upvoted 31 times

✉  **IRobot** 2 years, 2 months ago

Hope the best for my exam tomorrow

upvoted 4 times

✉  **[Removed]** 2 years, 2 months ago

how did it go for you? I sit mine tomorrow.

upvoted 3 times

✉  **MoreOps** 1 year ago

Mine is tommorow lol.

upvoted 2 times

✉  **Pranad** 1 year ago

My exam is tomorrow, hope for the best.

upvoted 2 times

✉  **moses101** 12 months ago

how did things go?

upvoted 1 times

✉  **arsh1916** Highly Voted 1 year, 5 months ago

A is best answer

upvoted 8 times

✉  **PKookNN** Most Recent 3 weeks ago

Selected Answer: A

A is more correct as you can show data from multiple billing accounts as well as different projects.

upvoted 1 times

✉  **AwesomeGCP** 3 weeks, 6 days ago

A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.

upvoted 1 times

✉  **Charumathi** 1 month ago

Selected Answer: D

D is the correct Answer,

Use the Reports in the Cloud Billing console to view the billing information for the resources in multiple projects.

<https://cloud.google.com/billing/docs/how-to/view-linked>

upvoted 1 times

✉  **KapilDhamija** 2 months, 4 weeks ago

Selected Answer: A

Would go with A (Configure Billing Data Export to BigQuery and visualize the data in Data Studio)

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is the best ..

upvoted 1 times

 **Tirthankar17** 5 months ago

Selected Answer: A

A is correct

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct

upvoted 6 times

 **GCP_Student1** 1 year, 8 months ago

A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.

upvoted 6 times

 **nitinz** 1 year, 8 months ago

A, Any time you see question about billing and has BQ - 99.99% of time correct answer is that has BQ in it.

upvoted 7 times

 **victory108** 1 year, 9 months ago

A - Configure Billing Data Export to BigQuery and visualize the data in Data Studio.

upvoted 3 times

 **sarahf** 1 year, 10 months ago

I was thinking A first.

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

"Cloud Billing export to BigQuery enables you to export detailed Google Cloud billing data (such as usage, cost estimates, and pricing data) automatically throughout the day to a BigQuery dataset that you specify."

It also says that:

"Be aware that your BigQuery dataset only reflects Google Cloud billing data incurred from the date you set up Cloud Billing export, and after. That is, Google Cloud billing data is not added retroactively, so you won't see Cloud Billing data from before you enable export."

But I am leaning towards alternative D after reading this page: <https://cloud.google.com/billing/docs/how-to/reports>.

You don't have to export/import to see data and you can see costs from different projects. All you need is the permission:

`billing.accounts.getSpendingInformation`. It can forecast future costs too. According to the video at the same page, BigQuery exports to Data Studio is for deeper cost analysis.

upvoted 6 times

 **raramos96** 1 year, 10 months ago

As we can read in the 1st paragraph at <https://cloud.google.com/billing/docs/how-to/reports>, Reports View from Cloud Billing Console shows info about a single Billing Account.

This question says:

- "all projects are linked to different billing accounts".
- "we want to have a single visual representation".

So, in my opinion, the best answer is [A] (BQ+data studio)

upvoted 7 times

 **don_v** 1 year, 11 months ago

I guess it's A.

I'm not sure how to get that info for multiple billing accounts with the billing report as they say "All projects are linked to different billing accounts."

upvoted 1 times

 **Bhagirathi** 1 year, 11 months ago

A or D which one to choose ?

upvoted 2 times

 **Lieder** 1 year, 11 months ago

Why not make use of Cloud Billing console? In the reports you get a dashboard with all the information needed (also forecasts) - across all billing accounts. And it is up to date - new resources are billed immediately. Everything is in place. No additional effort - like creating a dashboard in Data Studio.

I would go for D.

upvoted 4 times

 **[Removed]** 2 years, 1 month ago

I am also favour A, where you can use visualise the cost. But exporting and visualizing will show past data. We are talking about "estimating the further charges", that can be done by AI on big query else the closest will be pricing calculator so C. But since we don't know about the variable that will add on the cost that is you, for example, your data traffic is going to increase we can still go with A as answer here.

upvoted 1 times

Question #142

Your company has workloads running on Compute Engine and on-premises. The Google Cloud Virtual Private Cloud (VPC) is connected to your WAN over a Virtual Private Network (VPN). You need to deploy a new Compute Engine instance and ensure that no public Internet traffic can be routed to it. What should you do?

- A. Create the instance without a public IP address.
- B. Create the instance with Private Google Access enabled.
- C. Create a deny-all egress firewall rule on the VPC network.
- D. Create a route on the VPC to route all traffic to the instance over the VPN tunnel.

Correct Answer: B

Get private access to Google services, such as storage, big data, analytics, or machine learning, without having to give your service a public IP address.

Reference:

<https://cloud.google.com/vpc>

Community vote distribution

A (86%)	14%
---------	-----

 [Removed] Highly Voted 2 years, 2 months ago

A for sure

B - this allows internal communication, but does nothing to limit public traffic

C - deny all is nice, but it's for egress -- we're looking for ingress

D - this is way to invasive and doesn't explicitly address the issue of preventing public internet traffic from reaching your instance -- if it does, someone let me know how.

upvoted 27 times

 **MohammedGhouse** Highly Voted 2 years, 2 months ago

A: answer looks right

upvoted 11 times

 **Sam98845** Most Recent 2 weeks ago

should be A. VMs cannot communicate over the internet without a public IP address. Private Google Access permits access to Google APIs and services in Google's production infrastructure.

<https://cloud.google.com/vpc/docs/private-google-access>

upvoted 1 times

 **kailash** 4 weeks ago

Selected Answer: A

Elimination

upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: D

D is the right answer, with private google access for on-premises host,

Private Google Access for on-premises hosts

On-premises hosts with or without external IP addresses. Connect to Google APIs and services, from your on-premises network, through a Cloud VPN tunnel or Cloud Interconnect by using one of the Private Google Access-specific domains and VIPs. The Google services that you can access depend on which Private Google Access-specific domain you use. Use this option to connect to Google APIs and services through a VPC network. This method doesn't require your on-premises hosts to have external IP addresses.

please refer to the link below for more insights,

<https://cloud.google.com/vpc/docs/private-google-access-hybrid>

upvoted 1 times

 **tomis2** 4 months ago

Selected Answer: A

Through elimination - A

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right

upvoted 1 times

 **badrik** 6 months ago

Selected Answer: A

AAAAAAA

upvoted 1 times

 **HansKloss611** 9 months, 1 week ago

Selected Answer: A

A is correct

upvoted 3 times

 **Yaseed** 11 months, 2 weeks ago

PRIVATE INSTANCE WITH A PUBLIC LOADBALANCER WOULD GO PUBLIC! IT'S CONFUSING

upvoted 2 times

 **tvinay** 11 months ago

Why are you shouting here?

upvoted 10 times

 **lxs** 1 year ago

The question is about ingress traffic from Internet

- A - If the VM does not have public IP it is not routable from Internet. Correct answer
- B - It is about how to access Google Services API. It does not tell about ingress Internet traffic
- C - It is about egress traffic

D - It could be but we do not know anything about Internet ingress traffic to on prem. What's more default route tells about egress traffic to Internet. Nothing how Internet can access Compute instance.

Correct answer is A.

upvoted 4 times

 **arsh1916** 1 year, 5 months ago

A is correct

upvoted 4 times

 **Devgela** 1 year, 7 months ago

A for me

upvoted 3 times

 **GCP_Student1** 1 year, 8 months ago

A. Create the instance without a public IP address.

upvoted 3 times

 **victory108** 1 year, 9 months ago

A - Create the instance without a public IP address.

upvoted 2 times

 **DucSiu** 1 year, 9 months ago

A. Create the instance without a public IP address.

upvoted 1 times

 **guid1984** 1 year, 9 months ago

Seems like a multi-select option questions in this case its A+D

upvoted 2 times

Question #143

Your team maintains the infrastructure for your organization. The current infrastructure requires changes. You need to share your proposed changes with the rest of the team. You want to follow Google's recommended best practices. What should you do?

- A. Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket.
- B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories.
- C. Apply the changes in a development environment, run gcloud compute instances list, and then save the output in a shared Storage bucket.
- D. Apply the changes in a development environment, run gcloud compute instances list, and then save the output in Cloud Source Repositories.

Correct Answer: B*Community vote distribution*

B (58%)

A (42%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Connecting to Cloud Storage buckets

Cloud Storage is a flexible, scalable, and durable storage option for your virtual machine instances. You can read and write files to Cloud Storage buckets from almost anywhere, so you can use buckets as common storage between your instances, App Engine, your on-premises systems, and other cloud services.

<https://cloud.google.com/compute/docs/disks/gcs-buckets>

Why not (B)?

Caution

Cloud Source Repositories are intended to store only the source code for your app and not user or personal data. Don't store any Core App Engine Customer Data (as defined in your License Agreement) in Cloud Source Repositories.

<https://cloud.google.com/source-repositories/docs/features>

upvoted 29 times

 **ESP_SAP** 1 month, 4 weeks ago

Correction. 2 years later.

Correct Answer is (B):

Using Recommendations for Infrastructure as Code:

This tutorial uses Hashicorp Terraform as the IaC tool, however the architectural patterns and components used in the described automation pipeline can be leveraged even if you are using a different IaC management tool such as Deployment Manager.

Cloud Build

Purpose: Google Cloud Build automates the deployment of infrastructure based on the changes made to the IaC manifests per policy intelligence recommendations.

Access control: The Cloud Build service account must have the right set of permissions to interact with resources in your test project.

See the documentation for configuring a Cloud Build service account.

GitHub

Purpose: The IaC repository uses GitHub for source control. The IaC repository in GitHub is integrated with Cloud Build. When commits are made to the master branch, a Cloud Build job is triggered to run a set of preconfigured tasks.

<https://cloud.google.com/recommender/docs/tutorial-iac#overview>

upvoted 11 times

 **bobthebuilder55110** 3 months, 1 week ago

So, what do you think a deployment manager template will have? It will definitely have the source code for the source code/config for the template and if you are sharing with the team that is even better since that means that now people can collaborate that makes more sense to go with B over A here.

upvoted 1 times

 **mav3r1ck** 3 months ago

In reality, will you proposed the answer in "A" to your boss? I bet he will show you the exit door. Storing IaC in Cloud Storage? It's like telling to store CloudFormation in S3 instead of CodeCommit.

upvoted 1 times

 **gcpengineer** 1 year, 3 months ago

B is the ans

upvoted 3 times

 **SSPC** Highly Voted 2 years, 2 months ago

B is correct. <https://cloud.google.com/source-repositories/docs/features>

upvoted 23 times

 **AmitKM** 2 years, 2 months ago

Using Cloud Storage Repos, you can add comments and describe your changes to the team. Hence this might be a better option.

upvoted 3 times

 **magistrum** 1 year, 10 months ago

I don't see how you can do this when I tried creating:

Add code to your repository

info

Your repository is currently empty. Add some code using a selected method and then refresh your browser. Contents added to this repository can take some time to show up in search results. Learn more.

Select an option to push code to your repository:

Push code from a local Git repository

Clone your repository to a local Git repository

upvoted 2 times

 **diasporabro** Most Recent 2 weeks, 6 days ago

Selected Answer: B

Thought it was A for a sec, the realized Cloud Source is similar to Github/CodeCommit... So, B is a better choice

upvoted 1 times

 **AwesomeGCP** 3 weeks, 6 days ago

B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories

upvoted 1 times

 **KapilDhamija** 2 months, 4 weeks ago

Selected Answer: A

A is the right choice, Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket

upvoted 1 times

 **tomis2** 4 months ago

Selected Answer: B

Don't get confused - it is B. IaC is a matter of version control system like GIT.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right .. Showing Deployment Manager templates to your team will allow you to define the changes you want to implement in your cloud infrastructure. You can use Cloud Source Repositories to store Deployment Manager templates and collaborate with your team. Cloud Source Repositories are fully-featured, scalable, and private Git repositories you can use to store, manage and track changes to your code.

Hence, the correct answer is: Create Deployment Manager templates to define the proposed changes and save them into Cloud Source Repositories.

upvoted 2 times

 **SergijF** 5 months, 1 week ago

Selected Answer: B

Its B.

upvoted 1 times

 **Rutu_98** 5 months, 2 weeks ago

Selected Answer: A

A is correct Answer

Connecting to Cloud Storage buckets

Cloud Storage is a flexible, scalable, and durable storage option for your virtual machine instances. You can read and write files to Cloud Storage buckets from almost anywhere, so you can use buckets as common storage between your instances, App Engine, your on-premises systems, and other cloud services.

upvoted 1 times

 **nhadi82** 5 months, 3 weeks ago

Selected Answer: A

I would go for A as following

[#hosting_external_templates](https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally?hl=en)

upvoted 1 times

 **vruizrob** 8 months, 3 weeks ago

Correct answer: A. <https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally?hl=en>. You can host external templates in Google Cloud Storage or in a publicly-accessible location, such as GitHub.

upvoted 2 times

 **oracle111** 8 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally>

"Hosting templates in Google Cloud Storage let you control who has access to the templates because you can set access control options on the respective Cloud Storage bucket or the specific template."

In contrast, hosting your templates on other publicly-accessible URLs requires that you make the templates widely available."

It's all about access control. Using a publicly-accessible URL might not be ok.
Versioning we have on buckets, and as well on git (not comparable really).

Otherwise I understand the need of having the IAC on repo (eg. Terraform).

upvoted 1 times

✉ **[Removed]** 9 months ago

1. You can host external templates in Google Cloud Storage or in a publicly-accessible location, such as GitHub. As long as there is a URL to the file and you have access to the templates, Deployment Manager will be able to download it.

Ref:

https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally#hosting_external_templates

2. Cloud Source Repositories are private Git repositories hosted on Google Cloud. These repositories let you develop and deploy an app or service in a space that provides collaboration and version control for your code.

Ref:

<https://cloud.google.com/source-repositories/docs/quickstart>

Based on above 2 points, and "You need to share your proposed changes with the rest of the team." mentioned in question.
Option A should be better.

upvoted 1 times

✉ **raaj_p** 9 months, 1 week ago

Selected Answer: B

Just as Terraform, the source code needs to be in a Git style repository. This is Infra as code question. B is correct

upvoted 3 times

✉ **HansKloss611** 9 months, 1 week ago

B is correct. Cloud source repository is the first thing that you can think about

upvoted 1 times

✉ **RazOr** 9 months, 2 weeks ago

Selected Answer: B

B is right because in Cloud Source Repo you store code and the state can be stored in a bucket.

The question aims at code management and version control so B is the right way for that.

upvoted 1 times

✉ **mrNicky** 9 months, 3 weeks ago

Why people choose A. Using cloud repo is perfect for that

- Sharing code (yes deployment manager template is code!, IAC like terraform)
- We use google best practice, we share code through repo not putting this in a bucket and share the URI
- My company works with terraform, we store all code in git repo, so this is the same logic in google cloud: store and share code.

A it will be perfect if the question was "the fastest solution"

upvoted 3 times

Question #144

You have a Compute Engine instance hosting an application used between 9 AM and 6 PM on weekdays. You want to back up this instance daily for disaster recovery purposes. You want to keep the backups for 30 days. You want the Google-recommended solution with the least management overhead and the least number of services. What should you do?

- A. 1. Update your instances' metadata to add the following value: snapshot.schedule: 0 1 * * * 2. Update your instances' metadata to add the following value: snapshot.retention: 30
- B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: - Schedule frequency: Daily - Start time: 1:00 AM 2:00 AM - Autodelete snapshots after: 30 days
- C. 1. Create a Cloud Function that creates a snapshot of your instance's disk. 2. Create a Cloud Function that deletes snapshots that are older than 30 days. 3. Use Cloud Scheduler to trigger both Cloud Functions daily at 1:00 AM.
- D. 1. Create a bash script in the instance that copies the content of the disk to Cloud Storage. 2. Create a bash script in the instance that deletes data older than 30 days in the backup Cloud Storage bucket. 3. Configure the instance's crontab to execute these scripts daily at 1:00 AM.

Correct Answer: B*Community vote distribution*

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Creating scheduled snapshots for persistent disk

This document describes how to create a snapshot schedule to regularly and automatically back up your zonal and regional persistent disks. Use snapshot schedules as a best practice to back up your Compute Engine workloads. After creating a snapshot schedule, you can apply it to one or more persistent disks.

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>

upvoted 37 times

 **[Removed]** 2 years, 2 months ago

Definitely B.

With something like this, you should not have to write any custom scripts, custom functions, or cron jobs. This is google's way of saying 'hey, we've already built that stuff in to our snapshot schedules feature.'

upvoted 7 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

B is more appropriate

upvoted 1 times

 **nhadi82** 5 months, 3 weeks ago

Selected Answer: B

Correct Answer B

upvoted 1 times

 **Ridhanya** 11 months ago

it is b. we cannot define snapshot config in instance metadata.

VM instance metadata is used only for:

startup and shutdown scripts

host maintenance

guest attributes

upvoted 4 times

 **gcpengineer** 1 year, 3 months ago

why not C?

upvoted 2 times

 **Gianfry** 11 months, 2 weeks ago

The question calls for "Google-recommended solution with the least management overhead and the least number of services"

upvoted 2 times

 **arsh1916** 1 year, 5 months ago

B is correct

upvoted 2 times

✉ **Hi2ALL** 1 year, 7 months ago

B is more realistic approach

upvoted 3 times

✉ **GCP_user** 1 year, 7 months ago

B is the best option so far. However just wonder this: Schedule frequency: Daily " Start time: 1:00 AM " 2:00 AM" Autodelete snapshots: after 30 days; For Saturday and Sunday it will be a waste of resource to create snapshots since the instance is running during weekdays.

upvoted 2 times

✉ **GCP_Student1** 1 year, 8 months ago

B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 2 times

✉ **DucSiu** 1 year, 9 months ago

It's B

upvoted 1 times

✉ **Bhagirathi** 1 year, 11 months ago

B for sure, any doubt?

upvoted 1 times

✉ **swatititame** 1 year, 11 months ago

• B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 2 times

✉ **hiteshrup** 2 years, 1 month ago

Has to B...

- scripting which means overhead and maintenance So Option D neglected.
- Cloud function, adding another service which not require.
- MetaData, I am not sure if meta data have something to define as cron job schedule. So not be an option.

upvoted 1 times

✉ **stepkurniawan** 2 years, 2 months ago

Either A or B, I am not sure.

I have tested B and it works, but I dont know the name of the key to do A.

upvoted 1 times

✉ **juliandm** 2 years, 2 months ago

B is right <https://www.cloudbooklet.com/backup-google-cloud-vm-automatically-with-snapshot-schedules>

upvoted 2 times

✉ **baonguyen** 2 years, 2 months ago

The Ans is B.

Checked on my GCP account.

upvoted 4 times

✉ **SSPC** 2 years, 2 months ago

I would go with the answer "D" but I don't totally sure

upvoted 1 times

✉ **[Removed]** 2 years, 1 month ago

You need to write the script and test it that is a big step.

upvoted 2 times

✉ **[Removed]** 2 years, 1 month ago

So will go with B

upvoted 3 times

Question #145

Your existing application running in Google Kubernetes Engine (GKE) consists of multiple pods running on four GKE n1 "standard" 2 nodes. You need to deploy additional pods requiring n2 "highmem" 16 nodes without any downtime. What should you do?

- A. Use gcloud container clusters upgrade. Deploy the new services.
- B. Create a new Node Pool and specify machine type n2 "highmem16". Deploy the new pods.
- C. Create a new cluster with n2 "highmem16" nodes. Redeploy the pods and delete the old cluster.
- D. Create a new cluster with both n1 "standard2" and n2 "highmem16" nodes. Redeploy the pods and delete the old cluster.

Correct Answer: B

Community vote distribution

B (100%)

 **GCP_Student1** Highly Voted 1 year, 7 months ago
B is correct answer, read below from google docs;

This tutorial demonstrates how to migrate workloads running on a Google Kubernetes Engine (GKE) cluster to a new set of nodes within the same cluster without incurring downtime for your application. Such a migration can be useful if you want to migrate your workloads to nodes with a different machine type.

Background

A node pool is a subset of machines that all have the same configuration, including machine type (CPU and memory) authorization scopes. Node pools represent a subset of nodes within a cluster; a container cluster can contain one or more node pools.

When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.

To migrate your workloads without incurring downtime, you need to:

- Mark the existing node pool as unschedulable.
- Drain the workloads running on the existing node pool.
- Delete the existing node pool.

https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool#creating_a_node_pool_with_large_machine_type
upvoted 22 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago
B makes perfect sense.
upvoted 1 times

 **Tirthankar17** 5 months ago
Selected Answer: B
B is correct
upvoted 1 times

 **arsh1916** 1 year, 5 months ago
B is correct
upvoted 4 times

 **Jacky_YO** 1 year, 7 months ago
ANS : B

1. The title did not say to delete four GKE n1.

upvoted 4 times

 **pondai** 1 year, 7 months ago
B, You need to create new node pool for cluster
upvoted 3 times

 **dunhill** 1 year, 8 months ago
I guess it's B. I couldn't find resize parameter under cluster upgrade. C and D are incorrect because it's no need to create new cluster.
upvoted 2 times

 **GCP_Student1** 1 year, 8 months ago
A. Use gcloud container clusters upgrade. Deploy the new services.
upvoted 1 times

✉️ **GCP_Student1** 1 year, 7 months ago

I take it back, the correct answer is "B"

B. Create a new Node Pool and specify machine type n2highmem16. Deploy the new pods.

upvoted 2 times

✉️ **Jamaal_a** 1 year, 8 months ago

Answer is B - When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.

upvoted 3 times

Question #146

You have an application that uses Cloud Spanner as a database backend to keep current state information about users. Cloud Bigtable logs all events triggered by users. You export Cloud Spanner data to Cloud Storage during daily backups. One of your analysts asks you to join data from Cloud Spanner and Cloud Bigtable for specific users. You want to complete this ad hoc request as efficiently as possible. What should you do?

- A. Create a dataflow job that copies data from Cloud Bigtable and Cloud Storage for specific users.
- B. Create a dataflow job that copies data from Cloud Bigtable and Cloud Spanner for specific users.
- C. Create a Cloud Dataproc cluster that runs a Spark job to extract data from Cloud Bigtable and Cloud Storage for specific users.
- D. Create two separate BigQuery external tables on Cloud Storage and Cloud Bigtable. Use the BigQuery console to join these tables through user fields, and apply appropriate filters.

Correct Answer: B*Community vote distribution*

D (73%)	B (27%)
---------	---------

 **AmitKM** Highly Voted 2 years, 2 months ago

I think it should be D. <https://cloud.google.com/bigquery/external-data-sources>

upvoted 33 times

 **SSPC** 2 years, 2 months ago

The question says: " Join data from Cloud Spanner and Cloud Bigtable for specific users" You can see the Google documentation in the link <https://cloud.google.com/spanner/docs/export>

upvoted 3 times

 **Eshkrkrkr** 1 year, 12 months ago

Oh my god, SSPC read you your links!

The process uses Dataflow and exports data to a folder in a Cloud Storage bucket. The resulting folder contains a set of Avro files and JSON manifest files. And what next? I will tell - next you read below: Compute Engine: Before running your export job, you must set up initial quotas for Recommended starting values are:

CPUs: 200

In-use IP addresses: 200

Standard persistent disk: 50 TB

Still think its A?

upvoted 4 times

 **punjabishiva123** 1 year, 11 months ago

Hi,

Could u pls mail me all the correct answer to shivaasingh1104@gmail.com,

Really appreciated

Anyone pls help

upvoted 1 times

 **ash500** 1 year, 10 months ago

can you pls also share the correct answers with me

upvoted 1 times

 **kp0916** 1 year, 8 months ago

Does anyone have list of correct answers,let me know

upvoted 1 times

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (D):

Introduction to external data sources

This page provides an overview of querying data stored outside of BigQuery.

<https://cloud.google.com/bigquery/external-data-sources>

upvoted 24 times

 **ESP_SAP** 2 years, 2 months ago

BigQuery offers support for querying data directly from:

Bigtable

Cloud Storage
Google Drive
Cloud SQL (beta)
upvoted 5 times

 **djgodzilla** 1 year, 4 months ago

but here we're not talking about joining Cloud Storage and Cloud Bigtable external tables.
the join happens between a distributed relational database (Spanner) and key-value NoSQL Database (BigTable). how's converting Spanner to cloud storage an implicit and trivial step.

upvoted 1 times

 **djgodzilla** 1 year, 4 months ago

"The Cloud Spanner to Cloud Storage Text template is a batch pipeline that reads in data from a Cloud Spanner table, optionally transforms the data via a JavaScript User Defined Function (UDF) that you provide, and writes it to Cloud Storage as CSV text files."
<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#cloudspannertogcstext>

"The Dataflow connector for Cloud Spanner lets you read data from and write data to Cloud Spanner in a Dataflow pipeline"
<https://cloud.google.com/spanner/docs/dataflow-connector>

upvoted 2 times

 **ryzior** 8 months ago

update:
BigQuery supports the following external data sources:
Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive
upvoted 4 times

 **[Removed]** 2 years, 1 month ago

As per your comment D is the answer.
I also agree.
But can BigQuery read backed up data? , as we have backup data on Cloud storage, did not get any evidence in the link you shared.
upvoted 2 times

 **anolive** Most Recent 6 days, 18 hours ago

Selected Answer: D

I think is D, but not 100% sure, because D does not have any information about the specific user like others options.
upvoted 1 times

 **Charumathi** 1 month ago

Selected Answer: D

D is the correct answer,
An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Amazon S3
Azure Storage
Cloud Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive
upvoted 1 times

 **DualCore573** 1 month ago

Selected Answer: D

D makes sense as the BigQuery external tables are made for such use cases. and "efficient" keyword makes sense to use this way as resources used are less.
upvoted 1 times

 **soaresleo** 3 months ago

First of all, using Dataflow can perhaps be effective, but NOT efficient, specially because of costs.

Second:

"To query Cloud Bigtable data using a permanent external table, you: Create a table definition file (for the API or bq command-line tool); Create a table in BigQuery linked to the external data source; Query the data using the permanent table."

Source: <https://cloud.google.com/bigquery/docs/external-data-bigtable#:~:text=To%20query%20Cloud%20Bigtable%20data,data%20using%20the%20permanent%20table>

Third:

"To query a Cloud Storage external data source, provide the Cloud Storage URI path to your data and create a table that references the data source."

Source:
<https://cloud.google.com/bigquery/docs/external-data-cloud-storage>

Correct answer: D.
upvoted 1 times

✉ **jeffangel28** 3 months, 2 weeks ago

Selected Answer: D

"efficiently as possible" -> use the least amount of resources and achieve the same result... so I think it's D
upvoted 1 times

✉ **tomis2** 4 months ago

Selected Answer: D

Most "cloud" solution is D
upvoted 1 times

✉ **sabbella** 7 months, 3 weeks ago

Selected Answer: D

option d
upvoted 1 times

✉ **sabbella** 7 months, 3 weeks ago

Selected Answer: D

Option is D
upvoted 1 times

✉ **rlijhk** 9 months, 1 week ago

Selected Answer: B

I think it is B.
The data in Cloud storage is not up to date as backup window is daily. SO, there are chances is missing one day worth of data.
As it is mentioned as "efficiently" instead of quickly, I would choose "B".
upvoted 3 times

✉ **obeythefist** 8 months, 1 week ago

How does this create a "join" between the two tables?

upvoted 1 times

✉ **BigQuery** 7 months, 4 weeks ago

why do you think one cannot join 2 subsets of data in dataflow Its meant for processing sets of data.
upvoted 1 times

✉ **sasithra** 9 months, 1 week ago

B is correct
upvoted 1 times

✉ **raaj_p** 9 months, 1 week ago

Selected Answer: D

An external data source is a data source that you can query directly from BigQuery, even though the data is not stored in BigQuery storage.

BigQuery supports the following external data sources:

Bigtable
Cloud Spanner
Cloud SQL
Cloud Storage
Drive
upvoted 1 times

✉ **Ridhanya** 11 months ago

I think it has to be option A because Dataflow needs to be used for bigtable export and cloud spanner data is already backed up in cloud storage on a daily basis
upvoted 1 times

✉ **Ridhanya** 11 months ago

Um, the questions says "join" so D seems right. I take back my previous
upvoted 2 times

✉ **maggieli** 1 year ago

I vote D.
upvoted 1 times

✉ **vamgcp** 1 year, 1 month ago

Option B - <https://cloud.google.com/bigquery/external-data-sources>

upvoted 2 times

 **BalSan** 1 year, 4 months ago

One of the key word here is "ad hoc" request.... It doesn't make sense to create Dataflow Job for an adhoc request, ruling out A & B. C doesn't make sense for this use case and hence the answer is D

upvoted 3 times

Question #147

You are hosting an application from Compute Engine virtual machines (VMs) in us`"central1`"a. You want to adjust your design to support the failure of a single Compute Engine zone, eliminate downtime, and minimize cost. What should you do?

- A. " Create Compute Engine resources in us`"central1`"b. " Balance the load across both us`"central1`"a and us`"central1`"b.
- B. " Create a Managed Instance Group and specify us`"central1`"a as the zone. " Configure the Health Check with a short Health Interval.
- C. " Create an HTTP(S) Load Balancer. " Create one or more global forwarding rules to direct traffic to your VMs.
- D. " Perform regular backups of your application. " Create a Cloud Monitoring Alert and be notified if your application becomes unavailable. " Restore from backups when notified.

Correct Answer: A*Community vote distribution*

A (100%)

 **GCP_Student1** Highly Voted 1 year, 8 months ago

A. Create Compute Engine resources in us "central1" "b. " Balance the load across both us "central1" "a and us "central1" "b.
upvoted 14 times

 **arsh1916** Highly Voted 1 year, 5 months ago

A is best option
upvoted 6 times

 **Angel_99** Most Recent 2 months, 3 weeks ago

Selected Answer: A
A is best option
upvoted 1 times

 **abirroy** 3 months ago

Selected Answer: A
A is the best option
upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

A is fine.
upvoted 1 times

 **obeythefist** 8 months, 1 week ago

This seems straightforward. "A" is the only answer that involves putting instances in more than one zone!

- A. Yes, creating instances in another zone and balancing the loads will fix this problem
- B. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
- C. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
- D. Wrong. This keeps all the instances in one zone, but the question says we want to protect against zone failures.
upvoted 5 times

 **[Removed]** 8 months, 4 weeks ago

Option A.
Create VMs across more than one region and zone so that you have alternative VMs to point to if a zone or region containing one of your VMs is disrupted. If you host all your VMs in the same zone or region, you won't be able to access any of those VMs if that zone or region becomes unreachable.
<https://cloud.google.com/compute/docs/tutorials/robustsystems#distribute>
upvoted 2 times

 **Ridhanya** 11 months ago

A is correct because we have to eliminate single zone failure problem
upvoted 1 times

 **Gianfry** 11 months, 2 weeks ago

Why not "B" selecting "Regional (multi zone)" ?
"Regional (multiple zone) coverage. Regional MIGs let you spread app load across multiple zones. This replication protects against zonal failures. If that happens, your app can continue serving traffic from instances running in the remaining available zones in the same region."
<https://cloud.google.com/compute/docs/instance-groups/>
upvoted 4 times

 **kimharsh** 11 months, 2 weeks ago

it should be B , but because it specify the one Zone we can't pick this answer , the closest other option is A
upvoted 1 times

 **mj98** 1 year, 7 months ago

Can someone explain how A?

upvoted 2 times

 **jabrrJ68w02ond1** 11 months, 2 weeks ago

Other options do not prepare you for zonal outages

upvoted 1 times

 **tifo16** 1 year, 6 months ago

in order to remediate to the problem of single point of failure, we have to replicate VMs within multiple zones. Only A choice consider this concern

upvoted 10 times

 **nzexamtopics** 1 year, 7 months ago

A? Really? how?

upvoted 1 times

 **NARWAL** 1 year, 7 months ago

A is correct.

upvoted 4 times

 **victory108** 1 year, 8 months ago

A - " Create Compute Engine resources in us-central1-a. " Balance the load across both us-central1-a and us-central1-b.

upvoted 2 times

Question #148

A colleague handed over a Google Cloud Platform project for you to maintain. As part of a security checkup, you want to review who has been granted the Project Owner role. What should you do?

- A. In the console, validate which SSH keys have been stored as project-wide keys.
- B. Navigate to Identity-Aware Proxy and check the permissions for these resources.
- C. Enable Audit Logs on the IAM & admin page for all resources, and validate the results.
- D. Use the command `gcloud projects get-iam-policy` to view the current role assignments.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>*Community vote distribution*

D (100%)

ESP_SAP Highly Voted 2 years, 2 months ago

Correct Answer is (D):

A simple approach would be to use the command flags available when listing all the IAM policy for a given project. For instance, the following command:

```
'gcloud projects get-iam-policy $PROJECT_ID --flatten="bindings[].members" --format="table(bindings.members)" --filter="bindings.role:roles/owner'"
```

outputs all the users and service accounts associated with the role ‘roles/owner’ in the project in question.

<https://groups.google.com/g/google-cloud-dev/c/Z6sZs7TvygQ?pli=1>

upvoted 36 times

MohammedGhouse Highly Voted 2 years, 2 months ago

D: is the answer

upvoted 13 times

SSPC 2 years, 2 months ago

D is the correct.

upvoted 3 times

yurstev 1 year, 10 months ago

D IS THE ANSWER

upvoted 4 times

tomis2 Most Recent 4 months ago**Selected Answer: D**

gcloud iam get-iam-policy

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

D is right

upvoted 1 times

Rutu_98 5 months, 2 weeks ago**Selected Answer: D**

Answer is D

upvoted 2 times

somenick 7 months, 2 weeks ago**Selected Answer: D**

gcloud projects get-iam-policy \$PROJECT_ID

upvoted 1 times

obeythefist 8 months, 1 week ago

I chose D by a process of elimination. Here's my take:

A. There's more than one way to access an instance than just the SSH keys, and SSH keys have nothing to do with Project Owner role.

- B. Barking up the wrong tree here, Identity-Aware Proxy is more for remotely accessing resources, rather than Project Owner IAM roles.
C. This will only work if everyone who is a Project Owner accesses the system so you can see them in the logs. What if a Project Owner doesn't access the Project for a while? How long will you wait? Nope.
D. By elimination, this is the best result.

upvoted 5 times

-  **BigQuery** 7 months, 4 weeks ago

NICE EXPLANATION; WAY TO GO D

upvoted 1 times

-  **HansKloss611** 9 months, 1 week ago

Selected Answer: D

D is correct

upvoted 1 times

-  **PR0704** 11 months, 3 weeks ago

how can the admin be so inconsistent throughout with the answers..not good ..its so confusing

upvoted 2 times

-  **tvinay** 11 months ago

Confusion!! that's the main goal here so that we all go to the docs and Study hard xD

upvoted 1 times

-  **mohamedmahmoudf97** 1 year, 1 month ago

D is the correct answer

upvoted 1 times

-  **arsh1916** 1 year, 5 months ago

D is correct

upvoted 2 times

-  **GCP_Student1** 1 year, 8 months ago

D. Use the command gcloud projects get "iam" policy to view the current role assignments.

upvoted 3 times

-  **Bhagirathi** 1 year, 10 months ago

D 200%

upvoted 2 times

-  **Bhagirathi** 1 year, 11 months ago

anyone will be confused - solution says one answer

same time, all you guys have different choices here. what to take from this ?

upvoted 1 times

-  **swatititame** 1 year, 11 months ago

D. Use the command gcloud projects get "iam" policy to view the current role assignments.

upvoted 1 times

-  **adeyemi5700** 2 years ago

C is the likely answer. With D you see the current users with permission. With audit log you see those with prev. and current permission.

upvoted 1 times

-  **[Removed]** 2 years, 1 month ago

D is the answer but the command is wrong

gcloud projects get-iam-policy

upvoted 2 times

Question #149

You are running multiple VPC-native Google Kubernetes Engine clusters in the same subnet. The IPs available for the nodes are exhausted, and you want to ensure that the clusters can grow in nodes when needed. What should you do?

- A. Create a new subnet in the same region as the subnet being used.
- B. Add an alias IP range to the subnet used by the GKE clusters.
- C. Create a new VPC, and set up VPC peering with the existing VPC.
- D. Expand the CIDR range of the relevant subnet for the cluster.

Correct Answer: D*Community vote distribution*

D (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (D):

gcloud compute networks subnets expand-ip-range

NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 25 times

 **magistrum** 1 year, 10 months ago

Ok D it is, here's the GKE specific documentation

<https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips>

Every subnet must have a primary IP address range. You can expand the primary IP address range at any time, even when Google Cloud resources use the subnet; however, you cannot shrink or change a subnet's primary IP address scheme after the subnet has been created. The first two and last two IP addresses of a primary IP address range are reserved by Google Cloud.

upvoted 5 times

 **MohammedGhouse** Highly Voted 2 years, 2 months ago

D: is the answer

upvoted 12 times

 **SSPC** 2 years, 2 months ago

I agree with you. https://cloud.google.com/vpc/docs/configure-alias-ip-ranges#gcloud_1

upvoted 2 times

 **AwesomeGCP** Most Recent 3 weeks, 6 days ago

Selected Answer: D

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

Selected Answer: D

D. Expanding CIDR range is enough.

upvoted 1 times

 **sonuricky** 3 months, 1 week ago

C is the right answer

upvoted 1 times

 **ryumada** 3 months ago

Please provide the reason why you choose C as the right answer. ESP_SAP explains clearly about the reason why he choose D as the right answer even he add Google Documentation link too to prove his answer.

upvoted 1 times

 **Bumbah** 4 months ago

Selected Answer: D

Correct answer is D:

<https://cloud.google.com/vpc/docs/create-modify-vpc-networks#expand-subnet>

Just expand your subnet.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is right

upvoted 1 times

✉ **GCP_Student1** 1 year, 8 months ago

This might help

Node limiting ranges

The maximum number of Pods and Services for a given GKE cluster is limited by the size of the cluster's secondary ranges. The maximum number of nodes in the cluster is limited by the size of the cluster's subnet's primary IP address range and the cluster's Pod address range.

The Cloud Console shows error messages like the following to indicate that either the subnet's primary IP address range or the cluster's Pod IP address range (the subnet's secondary IP address range for Pods) has been exhausted:

Instance [node name] creation failed: IP space of [cluster subnet] is exhausted

Note: Secondary subnets are not visible in Cloud Console. If you can't find the [cluster subnet] reported by the above error message it means that the error is caused by IP exhaustion in a secondary subnet. In this case check the secondary ranges of the primary subnet.

https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips#node_limiters

upvoted 6 times

✉ **GCP_Student1** 1 year, 8 months ago

By the way the answer is;

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 3 times

✉ **Ozymandiax** 1 year, 10 months ago

UHmmm, 1 question. The description of the problem says that the ip's are EXHAUSTED. So, no more IP's available in this subnet.

It also states that we're having a multi-VPC environment... as always we should not interpret, just take the questions literally.

If we do not know the actual size of the deployment it can be ANY size, and if IP's are EXHAUSTED, it should BE, BIG as Galactic sized or so....

With all this I wonder if the right answer is not C...

upvoted 3 times

✉ **akshaym87** 9 months, 3 weeks ago

Same doubt!

VPC peering seems correct to me.

upvoted 1 times

✉ **Bhagirathi** 1 year, 11 months ago

D best option to think here.

upvoted 2 times

✉ **hicham** 1 year, 8 months ago

totally agree

upvoted 1 times

✉ **swatititame** 1 year, 11 months ago

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 1 times

Question #150

You have a batch workload that runs every night and uses a large number of virtual machines (VMs). It is fault-tolerant and can tolerate some of the VMs being terminated. The current cost of VMs is too high. What should you do?

- A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.
- B. Run a test using simulated maintenance events. If the test is successful, use N1 Standard VMs when running future jobs.
- C. Run a test using a managed instance group. If the test is successful, use N1 Standard VMs in the managed instance group when running future jobs.
- D. Run a test using N1 standard VMs instead of N2. If the test is successful, use N1 Standard VMs when running future jobs.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/vm-instance-pricing>*Community vote distribution*

A (100%)

  **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Creating and starting a preemptible VM instance

This page explains how to create and use a preemptible virtual machine (VM) instance. A preemptible instance is an instance you can create and run at a much lower price than normal instances. However, Compute Engine might terminate (preempt) these instances if it requires access to those resources for other tasks. Preemptible instances will always terminate after 24 hours. To learn more about preemptible instances, read the preemptible instances documentation.

Preemptible instances are recommended only for fault-tolerant applications that can withstand instance preemptions. Make sure your application can handle preemptions before you decide to create a preemptible instance. To understand the risks and value of preemptible instances, read the preemptible instances documentation.

<https://cloud.google.com/compute/docs/instances/create-start-preemptible-instance>

upvoted 35 times

  **MohammedGhouse** Highly Voted 2 years, 2 months ago

A: is the answer

upvoted 16 times

  **SSPC** 2 years, 2 months ago

"A" is correct

upvoted 3 times

  **juliandm** 2 years, 2 months ago

What about a mixture of preemptible N1 and normal N1 instances? i can't believe just having preemptible is a good practice

upvoted 1 times

  **Ale1973** 2 years, 2 months ago

Good point, in real-world your solution, is the best. For this scenario, the answer is A.

upvoted 8 times

  **Charumathi** Most Recent 1 month ago**Selected Answer: A**

A is correct, preemptible VMs reduce cost, and this is recommended to run batch jobs which run less than 24 hours

upvoted 1 times

  **gcpreviewer** 1 month, 1 week ago**Selected Answer: A**

I Vote A as it is clearly correct. Whenever something runs in under 24 hours and is fault tolerant we should be looking at preemptible VMs to save costs.

upvoted 1 times

  **ccieman2016** 2 months, 1 week ago**Selected Answer: A**

A is correct

upvoted 1 times

  **gcpj** 4 months, 1 week ago

Selected Answer: A

Answer should be A: preemptible VM instances. Because the workload is fault-tolerant and can tolerate some of the VMs being terminated.
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

A is right ..

upvoted 1 times

 **NBR1** 5 months ago

Selected Answer: A

I believe it is A

upvoted 1 times

 **Majkl93** 8 months, 3 weeks ago

Selected Answer: A

as per the comments

upvoted 1 times

 **maggieli** 1 year ago

I vote A. Preempt VM can costdown more.

upvoted 2 times

 **vamgcp** 1 year, 1 month ago

A is correct because preemptible VMs can provide up to 80% discount over normal VMs if the workloads are fault-tolerant

upvoted 2 times

 **AD_0525** 1 year, 4 months ago

Keyword- Fault tolerant, so answer should be pre emptible VMs, option A

upvoted 2 times

 **GCP_Student1** 1 year, 8 months ago

A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs

upvoted 3 times

 **swatititame** 1 year, 11 months ago

• A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.

upvoted 2 times

Question #151

You are working with a user to set up an application in a new VPC behind a firewall. The user is concerned about data egress. You want to configure the fewest open egress ports. What should you do?

- A. Set up a low-priority (65534) rule that blocks all egress and a high-priority rule (1000) that allows only the appropriate ports.
- B. Set up a high-priority (1000) rule that pairs both ingress and egress ports.
- C. Set up a high-priority (1000) rule that blocks all egress and a low-priority (65534) rule that allows only the appropriate ports.
- D. Set up a high-priority (1000) rule to allow the appropriate ports.

Correct Answer: C*Community vote distribution*

A (75%)

D (25%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (A):

Implied rules

Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:

Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see Internet access requirements.

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.

https://cloud.google.com/vpc/docs/firewalls#default_firewall_rules

upvoted 32 times

 **patashish** 4 months, 1 week ago

The correct answer is C

upvoted 1 times

 **ryumada** 3 months ago

You should visit the documentation link he attached. He's copy those statements from the Google Docs.

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Listen that guy because he is right

upvoted 1 times

 **MohammedGhouse** Highly Voted 2 years, 2 months ago

A: is the answer

upvoted 11 times

 **ryumada** Most Recent 3 months ago

Selected Answer: A

Read ESP_SAP comment for the explanation. He explains it clearly.

upvoted 1 times

 **bobthebuilder55110** 3 months, 1 week ago

Selected Answer: A

Answer is (A) :

First I was going with C but then I read the question again, let's try to understand both options here, the goal is to deny egress and only allow some ports for some functions to perform. If we go with C, lower the number higher the priority (1000) so the rule with this priority 1000 will overwrite (65534), so If we allow only appropriate ports it will be overwritten with the high-priority (1000) rule and all the egress traffic will be blocked.

Remember the goal here is to block egress but not all of it since we still want to configure the fewest open ports and this is statefull meaning for open ports traffic will be both ways.

A fits this condition where it is saying we block all traffic but the required ports are kept open with higher priority which will only allow the required traffic to leave the network.

upvoted 3 times

 **sonuricky** 3 months, 1 week ago

C is the correct answer

upvoted 1 times

 **gscharly** 3 months, 2 weeks ago

Selected Answer: A

A: is the answer

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: A

Correct answer is A

upvoted 1 times

 **patashish** 4 months, 1 week ago

Correct Answer is C

upvoted 1 times

 **patashish** 4 months, 1 week ago

Hint : All rules are stateful.

VPC firewall rules are stateful. When a connection is allowed through the firewall in either direction, return traffic matching this connection is also allowed. You cannot configure a firewall rule to deny associated response traffic.

As per question , we want to restrict egress traffic.

So focus to restrict egress traffic based on priority of rules.

Allow incoming traffic for appropriate traffic and block all traffic and allow only which are required.

Hence , as per my view C should be correct answer

upvoted 3 times

 **mani098** 4 months, 2 weeks ago

Selected Answer: D

A incorrect 65534 that blocks all ingress, not egress (except few default ports)

D is correct.

upvoted 3 times

 **patashish** 4 months, 1 week ago

But why D is correct ? Why not C ?

D is more generic , As per question , need to focus on egress traffic

upvoted 1 times

 **pnVino27** 11 months ago

Selected Answer: A

Correct Answer is A

upvoted 3 times

 **maggieli** 1 year ago

I vote A is correct. Block all port in gress and set low-priority.

upvoted 2 times

 **aamirahal** 1 year ago

A is correct

upvoted 2 times

 **vvkds** 1 year, 2 months ago

Correct answer is A. Firewall rules are executed based on the priority.

upvoted 2 times

 **tanito83** 1 year, 5 months ago

The correct answer is A. Please, modify it.

upvoted 3 times

 **kopper2019** 1 year, 6 months ago

A, lower is high priority

Implied rules

Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:

Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see Internet access requirements.

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.

upvoted 2 times

 **Devgela** 1 year, 7 months ago

A for me

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

A is correct.

upvoted 1 times

Question #152

Your company runs its Linux workloads on Compute Engine instances. Your company will be working with a new operations partner that does not use Google Accounts. You need to grant access to the instances to your operations partner so they can maintain the installed tooling. What should you do?

- A. Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User.
- B. Tag all the instances with the same network tag. Create a firewall rule in the VPC to grant TCP access on port 22 for traffic from the operations partner to instances with the network tag.
- C. Set up Cloud VPN between your Google Cloud VPC and the internal network of the operations partner.
- D. Ask the operations partner to generate SSH key pairs, and add the public keys to the VM instances.

Correct Answer: B

Reference:

<https://cloud.google.com/vpc/docs/firewalls>*Community vote distribution*

A (100%)

✉  **kulikBro** Highly Voted 1 year, 7 months ago

A - <https://cloud.google.com/iap/docs/external-identities>

upvoted 24 times

✉  **Bhagirathi** Highly Voted 1 year, 11 months ago

full of confusions for any reader....

You guys all say A, B, C & D but which one is correct ?

upvoted 21 times

✉  **yc25744** 1 year, 3 months ago

nothing

upvoted 4 times

✉  **Charumathi** Most Recent 1 month ago

Selected Answer: A

A is the correct answer,

IAP controls access to your App Engine apps and Compute Engine VMs running on Google Cloud. It leverages user identity and the context of a request to determine if a user should be allowed access. IAP is a building block toward BeyondCorp, an enterprise security model that enables employees to work from untrusted networks without using a VPN.

By default, IAP uses Google identities and IAM. By leveraging Identity Platform instead, you can authenticate users with a wide range of external identity providers, such as:

Email/password

OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)

SAML

OIDC

Phone number

Custom

Anonymous

This is useful if your application is already using an external authentication system, and migrating your users to Google accounts is impractical.

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A is right. IAP will allow you to connect compute engine without GCP account .

upvoted 2 times

✉  **pfabio** 5 months, 1 week ago

Selected Answer: A

A - To control which users and groups are allowed to use IAP TCP forwarding and which VM instances they're allowed to connect to, configure Identity and Access Management (IAM) permissions.

How to:

Open the IAP admin page and select the SSH and TCP Resources tab.

Open the IAP admin page

Select the VM instances that you want to configure.

Click Show info panel if the info panel is not visible.

Click Add member and configure the following:

New members: Specify the user or group you want to grant access.
Select a role Select Cloud IAP > IAP-Secured Tunnel User.

<https://cloud.google.com/iap/docs/using-tcp-forwarding#grant-permission>

upvoted 3 times

✉ **SleepyHitman** 9 months, 2 weeks ago

Selected Answer: A

The answer is A:
Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User as per [1] and [2].

[1]:<https://cloud.google.com/iap/docs/tcp-forwarding-overview#:~:text=To%20learn%20how%20to%20grant%20principals%20access%20to%20tunneled%20resources%20and%20how%20to%20create%20tunnels%20that%20route%20TCP%20traffic%2C%20see%20Using%20IAP%20for%20TCP%20forwarding>.

[2]: <https://cloud.google.com/iap/docs/tcp-forwarding-overview#:~:text=IAP%27s%20TCP%20forwarding%20feature%20lets%20you%20control%20who%20can%20access%20administrative%20services%20like%20SSH%20and%20RDP%20on%20your%20backends%20from%20the%20public%20internet>.

upvoted 4 times

✉ **emv** 9 months, 3 weeks ago

IAP controls access to your App Engine apps and Compute Engine VMs running on Google Cloud. It leverages user identity and the context of a request to determine if a user should be allowed access. IAP is a building block toward BeyondCorp, an enterprise security model that enables employees to work from untrusted networks without using a VPN.

By default, IAP uses Google identities and IAM. By leveraging Identity Platform instead, you can authenticate users with a wide range of external identity providers, such as:

Email/password

OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)

SAML

OIDC

Phone number

Custom

Anonymous

This is useful if your application is already using an external authentication system, and migrating your users to Google accounts is impractical.

upvoted 4 times

✉ **Surat** 10 months ago

Very confusing question, what's final answer?

upvoted 1 times

✉ **kimharsh** 10 months, 3 weeks ago

People THINK SIMPLE

the exam want to test you if you know about the IAP or not, the IAP allow you to access compute engine from the internet without having to have a GCP account,

the answer is A

upvoted 9 times

✉ **MariGK** 11 months ago

both (CD) combined seems to be the right answer

upvoted 1 times

✉ **wh1t4k3r** 11 months ago

B and C would make the VMs reachable, but it does not mention user access/auth to it, so I would rule them out.

D is viable, but as per google documentation, there is a risk on managing SSH keys manually:

<https://cloud.google.com/compute/docs/instances/access-overview>

"Risks of manual key management

If you create and manage public SSH keys yourself through the Cloud Console, the gcloud command-line tool, or the API, you must keep track of the used keys and delete the public SSH keys for users who no longer have access. For example, if a team member leaves your project, remove their public SSH keys from metadata, so they can't continue to access your instances.

Additionally, specifying your gcloud tool or API calls incorrectly can potentially wipe out all of the public SSH keys in your project or on your instances, which disrupts connections for your project members.

If you aren't sure that you want to manage your own keys, use Compute Engine tools to connect to your instances instead."

A seems to be the best choice in terms of security and administration.

upvoted 1 times

✉ **Ridhanya** 11 months ago

D seems right

upvoted 2 times



👤 **ME_MYSELF** 10 months, 3 weeks ago

Without goggle account...is that possible
upvoted 1 times

✉️ 🚩 **dark_3k03r** 5 months, 2 weeks ago

Yes, you can use the following:
Email/password
OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)
SAML
OIDC
Phone number
Custom
Anonymous
Per Google Documentation.
<https://cloud.google.com/iap/docs/external-identities>
upvoted 1 times

✉️ 🚩 **maggieli** 1 year ago

I vote B. Open firewall 22 port for maintain using.
upvoted 5 times

✉️ 🚩 **jabrrJ68w02ond1** 11 months, 2 weeks ago

And restrict IP range so only the operations team can login
upvoted 1 times

✉️ 🚩 **maggieli** 1 year ago

I vote B.
upvoted 1 times

✉️ 🚩 **aamirahal** 1 year ago

B is correct and i practically implemented this.
upvoted 2 times

✉️ 🚩 **bubblegumbeach** 1 year ago

A is incorrect because the operations partner does not have a google account.
B is incorrect because creating a firewall rule does not grant access to the VMs, as authentication is also required.
C is incorrect because setting up VPN does not grant access to VMs.
D is correct because the operations partner can use SSH keys to SSH into the VMs.
upvoted 9 times

✉️ 🚩 **vamgcp** 1 year, 1 month ago

correct answer is to Ask the operations partner to generate SSH key pairs and add the public keys to the VM instances.(option D)
upvoted 6 times

Question #153

You have created a code snippet that should be triggered whenever a new file is uploaded to a Cloud Storage bucket. You want to deploy this code snippet. What should you do?

- A. Use App Engine and configure Cloud Scheduler to trigger the application using Pub/Sub.
- B. Use Cloud Functions and configure the bucket as a trigger resource.
- C. Use Google Kubernetes Engine and configure a CronJob to trigger the application using Pub/Sub.
- D. Use Dataflow as a batch job, and configure the bucket as a data source.

Correct Answer: B*Community vote distribution*

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Google Cloud Storage Triggers

Cloud Functions can respond to change notifications emerging from Google Cloud Storage. These notifications can be configured to trigger in response to various events inside a bucket—object creation, deletion, archiving and metadata updates.

Note: Cloud Functions can only be triggered by Cloud Storage buckets in the same Google Cloud Platform project.

Event types

Cloud Storage events used by Cloud Functions are based on Cloud Pub/Sub Notifications for Google Cloud Storage and can be configured in a similar way.

Supported trigger type values are:

google.storage.object.finalize

google.storage.object.delete

google.storage.object.archive

google.storage.object.metadataUpdate

Object Finalize

Trigger type value: google.storage.object.finalize

This event is sent when a new object is created (or an existing object is overwritten, and a new generation of that object is created) in the bucket.

https://cloud.google.com/functions/docs/calling/storage#event_types

upvoted 35 times

 **francisco_guerra** Highly Voted 2 years, 2 months ago

The answer is B

upvoted 19 times

 **SSPC** 2 years, 2 months ago

Sure B? Please you could share the link with the Google documentation

upvoted 1 times

 **Ale1973** 2 years, 2 months ago

<https://cloud.google.com/functions/docs/calling/storage>

upvoted 3 times

 **Angel_99** Most Recent 2 months, 3 weeks ago

Selected Answer: B

The answer is B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct, it is required on demand when upload happens

upvoted 1 times

 **arvsrv** 9 months, 2 weeks ago

Selected Answer: B

The answer is B

upvoted 2 times

 **Surat** 10 months ago

I vote for B

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: B

Vote For B

upvoted 4 times

 **vamgcp** 1 year, 1 month ago

Correct Answer is B - Use Cloud Functions and configure the bucket as a trigger resource.

upvoted 2 times

 **Zimcruza** 1 year, 4 months ago

Question asks:

You want to deploy this code snippet. What should you do?

To me, none of the answers is relevant to DEPLOYMENT - they all are about how you get the trigger to run ...

upvoted 1 times

 **GoCloud** 1 year, 7 months ago

vote for B

upvoted 2 times

 **marialix87** 1 year, 7 months ago

I think is B

upvoted 3 times

 **GCP_Student1** 1 year, 7 months ago

B. Use Cloud Functions and configure the bucket as a trigger resource.

upvoted 3 times

 **lutoa** 1 year, 8 months ago

Answer is B, app engine is for applications, in this case it's just a code snippet which cloud functions is more suitable

upvoted 2 times

 **victory108** 1 year, 9 months ago

B - Use Cloud Functions and configure the bucket as a trigger resource.

upvoted 3 times

 **Bhagirathi** 1 year, 11 months ago

Most of you suggest B.

someone chose A - I will choose B

upvoted 1 times

 **swatititame** 1 year, 11 months ago

• B. Use Cloud Functions and configure the bucket as a trigger resource.

upvoted 2 times

 **ritzheg** 2 years ago

B.

<https://cloud.google.com/functions/docs/tutorials/storage#objectives>

upvoted 2 times

Question #154

You have been asked to set up Object Lifecycle Management for objects stored in storage buckets. The objects are written once and accessed frequently for 30 days. After 30 days, the objects are not read again unless there is a special need. The objects should be kept for three years, and you need to minimize cost.

What should you do?

- A. Set up a policy that uses Nearline storage for 30 days and then moves to Archive storage for three years.
- B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.
- C. Set up a policy that uses Nearline storage for 30 days, then moves the Coldline for one year, and then moves to Archive storage for two years.
- D. Set up a policy that uses Standard storage for 30 days, then moves to Coldline for one year, and then moves to Archive storage for two years.

Correct Answer: A

Reference:

https://books.google.com.pk/books?id=q0nhDwAAQBAJ&pg=PA52&lpg=PA52&dq=Set+up+a+policy+that+uses+Nearline+storage+for+30+days+and+then+moves+to+Archive+storage+for+three+years.&source=bl&ots=kYLZN1ymA8&sig=ACfU3U2XLmzQ39cmPDwjfWxRbNtDNLc_6g&hl=en&sa=X&ved=2ahUKEwjZmef0pr7qAhVzQkEAHTUGASYQ6AEwAHoECAoQAQ#v=onepage&q=Set%20up%20a%20policy%20that%20uses%20Nearline%20storage%20for%2030%20days%20and%20then%20moves%20to%20Archive%20storage%20for%20three%20years.&f=false

Community vote distribution

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

The key to understand the requirement is : "The objects are written once and accessed frequently for 30 days"

Standard Storage

Standard Storage is best for data that is frequently accessed ("hot" data) and/or stored for only brief periods of time.

Archive Storage

Archive Storage is the lowest-cost, highly durable storage service for data archiving, online backup, and disaster recovery. Unlike the "coldest" storage services offered by other Cloud providers, your data is available within milliseconds, not hours or days. Archive Storage is the best choice for data that you plan to access less than once a year.

<https://cloud.google.com/storage/docs/storage-classes#standard>

upvoted 47 times

 **naveedpk00** 2 years, 1 month ago

What if we chose option D to minimize the cost as asked in the question? What do you think?

upvoted 1 times

 **gcper** 2 years ago

It doesn't minimize the costs. Check the costs of coldline vs archival

upvoted 3 times

 **SSPC** Highly Voted 2 years, 2 months ago

I think the correct one is B. Because Nearline has a 30-day minimum storage duration.

<https://cloud.google.com/storage/docs/storage-classes>

upvoted 14 times

 **pepepy** 2 years, 2 months ago

The object should be kept for three years, and you need to minimize cost, after 30 days it will be moved to archive, ans A

upvoted 1 times

 **pepepy** 2 years, 2 months ago

Sorry you are right accessed frequently for 30 days, its B

upvoted 6 times

 **Charumathi** Most Recent 1 month ago

Selected Answer: B

B is the correct Answer,
 Frequently accessed data 'Hot Data' should be stored in Standard Storage for 30 days,
 Then this can be moved to Archive after 30 days for period of three years which is accessed only when a special need arises, to reduce cost.
 upvoted 1 times

 **taiyi078** 4 months, 2 weeks ago

<https://cloud.google.com/storage/docs/storage-classes#nearline>

Nearline storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline storage is a great choice.

Nearline storage is also appropriate for data backup, long-tail multimedia content, and data archiving. Note, however, that for data accessed less frequently than once a quarter, Coldline storage or Archive storage are more cost-effective, as they offer lower storage costs.

upvoted 1 times

 **taiyi078** 4 months, 2 weeks ago

Nearline storage is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline storage is a better choice than Standard storage in scenarios where slightly lower availability, a 30-day minimum storage duration, and costs for data access are acceptable trade-offs for lowered at-rest storage costs.

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

I am changing it to D. Set up a policy that uses Standard storage for 30 days, then moves to Coldline for one year, and then moves to Archive storage for two years.

upvoted 2 times

 **pluidust** 8 months ago

Selected Answer: B

B for sure

upvoted 1 times

 **SleepyHitman** 9 months, 2 weeks ago

Selected Answer: B

The answer is: B

Standard storage description:

[https://cloud.google.com/storage/docs/storage-classes#:~:text=Standard%20Storage%20is%20best%20for%20data%20that%20is%20frequently%20accessed%20\(%22hot%22%20data\)%20and/or%20stored%20for%20only%20brief%20periods%20of%20time](https://cloud.google.com/storage/docs/storage-classes#:~:text=Standard%20Storage%20is%20best%20for%20data%20that%20is%20frequently%20accessed%20(%22hot%22%20data)%20and/or%20stored%20for%20only%20brief%20periods%20of%20time).

Nearline storage imp description:

<https://cloud.google.com/storage/docs/storage-classes#:~:text=storage%20service%20for-,storing%20infrequently%20accessed%20data,-Nearline%20Storage%20is>

Pricing for frequent access:

<https://cloud.google.com/storage/pricing#:~:text=Free%20operations-,Standard%20Storage,Free,-Coldline%20Storage>

upvoted 4 times

 **Surat** 10 months ago

Standard – Frequently access and short period

Nearline - Low cost, highly durable for infrequent data access, lower availability

Coldline –Very low cost, highly durable for infrequently accessed, 90 days minimum storage

Archive - Lowest cost, highly durable for archiving, backup and DR, lower availability

I will also go for B

upvoted 3 times

 **fazalmf** 10 months, 3 weeks ago

B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.

_ clearly mentioned data access frequently for 30days and then object not read again until special need (archive suitable).

upvoted 1 times

 **Naren080914** 10 months, 3 weeks ago

Selected Answer: B

If you access frequently in Nearline, it will cost you more. There is no retrieval cost for Standard. So for frequently accessed object, go with standard. Also in Standard, there is no minimum period to store the object, so don't get confused by 30 days. Ans is B.

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Vote For B

upvoted 1 times

 **AD_0525** 1 year, 4 months ago

Answer should be D. object will be frequently access for first 30 days. cost of data access from standard storage is less than Nearline. Then after 30 days the objects can be moved to archive.

upvoted 2 times

 **AD_0525** 1 year, 4 months ago

Typo, answer should be B. Description remains the same.

upvoted 2 times

 **jahnu** 1 year, 6 months ago

My Answer is B. why because Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. so we use standard storage for frequent access.

upvoted 1 times

 **kopper2019** 1 year, 6 months ago

B, Standard since they frequently access for 30 days and after that rarely accessed and you need to reduce cost so Archive

upvoted 2 times

 **pondai** 1 year, 7 months ago

B is my answer after 30 day no read again ,so put in archive is the way to min cost

upvoted 1 times

 **Devgela** 1 year, 7 months ago

B for sure

upvoted 1 times

 **neerajgoyal** 1 year, 8 months ago

I will go with B

upvoted 1 times

Question #155

You are storing sensitive information in a Cloud Storage bucket. For legal reasons, you need to be able to record all requests that read any of the stored data. You want to make sure you comply with these requirements. What should you do?

- A. Enable the Identity Aware Proxy API on the project.
- B. Scan the bucket using the Data Loss Prevention API.
- C. Allow only a single Service Account access to read the data.
- D. Enable Data Access audit logs for the Cloud Storage API.

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/audit-logs>*Community vote distribution*

D (100%)

ESP_SAP Highly Voted 2 years, 2 months ago

Correct Answer is (D):

Logged information

Within Cloud Audit Logs, there are two types of logs:

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

<https://cloud.google.com/storage/docs/audit-logs#types>

upvoted 22 times

francisco_guerra Highly Voted 2 years, 2 months ago

D is the correct one

upvoted 19 times

SSPC 2 years, 2 months ago

Yes D is the correct

upvoted 6 times

AzureDP900 Most Recent 4 months, 2 weeks ago

D is right for this use case

upvoted 1 times

Akash7 5 months, 3 weeks ago

D is correct as Data Access logs pertaining to Cloud Storage operations are not recorded by default. You have to enable them ...

<https://cloud.google.com/storage/docs/audit-logging>

upvoted 2 times

wael_tn 6 months, 2 weeks ago**Selected Answer: D**

I think it's D

upvoted 1 times

Surat 10 months ago

I also vote for D

upvoted 2 times

Vinodh9289 1 year, 2 months ago

D is the correct Answer

upvoted 2 times

WakandaF 1 year, 5 months ago

seems that B is the right!

Cloud Data Loss Prevention (DLP) helps you to understand and manage such sensitive data. It provides fast, scalable classification and redaction for sensitive data elements. Using the Data Loss Prevention API and Cloud Functions, you can automatically scan this data before it is uploaded to the shared storage bucket.

upvoted 1 times

 **YAS007** 1 year, 3 months ago

the question doesn't ask you to manage or understand sensitive data :
" you need to be able to record all requests that read any of the stored data"

upvoted 3 times

 **victory108** 1 year, 8 months ago

D - Enable Data Access audit logs for the Cloud Storage API.

upvoted 1 times

 **EABDAJA** 1 year, 8 months ago

D is correct

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

D. Enable Data Access audit logs for the Cloud Storage API.

upvoted 2 times

 **swatititame** 1 year, 11 months ago

• D. Enable Data Access audit logs for the Cloud Storage API.

upvoted 1 times

 **RockAJ** 2 years, 1 month ago

Ans is D

upvoted 2 times

 **pepepy** 2 years, 2 months ago

agree with D

upvoted 2 times

Question #156

You are the team lead of a group of 10 developers. You provided each developer with an individual Google Cloud Project that they can use as their personal sandbox to experiment with different Google Cloud solutions. You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment. What should you do?

- A. Create a single budget for all projects and configure budget alerts on this budget.
- B. Create a separate billing account per sandbox project and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per billing account.
- C. Create a budget per project and configure budget alerts on all of these budgets.
- D. Create a single billing account for all sandbox projects and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per project.

Correct Answer: C

Reference:

<https://cloud.google.com/billing/docs/how-to/budgets>*Community vote distribution*

C (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (C):

Set budgets and budget alerts

Overview

Avoid surprises on your bill by creating Cloud Billing budgets to monitor all of your Google Cloud charges in one place. A budget enables you to track your actual Google Cloud spend against your planned spend. After you've set a budget amount, you set budget alert threshold rules that are used to trigger email notifications. Budget alert emails help you stay informed about how your spend is tracking against your budget.

2. Set budget scope

Set the budget Scope and then click Next.

In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all the projects in the Cloud Billing account, choose Select all.

<https://cloud.google.com/billing/docs/how-to/budgets#budget-scop>

upvoted 40 times

 **bobthebuilder55110** 3 months ago

wait a minute, why not A ?

As you said that

" In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all the projects in the Cloud Billing account, choose Select all. "

As per this I should be able to create single budget for all the projects and should be able to set alert on that, why create separate budget for all 10 projects ?

upvoted 1 times

 **Priyanka109** 1 month, 1 week ago

It will be a combined budget that's why it's C

upvoted 1 times

 **dang1986** 9 months ago

You're the only answer I take seriously "Thumbs up"

upvoted 3 times

 **Hjameel** Highly Voted 2 years, 2 months ago

I think C is the best answer.

upvoted 10 times

 **Angel_99** Most Recent 2 months, 3 weeks ago

Selected Answer: C

Key is anyone goes above \$500 means it requires project level

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

Key is anyone goes above \$500 means it requires project level so C is right

upvoted 1 times

✉  **wael_tn** 6 months, 2 weeks ago

Selected Answer: C

Clearly C is the answer

upvoted 1 times

✉  **pondai** 1 year, 7 months ago

Does anyone know Data Studio can be alert to email? If it can't I'll pick C

upvoted 5 times

✉  **GCP_Student1** 1 year, 8 months ago

C. Create a budget per project and configure budget alerts on all of these budgets.

upvoted 4 times

✉  **andregjp** 1 year, 9 months ago

I believe is A because of this...

Projects: In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all of the projects in the Cloud Billing account, choose Select all.

Some costs are not related to a project, such as the costs of subscriptions or Support costs.

In the budget's project scope, in the list of projects you can filter on, [Charges not specific to a project] is not an option you can select.

If you choose Select all, then the costs in all projects, including Charges not specific to a project, are included in the budget and cost trend chart cost calculations.

If you select one or more projects - but not all projects - then the Charges not specific to a project are not included in the budget and cost trend chart cost calculations.

You can view your costs that are not related to a project in the billing reports. Using the projects filter in the reports page, you can select and view [Charges not specific to a project].

URL: <https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 3 times

✉  **Ridhanya** 11 months ago

but how will you know who crossed the limit. what if the summation of their usage exceeds 500? the corresponding alert would be a false alarm

upvoted 5 times

✉  **Bhagirathi** 1 year, 11 months ago

which one correct?

upvoted 1 times

✉  **swatititame** 1 year, 11 months ago

• C. Create a budget per project and configure budget alerts on all of these budgets.

upvoted 1 times

✉  **xtian2900** 2 years, 1 month ago

is C, not A

with A, i guess if you create a single budget for all projects, together they can easily beat the \$500 mark and you need to know if "one developer" did it.

so one budget per project is the solution.

upvoted 4 times

✉  **[Removed]** 2 years, 1 month ago

Yes C is correct, I would have gone with B and D as billing export is the crucial element for billing, but both the options don't talk about notifying about spending.

Also, you don't want to combine the billing for all as each one can spend up to 500 so it will be better if they are individual so A is out.

upvoted 2 times

✉  **pepey** 2 years, 2 months ago

You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment, so Answer is C

upvoted 2 times

✉  **tothecloud** 2 years, 2 months ago

C is correct

upvoted 4 times

✉  **SSPC** 2 years, 2 months ago

Do you think "C" is correct?

upvoted 2 times

✉  **SSPC** 2 years, 2 months ago

A budget per project?

upvoted 1 times

✉  **Ale1973** 2 years, 2 months ago

Yes, when you are creating a budget, the budget can be scoped to focus on a specific set of resources (Projects, Products or Labels)

upvoted 3 times

Question #157

You are deploying a production application on Compute Engine. You want to prevent anyone from accidentally destroying the instance by clicking the wrong button. What should you do?

- A. Disable the flag `Delete boot disk when instance is deleted`.
- B. Enable delete protection on the instance.
- C. Disable Automatic restart on the instance.
- D. Enable Preemptibility on the instance.

Correct Answer: A

Reference:

<https://googlecloudplatform.uservoice.com/forums/302595-compute-engine/suggestions/14227521-set-delete-boot-disk-when-instance-is-deleted-to>

Community vote distribution

B (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answer is (B):

Preventing Accidental VM Deletion

This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an Instance resource. To learn more about VM instances, read the Instances documentation.

As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.

By setting the `deletionProtection` flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the `deletionProtection` flag, the request fails. Only a user that has been granted a role with `compute.instances.create` permission can reset the flag to allow the resource to be deleted.

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 32 times

 **MohammedGhouse** Highly Voted 2 years, 2 months ago

"B" is the answer

upvoted 11 times

 **AzureDP900** Most Recent 4 months, 2 weeks ago

This is straight forward question, enable delete protection. B is right

upvoted 1 times

 **Himadhar1997** 5 months, 2 weeks ago

Selected Answer: B

Preventing Accidental VM Deletion

This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an Instance resource. To learn more about VM instances, read the Instances documentation.

upvoted 2 times

 **Surat** 10 months ago

B seems right option

upvoted 3 times

 **kped21** 10 months, 1 week ago

B - on VM Enable delete protection

upvoted 2 times

 **jaffarali** 11 months ago

Selected Answer: B

Answer is B. there is an Option in VM instance while creating

upvoted 3 times

 **Sreedharveluru** 1 year, 3 months ago

Option A would not prevent , It can be used only after the damage is done. Hence B

upvoted 1 times

 **NARWAL** 1 year, 7 months ago

B should be the answer.

upvoted 1 times

 **GCP_Student1** 1 year, 8 months ago

B. Enable delete protection on the instance.

upvoted 2 times

 **Lomy** 1 year, 8 months ago

B

The ans is B

upvoted 1 times

 **RockAJ** 2 years, 1 month ago

B for me

upvoted 2 times

 **muk5658** 2 years, 1 month ago

Correct Answer is 'B'

upvoted 3 times

 **SSPC** 2 years, 2 months ago

B is correct answer. <https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>

upvoted 5 times

 **francisco_guerra** 2 years, 2 months ago

I think is B

upvoted 3 times

Question #158

Your company uses a large number of Google Cloud services centralized in a single project. All teams have specific projects for testing and development. The

DevOps team needs access to all of the production services in order to perform their job. You want to prevent Google Cloud product changes from broadening their permissions in the future. You want to follow Google-recommended practices. What should you do?

- A. Grant all members of the DevOps team the role of Project Editor on the organization level.
- B. Grant all members of the DevOps team the role of Project Editor on the production project.
- C. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the production project.
- D. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the organization level.

Correct Answer: B

Community vote distribution

C (100%)

ESP_SAP [Highly Voted] 2 years, 2 months ago

Correct Answer is (C):

Understanding IAM custom roles

Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.

Basic concepts

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

https://cloud.google.com/iam/docs/understanding-custom-roles#basic_concepts

upvoted 44 times

SSPC [Highly Voted] 2 years, 2 months ago

"You want to prevent Google Cloud product changes from broadening their permissions in the future." then CUSTOM ROLE

upvoted 18 times

Rothmansua 1 year ago

Great hint, thanks!

upvoted 2 times

Charumathi [Most Recent] 1 month ago

Selected Answer: C

C is the correct answer, give the devops team the least privileged role, only the required permissions to access the production services, as the question states 'to prevent product changes' for which editor role is not recommended either at Project or organizational level, organizational level access gives broad scope to all the projects in the organization, this role cannot be given to the devops team.

- A. Editor has privilege to change the products, and the scope is broad
- B. Editor has privilege to change the products
- C. Recommended, as this will give only required permission at project level to devops team.
- D. They require only project level access. This gives access to all project in the organization.

upvoted 1 times

ale_brd_ 1 month, 1 week ago

Selected Answer: C

Correct Answer is (C):

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

upvoted 1 times

sonuricky 3 months, 1 week ago

correct answer is A

upvoted 2 times

Roro_Brother 4 months, 1 week ago

Selected Answer: C

There is no doubt, the correct answer is C
upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is right
upvoted 1 times

 **DaveNZ** 8 months, 4 weeks ago

Selected Answer: C

C seems to be the popular answer, and it makes sense because the generic roles are not sufficient for these specific requirements. I added this voting comment because the community answers are not currently visible.
upvoted 2 times

 **Surat** 10 months ago

I vote for C
upvoted 2 times

 **rachee** 11 months, 1 week ago

I initially thought C. But I think this may be a trick question. "The DevOps team needs access to ALL of the PRODUCTION services..." which are in a "single" project. If "Project Editor" is assigned at on the "production" project it gives them access to "ALL" production services including product changes in the "production" project. A custom role would have to be modified to get access to product changes in the production project that required additional permissions, so the DevOps team would not have access to "ALL" services until the custom role is modified. I am changing my choice to B.

upvoted 4 times

 **theBestStudent** 2 months ago

Your choice doesn't follow the rule of least privilege. So correct answer is C.

upvoted 1 times

 **jahnu** 1 year, 6 months ago

Answer C: custom role enable privileges.
upvoted 2 times

 **GCP_Student1** 1 year, 7 months ago

C. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the production project.
upvoted 2 times

 **swatititame** 1 year, 12 months ago

A single folder contains multiple projects. In the below line, it should be a single organization or folder. If we open the GCP console we can see all the services without selecting the project.
Your company uses a large number of Google Cloud services centralized in a single project.

upvoted 2 times

 **anony1234567** 2 years, 2 months ago

Question looks incorrect. In 1st line they are saying one project and in the second line they are saying multiple projects. I don't know if I understood that correctly..

upvoted 8 times

 **rjperdigao** 1 year, 8 months ago

please look at the following link:
https://console.cloud.google.com/getting-started/enterprise;step=RESOURCE_MANAGEMENT
Imagine you have only a shared project for development, testing and production

upvoted 1 times

 **omunoz** 1 year, 7 months ago

I believe in the second line they are referring to company projects and no GCP projects... but yes, is confuse..
upvoted 1 times

 **SSPC** 2 years, 2 months ago

C better yes
upvoted 7 times

 **SSPC** 2 years, 2 months ago

What do you think C or D?
upvoted 2 times

 **[Removed]** 2 years, 2 months ago

I think it's C because the question says "The DevOps team needs access to all of the production services in order to perform their job" and so to give them the custom permissions across the whole organization would be too permissive for what they need.
upvoted 3 times

 **francisco_guerra** 2 years, 2 months ago

Maybe is C because is only one project

upvoted 4 times

 **SSPC** 2 years, 2 months ago

I think that the answer is D, because It says "DevOps team"

upvoted 2 times

Question #159

You are building an application that processes data files uploaded from thousands of suppliers. Your primary goals for the application are data security and the expiration of aged data. You need to design the application to:

- * Restrict access so that suppliers can access only their own data.
- * Give suppliers write access to data only for 30 minutes.
- * Delete data that is over 45 days old.

You have a very short development cycle, and you need to make sure that the application requires minimal maintenance. Which two strategies should you use?

(Choose two.)

- A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.
- B. Use signed URLs to allow suppliers limited time access to store their objects.
- C. Set up an SFTP server for your application, and create a separate user for each supplier.
- D. Build a Cloud function that triggers a timer of 45 days to delete objects that have expired.
- E. Develop a script that loops through all Cloud Storage buckets and deletes any buckets that are older than 45 days.

Correct Answer: AE

Community vote distribution

AB (100%)

 **ESP_SAP** Highly Voted 2 years, 2 months ago

Correct Answers are: (AB):

(A) Object Lifecycle Management

Delete

The Delete action deletes an object when the object meets all conditions specified in the lifecycle rule.

Exception: In buckets with Object Versioning enabled, deleting the live version of an object causes it to become a noncurrent version, while deleting a noncurrent version deletes that version permanently.

<https://cloud.google.com/storage/docs/lifecycle#delete>

(B) Signed URLs

This page provides an overview of signed URLs, which you use to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account

<https://cloud.google.com/storage/docs/access-control/signed-urls>

upvoted 38 times

 **francisco_guerra** Highly Voted 2 years, 2 months ago

AB is the answer

upvoted 14 times

 **diasporabro** Most Recent 2 weeks, 4 days ago

Selected Answer: AB

AB achieves this objective

upvoted 1 times

 **olme59** 2 months, 2 weeks ago

Selected Answer: AB

its clearly AB, life cycle and provider private url

upvoted 1 times

 **Angel_99** 2 months, 3 weeks ago

Selected Answer: AB

Correct Answer Combo: (AB)

upvoted 1 times

 **abirroy** 3 months, 1 week ago

Selected Answer: AB

Correct Answers are: (AB)

upvoted 1 times

✉  **patashish** 4 months, 1 week ago

Correct Answers are: A and B

upvoted 1 times

✉  **Roro_Brother** 4 months, 1 week ago

Selected Answer: AB

A and B is the right answer

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A and B right answer

upvoted 1 times

✉  **DaveNZ** 8 months, 4 weeks ago

Selected Answer: AB

The website answer says A and E, but these do the same thing. The question also asks about setting up security, so the other part of the answer must be B

upvoted 1 times

✉  **jahnu** 1 year, 6 months ago

My Answer: A B

upvoted 2 times

✉  **GCP_Student1** 1 year, 7 months ago

- A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.
- B. Use signed URLs to allow suppliers limited time access to store their objects.

upvoted 3 times

✉  **Bhagirathi** 1 year, 10 months ago

A & B helps.

upvoted 1 times

✉  **tuta** 1 year, 10 months ago

did you pass exam ?

upvoted 4 times

✉  **RockAJ** 2 years, 1 month ago

for me A and B

upvoted 2 times

✉  **[Removed]** 2 years, 2 months ago

100% A and B.

- A. lifecycle policies are designed for situations just like these. It can cleanly schedule deletion of objects that are past 45 days old.
- B. signed URIs also can allow the user to 1) access exclusively their own data and 2) access it for a set period of time, in this case 30 minutes
- C. an SFTP server? Really? This is GCP, not circa 2000. Plus it's high maintenance.
- D. Ok, though it's better than C, it's still clunky. Why not use a pre-built lifecycle policy? Plus this doesn't tell us how objects are marked as expired anyways. Sounds like extra work.
- E. Now this one is as bad as C. All that looping isn't sustainable as the dataset grows. Plus it's overly complex, and doesn't even mention how the script gets invoked in the first place. Another can of worms.

A and B for sure.

upvoted 9 times

✉  **stepkurniawan** 2 years, 2 months ago

But if we implement A: Storage lifecycle -> it will affect all the files, including the files that are not created by the supplier...

upvoted 1 times

✉  **tothecloud** 2 years, 2 months ago

AB is correct

upvoted 3 times

Question #160

Your company wants to standardize the creation and management of multiple Google Cloud resources using Infrastructure as Code. You want to minimize the amount of repetitive code needed to manage the environment. What should you do?

- A. Develop templates for the environment using Cloud Deployment Manager.
- B. Use curl in a terminal to send a REST request to the relevant Google API for each individual resource.
- C. Use the Cloud Console interface to provision and manage all related resources.
- D. Create a bash script that contains all requirement steps as gcloud commands.

Correct Answer: A

Reference:

<https://cloud.google.com/deployment-manager/docs/fundamentals>
(see templates)

Community vote distribution

A (100%)

✉  **jmgf** Highly Voted 1 year, 8 months ago

A

You can use Google Cloud Deployment Manager to create a set of Google Cloud resources and manage them as a unit, called a deployment. For example, if your team's development environment needs two virtual machines (VMs) and a BigQuery database, you can define these resources in a configuration file, and use Deployment Manager to create, change, or delete these resources. You can make the configuration file part of your team's code repository, so that anyone can create the same environment with consistent results.

<https://cloud.google.com/deployment-manager/docs/quickstart>

upvoted 30 times

✉  **GCP_Student1** Highly Voted 1 year, 7 months ago

A. Develop templates for the environment using Cloud Deployment Manager.

upvoted 11 times

✉  **AzureDP900** Most Recent 4 months, 2 weeks ago

A is right

upvoted 1 times

✉  **POOJA3808** 7 months, 3 weeks ago

Selected Answer: A

Develop templates for the environment using Cloud Deployment Manager.

upvoted 1 times

✉  **look1** 11 months, 1 week ago

Selected Answer: A

Templates only

upvoted 2 times

✉  **kartz14** 11 months, 2 weeks ago

Its A.

upvoted 2 times

✉  **TAvenger** 1 year, 7 months ago

A.

According to this document <https://cloud.google.com/solutions/infrastructure-as-code>

IaC (Infrastructure as code) tools for Google Cloud:

Deployment Manager, Terraform, Puppet, Chef ...

upvoted 3 times

✉  **lio123** 1 year, 8 months ago

I think A

upvoted 2 times

Question #161

You are performing a monthly security check of your Google Cloud environment and want to know who has access to view data stored in your Google Cloud Project. What should you do?

- A. Enable Audit Logs for all APIs that are related to data storage.
- B. Review the IAM permissions for any role that allows for data access.
- C. Review the Identity-Aware Proxy settings for each resource.
- D. Create a Data Loss Prevention job.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/docs/access>*Community vote distribution*

B (94%) 6%

 **JelloMan** Highly Voted 5 months, 2 weeks ago

Selected Answer: B

Only use audit logs to look at history (PAST)
 If you need current, up-to-date, info regarding permissions always go to IAM
 upvoted 6 times

 **Alejondri** Highly Voted 6 months ago

Selected Answer: B

B is the one:

- A. Enable Audit Logs for all APIs that are related to data storage. --> That is not the correct answer, if someone with permissions has not accessed or does not access, it will not be listed.
- B. Review the IAM permissions for any role that allows for data access. --> That's correct
- C. Review the Identity-Aware Proxy settings for each resource. --> Nothing relevant, Proxy? Is configured? The question don't ask or tell something about if it is configured.
- D. Create a Data Loss Prevention job. --> Data Loss Prevention nothing to see here.

upvoted 5 times

 **snkhatri** Most Recent 2 months ago

B "WHO HAS ACCESS"

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct

upvoted 1 times

 **akshaychavan7** 5 months, 2 weeks ago

Selected Answer: B

Without any doubt, it's B.

upvoted 3 times

 **Terzlightyear** 6 months ago

Selected Answer: B

B is the one

upvoted 3 times

 **sdflkds** 6 months ago

B. 'Audit logs help you answer "who did what, where, and when?"'(from <https://cloud.google.com/logging/docs/audit>). So, not who has access, but rather who accessed.

upvoted 1 times

 **Malib** 6 months, 1 week ago

Selected Answer: A

La réponse A.

upvoted 1 times

Question #162

Your company has embraced a hybrid cloud strategy where some of the applications are deployed on Google Cloud. A Virtual Private Network (VPN) tunnel connects your Virtual Private Cloud (VPC) in Google Cloud with your company's on-premises network. Multiple applications in Google Cloud need to connect to an on-premises database server, and you want to avoid having to change the IP configuration in all of your applications when the IP of the database changes.

What should you do?

- A. Configure Cloud NAT for all subnets of your VPC to be used when egressing from the VM instances.
- B. Create a private zone on Cloud DNS, and configure the applications with the DNS name.
- C. Configure the IP of the database as custom metadata for each instance, and query the metadata server.
- D. Query the Compute Engine internal DNS from the applications to retrieve the IP of the database.

Correct Answer: A

Community vote distribution

B (100%)

 **kopper2019**  1 year, 6 months ago

B,

Forwarding zones

Cloud DNS forwarding zones let you configure target name servers for specific private zones. Using a forwarding zone is one way to implement outbound DNS forwarding from your VPC network.

A Cloud DNS forwarding zone is a special type of Cloud DNS private zone. Instead of creating records within the zone, you specify a set of forwarding targets. Each forwarding target is an IP address of a DNS server, located in your VPC network, or in an on-premises network connected to your VPC network by Cloud VPN or Cloud Interconnect.

A does not apply, that is to provide internet access to resources

C, does not apply

D, I don't get it

so B

upvoted 25 times

 **MacFreak** 2 months, 1 week ago

"A does not apply, that is to provide internet access to resources" - do you really think NAT is only being used between public and private? Well...it's not! :)

upvoted 1 times

 **meh009** 1 year, 6 months ago

Agreed, It's B although I chose A initially. After some careful consideration and understanding how Cloud NAT works, I'm sticking with B
<https://cloud.google.com/nat/docs/overview>

upvoted 3 times

 **meh009** 1 year, 6 months ago

Further clarification:

"On-premises clients can resolve records in private zones, forwarding zones, and peering zones for which the VPC network has been authorized. On-premises clients use Cloud VPN or Cloud Interconnect to connect to the VPC network."

upvoted 1 times

 **djgodzilla** 1 year, 4 months ago

this is talking about On-premises client resolving nodes outside their network . the question is about how would the application tier within the VPC would resolve the database server . you're confusing the resolution direction my friend

upvoted 2 times

 **djgodzilla** 1 year, 4 months ago

It is still B , but it's rather outbound forward that's needed here :

DNS outbound Forwarding :

- Set up outbound forwarding private zones to query on-premises servers (On-prem Authoritative Zone: corp.example.com)
- In Cloud Router , add a custom route advertisement for GCP DNS proxies range 35.199.192.0/19 to the on-premises environment.
- Make sure inbound DNS traffic from 35.199.192.0/19 is allowed on on-prem firewall
- Cloud Router should be learning on-prem network route from On-prem Router

https://youtu.be/OH_Jw8NhEGU?t=1283

https://cloud.google.com/dns/docs/best-practices#use_forwarding_zones_to_query_on-premises_servers

upvoted 3 times

 **pondai**  1 year, 7 months ago

https://cloud.google.com/dns/docs/best-practices#best_practices_for_dns_forwarding_zones_and_server_policies

Cloud DNS offers DNS forwarding zones and DNS server policies to allow lookups of DNS names between your on-premises and Google Cloud environment. You have multiple options for configuring DNS forwarding. The following section lists best practices for hybrid DNS setup. These best practices are illustrated in the Reference architectures for hybrid DNS.

So I think B is correct

upvoted 6 times

 **Charumathi** Most Recent 1 month ago

Selected Answer: B

B is correct answer,
Configure Private Google Access for on-premises hosts,

DNS configuration

Your on-premises network must have DNS zones and records configured so that Google domain names resolve to the set of IP addresses for either private.googleapis.com or restricted.googleapis.com. You can create Cloud DNS managed private zones and use a Cloud DNS inbound server policy, or you can configure on-premises name servers. For example, you can use BIND or Microsoft Active Directory DNS.

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid#config-domain>

upvoted 1 times

 **aforolt** 1 month, 1 week ago

Ans is D, looks like there is typo

upvoted 1 times

 **habros** 2 months, 3 weeks ago

B. DNS works best with dynamic IPs.

upvoted 1 times

 **patashish** 4 months, 1 week ago

Correct Ans is B

Ref - https://cloud.google.com/dns/docs/best-practices#best_practices_for_private_zones

upvoted 1 times

 **deadlydeb** 4 months, 1 week ago

Selected Answer: B

B Cloud DNS

upvoted 2 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct

upvoted 1 times

 **deadlydeb** 4 months, 3 weeks ago

Selected Answer: B

B: DNS

upvoted 1 times

 **Mike_Malone** 10 months ago

ans D

<https://cloud.google.com/compute/docs/internal-dns>

upvoted 3 times

 **Ridhanya** 11 months ago

option B

upvoted 1 times

 **FernandoJ** 11 months, 2 weeks ago

Selected Answer: B

<https://cloud.google.com/dns/docs/overview#:~:text=Create%20an%20inbound,the%20VPC%20network.>

upvoted 1 times

 **gloisv** 1 year, 3 months ago

IT's D, because:

- A) Cloud NAT direction will be from the cloud resources to the on prem, an the DB is on prem (It will not work if the IP of the database on prem changes, and you have an VPN you should traffic your data through it).
- B) If you create a private zone and configure the applications, if your servers doesn't use the dns will not work.
- C) I think is not suitable. I guess you should re-deploy all your Apps with the new conf if the address change.
- D) If you have an A record ip of your DB HOST(which is on prem) in Engine internal DNS and if it changes, you can update the registry quickly to change it to the new IP address, so it will be the best option for sure.

upvoted 3 times

 **iloveexam** 1 year, 3 months ago

I choose D.

It means setting up a compute engine as internal DNS that contains a domain that resolves IP address of the database.

upvoted 2 times

 **AkshayKalbhor** 1 year, 4 months ago

Options A, B, C do not make sense to me, I think there is a typo in option D:

D: Query the [Compute Engine] internal DNS from the applications to retrieve the IP of the database.

D: Query the ON-PREMISE internal DNS from the applications to retrieve the IP of the database.

upvoted 3 times

 **crazyaboutazure** 1 year, 4 months ago

Must be A

When you create a Cloud NAT gateway, you can choose to have the gateway automatically allocate regional external IP addresses. Alternatively, you can manually assign a fixed number of regional external IP addresses to the gateway. For details about each method, see NAT IP addresses.

You can configure the number of source ports that each Cloud NAT gateway reserves to each VM for which it should provide NAT services. The VMs for which NAT should be provided are determined by the subnet IP address ranges that the gateway is configured to serve.

upvoted 2 times

 **crazyaboutazure** 1 year, 4 months ago

Sorry it has to be B

<https://gcloud.devoteam.com/blog/google-cloud-platform-dns-forwarding-big-thing-enterprises>

upvoted 3 times

 **AD_0525** 1 year, 4 months ago

Answer is A

upvoted 1 times

Question #163

You have developed a containerized web application that will serve internal colleagues during business hours. You want to ensure that no costs are incurred outside of the hours the application is used. You have just created a new Google Cloud project and want to deploy the application. What should you do?

- A. Deploy the container on Cloud Run for Anthos, and set the minimum number of instances to zero.
- B. Deploy the container on Cloud Run (fully managed), and set the minimum number of instances to zero.
- C. Deploy the container on App Engine flexible environment with autoscaling, and set the value min_instances to zero in the app.yaml.
- D. Deploy the container on App Engine flexible environment with manual scaling, and set the value instances to zero in the app.yaml.

Correct Answer: C*Community vote distribution*

B (100%)

 **crysk** Highly Voted 1 year, 8 months ago

I think that is B the correct answer, because Cloud Run can scale to 0:

<https://cloud.google.com/run/docs/about-instance-autoscaling>

And App Engine Flexible can't scale to 0, the minimum instance number is 1:

https://cloud.google.com/appengine/docs/the-appengine-environments#comparing_high-level_features

upvoted 30 times

 **ryumada** 3 months ago

No for the App Engine Flexible Environment, but App Engine Standard can also scale to zero.

upvoted 1 times

 **pca2b** Highly Voted 1 year, 7 months ago

B:

not A because Anthos is an add-on to GKE clusters, 'new project' means we don't have a GKE cluster to work with

https://cloud.google.com/kuberun/docs/architecture-overview#components_in_the_default_installation

upvoted 13 times

 **Charumathi** Most Recent 1 month ago

Selected Answer: B

B is the correct answer,

Cloud Functions can scale to zero, whereas App Engine will not be able to scale to zero, it should have at least one instance.

Add-on Info,

App-Engine Standard can scale to zero, whereas App-Engine Flexible couldn't scale down to zero.

upvoted 1 times

 **sylva91** 1 month, 3 weeks ago

Selected Answer: B

B is the answer since we can scale to 0 and the other key word is "containerized"

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: B

B is the right answer

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct

upvoted 1 times

 **tommatali75** 5 months, 4 weeks ago

<https://cloud.google.com/appengine/docs/standard/python/config/appref>
min_instances

Warning: For this feature to function properly, you must make sure that warmup requests are enabled and that your application handles warmup requests.

Note: This setting applies only if the version of the app defined by this app.yaml file is configured to receive traffic. To learn more about routing traffic to different versions of an app, see Splitting Traffic.

Optional. The minimum number of instances for App Engine to create for this module version. These instances serve traffic when requests arrive, and continue to serve traffic even when additional instances are started up as required to handle traffic.

Specify a value from 0 to 1000. You can set the parameter to the value 0 to allow scaling to 0 instances to lower costs when no requests are being served. Note that you are charged for the number of instances specified whether they are receiving traffic or not.

So C

upvoted 1 times

 **JelloMan** 6 months ago

Selected Answer: B

along with the reason that most have stated (only Cloud Run can scale down to 0 instances) another reason is that Cloud Run is pay-per-use. App-engine flexible is paid based on usage of vCPU, memory, and persistent disks, so you will be racking up cost quickly just because your VM's are created (regardless if they're running or not)

upvoted 2 times

 **jblima** 6 months, 2 weeks ago

Selected Answer: B

B is the correct

upvoted 1 times

 **nshah68** 8 months ago

Selected Answer: B

Flex cannot scale down to 0. Standard can but that is not relevant here.

Therefore, Cloud Run is the best answer since it can scale down to 0 when there is no traffic

upvoted 2 times

 **luciorifa** 8 months, 2 weeks ago

Selected Answer: B

Correct answer is B, Cloud Run is a serverless solution that, same as Cloud Functions, can be activated based on events.

upvoted 1 times

 **DaveNZ** 8 months, 2 weeks ago

Selected Answer: B

The majority vote here seems to be B so I'm posting this as a voting comment to make that visible. I am not entirely sure, but suspect that the "min-instances" issue is a red herring - that is, I understand why some comments raise this as being a reason why B may not be correct but I think maybe it's not the intention of the question to focus on that rather tricky aspect.

upvoted 1 times

 **mahesh0049** 9 months, 3 weeks ago

Option B is correct. After Business hours there will be no colleagues to work on that application then cloud run can scale to zero instances

upvoted 2 times

 **gcpengineer** 1 year, 3 months ago

Ans is A. as we seek the solution for internal users, Anthos makes vpn/interconnect to cloud Run, only from internal ip address

upvoted 3 times

 **umashankar_a** 1 year, 4 months ago

It's Option B... Folks !!

"By default, container instances have min-instances turned off, with a setting of 0"

This statement means - By default min instances is set to Zero, unless you want to change it and keep your payload warm for reducing the latency.

App Engine Flexible documentation clearly stated that min value cannot be less than 1. Crysk is absolutely right and that leaves us with option B - Cloud Run (Fully Managed.)

Please read through the doc supplied by Crysk.

upvoted 3 times

 **AkshayKalbhor** 1 year, 4 months ago

B IS WRONG,

Docs state the below line.

"By default, container instances have min-instances turned off, with a setting of 0" which means that the "minimum instances" value cannot have a value of '0'.

Reference : <https://cloud.google.com/run/docs/configuring/min-instances>

C IS CORRECT,

As we can deploy containerized web application and we can set the "min_instances" value to zero in the app.xml as mentioned in the answer.

upvoted 1 times

 **Ridhanya** 11 months ago

app engine flexible instances can never have 0 instances, so your argument is wrong

upvoted 3 times

 **mchaconr** 10 months, 3 weeks ago

This is wrong.

From same link you provided "By default, container instances have min-instances turned off, with a setting of 0. You can change this default using the Cloud Console, the gcloud command line, or a YAML file when you create a new service or deploy a new revision".

By default, cloud run scales to zero. That means, nothing is running until requested (reason why "minimum instances" is already 0 by default). The minimum instances is an optional value to Cloud Run, in case you want to have a warn (at least) instance running, an speed up that first

request (reduce latency for it), yet not mandatory. You are charged for the resources required to have that instance running (not to the full or same extend of a compute instance, of course).

upvoted 2 times

 **AD_0525** 1 year, 4 months ago

App engine Flexible does not support basic scaling i.e. it can't scale to 0. So answer should be B

upvoted 3 times

Question #164

You have experimented with Google Cloud using your own credit card and expensed the costs to your company. Your company wants to streamline the billing process and charge the costs of your projects to their monthly invoice. What should you do?

- A. Grant the financial team the IAM role of `Billing Account User` on the billing account linked to your credit card.
- B. Set up BigQuery billing export and grant your financial department IAM access to query the data.
- C. Create a ticket with Google Billing Support to ask them to send the invoice to your company.
- D. Change the billing account of your projects to the billing account of your company.

Correct Answer: D*Community vote distribution*

D (100%)

 **j_mrnr** Highly Voted 1 year, 8 months ago

1000% Ans D

upvoted 23 times

 **Cornholio_LMC** Most Recent 1 month, 2 weeks ago

had this one today

upvoted 2 times

 **deadlydeb** 4 months, 1 week ago

Selected Answer: D

D D D D D

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D is perfect

upvoted 1 times

 **rsuresh27** 6 months, 2 weeks ago

Please do not overthink the question. The question does not mention anything about finance teams, so A cannot be correct. D is the only one that makes sense out of the remaining options.

upvoted 3 times

 **Surat** 10 months ago

D seems correct ans

upvoted 2 times

 **walkwolf3** 11 months ago

Either A or D are incomplete solution.

Soution A

Grant financial team the IAM role of Project Owner or Project Billing Manager on your project, then let financial team change the billing account.

Solution B

Grant you IAM role of Billing Account Administrator or Billing Account User on company's project, then you can change the billing account.

I hope this question won't be in the exam, or has a more accurate answer in the exam.

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

D -- is the answer

cant be A. you dont want to share your personnel credit card data with company billing account

upvoted 1 times

 **vamgcp** 1 year, 1 month ago

correct answer D- Change the billing account of your projects to the billing account of your company

upvoted 2 times

 **arsh1916** 1 year, 5 months ago

D is correct

upvoted 2 times

 **sivre** 1 year, 7 months ago

Does the user have access to the company billing account? Not sure of D, I think more A.

To change the Cloud Billing account for a project, you need to be able to move a project from one Cloud Billing account to another. To accomplish this task, you need permissions adequate to unlink the project from the existing Cloud Billing account AND to link the project to the target Cloud Billing account.

Roles with adequate permissions to perform this task: Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account

https://cloud.google.com/billing/docs/how-to/modify-project#change_the_billing_account_for_a_project

upvoted 2 times

 **GoCloud** 1 year, 7 months ago

A would be ok if assigning Project Billing Manager IAM role to the finance team.

So the answer will be D.

upvoted 1 times

 **muneebarshad** 1 year, 4 months ago

"Billing Account User" Role Link projects to billing accounts. Since It does not say anything about user permissions , lets assume if user does not have sufficient permission then Granting the financial team the ""Billing Account User" " role will do the job

upvoted 1 times

 **wh1t4k3r** 11 months ago

"Grant the financial team the IAM role of `Billing Account User` on the billing account linked to your credit card. "

Giving billing user rights to the financial team to the billing account binded to your credit card wont give them any means to change it or maybe direct/export to the company's billing.

Billing user gives the following permissions on a billing account:

" This role has very restricted permissions, so you can grant it broadly, typically in combination with Project Creator. These two roles allow a user to create new projects linked to the billing account on which the role is granted. "

upvoted 1 times

 **wh1t4k3r** 11 months ago

Still, you do make a valid point regarding the access. How will i add my projects to the company's wihtout access? Its not possible indeed.

That said, i would say that the best option is B (and when i say best i mean for the given options), since you can explort billing data in bigquery:

<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

in B option we can get only billing / budget data, but question asks for streamlining of billing process, mean they want to pay bills using company invoice.

upvoted 1 times

 **GCP_Student1** 1 year, 7 months ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 4 times

 **TAvenger** 1 year, 7 months ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 2 times

 **lio123** 1 year, 8 months ago

yes it is D

upvoted 2 times

Question #165

You are running a data warehouse on BigQuery. A partner company is offering a recommendation engine based on the data in your data warehouse. The partner company is also running their application on Google Cloud. They manage the resources in their own project, but they need access to the BigQuery dataset in your project. You want to provide the partner company with access to the dataset. What should you do?

- A. Create a Service Account in your own project, and grant this Service Account access to BigQuery in your project.
- B. Create a Service Account in your own project, and ask the partner to grant this Service Account access to BigQuery in their project.
- C. Ask the partner to create a Service Account in their project, and have them give the Service Account access to BigQuery in their project.
- D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.

Correct Answer: D

Community vote distribution

D (100%)

✉️ **GCP_Student1** Highly Voted 1 year, 7 months ago

D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.
upvoted 28 times

✉️ **pondai** Highly Voted 1 year, 7 months ago

BigQuery is in our project, so we need to create a service account and grant it access BigQuery role. That can make partner company to use this account to use it to access our project's BigQuery. So I vote A.

upvoted 12 times

✉️ **tavva_prudhvi** 1 year, 7 months ago

See, the ones who want our access needs to create a service account (in our case it's the partner company), then we give access to the service account with the user permissions. Clearly, D says the same thing!
upvoted 7 times

✉️ **akshaychavan7** 5 months, 2 weeks ago

Your understanding is bit wrong here, my friend!
upvoted 1 times

✉️ **Aninina** Most Recent 1 week ago

Selected Answer: D

"Service accounts are both identities and resources. Because service accounts are identities, you can let a service account access resources in your project by granting it a role, just like you would for any other principal."
upvoted 1 times

✉️ **Cornholio_LMC** 1 month, 2 weeks ago

had this one today
upvoted 1 times

✉️ **AzureDP900** 4 months, 2 weeks ago

D is right
upvoted 1 times

✉️ **somenick** 7 months, 2 weeks ago

Selected Answer: D
<https://gtseres.medium.com/using-service-accounts-across-projects-in-gcp-cf9473fef8f0#:~:text=Go%20to%20the%20destination%20project,Voila!>
upvoted 2 times

✉️ **mk1708** 8 months, 1 week ago

I thought it was A. But when I quickly did some research I found this:
"Service accounts are both identities and resources. Because service accounts are identities, you can let a service account access resources in your project by granting it a role, just like you would for any other principal."
Thus, the answer is D.
upvoted 3 times

✉️ **wh1t4k3r** 11 months ago

D per my understanding: if the need is to authenticate the application to access your dataset, it's the application's service account that will be provided during the authentication, so the service account is to be created at their side to run the application, not the other way around.
upvoted 3 times

✉️ **wh1t4k3r** 11 months ago

Another insight:

A is too broad. The question states: "provide the partner company with access to the dataset"

A states: "grant this Service Account access to BigQuery in your project"

I think D is a more granular option, given that A would give access to all datasets in your bigquery data warehousing.

upvoted 2 times

 **jabrrJ68w02ond1** 11 months, 2 weeks ago

How is it D? I want to give access to my BigQuery data, so I need to provide the ServiceAccount. I create it, put some decent predefined roles on it, and whenever I stop working with the other company, I either invalidate the JSON key of the SA or I simply delete the SA. For me, it is A.

upvoted 3 times

 **JelloMan** 6 months, 3 weeks ago

Righ but how will the other project ever gain access to resources on YOUR project? Key thing here is that your sharing ACROSS different projects so you need a bridge between them. D provides that bridge by connecting THEIR service account with YOUR resource (big query)

upvoted 1 times

 **AD_0525** 1 year, 4 months ago

Answer should be D, as the other company project needs access in your project.

upvoted 6 times

 **EmreA** 1 year, 6 months ago

It's A

upvoted 4 times

 **kopper2019** 1 year, 6 months ago

D is the answer

upvoted 6 times

 **TAvenger** 1 year, 7 months ago

I think it is "D" but I have not found such usecase when you share dataset with another organization via service account

upvoted 3 times

 **lio123** 1 year, 8 months ago

it is D

upvoted 4 times

 **nshtnrt** 1 year, 8 months ago

D looks correct to me.

upvoted 6 times

Question #166

Your web application has been running successfully on Cloud Run for Anthos. You want to evaluate an updated version of the application with a specific percentage of your production users (canary deployment). What should you do?

- A. Create a new service with the new version of the application. Split traffic between this version and the version that is currently running.
- B. Create a new revision with the new version of the application. Split traffic between this version and the version that is currently running.
- C. Create a new service with the new version of the application. Add an HTTP Load Balancer in front of both services.
- D. Create a new revision with the new version of the application. Add an HTTP Load Balancer in front of both revisions.

Correct Answer: A*Community vote distribution*

B (100%)

 **crysk** Highly Voted 1 year, 8 months ago

In my opinion correct answer is B:

https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration?utm_campaign=CDR_ahm_aap-severless_cloud-run-faq_&utm_source=external&utm_medium=web

Cloud Run can split traffic between revisions

upvoted 33 times

 **TAvenger** 1 year, 7 months ago

The google doc link is incorrect. You need to specify CloudRun for Anthos

<https://cloud.google.com/kuberun/docs/rollouts-rollbacks-traffic-migration>

Anyway principles for CloudRun and CloudRun for Anthos are the same. Traffic can be split between multiple revisions.

The answer is "B"

upvoted 6 times

 **GCP_Student1** Highly Voted 1 year, 7 months ago

B. Create a new revision with the new version of the application. Split traffic between this version and the version that is currently running.

upvoted 9 times

 **Untamables** Most Recent 1 week, 6 days ago

Selected Answer: B

The latest Document

<https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration>

upvoted 1 times

 **nonsense** 2 weeks, 2 days ago

Selected Answer: B

B. Create a new revision

upvoted 1 times

 **sylva91** 1 month, 3 weeks ago

Selected Answer: B

B is the correct answer

upvoted 1 times

 **thimai** 2 months, 1 week ago

Selected Answer: B

i think B

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

B is correct for this scenario, there is no need to create new services

upvoted 2 times

 **pfabio** 5 months, 2 weeks ago

Selected Answer: B

Cloud Run for Anthos allows you to specify which revisions should receive traffic and to specify traffic percentages that are received by a revision

upvoted 1 times

 **HansKloss611** 9 months, 1 week ago

Selected Answer: B

If course B

upvoted 1 times

 **AkshayKalbhor** 1 year, 4 months ago

Answer should be B.

upvoted 5 times

 **RishiAbhishek** 1 year, 4 months ago

Canary deployments are a method of releasing software to a subset of users or servers. The plan is to deliver the update to a small selection of servers first, test it, and then roll it out to the other servers.

So it's a Revision or new update or a new version, not a service.

upvoted 4 times

 **nana1995** 1 year, 5 months ago

Its B

<https://cloud.google.com/run/docs/managing/revisions>

upvoted 2 times

 **Ashii** 1 year, 5 months ago

Its B, revision or version, not service.

upvoted 2 times

Question #167

Your company developed a mobile game that is deployed on Google Cloud. Gamers are connecting to the game with their personal phones over the Internet. The game sends UDP packets to update the servers about the gamers' actions while they are playing in multiplayer mode. Your game backend can scale over multiple virtual machines (VMs), and you want to expose the VMs over a single IP address. What should you do?

- A. Configure an SSL Proxy load balancer in front of the application servers.
- B. Configure an Internal UDP load balancer in front of the application servers.
- C. Configure an External HTTP(s) load balancer in front of the application servers.
- D. Configure an External Network load balancer in front of the application servers.

Correct Answer: A

Reference:

<https://cloud.google.com/solutions/connecting-securely>*Community vote distribution*

D (100%)

 **kopper2019** Highly Voted 1 year, 6 months ago

Answer is D, cell phones are sending UDP packets and the only that can receive that type of traffic is a External Network TCP/UDP
<https://cloud.google.com/load-balancing/docs/network>

upvoted 24 times

 **ashraff** 1 year, 2 months ago

Google Cloud HTTP(S) Load Balancing is a global, proxy-based Layer 7 load balancer that enables you to run and scale your services worldwide behind a single external IP address. External HTTP(S) Load Balancing distributes HTTP and HTTPS traffic to backends hosted on Compute Engine and Google Kubernetes Engine (GKE).

<https://cloud.google.com/load-balancing/docs/https>

upvoted 2 times

 **ryumada** 3 months ago

All the load balancer products in GCP give you a single IP address for the backend servers you registered to it.

Also, External HTTP(s) load balancer only support the port that used by HTTP which is the port 80 and HTTPS which is the port 443.

And Google Cloud external TCP/UDP Network Load Balancing is referred to as "Network Load Balancing" which supports UDP packets.

- <https://cloud.google.com/load-balancing/docs/load-balancing-overview#about>
- <https://cloud.google.com/load-balancing/docs/network>
- <https://cloud.google.com/load-balancing/docs/https>

upvoted 2 times

 **patashish** 3 months, 4 weeks ago

what you are trying to say ? What is your answer ? A B C D ?

upvoted 3 times

 **JH86** Highly Voted 1 year, 5 months ago

Answer is D. there are so many confusion here, from B,C or D. For myself im eliminating all options except B,D due to the traffic type. which leaves me with B or D. Then next the traffic source either external or internal which in this case is an external traffic from the internet, therefore my final answer is D.

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

upvoted 10 times

 **BobbyFlash** 1 year ago

Following the diagram, there's no doubt about D. We have external clients connecting to our gaming service on google cloud that works using UDP traffic that results in using External Network Load Balancing. I feel that it's simple as it is. I also go with D.

upvoted 4 times

 **abirroy** Most Recent 3 months, 1 week ago

Selected Answer: D

External Network LB used for UDP

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

D seems correct..

upvoted 1 times

 **akshaychavan7** 5 months, 2 weeks ago

Selected Answer: D

I'm dead sure, it's D!

upvoted 1 times

 **somenick** 7 months, 2 weeks ago

Selected Answer: D

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer#lb-decision-tree>

upvoted 2 times

 **[Removed]** 7 months, 4 weeks ago

Selected Answer: D

D - Check <https://cloud.google.com/load-balancing/images/choose-lb.svg>

upvoted 1 times

 **RazOr** 9 months, 2 weeks ago

Selected Answer: D

D, because:

<https://cloud.google.com/load-balancing/docs/network#:~:text=Google%20Cloud%20external-,TCP/UDP,-Network%20Load%20Balancing>

upvoted 2 times

 **exam_war** 10 months ago

Answer is C. Players need to access through internet by HTTP(S) load balancing

upvoted 2 times

 **S00999** 4 months, 3 weeks ago

It is not specified whether the app protocol is HTTP(S) or not, only that it is UDP packets. Internet is not limited to the http protocol.

Answer D

upvoted 1 times

 **[Removed]** 10 months ago

D. Configure an External Network load balancer in front of the application servers.

"VM over single (external) IP address ->>> getting UDP packets through External LB

"

upvoted 2 times

 **Eben01** 10 months, 2 weeks ago

The key statement we need to watch out for ; "And you want to expose the Vms over a single IP address" .

Google Cloud external TCP/UDP Network Load Balancing (after this referred to as Network Load Balancing) is a regional, "pass-through" load balancer. A network load balancer "distributes external traffic among virtual machine (VM) instances" in the same region.

You can configure a network load balancer for TCP, UDP, ESP, and ICMP traffic.

<https://cloud.google.com/load-balancing/docs/network>

upvoted 2 times

 **dttnc1** 1 year ago

My answer here is A for the following reasons. Feel free to comment if I missed something.

1. For a mobile game to reach its maximum audience, it must be available globally. SSL proxy LB supports global load balancing service with the Premium Tier.
2. It is intended for non-HTTP(S) traffic. For HTTP(S) traffic, GCP recommends using HTTP(S) Load Balancing.
3. Preserves the original source IP addresses of incoming connections to the load balancer
4. SSL proxy LB allows you to expose a single IP address

https://cloud.google.com/network-tiers/docs/overview#configuring_standard_tier_for_https_lb_and_tcpssl_proxy
<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

 **pas77** 1 year, 1 month ago

I was also thinking about option D. Whoever, my problem with option D is that the scenario mentions the following:

"Your game backend can scale over multiple virtual machines (VMs)"

And unless I am reading the documentation wrong, Layer 4 Network LB does not support autoscaling.

https://cloud.google.com/load-balancing/docs/features#autoscaling_and_autohealing

upvoted 2 times

 **procloud** 1 year, 1 month ago

Answer is D ,<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

Load balancer type Traffic type

External HTTP(S) HTTP or HTTPS

External TCP/UDP Network

(also referred to as Network Load Balancing) TCP, UDP, ESP, or ICMP (Preview)

upvoted 2 times

 **associatecloudexamuser** 1 year, 3 months ago

Oh. Missed "Internal" term. So it should be External Load Balancer.

Answer is D.

upvoted 4 times

 **curious_kitty9** 1 year, 5 months ago

Answer is D, Reason- UDP and preserve client ID

upvoted 1 times

 **jahnu** 1 year, 6 months ago

Ans: B check in <https://cloud.google.com/load-balancing/docs/internal/multiple-forwarding-rules-same-ip>

upvoted 3 times

Question #168

You are working for a hospital that stores its medical images in an on-premises data room. The hospital wants to use Cloud Storage for archival storage of these images. The hospital wants an automated process to upload any new medical images to Cloud Storage. You need to design and implement a solution. What should you do?

- A. Create a Pub/Sub topic, and enable a Cloud Storage trigger for the Pub/Sub topic. Create an application that sends all medical images to the Pub/Sub topic.
- B. Deploy a Dataflow job from the batch template, ↗Datastore to Cloud Storage. Schedule the batch job on the desired interval.
- C. Create a script that uses the gsutil command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.
- D. In the Cloud Console, go to Cloud Storage. Upload the relevant images to the appropriate bucket.

Correct Answer: C*Community vote distribution*

C (100%)

 **GCP_Student1** Highly Voted  1 year, 7 months ago

C. Create a script that uses the gsutil command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.

upvoted 19 times

 **cserra** 2 months, 3 weeks ago

Where does it say that the on-premises images are already digitized? and even if they are, where does it say that we also keep the old images?

I think the correct answer is "A"

upvoted 1 times

 **theBestStudent** 1 month, 4 weeks ago

Tell yo yourself how the images would end up in the pubsub first of all. Also usually the process is in the other way around for pubsub notifications: Once an object lands in GCS the pubsub is notified of it.

Option A makes totally nonsense. Check the flow again.

From the options the only one that "makes more sense" is Option C

upvoted 1 times

 **TAvenger** Highly Voted  1 year, 7 months ago

From the question the key point is "upload ANY NEW medical images to Cloud Storage". So we are not interested in old images. That's why we need some trigger that will upload images. I think option "A" with PubSub is the best

upvoted 18 times

 **dunhill** 1 year, 7 months ago

I am not sure but the question also mentions that "wants to use Cloud Storage for archival storage of these images". It can create an application that sends all medical images to storage and no need via PubSub?

upvoted 1 times

 **pca2b** 1 year, 7 months ago

Pub/Sub will be good for all future files in in-prem data-storage.

we want to sync all + new, so a local on-prem server running a cron job (not GCE CronJob) to run gsutil to transfer files to Cloud Storage would work.

I vote for C

upvoted 5 times

 **yvinisiupacuando** 1 year, 6 months ago

Sorry you are wrong, the question clearly indicates "The hospital wants an automated process to upload ANY NEW medical images to Cloud Storage." It does not mention the need to upload the original stock of images, only the new ones. Then I think the right answer must be A, as you said "Pub/sub will be good for all future files in prem data-storage" which is exactly what the questions is pointing to.

upvoted 5 times

 **Priyanka109** 1 month, 1 week ago

In option C we are using a cron job, not dragging and dropping the images.

upvoted 1 times

 **gcpengineer** 1 year, 3 months ago

ans is C

upvoted 3 times

 **Charumathi** Most Recent ⓘ 1 month ago

Selected Answer: C

C is the correct answer.

Keyword, they require cloud storage for archival and the want to automate the process to upload new medical image to cloud storage, hence we go for gsutil to copy on-prem images to cloud storage and automate the process via cron job. whereas Pub/Sub listens to the changes in the Cloud Storage bucket and triggers the pub/sub topic, which is not required.

upvoted 1 times

 **zolthar_z** 3 months, 1 week ago

Selected Answer: C

The Hospital wants Cloud storage for archival of old images and also sync the new images, for this logic the answer is C

upvoted 1 times

 **AzureDP900** 4 months, 2 weeks ago

C is correct

upvoted 1 times

 **jblima** 6 months, 2 weeks ago

Selected Answer: C

C is correct.

upvoted 1 times

 **wael_tn** 6 months, 2 weeks ago

Selected Answer: C

Discarding (A) because "Cloud Storage trigger". So for option A the triggering event should be making a change in Cloud Storage, while in the real use case, the triggering should be adding a new medical image to the "on-premises data room"

upvoted 2 times

 **theBestStudent** 1 month, 4 weeks ago

Correct, and on top of that, how are they suppose to connect the on premises to the cloud. Nothing is mentioned. But either way, option A does not make sense at all (as you already explained too). Correct option: C

upvoted 1 times

 **somenick** 7 months, 2 weeks ago

Selected Answer: C

I would go with C

Not A, don't think you can send image files to Pub/Sub. Technically you can do so by converting image to some binary text, but then we don't know the size of the image and there is a limitation on message size. Not recommended.

Not B – there is only this template “Datastore to Cloud Storage Text”, as the name implies it is for text,

<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#datastore-to-cloud-storage-text>, and it reads from datastore which is definitely not where the medical images are stored, from the question "... stores its medical images in an on-premises data room".

Not D – it's not automated

upvoted 3 times

 **lilapause** 9 months ago

Selected Answer: C

Pub/Sub could make sense if you ignore the size limit. But the way it is described with the storage trigger would just not be working.

upvoted 1 times

 **exam_war** 10 months ago

A is correct. The key is automation whenever there is a new image, it needs to upload to cloud storage. Only pub/sub can make the automation work.

upvoted 1 times

 **[Removed]** 10 months ago

I think A(once the hospital receive the new images Cloud pub/sub will act on it) and C(creating a script with rsync command <https://stackoverflow.com/questions/37662416/how-to-sync-a-local-folder-with-a-folder-in-a-google-cloud-platform-bucket>) must be the Correct options

upvoted 1 times

 **RealEL40** 10 months, 3 weeks ago

" any new medical images" so "A" using Pub\Sub

upvoted 1 times

 **wh1t4k3r** 11 months ago

I'll go with A for he following reason:

Yes, the question says that they use onprem to store images and they want those images in gcloud storage from now on, BUT the solution for automation is requested for uploading new images only: "The hospital wants an automated process to upload any new medical images to Cloud Storage"

Using sync (A) for new images implies that you will continue to use your onprem and keep synchronizing it forever... Sync just once for the old images, new images go directly to gcloud via pubsub, and eventually get rid of the onprem.

upvoted 2 times

 **wh1t4k3r** 11 months ago

Correcting: "Using sync (C) for new images implies that you will continue to use your onprem and keep synchronizing it forever... Sync just once for the old images, new images go directly to gcloud via pubsub, and eventually get rid of the onprem."

upvoted 1 times

 **ankatsu2010** 1 year ago

'C' is correct. Question says 'use Cloud Storage for archival'. Pub/Sub is for real time processing and in addition, 'Cloud Storage trigger' does not make sense.

upvoted 3 times

 **jackdbd** 1 year, 3 months ago

I would go with C.

Here is my rationale for excluding option A.

Pub/Sub accepts a maximum of 1,000 messages in a batch, and the size of a batch can not exceed 10 megabytes. A single CT scan can be 35 MB, so unless we chunk it up in many smaller pieces and publish them as smaller messages, using a Pub/Sub topic is not a viable alternative.

upvoted 11 times

 **ryzior** 8 months ago

I think the question is: why to use pubsub when I can do it better using direct gcs upload with multipart option etc. and this approach is covered in C :)

upvoted 1 times

 **associatecloudexamuser** 1 year, 3 months ago

Answer is C. gsutil rsync <source_location> <destination_location>. This can sync content with Google cloud storage locations

upvoted 3 times

 **JieHeng** 1 year, 4 months ago

I would go with C

Not A, don't think you can send image files to Pub/Sub. Technically you can do so by converting image to some binary text, but then we don't know the size of the image and there is a limitation on message size. Not recommended.

Not B – there is only this template “Datastore to Cloud Storage Text”, as the name implies it is for text,

<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#datastore-to-cloud-storage-text>, and it reads from datastore which is definitely not where the medical images are stored, from the question “... stores its medical images in an on-premises data room”.

Not D – it's not automated

upvoted 2 times

Question #169

Your auditor wants to view your organization's use of data in Google Cloud. The auditor is most interested in auditing who accessed data in Cloud Storage buckets. You need to help the auditor access the data they need. What should you do?

- A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.
- B. Assign the appropriate permissions, and then create a Data Studio report on Admin Activity Audit Logs.
- C. Assign the appropriate permissions, and then use Cloud Monitoring to review metrics.
- D. Use the export logs API to provide the Admin Activity Audit Logs in the format they want.

Correct Answer: D

Reference:

<https://cloud.google.com/storage/docs/audit-logging>

Community vote distribution

A (100%)

✉  **iri_gcp**  1 year, 8 months ago

It should be A.
Data access log are not enabled by default due to the fact that it incurs costs.
So you need to enable it first.
And then you can filter it in the log viewer
upvoted 33 times

✉  **GCP_Student1**  1 year, 7 months ago

A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.
upvoted 10 times

✉  **Charumathi**  1 month ago

Selected Answer: A
A is the correct answer,
Since the auditor wants to know who accessed the cloud storage data, we need data acces logs for cloud storage.

Types of audit logs

Cloud Audit Logs provides the following audit logs for each Cloud project, folder, and organization:

Admin Activity audit logs
Data Access audit logs
System Event audit logs
Policy Denied audit logs

***Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.

<https://cloud.google.com/logging/docs/audit#types>

upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

A is right
upvoted 1 times

✉  **Jman007** 4 months, 3 weeks ago

Selected Answer: A
question says auditor is most interested in who accessED data in Cloud Storage. im not sure how auditing is done for those who answered A but this means they want the logs for past users who accessed the data from a sepecified time. Turning on the feature now is kind of too late. poorly written question and answers. No point in an auditor coming in and giving the company all the exact questions they are going to ask and come back and ask them in a few months time. A seems like the better choices though
upvoted 2 times

✉  **akshaychavan7** 5 months, 2 weeks ago

If it's A then how will we assign the permission for the auditor to view the logs?
I had chosen option A on the first place, but later changed it considering that the auditor won't have the access to view the logs.
upvoted 1 times

✉  **peugeotdude** 6 months, 1 week ago

Selected Answer: A
Based on how I read the question-
We want Data Access log, not Admin Activity Audit Logs.

upvoted 1 times

✉ **somenick** 7 months, 2 weeks ago

Selected Answer: A
Data access log are not enabled by default due to the fact that it incurs costs.
So you need to enable it first.
And then you can filter it in the log viewer
upvoted 1 times

✉ **[Removed]** 7 months, 4 weeks ago

<https://cloud.google.com/logging/docs/audit#data-access>

Cloud Storage: When Cloud Storage usage logs are enabled, Cloud Storage writes usage data to the Cloud Storage bucket, which generates Data Access audit logs for the bucket. The generated Data Access audit log has its caller identity redacted.
upvoted 1 times

✉ **DaveNZ** 8 months, 2 weeks ago

Selected Answer: A
The majority vote here is A, despite some confusion around the wording of the question. I tend to agree because it's the solution that most closely reflects the requirements of the question (buckets, cloud storage).
upvoted 1 times

✉ **wh1t4k3r** 11 months ago

- A. I could not find a way to enable audit logs in specific buckets, only on the whole storage level:
<https://cloud.google.com/logging/docs/audit/services>
- B. Admin activity audit logs cover admin actions, such as metadata or config changes:
<https://cloud.google.com/logging/docs/audit#admin-activity>
- C. Cloud monitoring is not for auditing: <https://cloud.google.com/monitoring>
- D. Again, Admin Activity Audit Logs should not be used to audit data access, specifically from buckets.

My conclusion: all these answers are wrong. My assumption: A is badly written. Specific buckets were not mentioned. I Vote A, but I think this Q&A is messed up. Maybe a correction? or deletion.

upvoted 3 times

✉ **MarcoDipa** 11 months ago

Actually, there is a different service named User Logs that permits to focus on a single bucket.
Refer to google page:
<https://cloud.google.com/storage/docs/access-logs>
Usage logs provide information for all of the requests made on a specified bucket
upvoted 1 times

✉ **obeythefist** 8 months, 2 weeks ago

The question just says "buckets" and hints that the audit should cover all org data, so I don't think there is any need to overanalyse, you are correct in choosing A
upvoted 1 times

✉ **ericyev** 11 months ago

I choose D. reason is here: Cloud Audit Logs generates the following audit logs for operations in Cloud Storage:

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

upvoted 4 times

✉ **kimharsh** 11 months, 1 week ago

Also A because it's the only one that mentions DATA ACCESS LOGS, which is the one that logs objects access, t

Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:

ADMIN_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.

DATA_READ: Entries for operations that read an object.

DATA_WRITE: Entries for operations that create or modify an object.

<https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times

✉ **jackwillis** 1 year ago

Question is about user activities log not about Data Access log.

upvoted 1 times

✉ **ankatsu2010** 1 year ago

Looks like there is no correct answer in this question. Because you can't turn on Data Access Logs for the specific buckets independently. What should you do???

upvoted 1 times

✉ **learnazureportal** 1 year, 5 months ago

based on the <https://cloud.google.com/storage/docs/audit-logging> link, Answer is D.

upvoted 4 times

✉ **JieHeng** 1 year, 4 months ago

from the link, it's the Data Access logs that contain "Entries for operations that read the configuration or metadata of a project, bucket, or object." and "Entries for operations that read an object."

Admin Activity logs do not have these.

upvoted 4 times

✉ **arsh1916** 1 year, 5 months ago

A is correct option

upvoted 1 times

Question #170

You received a JSON file that contained a private key of a Service Account in order to get access to several resources in a Google Cloud project. You downloaded and installed the Cloud SDK and want to use this private key for authentication and authorization when performing gcloud commands. What should you do?

- A. Use the command gcloud auth login and point it to the private key.
- B. Use the command gcloud auth activate-service-account and point it to the private key.
- C. Place the private key file in the installation directory of the Cloud SDK and rename it to `credentials.json`.
- D. Place the private key file in your home directory and rename it to `GOOGLE_APPLICATION_CREDENTIALS`.

Correct Answer: B

Reference:

<https://cloud.google.com/sdk/docs/authorizing>

Community vote distribution

B (100%)

GCP_Student1 Highly Voted 1 year, 7 months ago

- B. Use the command gcloud auth activate-service-account and point it to the private key.

Authorizing with a service account

gcloud auth activate-service-account authorizes access using a service account. As with gcloud init and gcloud auth login, this command saves the service account credentials to the local system on successful completion and sets the specified account as the active account in your Cloud SDK configuration.

https://cloud.google.com/sdk/docs/authorizing#authorizing_with_a_service_account

upvoted 37 times

TAvenger Highly Voted 1 year, 7 months ago

- B.
gcloud auth activate-service-account --help

NAME)

gcloud auth activate-service-account - authorize access to Google Cloud Platform with a service account

SYNOPSIS

```
gcloud auth activate-service-account [ACCOUNT] --key-file=KEY_FILE  
[--password-file=PASSWORD_FILE | --prompt-for-password]  
[GCLOUD_WIDE_FLAG ...]
```

DESCRIPTION

To allow gcloud (and other tools in Cloud SDK) to use service account credentials to make requests, use this command to import these credentials from a file that contains a private authorization key, and activate them for use in gcloud. gcloud auth activate-service-account serves the same function as gcloud auth login but uses a service account rather than Google user credentials.

upvoted 17 times

abirroy Most Recent 3 months ago**Selected Answer: B**

Use the command gcloud auth activate-service-account and point it to the private key

upvoted 1 times

skrjha20 4 months, 1 week ago

ANSWER is B

To activate your service account, run gcloud auth activate-service-account:

```
gcloud auth activate-service-account [ACCOUNT] --key-file=[KEY_FILE]
```

upvoted 1 times

AzureDP900 4 months, 2 weeks ago

I will go with B

upvoted 1 times

AzureDP900 4 months, 4 weeks ago

B is right Please refer

<https://cloud.google.com/storage/docs/authentication>

upvoted 1 times

 **Rukman** 7 months, 3 weeks ago

Selected Answer: B

Ans: B

upvoted 1 times

 **KunK** 1 year, 1 month ago

B, really straightforward

upvoted 5 times

 **ravixkumar** 1 year, 8 months ago

Ans : B

gcloud auth activate-service-account --key-file=/test-service-account.json

upvoted 10 times

Question #171

You are working with a Cloud SQL MySQL database at your company. You need to retain a month-end copy of the database for three years for audit purposes.

What should you do?

- A. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.
- B. Save the automatic first-of-the-month backup for three years. Store the backup file in an Archive class Cloud Storage bucket.
- C. Set up an on-demand backup for the first of the month. Write the backup to an Archive class Cloud Storage bucket.
- D. Convert the automatic first-of-the-month backup to an export file. Write the export file to a Coldline class Cloud Storage bucket.

Correct Answer: B

Community vote distribution

A (100%)

✉️  **TAvenger** Highly Voted 1 year, 7 months ago

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

not B: Automatic backups are made EVERY SINGLE DAY. You can set only the number of backups up to 365. Also you cannot choose your Archival storage as destination

not C: You cannot setup "on-demand" backup. User would have to make backups manually every month. Also you cannot choose your Archival storage as destination

not D: You cannot convert backup to export file. Also Coldline class is less cost-effective than Archival class.

The only option left is "A"

You can set up your job with any date/time schedule. You can export file to any storage with any storage class.

upvoted 38 times

✉️  **djgodzilla** 1 year, 4 months ago

from the same link :

Can I export a backup?

No, you can't export a backup. You can only export instance data. See Exporting data from Cloud SQL to a dump in Cloud storage.

upvoted 4 times

✉️  **JieHeng** Highly Voted 1 year, 4 months ago

First need to understand backup vs export, two different concepts. - <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

A – yes, you can export data from Cloud SQL to Cloud Storage- <https://cloud.google.com/sql/docs/mysql/import-export/exporting#cloud-sql>

Not B, C, D – be it automatic or on-demand backup, according to the doc “No, you can't export a backup. You can only export instance data.”

upvoted 14 times

✉️  **Charumathi** Most Recent 1 month ago

Selected Answer: A

A is correct answer,

Export the SQL month end data as a CSV file to cloud storage bucket, and move the data to Archival Storages for 3 years for audit purpose.

<https://cloud.google.com/sql/docs/mysql/import-export?authuser=1>

upvoted 1 times

✉️  **learn_GCP** 1 month, 1 week ago

Selected Answer: A

A. Although Cloud SQL doesn't provide a built-in way to automate database exports, you can build your own automation tool using several Google Cloud components.

https://cloud.google.com/sql/docs/mysql/import-export#automating_export_operations

upvoted 1 times

✉️  **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

✉️  **butki** 1 month, 2 weeks ago

B is correct

upvoted 1 times

✉️  **snkhatri** 2 months ago

A seems right to me key word: "month-end copy"

upvoted 2 times

 **patashish** 3 months, 4 weeks ago

Correct Ans - C
On demand backup

You can create a backup at any time (Here we need backup a month-end copy of the database for three years). You can create on-demand backups for any instance, whether the instance has automatic backups enabled or not.

Reason :

1) you can't export a backup. You can only export instance data so export option A is out from answer.
Backups encompass the entire database. Exports can select specific contents.

As per question You need to retain a month-end copy of the database not specific contents.

2) Automated backups are taken daily, within a 4-hour backup window. Up to seven most recent backups are retained, by default.
*Cost** to store all backups ..

3) Option D not applicable **Coldline class** *Cost*

upvoted 2 times

 **mk1471** 9 months ago

why not D? it is the only one that doesn't store it as Archive class, and since it is for Audit purposes this can't be used as Archive allows LESS than one access per year.

upvoted 2 times

 **look1** 11 months ago

Selected Answer: A

I would go with export not backup. Question stated that this is copy. In addition restore of three years old backup might be not possible in newer version. In addition you can retain max 365 backups.

upvoted 4 times

 **jackwillis** 1 year ago

As question itself states that we need month end backup, automatic backup runs daily which is required as per the question.
<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 3 times

 **vamgcp** 1 year, 1 month ago

A is incorrect because there is no need to create an export job as the export functionality is built-in with Cloud SQL - I would go with B
upvoted 4 times

 **arsh1916** 1 year, 5 months ago

B is correct

upvoted 1 times

 **jahnu** 1 year, 6 months ago

Ans: B Backups are managed by Cloud SQL according to retention policies, and are stored separately from the Cloud SQL instance. So take the backup and store in Bucket.

check: <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 2 times

 **kopper2019** 1 year, 6 months ago

Answer is A, there is no way to take the backups made by MySQL and send them to Cloud Storage so you have to make an export...
Default backup only allows 365 retention points so you need to make export I mean program exports and send them to Cloud Datastore so A.

Backups are managed by Cloud SQL according to retention policies, and are stored separately from the Cloud SQL instance. Cloud SQL backups differ from an export uploaded to Cloud Storage, where you manage the lifecycle. Backups encompass the entire database. Exports can select specific contents.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 1 times

 **GCP_Student1** 1 year, 7 months ago

A. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.

upvoted 2 times

 **victory108** 1 year, 8 months ago

B- Save the automatic first-of-the-month backup for three years. Store the backup file in an Archive class Cloud Storage bucket.
upvoted 3 times

 **victory108** 1 year, 7 months ago

A is correct one. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.
upvoted 1 times

 **gcpengineer** 1 year, 3 months ago

you can not export backup
upvoted 1 times

 **[Removed]** 8 months, 3 weeks ago

You can
<https://cloud.google.com/sql/docs/mysql/import-export/import-export-sql>

upvoted 2 times

Question #172

You are monitoring an application and receive user feedback that a specific error is spiking. You notice that the error is caused by a Service Account having insufficient permissions. You are able to solve the problem but want to be notified if the problem recurs. What should you do?

- A. In the Log Viewer, filter the logs on severity 'Error' and the name of the Service Account.
- B. Create a sink to BigQuery to export all the logs. Create a Data Studio dashboard on the exported logs.
- C. Create a custom log-based metric for the specific error to be used in an Alerting Policy.
- D. Grant Project Owner access to the Service Account.

Correct Answer: A

Reference:

<https://cloud.google.com/logging/docs/view/advanced-queries>*Community vote distribution*

C (100%)

✉  **GCP_Student1** Highly Voted 1 year, 7 months ago

C. Create a custom log-based metrics for the specific error to be used in an Alerting Policy.
upvoted 23 times

✉  **greatsam321** Highly Voted 1 year, 8 months ago

C seems to be the right answer.
upvoted 11 times

✉  **Charumathi** Most Recent 1 month ago

Selected Answer: C

C is the correct answer,
Since the problem is resolved, We need to monitor if the error recurs, hence we create a custom log based metrics to monitor only the particular service account.
upvoted 2 times

✉  **snkhatri** 2 months ago

C as Keyword "want to be notified if the problem recurs"
upvoted 1 times

✉  **AzureDP900** 4 months, 2 weeks ago

C right
upvoted 1 times

✉  **AzureDP900** 4 months, 4 weeks ago

C is correct.
upvoted 1 times

✉  **PAUGURU** 6 months, 2 weeks ago

Selected Answer: C

C - the only answer that outputs a notification
upvoted 2 times

✉  **RazOr** 9 months, 2 weeks ago

Selected Answer: C
"C" is right, the only answer that includes setting up an alert!
upvoted 2 times

✉  **Wolf13ts** 9 months, 3 weeks ago

Selected Answer: C
Answer should be C
upvoted 1 times

✉  **alaahakim** 11 months, 3 weeks ago

Ans: C
upvoted 2 times

✉  **AD_0525** 1 year, 4 months ago

You want to be alerted next time, so only option C meets that criteria.
upvoted 4 times

 **Enzo** 1 year, 5 months ago

C is correct : You are able to solve the problem but want to be notified if the problem recurs.
upvoted 4 times

 **ersh1916** 1 year, 5 months ago

A is correct
upvoted 1 times

 **kopper2019** 1 year, 6 months ago

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?

- A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.
- B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.
- C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team's internal data warehouse.
- D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

it's A

upvoted 2 times

 **mj98** 1 year, 6 months ago

I think it's D. Can anyone confirm?
upvoted 2 times

 **kopper2019** 1 year, 6 months ago

Question 178
upvoted 1 times

 **Petza** 1 year, 6 months ago

A. "Every time the dashboard is refreshed, it pulls new data from the view, which in turn dynamically reflects the latest data in BigQuery". Data Science on the Google Cloud Platform: Implementing End-to-End Real-Time (C)
upvoted 2 times

 **lxgywil** 1 year, 6 months ago

The question is saying "want to be notified if the problem recurs", I don't see how A meets that requirement.
upvoted 1 times

 **kopper2019** 1 year, 6 months ago

You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

- A. Create and deploy a Custom Resource Definition per microservice.
- B. Create and deploy a Docker Compose File.
- C. Create and deploy a Job per microservice.
- D. Create and deploy a Deployment per microservice.

I think is D

upvoted 3 times

 **AmineHM** 1 year, 6 months ago

I think B
upvoted 1 times

 **AmineHM** 1 year, 6 months ago

It's D, yes !
upvoted 1 times

 **noreen** 1 year, 6 months ago

its c; Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a YAML file to configure your application's services. Then, with a single command, you create and start all the services from your configuration.
upvoted 2 times

 **noreen** 1 year, 6 months ago

b * B. Create and deploy a Docker Compose File.
upvoted 1 times

 **kopper2019** 1 year, 6 months ago

You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.
- B. Provision Compute Engine instances with GPUs attached.
- C. Provision Compute Engine instances with local SSDs attached.
- D. Provision Compute Engine instances with M1 machine type.

Answer is D type M,

upvoted 3 times

✉  **AmineHM** 1 year, 6 months ago

it's D

upvoted 1 times

✉  **kopper2019** 1 year, 6 months ago

hey guys go Question number 1 for New Question from 173 to 179

upvoted 3 times

✉  **AmineHM** 1 year, 6 months ago

I don't find 178 and question for 179

upvoted 1 times

Question #173

You are developing a financial trading application that will be used globally. Data is stored and queried using a relational structure, and clients from all over the world should get the exact identical state of the data. The application will be deployed in multiple regions to provide the lowest latency to end users. You need to select a storage option for the application data while minimizing latency. What should you do?

- A. Use Cloud Bigtable for data storage.
- B. Use Cloud SQL for data storage.
- C. Use Cloud Spanner for data storage.
- D. Use Firestore for data storage.

Correct Answer: C

Reference:

<https://cloud.google.com/solutions/best-practices-compute-engine-region-selection>

Community vote distribution

C (86%)	14%
---------	-----

 **JieHeng** Highly Voted 1 year, 4 months ago

C, Cloud Spanner, keywords are globally, relational structure and lastly "clients from all over the world should get the exact identical state of the data" which implies strong consistency is needed.

upvoted 12 times

 **Charumathi** Most Recent 1 month ago

Selected Answer: C

C is the correct answer,

Keywords, Financial data (large data) used globally, data stored and queried using relational structure (SQL), clients should get exact identical copies(Strong Consistency), Multiple region, low latency to end user, select storage option to minimize latency.

upvoted 1 times

 **Charumathi** 1 month ago

Spanner powers business-critical applications in retail, financial services, gaming, media and entertainment, technology, healthcare and more.

Use cases for Cloud Spanner

<https://www.youtube.com/watch?v=1b4flZwAQfM&t=1s>

upvoted 1 times

 **ale_brd_** 1 month, 1 week ago

Selected Answer: C

it's C 100%

Guys come on, it's a pretty straight forward scenario.

if you have the keywords "relational DB" and the word "Globally" in a sentence always go for Cloud Spanner.

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

Selected Answer: C

C. is the answer

upvoted 1 times

 **sri333** 1 month, 1 week ago

Why not A. Big table as per keywords relational, global and low latency

upvoted 1 times

 **zellick** 1 month, 3 weeks ago

Selected Answer: C

C is the answer.

Cloud Spanner is a global relational database.

upvoted 1 times

 **Mr_MIXER007** 2 months ago

Selected Answer: C

Should be the correct answer.

upvoted 1 times

 **GayuSundar** 2 months, 1 week ago

C, Cloud Spanner. Globally and trading (tend to receive 1000s Records per second) are key here.

upvoted 1 times

 **VietmanOfficiel** 2 months, 2 weeks ago

Selected Answer: B

Why not Cloud SQL with replicas in multiple region to serve "global" ?

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

This is straight forward question, Answer is C.

upvoted 1 times

 **PAUGURU** 6 months, 2 weeks ago

Selected Answer: C

C is correct - Relational and Global

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

Choose - C

upvoted 2 times

 **associatecloudexamuser** 1 year, 3 months ago

Yes. C is right answer

upvoted 3 times

 **Fidget_** 1 year, 6 months ago

C) Spanner - Global, low latency, relational

upvoted 4 times

 **Rightsaidfred** 1 year, 6 months ago

Yes C, Cloud Spanner

upvoted 2 times

 **norrec9** 1 year, 6 months ago

C is the answer

upvoted 2 times

 **Biju1** 1 year, 6 months ago

C Cloud Spanner to store relational data

upvoted 2 times

Question #174

You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.
- B. Provision Compute Engine instances with GPUs attached.
- C. Provision Compute Engine instances with local SSDs attached.
- D. Provision Compute Engine instances with M1 machine type.

Correct Answer: C

Reference:

<https://cloud.google.com/compute/docs/disks/local-ssd>*Community vote distribution*

D (100%)

✉  **Rightsaidfred** Highly Voted 1 year, 6 months ago
Yes D, M1 Machine types for ERP i.e. SAP-HANA:
<https://cloud.google.com/compute/docs/machine-types>
upvoted 26 times

✉  **Charumathi** Most Recent 1 month ago

Selected Answer: D

D is the correct answer,
M1 machine series
Medium in-memory databases such as SAP HANA
Tasks that require intensive use of memory with higher memory-to-vCPU ratios than the general-purpose high-memory machine types.
In-memory databases and in-memory analytics, business warehousing (BW) workloads, genomics analysis, SQL analysis services.
Microsoft SQL Server and similar databases.
upvoted 1 times

✉  **snkhatri** 2 months ago
D, keyword "Full database in-memory"
upvoted 1 times

✉  **ryumada** 3 months ago

Selected Answer: D

Vote for D as the right answer. M1 machine type is the one of two Memory-Optimized machine types in GCP.

<https://cloud.google.com/compute/docs/machine-types>
upvoted 1 times

✉  **ryumada** 3 months ago
Read this also to see the difference of the two.

<https://cloud.google.com/compute/docs/memory-optimized-machines>
upvoted 1 times

✉  **abirroy** 3 months, 1 week ago

Selected Answer: D

D: M1 Machine types for ERP i.e. SAP-HANA

Medium-large in-memory databases such as SAP HANA
In-memory databases and in-memory analytics
Microsoft SQL Server and similar databases
upvoted 2 times

✉  **AzureDP900** 4 months, 4 weeks ago

D is right choice, when answer selected by author is C doesn't make any sense. User also need to understand the services well before attempting exam.

upvoted 1 times

✉  **akshaychavan7** 5 months, 2 weeks ago

I chose option D, because the first three didn't make any sense :D
upvoted 1 times

✉ **RazOr** 9 months, 2 weeks ago

D!!!

The "M1" VM type is right, it offers between 1.4TB and 3.75TB of RAM.

upvoted 3 times

✉ **[Removed]** 10 months, 2 weeks ago

<https://cloud.google.com/compute/docs/machine-types#:~:text=databases%20such%20as-,SAP%20HANA,-In%2Dmemory%20databases>

<https://www.sap.com/india/products/hana.html#:~:text=is%20SAP%20HANA-,in%2Dmemory,-database%3F>

upvoted 1 times

✉ **jaffarali** 11 months ago

Selected Answer: D

Correct Answer is D

upvoted 1 times

✉ **alaahakim** 11 months, 3 weeks ago

D is the Answer

upvoted 1 times

✉ **maggieli** 1 year ago

"m1-megamem-96" can attach local SSD. Correct is D.

upvoted 1 times

✉ **jackdbd** 1 year, 3 months ago

Note that VM instances m1-megamem-96 are both from the M1 family AND can have local SSDs attached to them.

https://cloud.google.com/compute/docs/memory-optimized-machines#m1_vms

upvoted 2 times

✉ **associatecloudexamuser** 1 year, 3 months ago

Answer is D.

Applications of Memory optimized VMs are,

1. Medium-large in-memory databases such as SAP HANA
2. In-memory databases and in-memory analytics
3. Microsoft SQL Server and similar databases

upvoted 4 times

✉ **JieHeng** 1 year, 4 months ago

D, "The application holds the full database in-memory for fast data access", so it'll be more appropriate to use memory-optimized machine types

- https://cloud.google.com/compute/docs/machine-types#m1_machine_types

upvoted 4 times

✉ **AD_0525** 1 year, 4 months ago

Here the question is about data access not work load type, hence local SSDs are best option. So it is C

upvoted 2 times

✉ **csrazdan** 4 months ago

Local SSD is used to improve IO. A memory store or GCE with a larger memory footprint will help with fast data access. D is the correct answer.

upvoted 2 times

✉ **sanhoo** 1 year, 5 months ago

C: All the options use compute engine. A is incorrect because it uses preemptible instances, which only runs for max 24 hours. Hence A is not an option. Then look at the key "fast data access". Local SSD is the fastest

upvoted 2 times

✉ **csrazdan** 4 months ago

Local SSD has faster IO

upvoted 1 times

Question #175

You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

- A. Create and deploy a Custom Resource Definition per microservice.
- B. Create and deploy a Docker Compose File.
- C. Create and deploy a Job per microservice.
- D. Create and deploy a Deployment per microservice.

Correct Answer: D*Community vote distribution*

D (100%)

 **Kollipara** Highly Voted 1 year, 6 months ago

D is the correct answer

upvoted 17 times

 **obeythefist** Highly Voted 8 months, 1 week ago

I was a little unsure about this question, here's how I understand why D is the best answer

- A. Custom Resource Definition... we have docker containers already, which is an established kind of resource for Kubernetes. We don't need to create a whole new type of resource, so this is wrong.
- B. Docker Compose is a wholly different tool from Kubernetes.
- C. A Kubernetes job describes a specific "task" which involves a bunch of pods and things. It makes no sense to have one job per microservice, a "Job" would be a bunch of different microservices executing together.
- D. is the leftover, correct answer. You can add scaling to each Deployment, an important aspect of the question.

upvoted 9 times

 **akshaychavan7** 5 months, 2 weeks ago

Thanks for your insights! Makes sense.

upvoted 2 times

 **snkhatri** Most Recent 2 months ago

D, keyword "each microservice can be scaled individually"!

upvoted 1 times

 **abirroy** 3 months, 1 week ago

Selected Answer: D

D is the correct answer

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

D is the best answer among other choices.

upvoted 1 times

 **RazOr** 9 months, 2 weeks ago

D is right!

It's one of Googles main ideas to distribute a complex system into microservices. They do it as well and encourage customers to do the same.

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Selected Answer: D

D is the Answer

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

Ans: D

upvoted 2 times

 **associatecloudexamuser** 1 year, 3 months ago

Yes. D is correct. Can deploy each service through
kubectl apply -f <deployment_config.yaml>

upvoted 5 times

 **pacman_user** 1 year, 3 months ago

D is the correct answer

upvoted 1 times

 **AD_0525** 1 year, 4 months ago

D is correct

upvoted 2 times

 **norrec9** 1 year, 6 months ago

D is the answer

upvoted 6 times

Question #176

You will have several applications running on different Compute Engine instances in the same project. You want to specify at a more granular level the service account each instance uses when calling Google Cloud APIs. What should you do?

- A. When creating the instances, specify a Service Account for each instance.
- B. When creating the instances, assign the name of each Service Account as instance metadata.
- C. After starting the instances, use gcloud compute instances update to specify a Service Account for each instance.
- D. After starting the instances, use gcloud compute instances update to assign the name of the relevant Service Account as instance metadata.

Correct Answer: C*Community vote distribution*

A (100%)

 **GoCloud** Highly Voted 1 year, 6 months ago

A .

upvoted 22 times

 **JieHeng** Highly Voted 1 year, 4 months ago

A, when you create an instance using the gcloud command-line tool or the Google Cloud Console, you can specify which service account the instance uses when calling Google Cloud APIs - https://cloud.google.com/compute/docs/access/service-accounts#associating_a_service_account_to_an_instance

upvoted 14 times

 **ryumada** Most Recent 3 months ago

Selected Answer: A

Vote for A, because there is no instance running yet. "You will have several applications running..."

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: A

A, there is no instance running yet

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

A is good option for given scenario.

upvoted 1 times

 **somenick** 7 months, 2 weeks ago

Selected Answer: A

You can set/update the service account only when the instance is not running

upvoted 2 times

 **Majkl93** 8 months, 2 weeks ago

Selected Answer: A

A - the instances are not running yet

upvoted 1 times

 **Razor** 9 months, 2 weeks ago

A: you can define which GCP service account is associated with a Compute Engine instance when creating one. It is still possible to change the service account later.

Link to the GCP docs: <https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#using>

upvoted 1 times

 **liyux21** 1 year ago

A is correct. You can change the assigned service account, use gcloud compute instances set-service-account, not the update

upvoted 3 times

 **techabhi2_0** 1 year, 1 month ago

Answer is C. The instances are already running. So you need to change the Service account

upvoted 2 times

 **techabhi2_0** 1 year, 1 month ago

My bad - "you will have" . Correct answer - A

upvoted 8 times

 **AD_0525** 1 year, 4 months ago

A should be correct

upvoted 4 times

 **norrec9** 1 year, 6 months ago

A is the answer

upvoted 4 times

 **Biju1** 1 year, 6 months ago

It should be A

upvoted 3 times

Question #177

You are creating an application that will run on Google Kubernetes Engine. You have identified MongoDB as the most suitable database system for your application and want to deploy a managed MongoDB environment that provides a support SLA. What should you do?

- A. Create a Cloud Bigtable cluster, and use the HBase API.
- B. Deploy MongoDB Atlas from the Google Cloud Marketplace.
- C. Download a MongoDB installation package, and run it on Compute Engine instances.
- D. Download a MongoDB installation package, and run it on a Managed Instance Group.

Correct Answer: C*Community vote distribution*

B (83%)

C (17%)

✉ **arsh1916** Highly Voted 1 year, 5 months ago

Simple it's B

upvoted 13 times

✉ **lxgywil** Highly Voted 1 year, 6 months ago

MongoDB Atlas is actually managed and supported by third-party service providers.

<https://console.cloud.google.com/marketplace/details/gc-launcher-for-mongodb-atlas/mongodb-atlas>

upvoted 9 times

✉ **lxgywil** 1 year, 6 months ago

I think that's it. The answer is B

upvoted 6 times

✉ **PKookNN** Most Recent 2 weeks, 1 day ago

Selected Answer: B

the best answer is B

upvoted 1 times

✉ **nonsense** 2 weeks, 2 days ago

Selected Answer: B

b. fast and simple

upvoted 1 times

✉ **11kc03** 2 weeks, 2 days ago

Selected Answer: C

Answer is B

upvoted 1 times

✉ **learn_GCP** 1 month, 1 week ago

Selected Answer: B

B. is the answer

upvoted 1 times

✉ **snkhatri** 2 months ago

B, keyword " support SLA"

upvoted 1 times

✉ **AzureDP900** 4 months, 4 weeks ago

<https://cloud.google.com/mongodb>

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

B is right

upvoted 1 times

✉ **AzureDP900** 4 months, 4 weeks ago

B is correct

upvoted 1 times

✉ **rsuresh27** 6 months, 2 weeks ago

Anytime the question mentions a third party software, always use the cloud marketplace. Answer is B.

upvoted 4 times

 **dishum** 9 months, 3 weeks ago

the question says "want to deploy a managed MongoDB environment" which means it should be managed by something, i.e compute engine or MIG.

Near Ans is C or D

I choose C - becoz no need of Mongo db running on MIG, GKE can easily handle mongoDb on compute engine.

upvoted 1 times

 **fazalmf** 10 months, 3 weeks ago

Deploy MongoDB Atlas its free Tier - does free tier provides a Support SLA?

upvoted 1 times

 **jaffarali** 11 months ago

Selected Answer: B

B is the right option to use managed services.

upvoted 2 times

 **alaahakim** 11 months, 3 weeks ago

I Pick B

upvoted 1 times

 **AD_0525** 1 year, 4 months ago

Answer should be B

upvoted 2 times

 **GoCloud** 1 year, 6 months ago

B .

upvoted 1 times

 **norrec9** 1 year, 6 months ago

Answer is b

upvoted 2 times

Question #178

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?

- A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.
- B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.
- C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team's internal data warehouse.
- D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

Correct Answer: C

Community vote distribution

D (100%)

 **ApaMokus** Highly Voted 1 year, 6 months ago

D is correct

roles/bigquery.user

When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.

When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

upvoted 26 times

 **blan_ak** Highly Voted 1 year, 2 months ago

Why on the earth would the answer be C? It has no relevance to the question. The answer is D, hands down

upvoted 6 times

 **anolive** Most Recent 6 days, 22 hours ago

Selected Answer: D

makes mor sense

upvoted 1 times

 **sylva91** 1 month, 3 weeks ago

Selected Answer: D

D is correct because google recommendations are always to privilege groups to individual accounts and this is what can make the users query the database unlike the Data Studio

upvoted 1 times

 **snkhatri** 2 months ago

Selected Answer: D

D is right

upvoted 1 times

 **snkhatri** 2 months ago

D is correct

upvoted 1 times

 **patashish** 3 months, 4 weeks ago

D is the answer

Hint - to **run the custom SQL queries** against the latest data in BigQuery

upvoted 1 times

 **TaniaMalfoy** 4 months ago

C is correct, data pipeline is the key:

https://cloud.google.com/dataflow/docs/guides/data-pipelines#create_a_batch_data_pipeline

To create this sample batch data pipeline, you must have access to the following resources in your project:

A Cloud Storage bucket to store input and output files

A BigQuery dataset where you will create a table.

upvoted 1 times

 **S00999** 4 months, 3 weeks ago

Selected Answer: D

Answer: D

The simplest.

It is not requested to automate the query. The BI team may also need to modify their query or have several different ones to meet the needs.

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

D is correct

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

 **RazOr** 9 months, 2 weeks ago

Selected Answer: D

D sounds perfect with minimal steps.

Quote from the GCP docs: "BigQuery User
(roles/bigquery.user)

When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.

When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A principal with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets."

upvoted 1 times

 **dishum** 9 months, 3 weeks ago

Answer is C

Option C says, there is a probability of an internal BI datawarehouse. Before providing the iam permissions, it is better to copy data to internal BI.

My view

upvoted 1 times

 **arvsrv** 9 months, 3 weeks ago

Selected Answer: D

agree With D

upvoted 2 times

 **gioresin1** 10 months ago

maybe D is not sufficient because, as per documentation:

"Note: For a user to be able to query the tables in a dataset, it is not sufficient for the user to have access to the dataset. A user must also have permission to run a query job in a project. If you want to give a user permission to run a query from your project, give the user the bigquery.jobs.create permission for the project. You can do this by assigning the user the roles/bigquery.jobUser role for your project. For more information, see Access control examples".

upvoted 1 times

 **alaahakim** 11 months, 3 weeks ago

I agree With D

upvoted 2 times

 **AD_0525** 1 year, 4 months ago

D should be the answer

upvoted 4 times

 **victory108** 1 year, 6 months ago

D - Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

upvoted 5 times

Question #179

Your company is moving its entire workload to Compute Engine. Some servers should be accessible through the Internet, and other servers should only be accessible over the internal network. All servers need to be able to talk to each other over specific ports and protocols. The current on-premises network relies on a demilitarized zone (DMZ) for the public servers and a Local Area Network (LAN) for the private servers. You need to design the networking infrastructure on

Google Cloud to match these requirements. What should you do?

- A. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- B. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.
- C. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- D. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.

Correct Answer: A*Community vote distribution*

A (100%)

 **perdigiorno** Highly Voted  1 year, 4 months ago

Passed the test today. About 80% of the questions are here.

upvoted 21 times

 **associatecloudexamuser** 1 year, 3 months ago

Congratulations!

upvoted 3 times

 **sumantrao1** 1 year ago

you got same questions from this examtopics

upvoted 2 times

 **yvinisiupacuando** Highly Voted  1 year, 6 months ago

A is the Right answer. You can discard B and C because they lack the need of creating Network Peering to communicate the DMZ VPC with the LAN VPC (LAN VPC is not exposed to public so they need to communicate via private addresses which cannot be achieved with 2 VPCs without Network Peering). Plus, you can discard B, as you don't need to enable the egress traffic, you always need to enable the ingress traffic as this is never enabled by default.

upvoted 21 times

 **Alela** 1 year, 5 months ago

A is wrong. You don't need to set up firewall rules between subnets of the same VPC. C is the answer

upvoted 11 times

 **gcpengineer** 1 year, 3 months ago

You need fw rules

upvoted 1 times

 **demnok_lannik** 9 months, 1 week ago

of course you do

upvoted 1 times

 **Ashii** 1 year, 5 months ago

C is Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ. Without peering 2 VPC's how this this be done ?

upvoted 4 times

 **BenKenGo6** 2 months, 1 week ago

and where do you have the VPC peering to communicate both VPCs?

upvoted 1 times

 **diasporabro** Most Recent  2 weeks, 5 days ago

Selected Answer: A

A is the right choice

upvoted 1 times

✉ **snkhatri** 2 months ago

Selected Answer: A

A seems right

upvoted 1 times

✉ **an0nym0us1** 2 months, 2 weeks ago

hi All what is the ans

upvoted 1 times

✉ **AzureDP900** 4 months, 2 weeks ago

1 VPC enough for LAN and DMZ , Need to open appropriate firewall rules. A is right.

upvoted 1 times

✉ **S00999** 4 months, 3 weeks ago

Selected Answer: A

Vote for A

By default traffic between subnets on a VPC network is not allowed (except on the "default" network).

(This blocks traffic between all instances, not just traffic between subnets => FW rules must be defined to allow communications between all instances, regardless the subnets)

2 VPC will not work without peering.

upvoted 2 times

✉ **somenick** 7 months, 2 weeks ago

Selected Answer: A

You can't explicitly create a FW rule for the subnet, but connections are allowed or denied on a per-instance basis. You can think of the VPC firewall rules as existing not only between your instances and other networks, but also between individual instances within the same network. C will not work without peering...

upvoted 2 times

✉ **akshaym87** 9 months, 2 weeks ago

Guys i cleared my exam last week. This question bank is must. 80% questions were from here.

upvoted 5 times

✉ **Boumer** 3 months ago

congratulation, i need to pass exam in end month .Can you give me some advise please?

upvoted 1 times

✉ **HansKloss611** 9 months, 2 weeks ago

Selected Answer: A

A - my vote. Two different vpc need vpc peering.

upvoted 2 times

✉ **gioresin1** 10 months ago

I don't understand why you say that the answer is A. If you have 2 subnets in the same network you won't have firewall between the 2 subnets. So you can't have a DMZ that can communicate with a private network. So the answer should be C.

upvoted 1 times

✉ **GCPdump** 1 year, 1 month ago

A is correct

upvoted 1 times

✉ **maan2935** 1 year, 1 month ago

Hi, I have an exam today. Are the questions still to some extent valid?

upvoted 3 times

✉ **jackwillis** 1 year ago

Yes it is. I passed my exam on 3rd Oct received certificate on 7th Oct. Exam topics and a study course in Udemy which I bought during their discount sale, helped me.

upvoted 1 times

✉ **gerhardbl** 1 year, 2 months ago

Textbook example of DMZ and private subnet topology, hence answer A. Anyone who thinks C or multiple VPCs or whatever I strongly suggest you do CCNA before coming here.

upvoted 4 times

✉ **associatecloudexamuser** 1 year, 3 months ago

Yes. Correct answer is A. No need to complicate the setup by creating two different VPC networks.

upvoted 2 times

✉ **shayanahmed** 1 year, 4 months ago

All questions are still valid. I cleared my paper yesterday (shayan18@live.com)

upvoted 5 times

associatecloudexamuser 1 year, 3 months ago

Congratulations!
upvoted 1 times

JieHeng 1 year, 4 months ago

Should be A.
Not B, allow public egress traffic for DMZ won't help anything (also by default there is already this allow public egress traffic rule)
Not C & D, Network Peering is needed to allow internal IP address connectivity across two Virtual Private Cloud (VPC)
<https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 5 times

AD_0525 1 year, 4 months ago

C should be the correct answer. Once you put both set of VMs in a single VPC, either all will be exposed to external traffic or none will be exposed.

upvoted 2 times

Question #180

You have just created a new project which will be used to deploy a globally distributed application. You will use Cloud Spanner for data storage. You want to create a Cloud Spanner instance. You want to perform the first step in preparation of creating the instance. What should you do?

- A. Enable the Cloud Spanner API.
- B. Configure your Cloud Spanner instance to be multi-regional.
- C. Create a new VPC network with subnetworks in all desired regions.
- D. Grant yourself the IAM role of Cloud Spanner Admin.

Correct Answer: B

Reference:

<https://cloud.google.com/spanner/docs/instances>*Community vote distribution*

A (71%)	14%	14%
---------	-----	-----

 **AzureDP900** Highly Voted  4 months, 4 weeks ago

A is right

<https://cloud.google.com/spanner/docs/getting-started/set-up>

upvoted 5 times

 **snkhatri** Most Recent  2 months ago

Selected Answer: A

A seems right

upvoted 1 times

 **bobthebuilder55110** 3 months ago

Selected Answer: B

Answer must be B , here is why?

I was confused between A and B but I tested this by creating a new project, when you go to spanner and click on create a spanner instance it automatically enables the API for you and you can all see this activity on the notification panel on the top right along with the visibility of this message clearly on that instance page as well, First it will auto-enable the API and then It will give you an option to select multi-region, Now the questions say your first step that has to be multi-region since enabling the API was done by google automatically and none other options makes sense here.

upvoted 2 times

 **bobthebuilder55110** 3 months ago

Correction : It's A

Since it does not specify if we are using command line tool or UI, if you are using command line tool then you will have to enable this.

upvoted 2 times

 **mani098** 4 months, 2 weeks ago

Selected Answer: C

because Api auto enabled when you click create new instance on cloud spanner UI

upvoted 2 times

 **mav3r1ck** 3 months ago

This is TRUE.

Verified this:

- Go to your GCP project and verify that "Cloud Spanner API" is NOT enabled.
- Go to Cloud Spanner. Click "Create Instance".
- Check back again the "Cloud Spanner API" , you will see that status is "API Enabled".

upvoted 2 times

 **pfabio** 5 months, 1 week ago

Selected Answer: A

If you click on Create instance, the message is show in bottom: Cloud Spanner API for your project has been enabled.

upvoted 3 times

 **Rutu_98** 5 months, 2 weeks ago

Selected Answer: A

Answer : A --> Tested

upvoted 3 times

 **bobthebuilder55110** 3 months ago

How ? I tested this as well and It auto enables the API. Unless this is a new feature by google and the question is old then I am not sure.

upvoted 1 times

 **adarsh4503** 1 month, 2 weeks ago

If you used CLI then the API would have to be enabled manually.

upvoted 1 times

 **Random_Mane** 6 months, 1 week ago

Selected Answer: A

Enabling API is the first step

upvoted 2 times

 **jblima** 6 months, 1 week ago

Selected Answer: A

A - I tested...

upvoted 1 times

 **parag09** 6 months, 2 weeks ago

A:-tested it need to Enable Api first

upvoted 1 times

 **sivasan** 6 months, 2 weeks ago

shoudnt the ans be enable cloud spanner API.Option A

upvoted 1 times

 **PAUGURU** 6 months, 2 weeks ago

A before anything else

upvoted 2 times

Question #181

You have created a new project in Google Cloud through the gcloud command line interface (CLI) and linked a billing account. You need to create a new Compute Engine instance using the CLI. You need to perform the prerequisite steps. What should you do?

- A. Create a Cloud Monitoring Workspace.
- B. Create a VPC network in the project.
- C. Enable the compute googleapis.com API.
- D. Grant yourself the IAM role of Computer Admin.

Correct Answer: D

Reference:

<https://isb-cancer-genomics-cloud.readthedocs.io/en/latest/sections/gcp-info/gcp-info2/LaunchVM.html>*Community vote distribution*

C (100%)

✉  **sylva91** 1 month, 3 weeks ago

Selected Answer: C

nothing can be done before activating the API

upvoted 1 times

✉  **zelick** 1 month, 3 weeks ago

Selected Answer: C

C is the obvious answer.

upvoted 1 times

✉  **III_bbb** 2 months ago

Selected Answer: C

api first

upvoted 2 times

✉  **snkhatri** 2 months ago

Selected Answer: C

C the compute googleapis.com API

upvoted 1 times

✉  **Nishanth222** 2 months, 1 week ago

Must be C

upvoted 1 times

Question #182

Your company has developed a new application that consists of multiple microservices. You want to deploy the application to Google Kubernetes Engine (GKE), and you want to ensure that the cluster can scale as more applications are deployed in the future. You want to avoid manual intervention when each new application is deployed. What should you do?

- A. Deploy the application on GKE, and add a HorizontalPodAutoscaler to the deployment.
- B. Deploy the application on GKE, and add a VerticalPodAutoscaler to the deployment.
- C. Create a GKE cluster with autoscaling enabled on the node pool. Set a minimum and maximum for the size of the node pool.
- D. Create a separate node pool for each application, and deploy each application to its dedicated node pool.

Correct Answer: C*Community vote distribution*

C (88%) 13%

✉ **diasporabro** 2 weeks, 5 days ago

Selected Answer: C

C is the right choice... See this for reference <https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler>
upvoted 1 times

✉ **Charumathi** 1 month ago

Selected Answer: C

C is the correct answer, you can enable the cluster autoscaling in node pool by specifying min and max node size.

https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-autoscaler#adding_a_node_pool_with_autoscaling
upvoted 1 times

✉ **ale_brd_** 1 month, 1 week ago

Selected Answer: C

it's mentioning "the cluster can scale" the answer is C

upvoted 1 times

✉ **hanskristian** 1 month, 3 weeks ago

I Think it should be C

upvoted 1 times

✉ **zellck** 1 month, 3 weeks ago

Selected Answer: C

C is my answer. You need to scale the cluster, not the pod.

upvoted 1 times

✉ **manjtrade2** 1 month, 3 weeks ago

Selected Answer: A

A

<https://cloud.google.com/kubernetes-engine/docs/concepts/horizontalpodautoscaler>

upvoted 1 times

✉ **manjtrade2** 1 month, 3 weeks ago

ITs A

<https://cloud.google.com/kubernetes-engine/docs/concepts/horizontalpodautoscaler>

upvoted 1 times

✉ **snkhatri** 2 months ago

Selected Answer: C

C looks fine to me

upvoted 3 times

Question #183

You need to manage a third-party application that will run on a Compute Engine instance. Other Compute Engine instances are already running with default configuration. Application installation files are hosted on Cloud Storage. You need to access these files from the new instance without allowing other virtual machines (VMs) to access these files. What should you do?

- A. Create the instance with the default Compute Engine service account. Grant the service account permissions on Cloud Storage.
- B. Create the instance with the default Compute Engine service account. Add metadata to the objects on Cloud Storage that matches the metadata on the new instance.
- C. Create a new service account and assign this service account to the new instance. Grant the service account permissions on Cloud Storage.
- D. Create a new service account and assign this service account to the new instance. Add metadata to the objects on Cloud Storage that matches the metadata on the new instance.

Correct Answer: A

Reference:

<https://cloud.google.com/compute/docs/access/service-accounts>*Community vote distribution*

C (100%)

 **gcpviewer** 1 month, 1 week ago

Selected Answer: C

C is the clear choice. Want to create a new service account instead of using the default and grant it permissions in cloud storage. Straightforward C.

upvoted 1 times

 **manjtrade2** 1 month, 3 weeks ago

Selected Answer: C

C is right

upvoted 1 times

 **snkhatri** 2 months ago

Selected Answer: C

C seems right to me

upvoted 1 times

 **VietmanOfficiel** 2 months, 2 weeks ago

Selected Answer: C

"without allowing other instances" , the other instances are created with default compute engine service account. So you must create a new independant service account

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

C

<https://cloud.google.com/iam/docs/best-practices-for-using-and-managing-service-accounts>

If an application uses third-party or custom identities and needs to access a resource, such as a BigQuery dataset or a Cloud Storage bucket, it must perform a transition between principals. Because Google Cloud APIs don't recognize third-party or custom identities, the application can't propagate the end-user's identity to BigQuery or Cloud Storage. Instead, the application has to perform the access by using a different Google identity.

upvoted 1 times

 **KRIV_1** 5 months, 2 weeks ago

Although C is the correct answer notice that, as Google recommend, you first need to grant the service account the required permission before attach it to a resource.

upvoted 1 times

 **JelloMan** 6 months, 1 week ago

Selected Answer: C

C all the way. Restricts access to other VMs since they won't have the new service account you have associated with your new VM

upvoted 4 times

 **amindbesideitself** 6 months, 1 week ago

Selected Answer: C

C, other VMs will run as default service account.

upvoted 2 times

 **Akash7** 6 months, 2 weeks ago

C is correct as the other vms have default service accounts.

upvoted 2 times

 **PAUGURU** 6 months, 2 weeks ago

Selected Answer: C

C, using Default account makes the storage visible to other machines

upvoted 2 times

Question #184

You need to configure optimal data storage for files stored in Cloud Storage for minimal cost. The files are used in a mission-critical analytics pipeline that is used continually. The users are in Boston, MA (United States). What should you do?

- A. Configure regional storage for the region closest to the users. Configure a Nearline storage class.
- B. Configure regional storage for the region closest to the users. Configure a Standard storage class.
- C. Configure dual-regional storage for the dual region closest to the users. Configure a Nearline storage class.
- D. Configure dual-regional storage for the dual region closest to the users. Configure a Standard storage class.

Correct Answer: B*Community vote distribution*

D (55%) B (45%)

 **akshaychavan7** Highly Voted  5 months, 2 weeks ago

Selected Answer: D

Mission critical is the keyword here which specifies that we need to have a multi-regional backup of the data to survive any regional failures. So option D is the correct choice here.

upvoted 13 times

 **mav3r1ck** 3 months ago

Keywords: minimal cost and mission-critical

Looks like people are just looking to be on the cost side. You need to meet both.

In this case, it needs to be "dual-region". This is much cheaper than storage in "multi-region" which is obviously not in the choices.

upvoted 1 times

 **Aninina** 1 week, 2 days ago

Dual region is expensive than multi-region. (Also mentioned in the documentation: <https://cloud.google.com/storage/docs/locations>) When we set objects to be multi-regional, we get to decide/shuffle the data around at will to meet our storage needs. When you take that control away from us, it reduces the flexibility of our systems, making it more expensive to operate.

upvoted 1 times

 **ryumada** 3 months ago

At the first point in this documentation says that dual-regional storage is used for business continuity and disaster recovery. Disaster can affect to a regional architecture. I think it's make sense to use dual-regional storage for this case. Also, dual-regional storage is cheaper than multi-regional.

<https://cloud.google.com/storage/docs/dual-regions#use-dual-region-storage>

upvoted 1 times

 **JelloMan** Highly Voted  6 months, 1 week ago

Selected Answer: B

Continuous access to data means Standard since all of the other options are for infrequently accessed storage (Nearline, Coldline, Archive). Since no other regions are mentioned, single region is best in this case

upvoted 12 times

 **KRIV_1** 5 months, 2 weeks ago

And beacuse single region is "costly-effective".

upvoted 6 times

 **AMIGOY** Most Recent  1 week ago

Selected Answer: B

The data is mission critical but it is still redundant across at least 2 zones with single region

upvoted 1 times

 **hiromi** 1 week, 1 day ago

Selected Answer: D

D is right.

Keywords:

- continually -> Standard
- mission-critical analytics -> dual-regional

upvoted 1 times

 **Aninina** 1 week, 2 days ago

It's B:

IMPORTANT NOTE: Dual region is expensive than multi-region. (Also mentioned in the documentation:)

[https://cloud.google.com/storage/docs/locations\)](https://cloud.google.com/storage/docs/locations)

When we set objects to be multi-regional, we get to decide/shuffle the data around at will to meet our storage needs. When you take that control away from us, it reduces the flexibility of our systems, making it more expensive to operate.

upvoted 1 times

✉ **PKookNN** 2 weeks, 1 day ago

Selected Answer: D

I was thinking B until I saw 'mission critical' + cost optimise + constant access = D is the best answer

upvoted 1 times

✉ **nonsense** 3 weeks ago

Selected Answer: D

d is correctly

upvoted 1 times

✉ **DualCore573** 1 month ago

Selected Answer: D

Mission critical is the keyword so D

upvoted 1 times

✉ **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

✉ **snkhatri** 2 months ago

Selected Answer: B

B, as there will be multiple zone and standard

upvoted 1 times

✉ **VietmanOfficiel** 2 months, 2 weeks ago

Selected Answer: B

Region contains multi zones so if one zone fail other zones have replicas

upvoted 2 times

✉ **its_another_samsun** 2 months, 4 weeks ago

Selected Answer: D

Key words -- "minimal cost", "mission-critical", "continually"

if u go with B, 2nd condition is not met & also if a region fails 3rd condition also fails

if u go with D, only thing is there will be slight difference in cost

upvoted 1 times

✉ **bobthebuilder55110** 3 months ago

Selected Answer: B

The data is mission critical but it is still redundant across at least 2 zones with single region storage, so this is sufficient for redundancy , the questions clearly mentions about lowering the cost as well, so B.

<https://cloud.google.com/storage/docs/locations>

upvoted 1 times

✉ **bobthebuilder55110** 3 months ago

Correction: D, since that same link talk about why dual-region is great for analytics and that's what we are doing in this questions along with it says "You need to configure optimal data storage for files stored in Cloud Storage for minimal cost" so, for minimal cost(it does not say that we want to reduce the cost, it asking the lowest as possible while maintaining the business needs) but it also has to be optimal for business at the same time.

upvoted 1 times

✉ **Sans_12450** 3 months, 2 weeks ago

Selected Answer: B

Keyword: minimal cost, used continually

upvoted 1 times

✉ **nginx_aws** 3 months, 3 weeks ago

It's B see the scenario here <https://cloud.google.com/architecture/building-and-orchestrating-data-analytics-and-machine-learning-pipeline#objectives>

upvoted 1 times

✉ **csrazdan** 3 months, 3 weeks ago

Selected Answer: B

Regional buckets can tolerate zonal failures and can be used for mission-critical deployments. The actual use case for multi-regional is when you want to reduce latency and move data closer to user.

upvoted 2 times

✉ **amenur** 3 months, 4 weeks ago

Selected Answer: B

"minimal cost". "The files are used in pipeline that is used continually" ."The users are in Boston, MA (United States)".
For these reasons, i think the correct answer is: B

upvoted 1 times

Question #185

You are developing a new web application that will be deployed on Google Cloud Platform. As part of your release cycle, you want to test updates to your application on a small portion of real user traffic. The majority of the users should still be directed towards a stable version of your application. What should you do?

- A. Deploy the application on App Engine. For each update, create a new version of the same service. Configure traffic splitting to send a small percentage of traffic to the new version.
- B. Deploy the application on App Engine. For each update, create a new service. Configure traffic splitting to send a small percentage of traffic to the new service.
- C. Deploy the application on Kubernetes Engine. For a new release, update the deployment to use the new version.
- D. Deploy the application on Kubernetes Engine. For a new release, create a new deployment for the new version. Update the service to use the new deployment.

Correct Answer: B

Reference:

<https://cloud.google.com/appengine/docs/admin-api/migrating-splitting-traffic>*Community vote distribution*

A (100%)

Charumathi 1 month ago**Selected Answer: A**

A is correct answer,
Keyword, Version, traffic splitting, App Engine supports traffic splitting for versions before releasing.
upvoted 1 times

snkhatri 2 months ago**Selected Answer: A**

A obvious choice
upvoted 1 times

KapilDhamija 2 months, 3 weeks ago**Selected Answer: A**

Vote goes to A
upvoted 1 times

AzureDP900 4 months, 4 weeks ago

It is no brainer questions, It is A.
upvoted 3 times

Tirthankar17 5 months ago

A obviously. No need to create a new service.
upvoted 1 times

pluidust 5 months, 1 week ago**Selected Answer: A**

Obviously A
upvoted 1 times

JelloMan 6 months, 1 week ago**Selected Answer: A**

A all the way. When deploying new versions you can split traffic for A/B testing to see how user base reacts to changes.
upvoted 4 times

amindbesideitself 6 months, 1 week ago**Selected Answer: A**

A is correct, there's no need to create a new service for each update.
upvoted 1 times

aabbcc321 6 months, 2 weeks ago

ans is A
upvoted 1 times

 **sivasan** 6 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

 **PAUGURU** 6 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 3 times

Question #186

You need to add a group of new users to Cloud Identity. Some of the users already have existing Google accounts. You want to follow one of Google's recommended practices and avoid conflicting accounts. What should you do?

- A. Invite the user to transfer their existing account.
- B. Invite the user to use an email alias to resolve the conflict.
- C. Tell the user that they must delete their existing account.
- D. Tell the user to remove all personal email from the existing account.

Correct Answer: A

Reference:

<https://support.google.com/cloudidentity/answer/7332836?hl=en>*Community vote distribution*

A (92%)	8%
---------	----

 **snkhatri** 2 months ago

Selected Answer: A

A obvious choice

upvoted 1 times

 **bobthebuilder55110** 3 months ago

Selected Answer: A

Correct Answer: A

upvoted 1 times

 **bobthebuilder55110** 3 months ago

Here is why ?

Question states "Some of the users already have existing Google accounts." Meaning they have personal account or any google account and what Option B is saying is to use aliases, as per google documentation this is only helpful when we want someone to receive emails in one inbox with 2 email names, meaning x@google.com and y@google.com goes to the same inbox BUT what you can't do is to have personal@google.com and company@google.com since the company wouldn't add you to their domain as that is not google recommended practice.

https://support.google.com/a/answer/33327?hl=en#when_to_use

upvoted 1 times

 **ryumada** 3 months ago

Selected Answer: A

Vote for A as the right answer. The docs in this link:

<https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

as provided by PAUGURU in his comment explains clearly about resolving account conflict. In the doc says nothing about to change email alias to resolve the conflict. So, following the documentation in that link means you are following the Googles Recommended Practices.

upvoted 2 times

 **zolthar_z** 3 months, 1 week ago

Selected Answer: A

A is the answer, for security reasons google best practices recommend transfer the account

upvoted 1 times

 **sai_learner** 3 months, 4 weeks ago

Selected Answer: B

Answer is B

<https://support.google.com/cloudidentity/answer/7062710>

upvoted 1 times

 **ryumada** 3 months ago

I am not sure if the link you provide explains the reason of the reason of your chosen answer. As in the documentation stated the email alias after what happen if you rename the email. Also, the documentation doesn't explain about account conflict.

Better with the docs link provided by PAUGURU: <https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

It's explains clearly about the conflicting email and best practices too.

upvoted 1 times

✉ **ryumada** 3 months ago

*explains the reason of your chosen answer.

sorry, messed up sentence

upvoted 1 times

✉ **bobthebuilder55110** 3 months ago

It should be A, I was confused with this as well but B is not relevant in this use case. Look at my Above comments.

upvoted 1 times

✉ **ramss** 4 months, 1 week ago

As per my understanding, B is the correct answer.

upvoted 2 times

✉ **AzureDP900** 4 months, 4 weeks ago

A is right and followed GCP documentation to get more information.

upvoted 1 times

✉ **ggupton1** 6 months ago

Selected Answer: A

<https://cloud.google.com/architecture/identity/assessing-existing-user-accounts>

If you want to maintain the access rights and some of the data associated with the Gmail account, you can ask the owner to remove Gmail from the user account so that you can then migrate them to Cloud Identity or Google Workspace.

upvoted 4 times

✉ **PAUGURU** 6 months, 2 weeks ago

Selected Answer: A

A looks right

<https://cloud.google.com/architecture/identity/migrating-consumer-accounts>

upvoted 3 times

Question #187

Topic 1

You need to manage a Cloud Spanner instance for best query performance. Your instance in production runs in a single Google Cloud region. You need to improve performance in the shortest amount of time. You want to follow Google best practices for service configuration. What should you do?

- A. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 45%. If you exceed this threshold, add nodes to your instance.
- B. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 45%. Use database query statistics to identify queries that result in high CPU usage, and then rewrite those queries to optimize their resource usage.
- C. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 65%. If you exceed this threshold, add nodes to your instance.
- D. Create an alert in Cloud Monitoring to alert when the percentage of high priority CPU utilization reaches 65%. Use database query statistics to identify queries that result in high CPU usage, and then rewrite those queries to optimize their resource usage.

Correct Answer: C

Reference:

<https://cloud.google.com/spanner/docs/instances>*Community vote distribution*

C (100%)

  **PAUGURU** Highly Voted 6 months, 2 weeks ago**Selected Answer: C**

C looks correct, increase instances on single region if CPU above 65%
<https://cloud.google.com/spanner/docs/cpu-utilization#recommended-max>

upvoted 6 times

  **Aninina** Most Recent 1 week, 2 days ago**Selected Answer: C**

<https://cloud.google.com/spanner/docs/cpu-utilization>

upvoted 1 times

  **snkhatri** 2 months ago**Selected Answer: C**

C looks correct
upvoted 1 times

  **AzureDP900** 4 months, 4 weeks ago

shortest timeframe is key here , I am going with C as my answer.
upvoted 1 times

  **Tirthankar17** 5 months ago**Selected Answer: C**

C is correct
upvoted 1 times

  **pfabio** 5 months, 1 week ago**Selected Answer: C**

Metric Maximum for single-region instances Maximum per region for multi-region instances
High priority total > 65% 45%
24-hour smoothed aggregate > 90% 90%
upvoted 1 times

  **rsuresh27** 6 months, 2 weeks ago**Selected Answer: C**

The correct answer is C
<https://cloud.google.com/spanner/docs/cpu-utilization#recommended-max>
upvoted 3 times

Question #188

Your company has an internal application for managing transactional orders. The application is used exclusively by employees in a single physical location. The application requires strong consistency, fast queries, and ACID guarantees for multi-table transactional updates. The first version of the application is implemented in PostgreSQL, and you want to deploy it to the cloud with minimal code changes. Which database is most appropriate for this application?

- A. BigQuery
- B. Cloud SQL
- C. Cloud Spanner
- D. Cloud Datastore

Correct Answer: C

Reference:

<https://cloud.google.com/spanner/docs/transactions>

Community vote distribution

B (57%) C (43%)

✉  **peugeotduke** Highly Voted 5 months, 3 weeks ago

Read the question :

The application is used exclusively by employees in a single physical location.

upvoted 10 times

✉  **theBestStudent** 1 month, 2 weeks ago

Correct. That is the key thing. I have no idea why some people ended up thinking cloud spanner is better. Definitely is alternative A.

upvoted 1 times

✉  **theBestStudent** 1 month, 2 weeks ago

Sorry I meant B, I had a typo.

upvoted 2 times

✉  **PAUGURU** Highly Voted 6 months, 2 weeks ago

Selected Answer: B

B -> minimal code changes

upvoted 8 times

✉  **cheeseburger12388** 5 months, 4 weeks ago

Cloud SQL for PostgreSQL is a fully-managed database service that helps you set up, maintain, manage, and administer your PostgreSQL relational databases on Google Cloud Platform.

<https://cloud.google.com/sql/docs/postgres>

upvoted 3 times

✉  **AwesomeGCP** Most Recent 3 weeks, 6 days ago

Selected Answer: B

B. Cloud SQL

and why C. Cloud Spanner NOT bcoz(The application is used exclusively by employees in a single physical location.)

upvoted 1 times

✉  **gcpreviewer** 1 month, 1 week ago

Selected Answer: B

Everything in the question screams SQL. Spanner is used for global scaling, not internally in a single location. B is the correct answer here.

upvoted 1 times

✉  **Cornholio_LMC** 1 month, 2 weeks ago

had this question today

upvoted 2 times

✉  **scjs** 1 month, 4 weeks ago

Selected Answer: B

Keywords: "a single physical location", "minimal code changes"

upvoted 1 times

✉  **snkhatri** 2 months ago

Selected Answer: B

B..seems right
upvoted 1 times

 **bobthebuilder55110** 3 months ago

Selected Answer: C

Answer is C:

Here is where google explains about different options and if you read closely you will see that cloud sql supports ACID but it does not say strong-consistency and you will see in cloud spanner it exclusively mentions strong-consistency. At first I thought why need a sword for needle since this questions mentions that it's only used in one location and exclusively by employees that fits cloud SQL easily but then I read the documentation on string-consistency + It does not say anything about keep the cost low so we can definitely pick cloud spanner here.
<https://cloud.google.com/blog/topics/developers-practitioners/your-google-cloud-database-options-explained>

upvoted 2 times

 **ryumada** 3 months ago

Selected Answer: C

ACID guarantee is only available at Cloud Spanner for relational database products on GCP.

- <https://cloud.google.com/sql>
- <https://cloud.google.com/spanner>

PostgreSQL also available as the interface for Cloud Spanner which could minimizes code changes. Just like akshaychavan7 has said in his comment.

- <https://cloud.google.com/spanner/docs/postgresql-interface>

upvoted 2 times

 **abirroy** 3 months ago

Selected Answer: B

Cloud spanner will be a over kill. Also, spanner is better for multiple geo.

upvoted 1 times

 **zolthar_z** 3 months, 1 week ago

Selected Answer: B

Answer is B, <https://cloud.google.com/spanner/docs/migrating-postgres-spanner>

upvoted 1 times

 **JasmineS1975** 3 months, 1 week ago

Selected Answer: C

Google docs only uses the term "strong consistency" with Cloud Spanner.

upvoted 1 times

 **yandieg** 3 months, 3 weeks ago

Selected Answer: B

It is B.

They want minimum change, use G Spanner required to change application code

upvoted 1 times

 **samsonakala** 3 months, 4 weeks ago

Cloud SQL...very easy and straightforward

upvoted 1 times

 **jitubaba80** 4 months ago

Selected Answer: C

I will go with Option C Cloud Spanner because the question is saying application requires strong consistency

upvoted 1 times

 **Roro_Brother** 4 months, 1 week ago

Selected Answer: C

The correct answer is C, for sure

upvoted 1 times

 **johnconnor** 4 months, 2 weeks ago

I am pretty sure it is B, all Google relational databases offer ACID consistency mode for the data, which means:

Atomic: All operations in a transaction succeed or the operation is rolled back.

Consistent: On the completion of a transaction, the database is structurally sound.

Isolated: Transactions do not contend with one another. Contentious access to data is moderated by the database so that transactions appear to run sequentially.

Durable: The results of applying a transaction are permanent, even in the presence of failures

It is for internal use, no requirements of anything global and we want something simply with minimum code changes.

upvoted 2 times

Question #189

You are assigned to maintain a Google Kubernetes Engine (GKE) cluster named 'dev' that was deployed on Google Cloud. You want to manage the GKE configuration using the command line interface (CLI). You have just downloaded and installed the Cloud SDK. You want to ensure that future CLI commands by default address this specific cluster. What should you do?

- A. Use the command gcloud config set container/cluster dev.
- B. Use the command gcloud container clusters update dev.
- C. Create a file called gke.default in the ~/.gcloud folder that contains the cluster name.
- D. Create a file called defaults.json in the ~/.gcloud folder that contains the cluster name.

Correct Answer: A

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-access-for-kubectl#gcloud-config>*Community vote distribution*

A (92%) 8%

snkhatri 2 months ago**Selected Answer: A**

A looks right to me

upvoted 1 times

73173v2 2 months, 1 week ago**Selected Answer: A**

To set a default cluster for gcloud commands, run the following command:

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters>

upvoted 2 times

AzureDP900 4 months, 4 weeks ago

A is right

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

ggupton1 6 months ago**Selected Answer: A**

Set a default cluster for gcloud

To set a default cluster for commands gcloud, run the following command:

Per <https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=fr>

gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

Akash7 6 months ago

Answer is A,

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=en>

upvoted 1 times

MadMikedD 6 months ago**Selected Answer: A**

To set a default cluster for gcloud commands, run the following command:

gcloud config set container/cluster CLUSTER_NAME

upvoted 3 times

cheeseburger12388 5 months, 4 weeks agohttps://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters#default_cluster_kubectl

upvoted 1 times

aswinachu 6 months ago**Selected Answer: B**

Correct Ans B

<https://cloud.google.com/sdk/gcloud/reference/container/clusters/update>

upvoted 1 times

 **VympelGRU** 6 months, 1 week ago

B sounds right.

upvoted 2 times

 **Malib** 6 months, 1 week ago

Selected Answer: A

<https://cloud.google.com/kubernetes-engine/docs/how-to/managing-clusters?hl=fr>

upvoted 3 times

Question #190

The sales team has a project named Sales Data Digest that has the ID acme-data-digest. You need to set up similar Google Cloud resources for the marketing team but their resources must be organized independently of the sales team. What should you do?

- A. Grant the Project Editor role to the Marketing team for acme-data-digest.
- B. Create a Project Lien on acme-data-digest and then grant the Project Editor role to the Marketing team.
- C. Create another project with the ID acme-marketing-data-digest for the Marketing team and deploy the resources there.
- D. Create a new project named Marketing Data Digest and use the ID acme-data-digest. Grant the Project Editor role to the Marketing team.

Correct Answer: A

Community vote distribution

C (100%)

 **gcpj** Highly Voted 4 months, 1 week ago

Selected Answer: C

Answer should be C because the resources for the marketing team should be independent from the Sales team. Resources are tied and separated by projects.

upvoted 5 times

 **Priyanka109** Most Recent 1 month, 1 week ago

Sales and Marketing resource has to be in separate project so C.

upvoted 1 times

 **snkhatri** 2 months ago

Selected Answer: C

C looks right

upvoted 1 times

 **ATANGA** 3 months ago

WhatsApp : +1(956)-520-4006 to obtain PMP, CISM, CCNA, CEH, PRINCE2, CISCO, ISTQB, PRINCE2, AWS/Azure/Sale force/ITIL Foundation/EC- COUNCIL...

Get Certified with 100% pass guarantee. Pay after exam.

all CISCO, ISACA & EC- COUNCIL certifications

For the Below certificates we offer 100% pass guarantee:

1. AWS Certification
2. Sales force
3. Scrum Master
4. Oracle Certification: OCA, OCP
5. Cisco Certification: CCNA, CCNP, CCIE
6. ITIL Foundation & Intermediate
7. Prince 2 Foundation and Practitioner
8. VMWARE Certification
9. Check Point Certification
10. EC-COUNCIL Certification (CEH V-9, CCISO, CND)
11. Cloud Certification
12. IBM Certification
13. HP Certification
14. Citrix Certification
15. Juniper certification
16. Azure
17. Skype 70-333/34
18. PMI (PMP/CAPM/ACP/PBA)
19. ISTQB
20. SAP
21. ISACA (CISA, CISM, CRISC, CGEIT, COBIT)

PAYMENT ONLY AFTER CERTIFICATION AND RESULT CONFIRMATION.

WhatsApp : +1(956)-520-4006

upvoted 1 times

 **Gautam_Thampy** 1 month, 2 weeks ago

what the heck is this?... good job mods

upvoted 1 times

 **XDcloud** 3 months, 1 week ago

Selected Answer: C

C for sure

upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

C is right answer based on given scenario..

upvoted 2 times

 **JelloMan** 6 months ago

C is straight to the point and addresses exactly all of the concerns

upvoted 3 times

 **jblima** 6 months, 1 week ago

C is the better choice

upvoted 3 times

 **amindbesideitself** 6 months, 1 week ago

Selected Answer: C

C, other options seem obviously wrong

upvoted 2 times

 **Terzlightyear** 6 months, 2 weeks ago

Selected Answer: C

C seems the right answer

upvoted 3 times

 **PAUGURU** 6 months, 2 weeks ago

Selected Answer: C

Looks more like C

upvoted 2 times

Question #191

You have deployed multiple Linux instances on Compute Engine. You plan on adding more instances in the coming weeks. You want to be able to access all of these instances through your SSH client over the internet without having to configure specific access on the existing and new instances. You do not want the Compute Engine instances to have a public IP. What should you do?

- A. Configure Cloud Identity-Aware Proxy for HTTPS resources.
- B. Configure Cloud Identity-Aware Proxy for SSH and TCP resources
- C. Create an SSH keypair and store the public key as a project-wide SSH Key.
- D. Create an SSH keypair and store the private key as a project-wide SSH Key.

Correct Answer: B*Community vote distribution*

B (67%)	C (25%)	8%
---------	---------	----

 **Akash7** Highly Voted 6 months, 2 weeks ago

B is correct as question say no public IP on the instance.
upvoted 8 times

 **Akash7** 6 months ago

Use IAP TCP to enable access to VM instances that do not have external IP addresses or do not permit direct access over the internet.
<https://cloud.google.com/iap/docs/using-tcp-forwarding>
upvoted 9 times

 **Untamables** Most Recent 1 week, 5 days ago

Selected Answer: B

Absolutely B
https://cloud.google.com/iap/docs/using-tcp-forwarding#tunneling_ssh_connections
upvoted 1 times

 **Gautam_Thampy** 1 month, 2 weeks ago

Selected Answer: B

b is right
upvoted 1 times

 **snkhatri** 2 months ago

Selected Answer: B

B looks right
upvoted 1 times

 **AzureDP900** 4 months, 4 weeks ago

B is correct, With TCP forwarding, IAP can protect SSH and RDP access to your VMs hosted on Google Cloud. Your VM instances don't even need public IP addresses.
upvoted 3 times

 **Rutu_98** 5 months, 2 weeks ago

Selected Answer: B

B is correct
upvoted 2 times

 **lixamec** 5 months, 3 weeks ago

Selected Answer: B

I think it is B
<https://medium.com/google-cloud/how-to-ssh-into-your-gce-machine-without-a-public-ip-4d78bd23309e>
upvoted 1 times

 **CloudAce7890** 5 months, 4 weeks ago

Selected Answer: B

B is correct as it uses IAP
upvoted 2 times

 **JelloMan** 6 months ago

Selected Answer: C

IAP lets you establish a central authorization layer for applications accessed by HTTPS. This statement immediately eliminates A & B since you would need to publicly access the instance. C is also incorrect because it uses a private SSH key. Private keys are only meant for the user themselves as proof of their identity. Public SSH keys are meant to be used for access within applications, so that is the most applicable in this case. D, final answer.

upvoted 1 times

 **JelloMan** 6 months ago

Selected Answer: D

IAP lets you establish a central authorization layer for applications accessed by HTTPS. This statement immediately eliminates A & B since you would need to publicly access the instance. C is also incorrect because it uses a private SSH key. Private keys are only meant for the user themselves as proof of their identity. Public SSH keys are meant to be used for access within applications, so that is the most applicable in this case. D, final answer.

upvoted 1 times

 **jeffangel28** 3 months, 2 weeks ago

<https://cloud.google.com/iap/docs/using-tcp-forwarding>
read pls!!

upvoted 1 times

 **JelloMan** 6 months ago

Scratch this. Made a mistake. C is the correct answer (thought they were flipped)

upvoted 1 times

 **dark_3k03r** 5 months, 2 weeks ago

The question states "You do not want the Compute Engine instances to have a public IP", so that knocks out C and D as both options require public access.

IAP however supports port forwarding for your client so the instances are never exposed. That leaves A and B. With the question explicitly stating "your ssh client", that means you need to configure the Cloud Identity Aware proxy for ssh (i.e. port forwarding. Touching the proxies is not "configure specific access on the existing and new instances" as all this occurs on IAP and not the compute engine api.

upvoted 5 times

 **PAUGURU** 6 months, 2 weeks ago

Selected Answer: C

C looks better

upvoted 2 times

Question #192

You have created an application that is packaged into a Docker image. You want to deploy the Docker image as a workload on Google Kubernetes Engine. What should you do?

- A. Upload the image to Cloud Storage and create a Kubernetes Service referencing the image.
- B. Upload the image to Cloud Storage and create a Kubernetes Deployment referencing the image.
- C. Upload the image to Container Registry and create a Kubernetes Service referencing the image.
- D. Upload the image to Container Registry and create a Kubernetes Deployment referencing the image.

Correct Answer: B

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>*Community vote distribution*

D (100%)

  **Aninina** 1 week, 2 days ago**Selected Answer: D**

A deployment is responsible for keeping a set of pods running. A service is responsible for enabling network access to a set of pods.
upvoted 1 times

  **Priyanka109** 1 month, 1 week ago

Upload your docker image on container registry then give a ref while creating deployment. So D!
upvoted 1 times

  **rimjhim09** 1 month, 2 weeks ago**Selected Answer: D**

I also vote for D. I passed my exam today and this question was there.
upvoted 2 times

  **Gautam_Thampy** 1 month, 2 weeks ago

its D , A and B are obviously incorrect
upvoted 2 times

  **pkmdb66** 1 month, 4 weeks ago**Selected Answer: D**

It's D
upvoted 1 times

  **snkhatri** 2 months ago**Selected Answer: D**

D is right
upvoted 1 times

  **Sakhi1234** 2 months, 1 week ago**Selected Answer: D**

I have hands on expereince to answer this question.
upvoted 1 times

  **learn_GCP** 2 months, 2 weeks ago**Selected Answer: D**

D is the Answer
upvoted 1 times

  **Bootsdale** 4 months, 3 weeks ago**Selected Answer: D**

D - not even a debate!
upvoted 1 times

  **AzureDP900** 4 months, 4 weeks ago

D perfect answer for given scenario.
upvoted 1 times

  **akshaychavan7** 5 months, 2 weeks ago

Selected Answer: D

D!!!!!!

upvoted 1 times

  **JelloMan** 6 months, 1 week ago**Selected Answer: D**
You can only create GKE instances for docker containers through container registry. D all the way

upvoted 4 times

  **amindbesideitself** 6 months, 1 week ago**Selected Answer: D**
D is best for storing and deploying containers to GKE

upvoted 2 times

  **Akash7** 6 months, 2 weeks ago

Correct answer is D.

upvoted 3 times

Question #193

You are using Data Studio to visualize a table from your data warehouse that is built on top of BigQuery. Data is appended to the data warehouse during the day.

At night, the daily summary is recalculated by overwriting the table. You just noticed that the charts in Data Studio are broken, and you want to analyze the problem. What should you do?

- A. Review the Error Reporting page in the Cloud Console to find any errors.
- B. Use the BigQuery interface to review the nightly job and look for any errors.
- C. Use Cloud Debugger to find out why the data was not refreshed correctly.
- D. In Cloud Logging, create a filter for your Data Studio report.

Correct Answer: C*Community vote distribution*

D (45%)	C (27%)	A (18%)	9%
---------	---------	---------	----

 **Untamables** 1 week, 5 days ago

Selected Answer: A

I vote A.
<https://cloud.google.com/error-reporting/docs/check-app-crash>
 upvoted 1 times

 **PKookNN** 2 weeks, 1 day ago

Selected Answer: D

not sure but I think D is the most appropriate
 upvoted 1 times

 **btimmons** 2 weeks, 2 days ago

The answer is D
 upvoted 1 times

 **Erie** 2 weeks, 4 days ago

Selected Answer: D

D because cloud logging is enabled by default. Cloud Debugger is something we need to configure manually in our code.
 upvoted 1 times

 **diasporabro** 2 weeks, 5 days ago

Selected Answer: A

Hmmm.... Between A and D, I'll just select A and move on...
 upvoted 1 times

 **diasporabro** 2 weeks, 5 days ago

I take that back, will select D... seems to be a more direct approach to analyze...
 upvoted 1 times

 **AwesomeGCP** 3 weeks, 6 days ago

Selected Answer: B

B. Use the BigQuery interface to review the nightly job and look for any errors.
 upvoted 1 times

 **theBestStudent** 1 month, 2 weeks ago

Selected Answer: D

On top of being almost deprecated, I'm not even sure if cloud debugger would be suitable for bigquery or cloud data studio stuff.
 Error reporting I'm not even sure if it supports error reporting for data studio stuff.

For my sounds more simpler to just create a filter to see the logs in

"cloud logging" seems simpler, faster are more suitable solution. I go with D.
 upvoted 3 times

 **arj4c3** 1 month, 2 weeks ago

Selected Answer: C

C->
 Cloud Debugger helps inspect the state of an application, at any code location, without stopping or slowing down the running app //

<https://cloud.google.com/debugger/docs>

A -> <https://cloud.google.com/error-reporting>

quote // Error Reporting counts, analyzes, and aggregates the crashes in your running cloud services. A centralized error management interface displays the results with sorting and filtering capabilities. //

Here the Data Studio charts seem to be broken this occurs because we overwrite the table at night. so each day the source data (from morning) is overwritten resulting in charts being broken. So the issue seems to be caused by the application that overwrites the table at night so i think answer should be C

B only gives you the data written at night so you cant find errors there
D seems irrelevant since it only gives you logs

So ans should be C, Cloud Debugger

upvoted 1 times

✉  **arj4c3** 1 month, 2 weeks ago

Datastudio errors -> <https://portermetrics.com/en/tutorial/google-data-studio-common-errors/>

upvoted 1 times

✉  **DualCore573** 1 month, 2 weeks ago

What is the correct answer?

upvoted 1 times

✉  **ace_2021_1** 1 month, 3 weeks ago

I don't agree with C.....Cloud Debugger is a feature of Google Cloud that lets you inspect the state of a running application, in real-time, without stopping or slowing it. the Answer is A. Error Reporting counts, analyzes, and aggregates the errors in your running cloud services.

upvoted 1 times

✉  **Kolonelpapa** 2 months ago

I would go for BQ interface: <https://cloud.google.com/bigquery/docs/managing-jobs#view-job>. Cloud Debugger is for inspecting apps: <https://cloud.google.com/debugger/docs>.

upvoted 4 times

✉  **snkhatri** 2 months ago

Selected Answer: C

C seems ok

upvoted 2 times

Question #194

You have been asked to set up the billing configuration for a new Google Cloud customer. Your customer wants to group resources that share common IAM policies. What should you do?

- A. Use labels to group resources that share common IAM policies.
- B. Use folders to group resources that share common IAM policies.
- C. Set up a proper billing account structure to group IAM policies.
- D. Set up a proper project naming structure to group IAM policies.

Correct Answer: B*Community vote distribution*

B (100%)

  **diasporabro** 2 weeks, 5 days ago**Selected Answer: B**

"Folders are used to group resources that share common IAM policies."

upvoted 2 times

  **Charumathi** 1 month ago**Selected Answer: B**

B is correct Answer,

Folders are nodes in the Cloud Platform Resource Hierarchy. A folder can contain projects, other folders, or a combination of both. Organizations can use folders to group projects under the organization node in a hierarchy. For example, your organization might contain multiple departments, each with its own set of Google Cloud resources. Folders allow you to group these resources on a per-department basis. Folders are used to group resources that share common IAM policies. While a folder can contain multiple folders or resources, a given folder or resource can have exactly one parent.

<https://cloud.google.com/resource-manager/docs/creating-managing-folders>

upvoted 3 times

  **zellck** 1 month, 3 weeks ago**Selected Answer: B**

B is the answer.

<https://cloud.google.com/resource-manager/docs/access-control-folders#best-practices-folders-iam>

upvoted 1 times

  **RockingRohit6** 1 month, 3 weeks ago

Folders are used to group resources that share common IAM policies

upvoted 1 times

  **thimai** 2 months ago**Selected Answer: B**

B for me

"Folders are used to group resources that share common IAM policies"

<https://cloud.google.com/resource-manager/docs/creating-managing-folders>

upvoted 3 times

  **snkhatri** 2 months ago**Selected Answer: B**

B seems right to me

upvoted 1 times

Question #195

You have been asked to create robust Virtual Private Network (VPN) connectivity between a new Virtual Private Cloud (VPC) and a remote site. Key requirements include dynamic routing, a shared address space of 10.19.0.1/22, and no overprovisioning of tunnels during a failover event. You want to follow Google-recommended practices to set up a high availability Cloud VPN. What should you do?

- A. Use a custom mode VPC network, configure static routes, and use active/passive routing.
- B. Use an automatic mode VPC network, configure static routes, and use active/active routing.
- C. Use a custom mode VPC network, use Cloud Router border gateway protocol (BGP) routes, and use active/passive routing.
- D. Use an automatic mode VPC network, use Cloud Router border gateway protocol (BGP) routes, and configure policy-based routing.

Correct Answer: D*Community vote distribution*

C (100%)

 **Charumathi** 1 month ago

Selected Answer: C

C . Choose a Cloud VPN gateway that uses dynamic routing and the Border Gateway Protocol (BGP). Google recommends using HA VPN and deploying on-premises devices that support BGP.

Choose the appropriate tunnel configuration

Choose the appropriate tunnel configuration based on the number of HA VPN gateways:

If you have a single HA VPN gateway, use an active/passive tunnel configuration.

If you have more than one HA VPN gateway, use an active/active tunnel configuration.

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/best-practices>

upvoted 2 times

 **ale_brd_** 1 month ago

Selected Answer: C

c is the correct one

upvoted 1 times

 **theBestStudent** 1 month, 2 weeks ago

Selected Answer: C

we need custom mode vpc so subnets are not created automatically (the ip range is mentioned in the question) also we will need active/passive HA VPN (as it is not mentioned we will have to use more than one HA VPN gateway).

Links : <https://cloud.google.com/network-connectivity/docs/vpn/concepts/best-practices>

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview#active>

<https://cloud.google.com/vpc/docs/vpc#subnet-ranges>

upvoted 1 times

 **theBestStudent** 1 month, 2 weeks ago

Also for dynamic routing we need HA VPN

Link: <https://cloud.google.com/network-connectivity/docs/vpn/concepts/choosing-networks-routing#dynamic-routing>

upvoted 1 times

 **Arulkumar** 1 month, 2 weeks ago

Selected Answer: C

Google Cloud Router

On Google Cloud, dynamic routing can be established using Cloud Router. It exchanges network topology information through Border Gateway Protocol (BGP). Cloud Router advertises subnets from its VPC network to another router or gateway via BGP. This is great for setting up VPN between the cloud and on-prem, as topology changes automatically propagate with no manual intervention and higher redundancy for your systems.

You now have:

Discovery of remote networks

Maintaining up-to-date routing information

Choosing the best path to destination networks

Ability to find a new best path if the current path is no longer available

And a great side effect can be lower latency because Cloud Router learns routes through BGP which allows for optimal data paths to reach its destination, whether that be another network or a VPN gateway to on-premise. Cloud Router is also how Dedicated Interconnect can give you 10 gbps bandwidth between your cloud VPC and your peered on-premise data center.

upvoted 2 times

 **manjtrade2** 1 month, 3 weeks ago

Selected Answer: C

C might be right

upvoted 1 times

 **snkhatri** 2 months ago

Selected Answer: C

I think it should be C as there is too much customisation.

upvoted 1 times

Question #196

Topic 1

You are running multiple microservices in a Kubernetes Engine cluster. One microservice is rendering images. The microservice responsible for the image rendering requires a large amount of CPU time compared to the memory it requires. The other microservices are workloads that are optimized for n1-standard machine types. You need to optimize your cluster so that all workloads are using resources as efficiently as possible. What should you do?

- A. Assign the pods of the image rendering microservice a higher pod priority than the other microservices.
- B. Create a node pool with compute-optimized machine type nodes for the image rendering microservice. Use the node pool with general-purpose machine type nodes for the other microservices.
- C. Use the node pool with general-purpose machine type nodes for the image rendering microservice. Create a node pool with compute-optimized machine type nodes for the other microservices.
- D. Configure the required amount of CPU and memory in the resource requests specification of the image rendering microservice deployment. Keep the resource requests for the other microservices at the default.

Correct Answer: B*Community vote distribution*

B (100%)

 **diasporabro** 2 weeks, 5 days ago

Selected Answer: B

B looks like the right choice here

upvoted 1 times

 **learn_GCP** 1 month, 1 week ago

Selected Answer: B

B. is the Answer

upvoted 1 times

 **adarsh4503** 1 month, 2 weeks ago

I agree B is the answer.

upvoted 2 times

 **zellck** 1 month, 2 weeks ago

Selected Answer: B

B is the answer.

upvoted 1 times

 **Gautam_Thampy** 1 month, 2 weeks ago

Selected Answer: B

C is not correct coz general purpose machine types will not suffice for image rendering.

B is the most suitable answer.

upvoted 2 times

 **osanchez** 1 month, 3 weeks ago

B is correct

upvoted 1 times

Question #197

Your organization has three existing Google Cloud projects. You need to bill the Marketing department for only their Google Cloud services for a new initiative within their group. What should you do?

- A. 1. Verify that you are assigned the Billing Administrator IAM role for your organization's Google Cloud Project for the Marketing department.
2. Link the new project to a Marketing Billing Account.
- B. 1. Verify that you are assigned the Billing Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Set the default key-value project labels to department:marketing for all services in this project.
- C. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Link the new project to a Marketing Billing Account.
- D. 1. Verify that you are assigned the Organization Administrator IAM role for your organization's Google Cloud account. 2. Create a new Google Cloud Project for the Marketing department. 3. Set the default key-value project labels to department:marketing for all services in this project.

Correct Answer: A*Community vote distribution*

A (53%)

B (47%)

 **moitsu** Highly Voted 2 months ago

Selected Answer: B

Between A& B, Billing Administrator IAM role is either at the organisation level not project level. Hence A is out. C & D doesn't make sense.
upvoted 5 times

 **Gautam_Thampy** 1 month, 2 weeks ago

The billing account administrator role can also be given at the project level. A is correct. Refer to this doc:
<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times

 **Gautam_Thampy** 1 month, 2 weeks ago

Correction I meant the billing account admin role can be given at the organisation or the billing account level.
upvoted 1 times

 **gneknurd** Most Recent 1 week, 4 days ago

This question was on the test which I passed. I chose A

upvoted 2 times

 **PKookNN** 2 weeks, 1 day ago

Selected Answer: B

it was either A or B but since Billing Admin is only granted at org level or billing level so A is out. Correct A (C & D are not correct)
upvoted 1 times

 **kishorparihar07** 3 weeks, 2 days ago

Most correct Ans is : A.

Other than A ,every remaining option is stating to create a new project with Role of Billing Account Administrator role which is not possible. You need Project Creator role to create new Project.

upvoted 1 times

 **ale_brd_** 1 month ago

Selected Answer: A

roles/billing.admin role is an owner role for a billing account. Use it to manage payment instruments, configure billing exports, view cost information, link and unlink projects and manage other user roles on the billing account.

nowhere is mentioned you need to create a new project, projects are already created.

upvoted 1 times

 **gcpreviewer** 1 month, 1 week ago

Selected Answer: A

I understand that the question implies the creation of a new project, however neither of the roles listed have that functionality. If you chose B you are choosing an answer that has a direct contradiction because the Billing Account Admin does not have the permissions to create a new project. Thus, I think it is better to assume the new initiative/project is already created or being created by someone else and your job is simply to link the project to the account which you do have the appropriate permissions to perform.

A is my choice.

upvoted 4 times

✉ **Pushpen_27** 1 month, 2 weeks ago

Selected Answer: A

A seems better choice

upvoted 1 times

✉ **arj4c3** 1 month, 2 weeks ago

Selected Answer: A

Option A seems the only valid answer as Billing Admin can link a new project to a billing acc

B & D assumes you have the resourcemanager.projects.create permission to create a new project. According to <https://cloud.google.com/resource-manager/docs/creating-managing-labels#permissions> resourcemanager.projects.update permission is needed to edit the labels

neither Billing Administrator nor Organization Adminstrator has resourcemanager.projects.update permission to edit the labels they only can see the labels of a project as both have only resourcemanager.projects.get permission so Options B & D are both wrong

C is wrong because neither does Organizational admin have resourcemanager.projects.create permission to create a project nor does it have permission to link a project to a billing account - > <https://cloud.google.com/iam/docs/understanding-roles#resourcemanager.organizationAdmin>

so my final answer is A

upvoted 1 times

✉ **arj4c3** 1 month, 2 weeks ago

We assume that the question means a new project by mentioning a "new initiatve"

upvoted 1 times

✉ **Gautam_Thamby** 1 month, 2 weeks ago

i disagree get rekt noob

upvoted 1 times

✉ **ace_2021_1** 1 month, 3 weeks ago

There is no billing administrator role.....Billing Account Administrator

(roles/billing.admin) Manage billing accounts (but not create them). Organization or billing account. This role is an owner role for a billing account. Use it to manage payment instruments, configure billing exports, view cost information, link and unlink projects and manage other user roles on the billing account.

upvoted 1 times

✉ **ace_2021_1** 1 month, 3 weeks ago

Billing Administrator IAM role creates a new project? in what universe?

upvoted 1 times

✉ **jaravu** 1 month, 3 weeks ago

B is correct...

upvoted 1 times

✉ **balamuthusamy** 1 month, 3 weeks ago

Question says: Your organization has three existing Google Cloud projects, then why do we need to create separate Project for Marketing department?

upvoted 1 times

✉ **alleinallein** 1 month, 4 weeks ago

Selected Answer: B

B in my opinion, A is wrong cause there is nothing said there about creation of new project

upvoted 1 times

✉ **anggamatt** 1 month, 4 weeks ago

Selected Answer: A

it should be A

upvoted 1 times